The syslog-ng Open Source Edition 3.6 Administrator Guide

Publication date December 09, 2014

Abstract
This manual is the primary documentation of the syslog-ng Open Source Edition 3.6 application.

Table of Contents

List of Examples

List of Procedures

Preface

Welcome to the syslog-ng Open Source Edition 3.6 Administrator Guide!

This document describes how to configure and manage syslog-ng. Background information for the technology and concepts used by the product is also discussed.

1. Summary of contents

Chapter 1, Introduction to syslog-ng (p. 1) describes the main functionality and purpose of syslog-ng OSE.

Chapter 2, The concepts of syslog-ng (p. 4) discusses the technical concepts and philosophies behind syslog-ng OSE.

Chapter 3, Installing syslog-ng (p. 23) describes how to install syslog-ng OSE on various UNIX-based platforms using the precompiled binaries.

Chapter 4, The syslog-ng OSE quick-start guide (p. 34) provides a briefly explains how to perform the most common log collecting tasks with syslog-ng OSE.

Chapter 5, The syslog-ng OSE configuration file (p. 40) discusses the configuration file format and syntax in detail, and explains how to manage large-scale configurations using included files and reusable configuration snippets.

Chapter 6, Collecting log messages — sources and source drivers (p. 50) explains how to collect and receive log messages from various sources.

Chapter 7, Sending and storing log messages — destinations and destination drivers (p. 116) describes the different methods to store and forward log messages.

Chapter 8, Routing messages: log paths and filters (p. 198) explains how to route and sort log messages, and how to use filters to select specific messages.

Chapter 9, Global options of syslog-ng OSE (p. 217) lists the global options of syslog-ng OSE and explains how to use them.

Chapter 10, TLS-encrypted message transfer (p. 229) shows how to secure and authenticate log transport using TLS encryption.

Chapter 11, Manipulating messages (p. 238) describes how to customize message format using templates and macros, how to rewrite and modify messages, and how to use regular expressions.

Chapter 12, Parsing and segmenting structured messages (p. 267) describes how to segment and process structured messages like comma-separated values.

Chapter 13, Processing message content with a pattern database (p. 275) explains how to identify and process log messages using a pattern database.

Chapter 14, Statistics of syslog-ng (p. 297) details the available statistics that syslog-ng OSE collects about the processed log messages.

Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300) describes how to configure syslog-ng OSE to use multiple processors, and how to optimize its performance.

Chapter 16, Troubleshooting syslog-ng (p. 303) offers tips to solving problems.

Chapter 17, Best practices and examples (p. 306) gives recommendations to configure special features of syslog-ng OSE.

Appendix A, The syslog-ng manual pages (p. 310) contains the manual pages of the syslog-ng OSE application.

Appendix C, GNU Lesser General Public License (p. 341) includes the text of the LGPLv2.1 license applicable to the core of syslog-ng Open Source Edition.

Appendix B, GNU General Public License (p. 335) includes the text of the GPLv2 license applicable to syslog-ng Open Source Edition.

Appendix D, Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License (p. 350) includes the text of the Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License applicable to The syslog-ng Open Source Edition 3.6 Administrator Guide.

Glossary (p. 355) defines the important terms used in this guide.

List of syslog-ng OSE parameters (p. 371) provides cross-references to the definitions of options, parameters, and macros available in syslog-ng OSE.

The *Index* provides cross-references to important terms used in this guide.

2. Target audience and prerequisites

This guide is intended for system administrators and consultants responsible for designing and maintaining logging solutions and log centers. It is also useful for IT decision makers looking for a tool to implement centralized logging in heterogeneous environments.

The following skills and knowledge are necessary for a successful syslog-ng administrator:

- At least basic system administration knowledge.
- An understanding of networks, TCP/IP protocols, and general network terminology.
- Working knowledge of the UNIX or Linux operating system.
- In-depth knowledge of the logging process of various platforms and applications.
- An understanding of the *legacy syslog (BSD-syslog) protocol*) and the *new syslog (IETF-syslog) protocol*) standard.

3. Products covered in this guide

This guide describes the use of the following products:

- syslog-ng Open Source Edition (syslog-ng OSE) 3.6.1 and later

4. Typographical conventions

Before you start using this guide, it is important to understand the terms and typographical conventions used in the documentation. For more information on specialized terms and abbreviations used in the documentation, see the Glossary at the end of this document.

The following kinds of text formatting and icons identify special information in the document.

Tip
Tips provide best practices and recommendations.

Note
Notes provide additional information on a topic, and emphasize important facts and considerations.

Warning
Warnings mark situations where loss of data or misconfiguration of the device is possible if the instructions are not obeyed.

`Command`	Commands you have to execute.
Emphasis	Reference items, additional readings.
`/path/to/file`	File names.
Parameters	Parameter and attribute names.
Label	GUI output messages or dialog labels.
Menu	A submenu or menu item in the menu bar.
Button	Buttons in dialog windows.

5. Contact and support information

This product is developed and maintained by BalaBit-Europe. We are located in Budapest, Hungary. Our address is:

BalaBit S.a.r.l.
2 Alíz Street
H-1117 Budapest, Hungary
Tel: +36 1 398-6700
Fax: +36 1 208-0875

E-mail: `<info@balabit.com>`
Web: *http://www.balabit.com/*

5.1. Sales contact

You can directly contact us with sales related topics at the e-mail address `<sales@balabit.com>`, or *leave us your contact information and we call you back*.

5.2. Support contact

In case you experience a problem that is not covered in this guide, *visit the syslog-ng wiki* or *post it on syslog-ng mailing list*.

To report bugs found in syslog-ng OSE, *visit this page*.

Precompiled binary packages are available for free from various third-parties. Most of these are listed on the *syslog-ng website*.

> **Note**
> BalaBit-Europe offers a Premium Edition of syslog-ng and a log management appliance syslog-ng Store Box for which professional support is available. For more information on which syslog-ng product version meets your needs contact our sales department at `<sales@balabit.com>`.
>
> For other products of BalaBit-Europe visit *http://www.balabit.com/*.

5.3. Training

BalaBit S.a.r.l. holds courses on using its products for new and experienced users. For dates, details, and application forms, visit the *http://www.balabit.com/support/trainings/* webpage.

6. About this document

This guide is a work-in-progress document with new versions appearing periodically.

The latest version of this document can be downloaded from the BalaBit website *here*.

6.1. Summary of changes

This section lists the changes of The syslog-ng Open Source Edition Administrator Guide.

6.1.1. Version 3.5 - 3.6

Changes in product:

Changes in documentation:

■ *Section 7.10, Monitoring your data with Riemann (p. 154)* has been added to the document.

- *Section 6.5, Receiving JSON messages from nodejs applications (p. 69)* has been added to the document.

- *Section 6.12, Collecting messages from the systemd-journal system log storage (p. 97)* has been added to the document.

- *Section 6.13, Collecting systemd messages using a socket (p. 99)* has been added to the document.

- *Section use-rcptid() (p. 228)* has been added to the document.

- *Section 11.2.4, Setting multiple message fields to specific values (p. 260)* has been added to the document.

- The `retries` and `throttle` options are available for the SMTP, MongoDB, AMQP, and Redis destinations.

- The description of the `multi-line-mode` option has been updated.

- *Section 6.15.1, UNIX credentials and other metadata (p. 109)* has been added to the document.

- *Section RUNID (p. 245)* has been added to *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*.

- The `extract-prefix` option has been added to *Section 12.3, The JSON parser (p. 272)*.

- The `graphite-output`, `or` and `padding` template functions have been added to *Section 11.1.7, Template functions of syslog-ng OSE (p. 248)*.

- PCRE is now a required dependency of syslog-ng OSE, and by default, syslog-ng OSE uses PCRE-style regular expressions. Therefore, the `--enable-pcre` compliation option has been removed.

- *Section 7.3, Sending metrics to Graphite (p. 128)* has been added to the document.

- *Section 7.8, pseudofile() (p. 152)* has been added to the document.

- The `custom-domain()` and `stats-lifetime()` options have been added to *Section 9.2, Global options (p. 217)*.

- The `retry_sql_inserts` option has been renamed to `retries` to increase consistency.

- *Section on-error() (p. 224)* can be set locally for MongoDB destinations as well. Also, MongoDB destinations support the `username` and `password` options, and connecting to the server using UNIX domain sockets. For details, see *Section 7.4, Storing messages in a MongoDB database (p. 129)*.

- *Procedure 7.4.1, How syslog-ng OSE connects the MongoDB server (p. 130)* has been added to the document.

- Several typos and syntax errors in examples have been corrected.

6.1.2. Version 3.4 - 3.5

Changes in product:

- *Section 7.13, Publishing messages using STOMP (p. 173)* has been added to the document.

- *Section 7.9, Storing name-value pairs in Redis (p. 152)* has been added to the document.
- The `replace()` option of `value-pairs()` has been renamed to `replace-prefix()` in *Section value-pairs() (p. 19)*.

- *Section 2.10.1, Specifying data types in value-pairs (p. 17)* has been added to the document.
- The `multi-line-mode()` option has been added to *Section 6.3.2, file() source options (p. 54)* and *Section 6.6.1, pipe() source options (p. 71)*.

- *Section inlist() (p. 213)* has been added to the document.
- *Section env (p. 248)* has been added to the document.
- *Section lowercase (p. 253)* has been added to the document.
- *Section uppercase (p. 256)* has been added to the document.
- *Section replace-delimiter (p. 254)* has been added to the document.
- The `ip-protocol()` option of the `network()` destination has been added to *Section 7.5.1, network() destination options (p. 135)*.

- *Section 6.2, Collecting internal messages (p. 52)* has been updated.
- The description of escaping special characters used in templates has been updated in *Section 11.1.2, Templates and macros (p. 239)*.

Changes in documentation:

- The synopsis and usage of the `facility()` and `priority()` filters have been corrected in *Section facility() (p. 212)* and *Section level() or priority() (p. 214)*.

- Incorrect references to escaping colons have been removed from *Section @ESTRING@ (p. 288)*.
- The default ports of the syslog drivers have been corrected.
- Erroneous references to RFC5427 and RFC5428 have been removed.

6.1.3. Version 3.3 - 3.4

Changes in product:

- New configuration objects called junctions and channels are available to improve the flexibility of configuring syslog-ng OSE. For details, see *Section 8.1.2, Junctions and channels (p. 201)* and *Section 5.5, Using channels in configuration objects (p. 44)*.
- syslog-ng OSE can publish messages using the AMQP (Advanced Message Queuing Protocol). For details, see *Section 7.1, Publishing messages using AMQP (p. 117)*.
- The `tcp, tcp6, udp, and udp6` drivers have been merged into a single `network()` driver. For details, see *Section 6.4, Collecting messages using the RFC3164 protocol (network() driver) (p. 61)* and *Section 7.5, Sending messages to a remote log server using the RFC3164 protocol (network() driver) (p. 134)*.

- The `tcp, tcp6, syslog, and network` source and destination drivers support the new `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options. For details, see respective source and driver descriptions, for example, *Section 6.10.1, syslog() source options (p. 88)*.

- Objects can be defined inline as well. This is useful if you use the object only once (for example, a filter). For details, see *Section 5.4, Defining configuration objects inline (p. 43)*.

- The `mongodb()` destination supports using replicasets, sockets, and safe-mode. For details, see *Section servers() (p. 133), Section path() (p. 133)*, and *Section safe-mode() (p. 133)*, respectively.

- The `value-pairs()` option can modify the names of the value-pairs using the `rekey` option. For details, see *Section value-pairs() (p. 19)*.

- A new parser is available to explicitly parse messages as syslog messages. For details, see *Section 12.1, Parsing syslog messages (p. 267)*.

- A new parser is available to parse JSON-formatted messages. For details, see *Section 12.3, The JSON parser (p. 272)*.

- The following new macros are available in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*: `AMPM, HOUR12, LOGHOST, MSEC, SYSUPTIME, USEC`.

- The following new template functions are available in *Section 11.1.7, Template functions of syslog-ng OSE (p. 248)*: `GEOIP, LENGTH, STRIP, SUBSTR, TFHASH, UUID`, and functions for various numerical operations.

- The following new parsers are available in *Section 13.5.1, Using pattern parsers (p. 286)*: `@EMAIL@, @HOSTNAME@, @MACADDR@, @LLADDR@, @PCRE@, @SET@`.

- Actions triggered from the pattern database can use the number of messages as a condition. For details, see *Section 13.4.1, Conditional actions (p. 284)*.

- Messages triggered from the pattern database can inherit the properties of the original message. For details, see *Example 13.8, Generating messages for pattern database matches (p. 283)*.

- Tags can be added and deleted using rewrite rules. For details, see *Section 11.2.6, Adding and deleting tags (p. 262)*.

- The `file-template()` and `proto-template()` global options have been documented. For details, see *Section file-template() (p. 219)* and *Section proto-template() (p. 225)*, respectively.

- The `dbd-option()` and `session-statements()` options are available for the `sql()` driver. For details, see *Section dbd-option() (p. 168)* and *Section session-statements() (p. 172)*.

- The tags and name-value pairs set in a pattern database file can be listed using the `pdbtool dictionary` command. For details, see *the section called "The dictionary command" (p. 315)*.

- The `expect-hostname` source flag has been documented.

- The `--caps` command-line option has been documented. For details, see *syslog-ng(8) (p. 321)*.

■ The `reload` option is available for the `syslog-ng-ctl` utility. For details, see *syslog-ng-ctl(1) (p. 332)*.

■ Wildcards can be used to include multiple configuration files. For details, see *Section 5.8.1, Including configuration files (p. 46)*.

■ *Section 7.11, Generating SMTP messages (e-mail) from logs (p. 158)* has been added to the document.

Changes in documentation:

■ Updated the documentation of the `system()` source in *Section 6.11, Collecting the system-specific log messages of a platform (p. 95)*.

■ *Section 13.4, Triggering actions for identified messages (p. 283)* has been split into several subsections.

■ Documented several missing `pdbtool` options in *Appendix A, The syslog-ng manual pages (p. 310)*.

■ The `--without-compile-date` compiling option has been documented in *Procedure 3.1, Compiling syslog-ng from source (p. 23)*.

■ The description of the `tags()` option has been added to *Section 6.2.1, internal() source options (p. 53)*.

■ Links and compiling option descriptions have been updated in *Procedure 3.1, Compiling syslog-ng from source (p. 23)*.

■ *Section retries (p. 171)* has been added to the document.

■ A few notes regarding kernel messages and file sources have been reorganized to *Section 6.3.1, Notes on reading kernel messages (p. 54)*.

■ Clarified the use of double-quotes and special characters in *Section 11.3, Regular expressions (p. 263)*.

■ Added a note about unsupported column types to *Section 7.12.2, Using the sql() driver with a Microsoft SQL database (p. 165)*.

■ The maximal number of worker-threads has been clarified in *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)*.

■ Corrected a note about persistent message contexts in *Section 13.3, Correlating log messages (p. 281)*.

■ Clarifications in *Section 7.12.1, Using the sql() driver with an Oracle database (p. 164)*.

■ A description of the `BOM` character has been added to *BOM (p. 355)*.

■ The description of the `delimiter` option of csv-parser() has been clarified. For details, see *Section delimiters (p. 270)*.

■ Missing facility names have been added to *Section facility() (p. 212)*.

■ The description of the `chain_hostnames` global option has been extended. For details, see *Section chain-hostnames() (p. 217)*.

■ The descriptions of statistics types have been clarified in *Chapter 14, Statistics of syslog-ng (p. 297)*.

■ The description of pipe() has been clarified.

- The manual pages of syslog-ng Open Source Edition have been relicensed under the GPLv2+ license. See *Appendix A, The syslog-ng manual pages (p. 310)* for details.
- Information about chrooting syslog-ng has been corrected in the `syslog-ng.8` manual page. See *syslog-ng(8) (p. 321)* for details.

- The example of *Section format-json (p. 249)* has been expanded and corrected.
- Lots of other corrections and clarifications.

6.2. Feedback

Any feedback is greatly appreciated, especially on what else this document should cover. General comments, errors found in the text, and any suggestions about how to improve the documentation is welcome at documentation@balabit.com.

6.3. Acknowledgments

BalaBit would like to express its gratitude to the syslog-ng users and the syslog-ng community for their invaluable help and support, including the community members listed at *syslog-ng Community Page*.

Chapter 1. Introduction to syslog-ng

This chapter introduces the syslog-ng Open Source Edition application in a non-technical manner, discussing how and why is it useful, and the benefits it offers to an existing IT infrastructure.

1.1. What syslog-ng is

The syslog-ng application is a flexible and highly scalable system logging application that is ideal for creating centralized and trusted logging solutions. The main features of syslog-ng are summarized below.

- *Reliable log transfer*: The syslog-ng application enables you to send the log messages of your hosts to remote servers using the latest protocol standards. The logs of different servers can be collected and stored centrally on dedicated log servers. Transferring log messages using the TCP protocol ensures that no messages are lost.

- *Secure logging using TLS*: Log messages may contain sensitive information that should not be accessed by third parties. Therefore, syslog-ng supports the Transport Layer Security (TLS) protocol to encrypt the communication. TLS also allows the mutual authentication of the host and the server using X.509 certificates.

- *Direct database access*: Storing your log messages in a database allows you to easily search and query the messages and interoperate with log analyzing applications. The syslog-ng application supports the following databases: MSSQL, MySQL, Oracle, PostgreSQL, and SQLite.

- *Heterogeneous environments*: The syslog-ng application is the ideal choice to collect logs in massively heterogeneous environments using several different operating systems and hardware platforms, including Linux, Unix, BSD, Sun Solaris, HP-UX, Tru64, and AIX.

- *Filter and classify*: The syslog-ng application can sort the incoming log messages based on their content and various parameters like the source host, application, and priority. Directories, files, and database tables can be created dynamically using macros. Complex filtering using regular expressions and boolean operators offers almost unlimited flexibility to forward only the important log messages to the selected destinations.

- *Parse and rewrite*: The syslog-ng application can segment log messages to named fields or columns, and also modify the values of these fields.

- *IPv4 and IPv6 support*: The syslog-ng application can operate in both IPv4 and IPv6 network environments; it can receive and send messages to both types of networks.

1.2. What syslog-ng is not

The syslog-ng application is not log analysis software. It can filter log messages and select only the ones matching certain criteria. It can even convert the messages and restructure them to a predefined format, or parse the messages and segment them into different fields. But syslog-ng cannot interpret and analyze the meaning behind the messages, or recognize patterns in the occurrence of different messages.

1.3. Why is syslog-ng needed?

Log messages contain information about the events happening on the hosts. Monitoring system events is essential for security and system health monitoring reasons.

The original syslog protocol separates messages based on the priority of the message and the facility sending the message. These two parameters alone are often inadequate to consistently classify messages, as many applications might use the same facility — and the facility itself is not even included in the log message. To make things worse, many log messages contain unimportant information. The syslog-ng application helps you to select only the really interesting messages, and forward them to a central server.

Company policies or other regulations often require log messages to be archived. Storing the important messages in a central location greatly simplifies this process.

For details on how can you use syslog-ng to comply with various regulations, see the *Regulatory compliance and system logging* whitepaper *available here*

1.4. What is new in syslog-ng Open Source Edition 3.6?

Version 3.6 of syslog-ng Open Source Edition includes the following main features:

- syslog-ng OSE can send messages and metrics to the Riemann monitoring system. For details, see *Section 7.10, Monitoring your data with Riemann (p. 154)*.
- syslog-ng OSE can send messages and metrics to a Graphite server. For details, see *Section 7.3, Sending metrics to Graphite (p. 128)*.
- syslog-ng OSE can directly read log messages from the journal file of platforms using systemd. For details, see *Section 6.12, Collecting messages from the systemd-journal system log storage (p. 97)*.
- syslog-ng OSE can read log messages from the syslog socket of platforms using systemd. For details, see *Section 6.13, Collecting systemd messages using a socket (p. 99)*.
- syslog-ng OSE can receive and automatically parse messages from nodejs application. For details, see *Section 6.5, Receiving JSON messages from nodejs applications (p. 69)*.
- syslog-ng OSE can rewrite multiple macros at the same time, making it easy to modify the values of fields parsed using patterndb or from JSON. For details, see *Section 11.2.4, Setting multiple message fields to specific values (p. 260)*.
- For a more detailed list of changes in syslog-ng OSE 3.6 and in The syslog-ng Open Source Edition 3.6 Administrator Guide, see *Section 6.1.1, Version 3.5 - 3.6 (p. xvi)*.

1.5. Who uses syslog-ng?

The syslog-ng application is used worldwide by companies and institutions who collect and manage the logs of several hosts, and want to store them in a centralized, organized way. Using syslog-ng is particularly advantageous for:

- Internet Service Providers;
- Financial institutions and companies requiring policy compliance;
- Server, web, and application hosting companies;

- Datacenters;
- Wide area network (WAN) operators;
- Server farm administrators.

1.6. Supported platforms

The syslog-ng Open Source Edition application is highly portable and is known to run on a wide range of hardware architectures (x86, x86_64, SUN Sparc, PowerPC 32 and 64, Alpha) and operating systems, including Linux, BSD, Solaris, IBM AIX, HP-UX, Mac OS X, Cygwin, Tru64, and others.

- The source code of syslog-ng Open Source Edition is released under the GPLv2 license and is *available here*.
- For a list of precompiled binary packages, see *http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/downloads/3rd-party*.

Chapter 2. The concepts of syslog-ng

This chapter discusses the technical concepts of syslog-ng.

2.1. The philosophy of syslog-ng

Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices — called syslog-ng clients — all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, which sorts and stores them.

2.2. Logging with syslog-ng

The syslog-ng application reads incoming messages and forwards them to the selected *destinations*. The syslog-ng application can receive messages from files, remote hosts, and other *sources*.

Log messages enter syslog-ng in one of the defined sources, and are sent to one or more *destinations*.

Sources and destinations are independent objects; *log paths* define what syslog-ng does with a message, connecting the sources to the destinations. A log path consists of one or more sources and one or more destinations; messages arriving from a source are sent to every destination listed in the log path. A log path defined in syslog-ng is called a *log statement*.

Optionally, log paths can include *filters*. Filters are rules that select only certain messages, for example, selecting only messages sent by a specific application. If a log path includes filters, syslog-ng sends only the messages satisfying the filter rules to the destinations set in the log path.

Other optional elements that can appear in log statements are *parsers* and *rewriting rules*. Parsers segment messages into different fields to help processing the messages, while rewrite rules modify the messages by adding, replacing, or removing parts of the messages.

2.2.1. Procedure – The route of a log message in syslog-ng

Purpose:

The following procedure illustrates the route of a log message from its source on the syslog-ng client to its final destination on the central syslog-ng server.
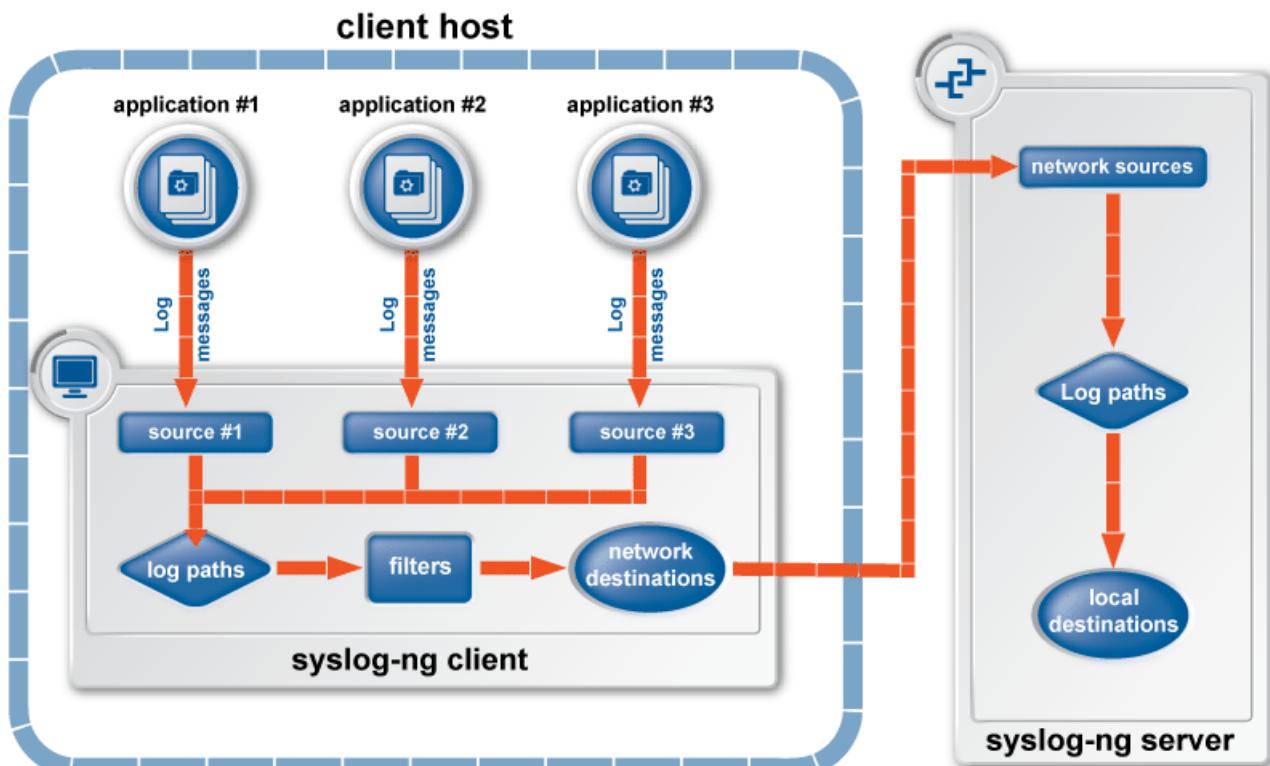
Figure 2.1. The route of a log message

Steps:

Step 1. A device or application sends a log message to a source on the syslog-ng client. For example, an Apache web server running on Linux enters a message into the `/var/log/apache` file.

Step 2. The syslog-ng client running on the web server reads the message from its `/var/log/apache` source.

Step 3. The syslog-ng client processes the first log statement that includes the `/var/log/apache` source.

Step 4. The syslog-ng client performs optional operations (message filtering, parsing, and rewriting) on the message; for example, it compares the message to the filters of the log statement (if any). If the message complies with all filter rules, syslog-ng sends the message to the destinations set in the log statement, for example, to the remote syslog-ng server.

Warning
Message filtering, parsing, and rewriting is performed in the order that the operations appear in the log statement.

Note
The syslog-ng client sends a message to *all* matching destinations by default. As a result, a message may be sent to a destination more than once, if the destination is used in multiple log statements. To prevent such situations, use the `final` flag in the destination statements. For details, see *Table 8.1, Log statement flags (p. 202)*.

Step 5.   The syslog-ng client processes the next log statement that includes the `/var/log/apache` source, repeating Steps 3-4.

Step 6.   The message sent by the syslog-ng client arrives from a source set in the syslog-ng server.

Step 7.   The syslog-ng server reads the message from its source and processes the first log statement that includes that source.

Step 8.   The syslog-ng server performs optional operations (message filtering, parsing, and rewriting) on the message; for example, it compares the message to the filters of the log statement (if any). If the message complies with all filter rules, syslog-ng sends the message to the destinations set in the log statement.

> **Warning**
> Message filtering, parsing, and rewriting is performed in the order that the operations appear in the log statement.

Step 9.   The syslog-ng server processes the next log statement, repeating Steps 7-9.

> **Note**
> The syslog-ng application can stop reading messages from its sources if the destinations cannot process the sent messages. This feature is called flow-control and is detailed in *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

2.3. Modes of operation

The syslog-ng Open Source Edition application has three typical operation scenarios: *Client*, *Server*, and *Relay*.

2.3.1. Client mode
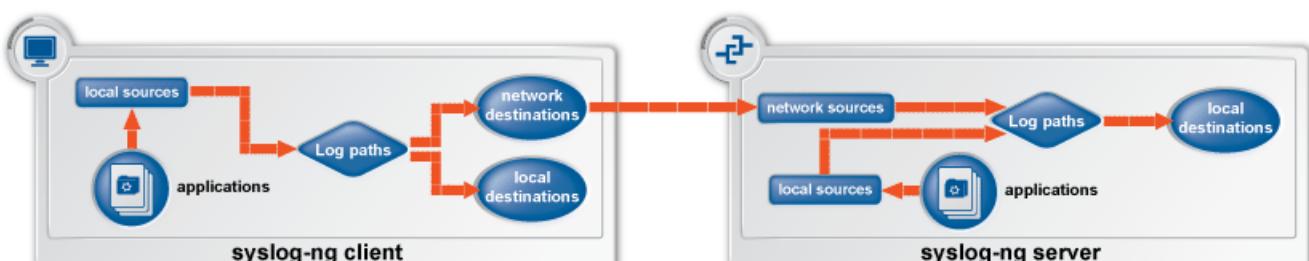
Figure 2.2. Client-mode operation

In client mode, syslog-ng collects the local logs generated by the host and forwards them through a network connection to the central syslog-ng server or to a relay. Clients often also log the messages locally into files.

2.3.2. Relay mode
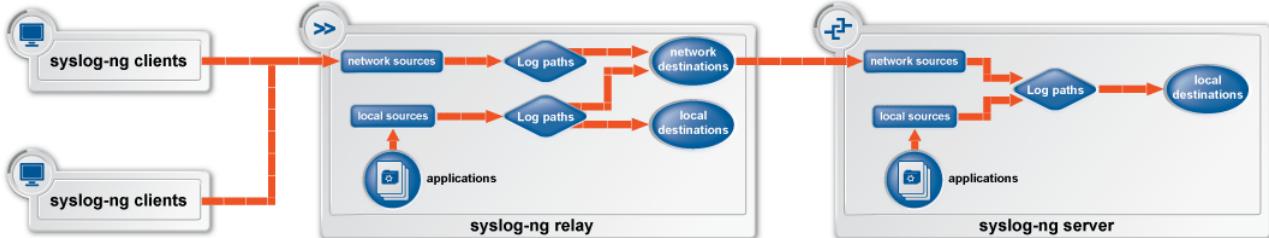
Figure 2.3. Relay-mode operation

In relay mode, syslog-ng receives logs through the network from syslog-ng clients and forwards them to the central syslog-ng server using a network connection. Relays also log the messages from the relay host into a local file, or forward these messages to the central syslog-ng server.

2.3.3. Server mode
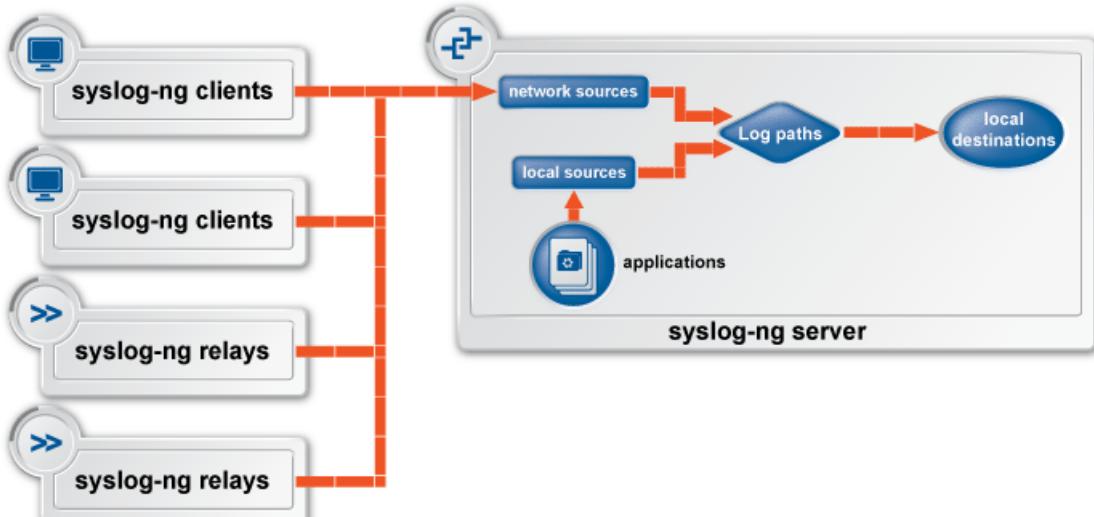
Figure 2.4. Server-mode operation

In server mode, syslog-ng acts as a central log-collecting server. It receives messages from syslog-ng clients and relays over the network, and stores them locally in files, or passes them to other applications, for example log analyzers.

2.4. Global objects

The syslog-ng application uses the following objects:

- *Source driver*: A communication method used to receive log messages. For example, syslog-ng can receive messages from a remote host via TCP/IP, or read the messages of a local application from a file. For details on source drivers, see *Chapter 6, Collecting log messages — sources and source drivers (p. 50)*.

- *Source*: A named collection of configured source drivers.

- *Destination driver*: A communication method used to send log messages. For example, syslog-ng can send messages to a remote host via TCP/IP, or write the messages into a file or database. For details on destination drivers, see *Chapter 7, Sending and storing log messages — destinations and destination drivers (p. 116)*.

- *Destination*: A named collection of configured destination drivers.

- *Filter*: An expression to select messages. For example, a simple filter can select the messages received from a specific host. For details, see *Section 11.1, Customizing message format (p. 238)*.

- *Macro*: An identifier that refers to a part of the log message. For example, the `${HOST}` macro returns the name of the host that sent the message. Macros are often used in templates and filenames. For details, see *Section 11.1, Customizing message format (p. 238)*.

- *Parser*: Parsers are objects that parse the incoming messages, or parts of a message. For example, the `csv-parser()` can segment messages into separate columns at a predefined separator character (for example a comma). Every column has a unique name that can be used as a macro. For details, see *Chapter 12, Parsing and segmenting structured messages (p. 267)* and *Chapter 13, Processing message content with a pattern database (p. 275)*.

- *Rewrite rule*: A rule modifies a part of the message, for example, replaces a string, or sets a field to a specified value. For details, see *Section 11.2, Modifying messages (p. 257)*.

- *Log paths*: A combination of sources, destinations, and other objects like filters, parsers, and rewrite rules. The syslog-ng application sends messages arriving from the sources of the log paths to the defined destinations, and performs filtering, parsing, and rewriting of the messages. Log paths are also called log statements. Log statements can include other (embedded) log statements and junctions to create complex log paths. For details, see *Chapter 8, Routing messages: log paths and filters (p. 198)*.

- *Template*: A template is a set of macros that can be used to restructure log messages or automatically generate file names. For example, a template can add the hostname and the date to the beginning of every log message. For details, see *Section 11.1, Customizing message format (p. 238)*.

- *Option*: Options set global parameters of syslog-ng, like the parameters of name resolution and timezone handling. For details, see *Chapter 9, Global options of syslog-ng OSE (p. 217)*.

For details on the above objects, see *Section 5.2, The configuration syntax in detail (p. 40)*.

2.5. Timezones and daylight saving

The syslog-ng application receives the timezone and daylight saving information from the operating system it is installed on. If the operating system handles daylight saving correctly, so does syslog-ng.

The syslog-ng application supports messages originating from different timezones. The original syslog protocol (RFC3164) does not include timezone information, but syslog-ng provides a solution by extending the syslog

protocol to include the timezone in the log messages. The syslog-ng application also enables administrators to supply timezone information for legacy devices which do not support the protocol extension.

2.5.1. Procedure – How syslog-ng OSE assigns timezone to the message

When syslog-ng OSE receives a message, it assigns timezone information to the message using the following algorithm.

Step 1. The sender application (for example the syslog-ng client) or host specifies the timezone of the messages. If the incoming message includes a timezone it is associated with the message. Otherwise, the local timezone is assumed.

Step 2. Specify the *time-zone()* parameter for the source driver that reads the message. This timezone will be associated with the messages only if no timezone is specified within the message itself. Each source defaults to the value of the *recv-time-zone()* global option. It is not possible to override only the timezone information of the incoming message; but setting the *keep-timestamp()* option to no allows syslog-ng OSE to replace the full timestamp (timezone included) with the time the message was received.

> **Note**
> When processing a message that does not contain timezone information, the syslog-ng OSE application will use the timezone and daylight-saving that was effective when the timestamp was generated. For example, the current time is 2011-03-11 (March 11, 2011) in the EU/Budapest timezone. When daylight-saving is active (summertime), the offset is +02:00. When daylight-saving is inactive (wintertime) the timezone offset is +01:00. If the timestamp of an incoming message is 2011-01-01, the timezone associated with the message will be +01:00, but the timestamp will be converted, because 2011-01-01 meant winter time when daylight saving is not active but the current timezone is +02:00.

Step 3. Specify the timezone in the destination driver using the *time-zone()* parameter. Each destination driver might have an associated timezone value; syslog-ng converts message timestamps to this timezone before sending the message to its destination (file or network socket). Each destination defaults to the value of the *send-time-zone()* global option.

> **Note**
> A message can be sent to multiple destination zones. The syslog-ng application converts the timezone information properly for every individual destination zone.

> **Warning**
> If syslog-ng OSE sends the message is to the destination using the legacy-syslog protocol (RFC3164) which does not support timezone information in its timestamps, the timezone information cannot be encapsulated into the sent timestamp, so syslog-ng OSE will convert the hour:min values based on the explicitly specified timezone.

Step 4. If the timezone is not specified, local timezone is used.

Step 5. When macro expansions are used in the destination filenames, the local timezone is used. (Also, if the timestamp of the received message does not contain the year of the message, syslog-ng OSE uses the local year.)

2.5.2. A note on timezones and timestamps

If the clients run syslog-ng, then use the ISO timestamp, because it includes timezone information. That way you do not need to adjust the *recv-time-zone()* parameter of syslog-ng.

If you want syslog-ng to output timestamps in Unix (POSIX) time format, use the *S_UNIXTIME* and *R_UNIXTIME* macros. You do not need to change any of the timezone related parameters, because the timestamp information of incoming messages is converted to Unix time internally, and Unix time is a timezone-independent time representation. (Actually, Unix time measures the number of seconds elapsed since midnight of Coordinated Universal Time (UTC) January 1, 1970, but does not count leap seconds.)

2.6. The license of syslog-ng OSE

Starting with version 3.2, the syslog-ng Open Source Edition application is licensed under a combined LGPL+GPL license. The core of syslog-ng OSE is licensed under the GNU Lesser General Public License Version 2.1 license, while the rest of the codebase is licensed under the GNU General Public License Version 2 license.

Note
Practically, the code stored under the `lib` directory of the source code package is under LGPL, the rest is GPL.

For details about the LGPL and GPL licenses, see *Appendix C, GNU Lesser General Public License (p. 341)* and *Appendix B, GNU General Public License (p. 335)*, respectively.

2.7. High availability support

Multiple syslog-ng servers can be run in fail-over mode. The syslog-ng application does not include any internal support for this, as clustering support must be implemented on the operating system level. A tool that can be used to create UNIX clusters is Heartbeat (for details, see *this page*).

2.8. The structure of a log message

The following sections describe the structure of log messages. Currently there are two standard syslog message formats:

- The old standard described in RFC 3164 (also called the BSD-syslog or the legacy-syslog protocol): see *Section 2.8.1, BSD-syslog or legacy-syslog messages (p. 11)*

- The new standard described in RFC 5424 (also called the IETF-syslog protocol): see *Section 2.8.2, IETF-syslog messages (p. 13)*

- How messages are represented in syslog-ng OSE: see *Section 2.9, Message representation in syslog-ng OSE (p. 16)*.

2.8.1. BSD-syslog or legacy-syslog messages

This section describes the format of a syslog message, according to the *legacy-syslog or BSD-syslog protocol*). A syslog message consists of the following parts:

- *PRI*

- *HEADER*

- *MSG*

The total message cannot be longer than 1024 bytes.

The following is a sample syslog message: *<133>Feb 25 14:09:07 webserver syslogd: restart*. The message corresponds to the following format: *<priority>timestamp hostname application: message*. The different parts of the message are explained in the following sections.

> **Note**
> The syslog-ng application supports longer messages as well. For details, see the `log-msg-size()` option in *Section 9.2, Global options (p. 217)*. However, it is not recommended to enable messages larger than the packet size when using UDP destinations.

2.8.1.1. The PRI message part

The PRI part of the syslog message (known as Priority value) represents the Facility and Severity of the message. Facility represents the part of the system sending the message, while severity marks its importance. The Priority value is calculated by first multiplying the Facility number by 8 and then adding the numerical value of the Severity. The possible facility and severity values are presented below.

> **Note**
> Facility codes may slightly vary between different platforms. The syslog-ng application accepts facility codes as numerical values as well.

Numerical Code	Facility
0	kernel messages
1	user-level messages
2	mail system
3	system daemons
4	security/authorization messages
5	messages generated internally by syslogd
6	line printer subsystem
7	network news subsystem

Numerical Code	Facility
8	UUCP subsystem
9	clock daemon
10	security/authorization messages
11	FTP daemon
12	NTP subsystem
13	log audit
14	log alert
15	clock daemon
16-23	locally used facilities (local0-local7)

Table 2.1. syslog Message Facilities

The following table lists the severity values.

Numerical Code	Severity
0	Emergency: system is unusable
1	Alert: action must be taken immediately
2	Critical: critical conditions
3	Error: error conditions
4	Warning: warning conditions
5	Notice: normal but significant condition
6	Informational: informational messages
7	Debug: debug-level messages

Table 2.2. syslog Message Severities

2.8.1.2. The HEADER message part

The HEADER part contains a timestamp and the hostname (without the domain name) or the IP address of the device. The timestamp field is the local time in the *Mmm dd hh:mm:ss* format, where:

- *Mmm* is the English abbreviation of the month: Jan, Feb, Mar, Apr, May, Jun, Jul, Aug, Sep, Oct, Nov, Dec.

- *dd* is the day of the month on two digits. If the day of the month is less than 10, the first digit is replaced with a space. (For example *Aug 7*.)

- *hh:mm:ss* is the local time. The hour (hh) is represented in a 24-hour format. Valid entries are between 00 and 23, inclusive. The minute (mm) and second (ss) entries are between 00 and 59 inclusive.

Note
The syslog-ng application supports other timestamp formats as well, like ISO, or the PIX extended format. For details, see the `ts-format()` option in *Section 9.2, Global options (p. 217)*.

2.8.1.3. The MSG message part

The MSG part contains the name of the program or process that generated the message, and the text of the message itself. The MSG part is usually in the following format: *program[pid]: message text*.

2.8.2. IETF-syslog messages

This section describes the format of a syslog message, according to the *IETF-syslog protocol*). A syslog message consists of the following parts:

- *HEADER* (includes the *PRI* as well)

- *STRUCTURED-DATA*

- *MSG*

The following is a sample syslog message:

```
<34>1 2003-10-11T22:14:15.003Z mymachine.example.com su - ID47 - BOM'su root' failed
 for lonvick on /dev/pts/8
```

The message corresponds to the following format:

```
<priority>VERSION ISOTIMESTAMP HOSTNAME APPLICATION PID MESSAGEID STRUCTURED-DATA
 MSG
```

In this example, the Facility has the value of 4, severity is 2, so PRI is 34. The VERSION is 1. The message was created on 11 October 2003 at 10:14:15pm UTC, 3 milliseconds into the next second. The message originated from a host that identifies itself as "mymachine.example.com". The APP-NAME is "su" and the PROCID is unknown. The MSGID is "ID47". The MSG is "'su root' failed for lonvick...", encoded in UTF-8. The encoding is defined by the *BOM* . There is no STRUCTURED-DATA present in the message, this is indicated by "-" in the STRUCTURED-DATA field. The MSG is "'su root' failed for lonvick...".

The HEADER part of the message must be in plain ASCII format, the parameter values of the STRUCTURED-DATA part must be in UTF-8, while the MSG part should be in UTF-8. The different parts of the message are explained in the following sections.

2.8.2.1. The PRI message part

The PRI part of the syslog message (known as Priority value) represents the Facility and Severity of the message. Facility represents the part of the system sending the message, while severity marks its importance. The Priority value is calculated by first multiplying the Facility number by 8 and then adding the numerical value of the Severity. The possible facility and severity values are presented below.

Source: http://tools.ietf.org/html/rfc5424
The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

Note
Facility codes may slightly vary between different platforms. The syslog-ng application accepts facility codes as numerical values as well.

Numerical Code	Facility
0	kernel messages
1	user-level messages
2	mail system
3	system daemons
4	security/authorization messages
5	messages generated internally by syslogd
6	line printer subsystem
7	network news subsystem
8	UUCP subsystem
9	clock daemon
10	security/authorization messages
11	FTP daemon
12	NTP subsystem
13	log audit
14	log alert
15	clock daemon
16-23	locally used facilities (local0-local7)

Table 2.3. syslog Message Facilities

The following table lists the severity values.

Numerical Code	Severity
0	Emergency: system is unusable
1	Alert: action must be taken immediately
2	Critical: critical conditions
3	Error: error conditions
4	Warning: warning conditions
5	Notice: normal but significant condition
6	Informational: informational messages

Numerical Code	Severity
7	Debug: debug-level messages

Table 2.4. syslog Message Severities

2.8.2.2. The HEADER message part

The HEADER part contains the following elements:

- *VERSION*: Version number of the syslog protocol standard. Currently this can only be 1.

- *ISOTIMESTAMP*: The time when the message was generated in the ISO 8601 compatible standard timestamp format (yyyy-mm-ddThh:mm:ss+-ZONE), for example: `2006-06-13T15:58:00.123+01:00`.

- *HOSTNAME*: The machine that originally sent the message.

- *APPLICATION*: The device or application that generated the message

- *PID*: The process name or process ID of the syslog application that sent the message. It is not necessarily the process ID of the application that generated the message.

- *MESSAGEID*: The ID number of the message.

Note
The syslog-ng application supports other timestamp formats as well, like ISO, or the PIX extended format. The timestamp used in the IETF-syslog protocol is derived from RFC3339, which is based on ISO8601. For details, see the `ts-format()` option in *Section 9.2, Global options (p. 217)*.

The syslog-ng OSE application will truncate the following fields:

- If *APP-NAME* is longer than 48 characters it will be truncated to 48 characters.

- If *PROC-ID* is longer than 128 characters it will be truncated to 128 characters.

- If *MSGID* is longer than 32 characters it will be truncated to 32 characters.

- If *HOSTNAME* is longer than 255 characters it will be truncated to 255 characters.

2.8.2.3. The STRUCTURED-DATA message part

The STRUCTURED-DATA message part may contain meta- information about the syslog message, or application-specific information such as traffic counters or IP addresses. STRUCTURED-DATA consists of data blocks enclosed in brackets (*[]*). Every block includes the ID of the block, and one or more *name=value* pairs. The syslog-ng application automatically parses the STRUCTURED-DATA part of syslog messages, which can be referenced in macros (for details, see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*). An example STRUCTURED-DATA block looks like:

```
[exampleSDID@O iut="3" eventSource="Application" eventID="1011"][examplePriority@O
 class="high"]
```

2.8.2.4. The MSG message part

The MSG part contains the text of the message itself. The encoding of the text must be UTF-8 if the *BOM* character is present in the message. If the message does not contain the BOM character, the encoding is treated as unknown. Usually messages arriving from legacy sources do not include the BOM character. CRLF characters will not be removed from the message.

2.9. Message representation in syslog-ng OSE

When the syslog-ng OSE application receives a message, it automatically parses the message. The syslog-ng OSE application can automatically parse log messages that conform to the RFC3164 (BSD or legacy-syslog) or the RFC5424 (IETF-syslog) message formats. If syslog-ng OSE cannot parse a message, it results in an error.

> **Tip**
> In case you need to relay messages that cannot be parsed without any modifications or changes, use the `flags(no-parse)` option in the source definition, and a template containing only the `${MSG}` macro in the destination definition.

A parsed message has the following parts.

- **Timestamps.** Two timestamps are associated with every message: one is the timestamp contained within the message (that is, when the sender sent the message), the other is the time when syslog-ng OSE has actually received the message.

- **Severity.** The severity of the message.

- **Facility.** The facility that sent the message.

- **Tags.** Custom text labels added to the message that are mainly used for filtering. None of the current message transport protocols adds tags to the log messages. Tags can be added to the log message only within syslog-ng OSE. The syslog-ng OSE application automatically adds the id of the source as a tag to the incoming messages. Other tags can be added to the message by the pattern database, or using the `tags()` option of the source.

- **IP address of the sender.** The IP address of the host that sent the message. Note that the IP address of the sender is a hard macro and cannot be modified within syslog-ng OSE but the associated hostname can be modified, for example, using rewrite rules.

- **Hard macros.** Hard macros contain data that is directly derived from the log message, for example, the ${MONTH} macro derives its value from the timestamp. The most important consideration with hard macros is that they are read-only, meaning they cannot be modified using rewrite rules or other means.

- **Soft macros.** Soft macros (sometimes also called name-value pairs) are either built-in macros automatically generated from the log message (for example, ${HOST}), or custom user-created macros generated by using the syslog-ng pattern database or a CSV-parser. The SDATA fields of RFC5424-formatted log messages become soft macros as well. In contrast with hard macros, soft macros are writable and can be modified within syslog-ng OSE, for example, using rewrite rules.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

Note
It is also possible to set the value of built-in soft macros using parsers, for example, to set the ${HOST} macro from the message using a column of a CSV-parser.

The data extracted from the log messages using named pattern parsers in the pattern database are also soft macros.

Tip
For the list of hard and soft macros, see *Section 11.1.4, Hard vs. soft macros (p. 241)*.

2.10. Structuring macros, metadata, and other value-pairs

Available in syslog-ng OSE 3.3 and later.

The syslog-ng OSE application allows you to select and construct name-value pairs from any information already available about the log message, or extracted from the message itself. You can directly use this structured information, for example, in the following places:

- *amqp()* destination

- *mongodb()* destination

- *stomp()* destination

- or in other destinations using the *format-json()* template function.

When using *value-pairs*, there are three ways to specify which information (that is, macros or other name-value pairs) to include in the selection.

- Select groups of macros using the *scope()* parameter, and optionally remove certain macros from the group using the *exclude()* parameter.

- List specific macros to include using the *key()* parameter.

- Define new name-value pairs to include using the *pair()* parameter.

These parameters are detailed in *Section value-pairs() (p. 19)*.

2.10.1. Specifying data types in value-pairs

By default, syslog-ng OSE handles every data as strings. However, certain destinations and data formats (for example, SQL, MongoDB, JSON, AMQP) support other types of data as well, for example, numbers or dates. The syslog-ng OSE application allows you to specify the data type in templates (this is also called type-hinting).

If the destination driver supports data types, it converts the incoming data to the specified data type. For example, this allows you to store integer numbers as numbers in MongoDB, instead of strings.

Warning
Hazard of data loss! If syslog-ng OSE cannot convert the data into the specified type, an error occurs, and syslog-ng OSE drops the message by default. To change how syslog-ng OSE handles data-conversion errors, see *Section on-error() (p. 224)*.

To use type-hinting, enclose the macro or template containing the data with the type: `<datatype>("<macro>")`, for example: `int("$PID")`.

Currently the *mongodb()* destination and the *format-json* template function supports data types.

Example 2.1. Using type-hinting
The following example stores the MESSAGE, PID, DATE, and PROGRAM fields of a log message in a MongoDB database. The DATE and PID parts are stored as numbers instead of strings.

```
mongodb(
    value-pairs(pair("date", datetime("$UNIXTIME"))
            pair("pid", int64("$PID"))
            pair("program", "$PROGRAM"))
            pair("message", "$MESSAGE"))
             )
);
```

The following example formats the same fields into JSON.

```
$(format-json date=datetime("$UNIXTIME") pid=int64("$PID") program="$PROGRAM"
message="$MESSAGE")
```

The syslog-ng OSE application currently supports the following data-types.

- *boolean*: Converts the data to a boolean value. Anything that begins with a `t` or `1` is converted to true, anything that begins with an `f` or `0` is converted to false.

- *datetime*: Use it only with UNIX timestamps, anything else will likely result in an error. This means that currently you can use only the *$UNIXTIME* macro for this purpose.

- *double*: A floating-point number.

- *literal*: The data as a literal string, without adding any quotes or escape characters.

- *int* or *int32*: 32-bit integer.

- *int64*: 64-bit integer.

- *string*: The data as a string.

value-pairs()

Type:	parameter list of the *value-pairs()* option
Default:	empty string

Description: The *value-pairs()* option allows you to select specific information about a message easily using predefined macro groups. The selected information is represented as name-value pairs and can be used formatted to JSON format, or directly used in a *mongodb()* destination.

Example 2.2. Using the *value-pairs()* option
The following example selects every available information about the log message, except for the date-related macros (R_* and S_*), selects the .SDATA.meta.sequenceId macro, and defines a new value-pair called MSGHDR that contains the program name and PID of the application that sent the log message.

```
value-pairs(
    scope(nv_pairs core syslog all_macros selected_macros everything)
    exclude("R_*")
    exclude("S_*")
    key(".SDATA.meta.sequenceId")
    pair("MSGHDR" "$PROGRAM[$PID]: ")
)
```

The following example selects the same information as the previous example, but converts it into JSON format.

```
$(format-json --scope nv_pairs,core,syslog,all_macros,selected_macros,everything \
  --exclude R_* --exclude S_* --key .SDATA.meta.sequenceId \
  --pair MSGHDR="$PROGRAM[$PID]: ")
```

Note
Every macro is included in the selection only once, but redundant information may appear if multiple macros include the same information (for example, including several date-related macros in the selection).

The *value-pairs()* option has the following parameters. The parameters are evaluated in the following order:

1. *scope()*

2. *exclude()*

3. *key()*

4. *pair()*

exclude()

Type: Space-separated list of macros to remove from the selection created using the *scope()* option.

Default: empty string

Description: This option removes the specified macros from the selection. Use it to remove unneeded macros selected using the *scope()* parameter.

For example, the following example removes the SDATA macros from the selection.

```
value-pairs(
        scope(rfc5424 selected_macros)
        exclude(".SDATA*")
    )
```

The name of the macro to remove can include glob expressions and wildcards (`*`, `?`, `[ab]`, `{foo,bar}`), as described in *Section glob (p. 265)*. Regular expressions are not supported.

key()

Type: A glob expression (or a comma-separated list of glob expressions) specifying the macros to be included in selection

Default: empty string

Description: This option selects the specified macros. The selected macros will be included as MACRONAME = MACROVALUE, that is using key("HOST") will result in HOST = $HOST. Multiple macros can be selected using glob expressions. For details on globs, see *Section glob (p. 265)*. For example:

```
value-pairs(
        scope(rfc3164)
        key("HOST"))
```

```
value-pairs(
        scope(rfc3164)
        key("HOST", "PROGRAM"))
```

pair()

Type: name value pairs in "<NAME>" "<VALUE>" format

Default: empty string

Description: This option defines a new name-value pair to be included in the message. The value part can include macros, templates, and template functions as well. For example:

```
value-pairs(
        scope(rfc3164)
        pair("TIME" "$HOUR:$MIN")
        pair("MSGHDR" "$PROGRAM[$PID]: "))
```

rekey()

Type: <pattern-to-select-names>, <list of transformations>

Default: empty string

Description: This option allows you to manipulate and modify the name of the value-pairs. You can define transformations, which are are applied to the selected name-value pairs. The first parameter of the *rekey()* option is a glob pattern that selects the name-value pairs to modify. If you omit the pattern, the transformations are applied to every key of the scope. For details on globs, see *Section glob (p. 265)*.

- If *rekey()* is used within a *key()* option, the name-value pairs specified in the glob of the *key()* option are transformed.

- If *rekey()* is used outside the *key()* option, every name-value pair of the *scope()* is transformed.

The following transformations are available:

`add-prefix("<my-prefix>")`	Adds the specified prefix to every name. For example, rekey(add-prefix("my-prefix."))
`replace-prefix("<prefix-to-replace>", "<new-prefix>")`	Replaces a substring at the beginning of the key with another string. Only prefixes can be replaced. For example, replace-prefix(".class", ",patterndb") changes the beginning tag `.class` to `.patterndb` This option was called *replace()* in syslog-ng OSE version 3.4.
`shift("<number>")`	Cuts the specified number of characters from the beginning of the name.

Example 2.3. Using the rekey() option
The following sample selects every value-pair that begins with `.cee.`, deletes this prefix by cutting 4 characters from the names, and adds a new prefix (`events.`).

```
value-pairs(
    key(".cee.*"
        rekey(
            shift(4)
            add-prefix("events.")
        )
    )
)
```

The *rekey()* option can be used with the *format-json* template-function as well, using the following syntax:

```
$(format-json --rekey .cee.* --add-prefix events.)
```

scope()

Type: space-separated list of macro groups to include in selection

Default: empty string

Description: This option selects predefined groups of macros. The following groups are available:

- *nv-pairs*: Every soft macro (name-value pair) associated with the message, except the ones that start with a dot (.) character. Macros starting with a dot character are generated within syslog-ng OSE and are not originally part of the message, therefore are not included in this group.

- *dot-nv-pairs*: Every soft macro (name-value pair) associated with the message which starts with a dot (.) character. For example, `.classifier.rule_id` and `.sdata.*`. Macros starting with a dot character are generated within syslog-ng OSE and are not originally part of the message.

- *all-nv-pairs*: Include every soft macro (name-value pair). Equivalent to using both `nv-pairs` and `dot-nv-pairs`.

- *rfc3164*: The macros that correspond to the RFC3164 (legacy or BSD-syslog) message format: `$FACILITY`, `$PRIORITY`, `$HOST`, `$PROGRAM`, `$PID`, `$MESSAGE`, and `$DATE`.

- *rfc5424*: The macros that correspond to the RFC5424 (IETF-syslog) message format: `$FACILITY`, `$PRIORITY`, `$HOST`, `$PROGRAM`, `$PID`, `$MESSAGE`, `$MSGID`, `$R_DATE`, and the metadata from the structured-data (SDATA) part of RFC5424-formatted messages, that is, every macro that starts with `.SDATA.`.

  The `rfc5424` group also has the following alias: `syslog-proto`. Note that the value of `$R_DATE` will be listed under the `DATE` key.

  The `rfc5424` group does not contain any metadata about the message, only information that was present in the original message. To include the most commonly used metadata (for example, the `$SOURCEIP` macro), use the `selected-macros` group instead.

- *all-macros*: Include every hard macro. This group is mainly useful for debugging, as it contains redundant information (for example, the date-related macros include the date-related information several times in various formats).

- *selected-macros*: Include the macros of the `rfc3164` groups, and the most commonly used metadata about the log message: the `$TAGS`, `$SOURCEIP`, and `$SEQNUM` macros.

- *sdata*: The metadata from the structured-data (SDATA) part of RFC5424-formatted messages, that is, every macro that starts with `.SDATA.`

- *everything*: Include every hard and soft macros. This group is mainly useful for debugging, as it contains redundant information (for example, the date-related macros include the date-related information several times in various formats).

For example:

```
value-pairs(
        scope(rfc3164 selected-macros))
```

Chapter 3. Installing syslog-ng

This chapter explains how to install syslog-ng Open Source Edition on various platforms.

- You can install syslog-ng OSE on many platforms using the package manager and official repositories of the platform. For a list of third-party packages available for various Linux, UNIX, and other platforms, see *third-party binaries page*.
- For instructions on compiling syslog-ng Open Source Edition from the source code, see *Procedure 3.1, Compiling syslog-ng from source (p. 23)*.

3.1. Procedure – Compiling syslog-ng from source

Purpose:

To compile syslog-ng Open Source Edition (OSE) from the source code, complete the following steps. Alternatively, you can use precompiled binary packages on several platforms. For a list of third-party packages available for various Linux, UNIX, and other platforms, see *third-party binaries page*.

Steps:

Step 1. Download the latest version of syslog-ng OSE from the *BalaBit website* or from *GitHub*. The source code is available as a tar.gz archive file.

Step 2. Download the latest version of the EventLog library *here*.

Step 3. Install the following packages that are required to compile syslog-ng. These packages are available for most UNIX/Linux systems. Alternatively, you can also download the sources and compile them.

- A version of the *gcc* C compiler that properly supports Thread Local Storage (TLS), for example, version 4.5 (at least version. For building recent gcc versions on Solaris, see *http://jblopen.com/node/16*.
- The *GNU flex* lexical analyser generator, *available here*.
- The *bison* parser generator, *available here*.
- The development files of the *glib* library, *available here*.
- The syslog-ng OSE application now uses PCRE-type regular expressions by default. It requires the `libpcre` library package, *available here*.

Step 4. If you want to use the spoof-source function of syslog-ng, install the development files of the *libnet* library, *available here*.

Step 5. If you want to send e-mails using the `smtp()` destination, install the development files of the *libesmtp* library, *available here*. This library is not needed if you use the `--disable-smtp` compile option.

Step 6. If you want to use the */etc/hosts.deny* and */etc/hosts.allow* for TCP access, install the development files of the *libwrap* (also called TCP-wrappers) library, _ available here.

Step 7. Uncompress the eventlog archive using the

```
$ tar xvfz eventlog-x.x.x.x.tar.gz
```

or the

```
$ gunzip -c eventlog-x.x.x.x.tar.gz | tar xvf -
```

command. A new directory containing the source code of eventlog will be created.

Step 8. By default, eventlog creates a file used by the syslog-ng configure script in the */usr/local/lib/pkgconfig* directory. Issue the following command to add this directory to your PKG_CONFIG_PATH:

```
PKG_CONFIG_PATH=/usr/local/lib/pkgconfig:$PKG_CONFIG_PATH
```

Step 9. Enter the new directory and issue the following commands:

```
$ ./configure
$ make
$ make install
```

Step 10. Uncompress the syslog-ng archive using the

```
tar xvfz syslog-ng-x.xx.tar.gz
```

or the

```
unzip -c syslog-ng-x.xx.tar.gz | tar xvf -
```

command. A new directory containing the source code of syslog-ng will be created.

Step 11. Enter the new directory and issue the following commands:

```
$ ./configure
$ make
$ make install
```

These commands will build syslog-ng using its default options.

> **Note**
> - On Solaris, use *gmake* (GNU make) instead of *make*.
>
> - To build syslog-ng OSE with less verbose output, use the make V=0 command. This results in shorter, less verbose output, making warnings and other anomalies easier to notice. Note that silent-rules support is only available in recent automake versions.

Step 12. If needed, use the following options to change how syslog-ng is compiled using the following command syntax:

```
$ ./configure --compile-time-option-name
```

Note

You can also use *--disable options*, to explicitly disable a feature and override autodetection. For example, to disable the TCP-wrapper support, use the *--disable-tcp-wrapper* option.

Warning

Starting with syslog-ng Open Source Edition 3.0.2, default linking mode of syslog-ng is `dynamic`. This means that syslog-ng might not be able to start up if the `/usr` directory is on NFS. On platforms where syslog-ng is used as a system logger, the `--enable-mixed-linking` is preferred.

- *--disable-json* Disable JSON support. It also disables `json-parser`, and the `format-json` template function. Also, it disables JSON support even if the `json-c` library is installed and detected (see *--enable-json*).

- *--disable-smtp* Disable SMTP support. By default, SMTP support is enabled if the libesmtp library is detected.

- *--enable-amqp* Enable the amqp destination (enabled by default). The source of the RabbitMQ client is included in the source code package of syslog-ng OSE. To use an external client instead, use the `--with-librabbitmq-client=system` compiling option. For details on using this destination, see *Section 7.1, Publishing messages using AMQP (p. 117)*.

- *--enable-debug* Include debug information.

- *--enable-dynamic-linking* Compile syslog-ng as a completely dynamic binary. If not specified syslog-ng uses mixed linking (`--enable-mixed-linking`): it links dynamically to system libraries and statically to everything else.

- *--enable-geoip* Enable GEOIP support, required for the `geoip` template function (enabled automatically if the `libgeoip` library is detected).

- *--enable-ipv6* Enable IPv6 support.

- *--enable-json* Enables JSON support (enabled automatically if the `json-c` 0.9 or newer library is installed and detected).

- *--enable-linux-caps* Enable support for capabilities on Linux. For details, see *syslog-ng(8) (p. 321)*.

- *--enable-mongodb* Enable the mongodb destination (enabled by default). The source of the MongoDB client is included in the source code package of syslog-ng OSE. To use an external MongoDB client instead, use the `--with-libmongo-client=system` compiling option. For details on using this destination, see *Section 7.4, Storing messages in a MongoDB database (p. 129)*.

- *--enable-pacct* Enable using the `pacct()` driver to collect process-accounting logs on Linux systems.

- *--enable-redis* Enable the redis destination (enabled by default). The source of the libhiredis client must be available. To specify the location of the library, use the

`--with-libhiredis=<path-to-libhiredis>` compiling option. For details on using this destination, see *Section 7.9, Storing name-value pairs in Redis (p. 152)*.

- *--enable-spoof-source* Enable spoof_source feature (disabled by default).

- *--enable-sql* Enables the sql() destination (enabled automatically if the `libdbi` library is installed and detected).

- *--enable-ssl* Enable SSL support, required for encrypted message transfer, as well as template functions that calculate hashes and UUIDs (enabled automatically if the `libopenssl` library is detected).

- *--enable-sun-door* Enable Sun door support even if not detected (autodetected by default).

- *--enable-sun-streams* Enable Sun STREAMS support even if not detected (autodetected by default).

- *--enable-systemd* Enable systemd support on Linux platforms (autodetected by default) (enabled automatically if the `libsystemd-daemon` library is detected).

- *--enable-tcp-wrapper* Enable using */etc/hosts.deny* and */etc/hosts.allow* for TCP access (enabled automatically if the `libwrap` libraries are detected).

- *--with-embedded-crypto* If this option is set, the crypto library is linked directly into libsyslog-ng: the sources of libsyslog-ng-crypto will be appended to the libsyslog-ng sources, and -crypto is not built.

- *--with-ivykis* Specifies which ivykis implementation to use (default value: internal). The source of ivykis is included in the source code package of syslog-ng OSE and is used by default. To use an external implementation instead, use the `--with-ivykis=system` compiling option.

- *--with-libhiredis* Specifies the path to the libhiredis library. For details on using this destination, see *Section 7.9, Storing name-value pairs in Redis (p. 152)*.

- *--with-libmongo-client* Specifies which MongoDB client to use (default value: auto). The source of the mongodb client is included in the source code package of syslog-ng OSE, but the compiler will use an external MongoDB client if it is installed. To force the compiler to use the internal client instead, use the `--with-libmongo-client=internal` compiling option. For details on using this destination, see *Section 7.4, Storing messages in a MongoDB database (p. 129)*.

- *--with-librabbitmq-client* Specifies which RabbitMQ client to use (default value: internal). The source of the rabbitmq client is included in the source code package of syslog-ng OSE and is used by default. To use an external client instead, use the `--with-librabbitmq-client=system` compiling option. For details on using this destination, see *Section 7.1, Publishing messages using AMQP (p. 117)*.

- *--with-module-dir* Specifies a single directory where the syslog-ng OSE Makefile will install the modules.

- *--with-module-path* Specifies a colon-separated (:) list of directories, where the syslog-ng OSE binary will search for modules.

- *--with-timezone-dir* Specifies the directory where syslog-ng looks for the timezone files to resolve the `time-zone()` and `local-time-zone()` options. If not specified, the `/opt/syslog-ng/share/zoneinfo/` and `/usr/share/zoneinfo/` directories are checked, respectively. Note that HP-UX uses a unique file format (`tztab`) to describe the timezone information; that format is currently not supported in syslog-ng. As a workaround, copy the zoneinfo files from another, non-HP-UX system to the `/opt/syslog-ng/share/zoneinfo/` directory of your HP-UX system.

- *--without-compile-date* Removes the compilation date from the binary. For example, as openSUSE checks if recompilation changes the binary to detect if dependent packages need to be rebuilt or not, and including the date changes the binary every time.

3.2. Uninstalling syslog-ng OSE

If you need to uninstall syslog-ng OSE for some reason, you have the following options:

- *If you have installed syslog-ng OSE from a .deb package*: Execute the `dpkg -r syslog-ng` command to remove syslog-ng; or the `dpkg -P syslog-ng` command to remove syslog-ng OSE and the configuration files as well. Note that removing syslog-ng OSE does not restore the syslog daemon used before syslog-ng.

- *If you have installed syslog-ng OSE from an .rpm package*: Execute the `rpm -e syslog-ng` command to remove syslog-ng OSE. Note that removing syslog-ng OSE does not restore the syslog daemon used before syslog-ng OSE.

3.3. Procedure – Configuring Microsoft SQL Server to accept logs from syslog-ng

Purpose:

Complete the following steps to configure your Microsoft SQL Server to enable remote logins and accept log messages from syslog-ng.

Steps:

Step 1. Start the SQL Server Management Studio application. Select **Start > Programs > Microsoft SQL Server 2005 > SQL Server Management Studio**.

Step 2. Create a new database.

Step a.
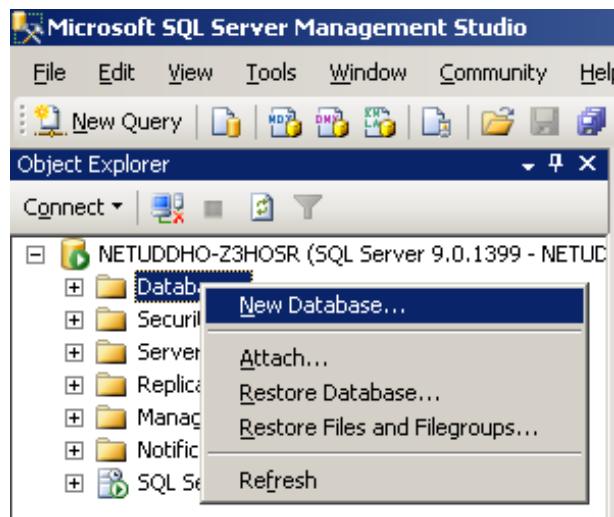
Figure 3.1. Creating a new MSSQL database 1.

In the Object Explorer, right-click on the **Databases** entry and select **New Database**.

Step b.
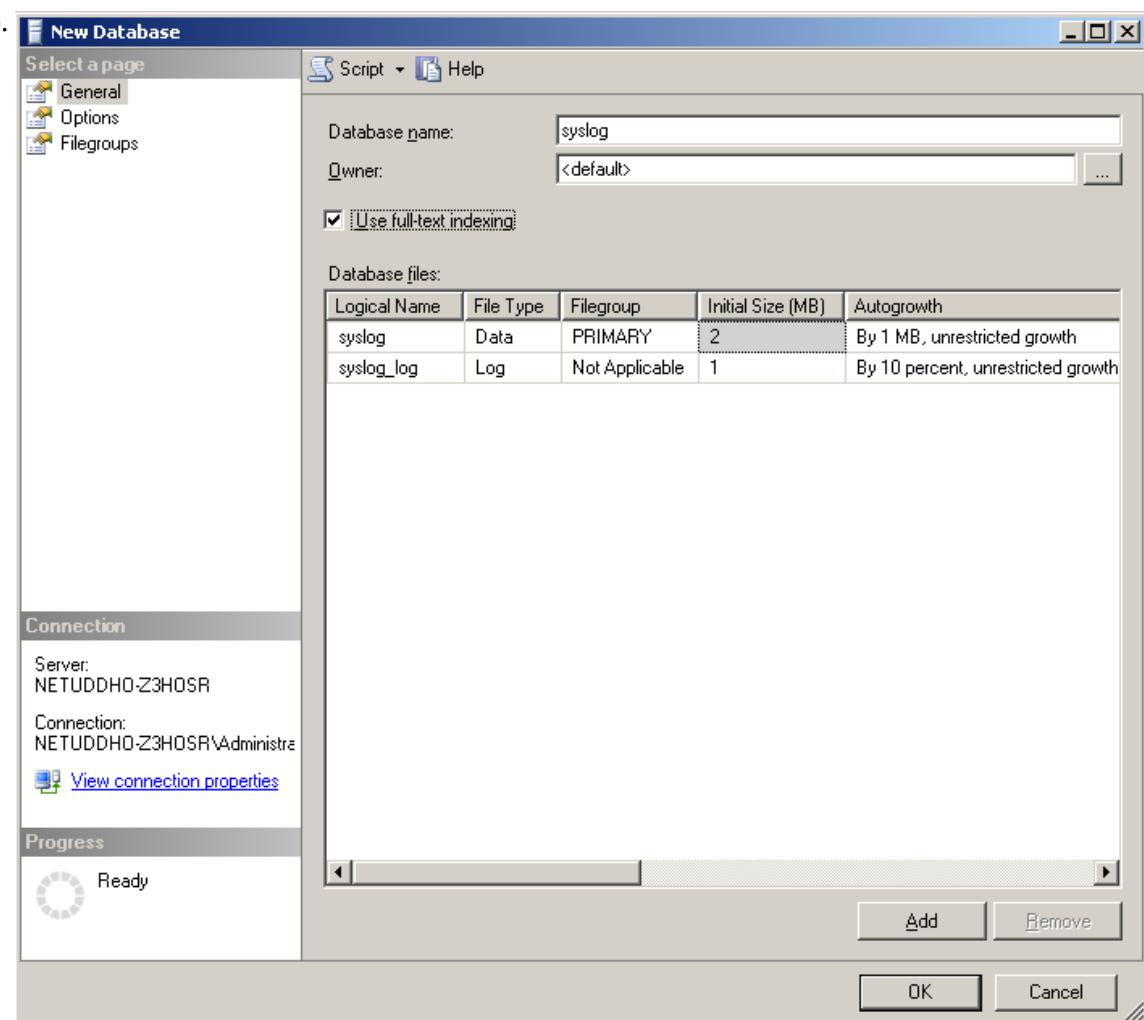
Figure 3.2. Creating a new MSSQL database 2.

Enter the name of the new database (for example `syslogng`) into the **Database name** field and click **OK**.

Step 3. Create a new database user and associate it with the new database.

Step a.
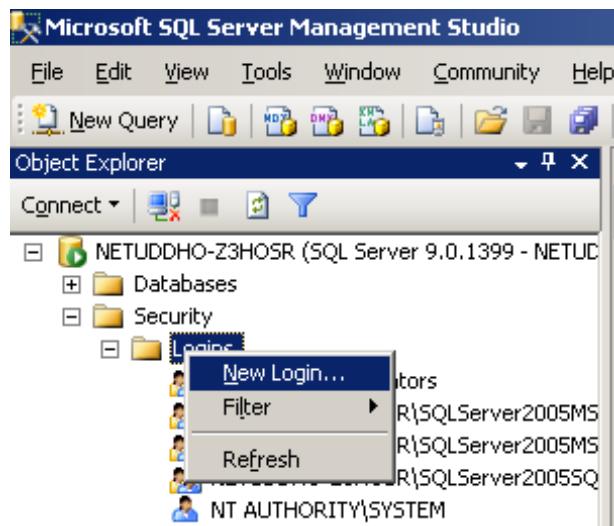
Figure 3.3. Creating a new MSSQL user 1.

In the Object Explorer, select **Security**, right-click on the **Logins** entry, then select **New Login**.

Step b.
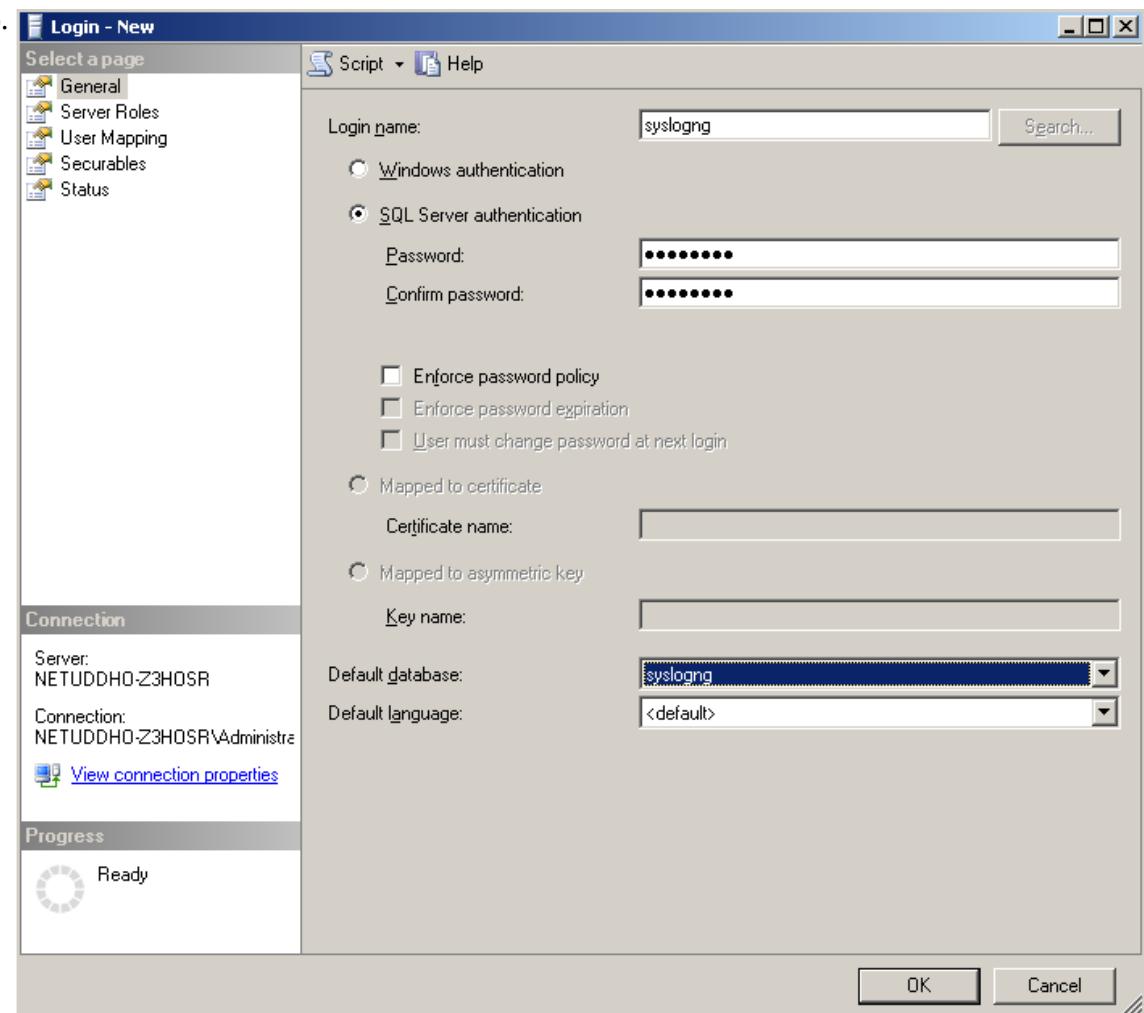
Figure 3.4. Creating a new MSSQL user 2.

Enter a name (for example `syslog-ng`) for the user into the **Login name** field.

Step c. Select the **SQL Server Authentication** option and enter a password for the user.

Step d. In the **Default database** field, select the database created in Step 2 (for example `syslogng`).

Step e. In the **Default language** field, select the language of log messages that you want to store in the database, then click **OK**.

Warning
Incorrect language settings may result in the database converting the messages to a different character-encoding format. That way the log messages may become unreadable, causing information loss.

Step f. In the Object Explorer, select **Security > Logins**, then right-click on the new login created in the previous step, and select **Properties**.

Step g.
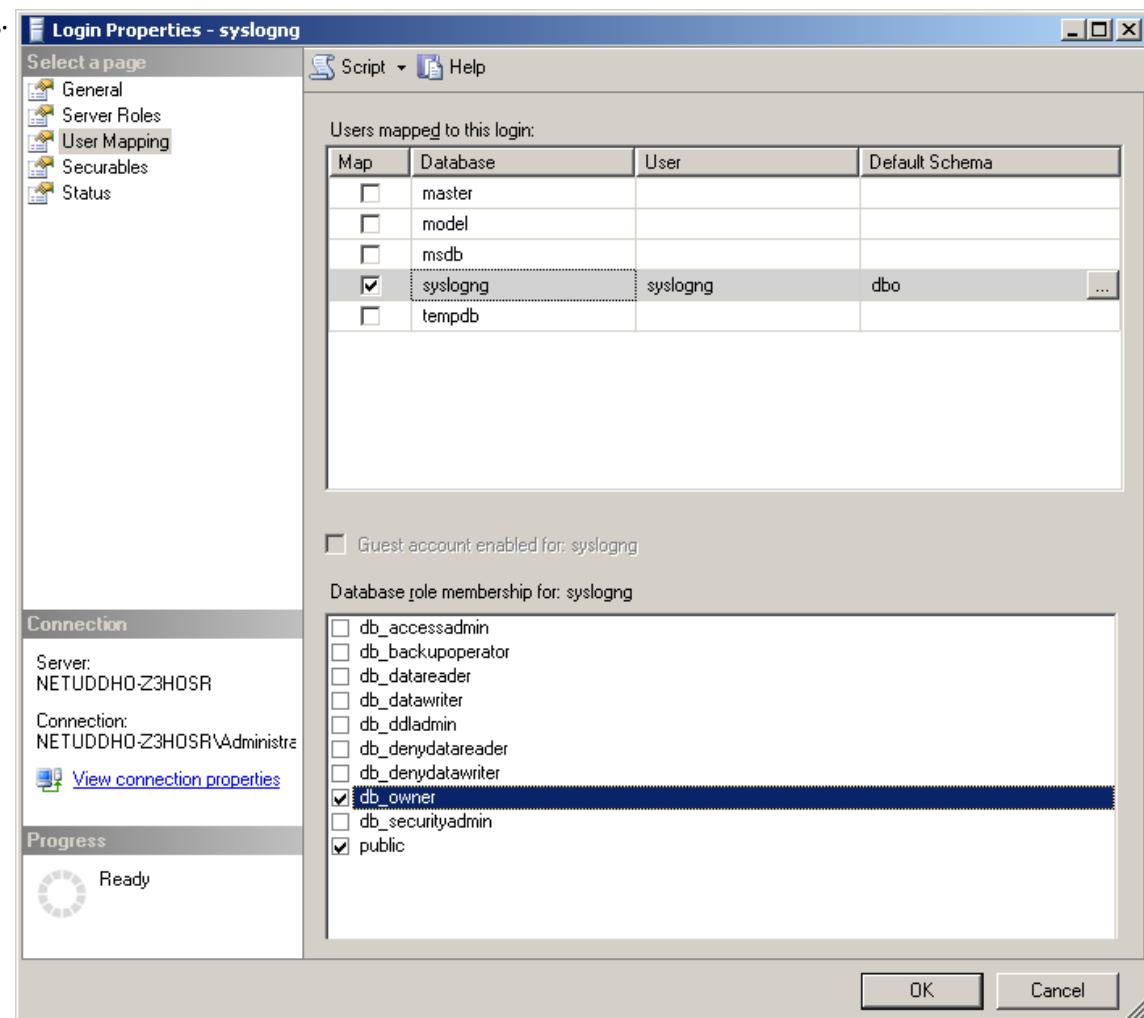
Figure 3.5. Associating database with the new user

Select **User Mapping**. In the **Users mapped to this login** option, check the line
corresponding to the new login (for example `syslogng`). In the **Database role
membership** field, check the **db_owner** and **public** options.

Step 4.
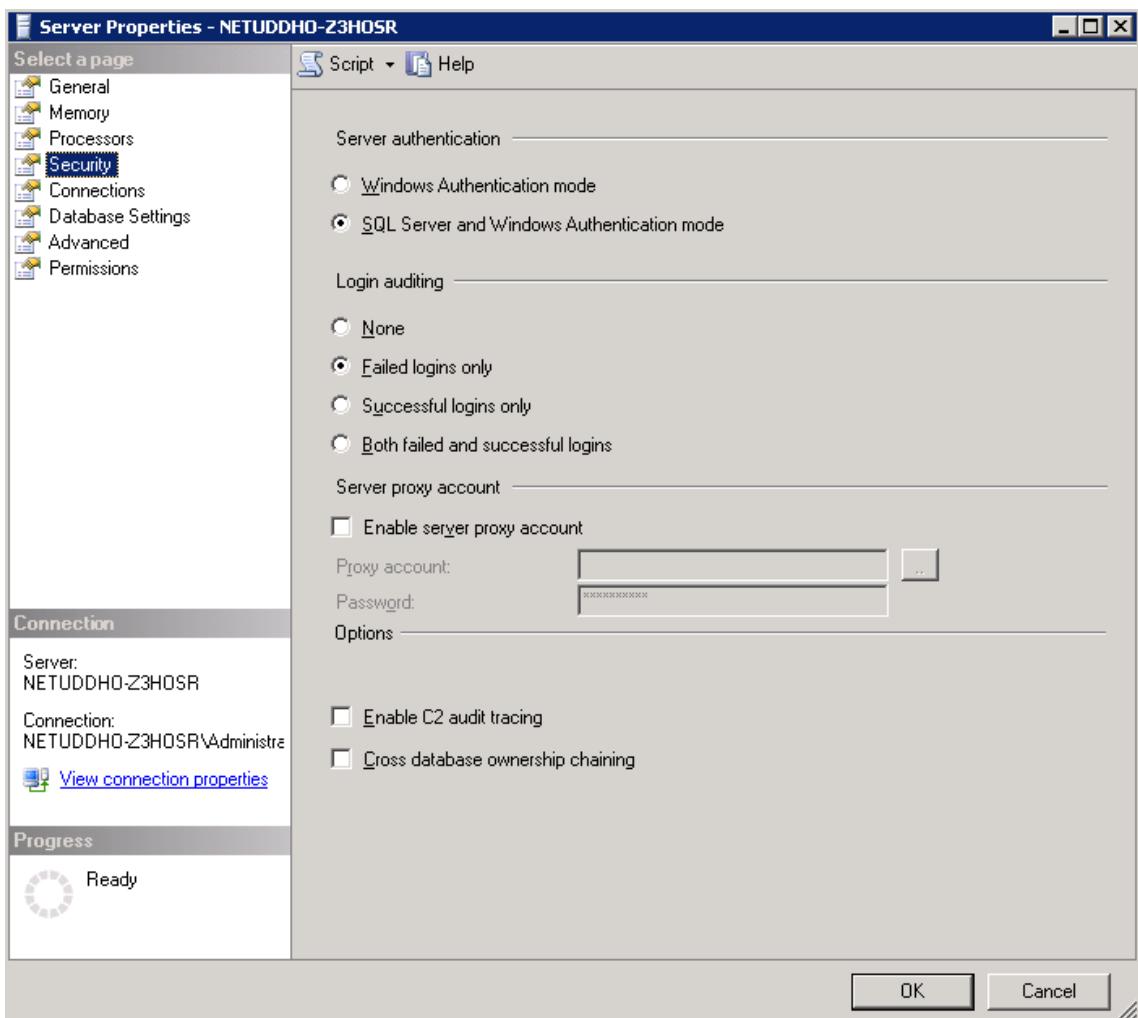
Figure 3.6. Associating database with the new user

Enable remote logins for SQL users.

In the Object Explorer right-click on your database server, and select **Properties > Security**, and set the **Server Authentication** option to **SQL Server and Windows Authentication mode**.

Chapter 4. The syslog-ng OSE quick-start guide

This chapter provides a very brief introduction into configuring the syslog-ng OSE application. For details on the format of the configuration file and how to configure sources, destinations, and other features, refer to the subsequent chapters.

- To configure syslog-ng OSE as a client that sends log messages to a central logserver, see *Procedure 4.1, Configuring syslog-ng on client hosts (p. 34)*.

- To configure syslog-ng OSE as a server that receives log messages from client hosts, see *Procedure 4.2, Configuring syslog-ng on server hosts (p. 36)*.

- To configure syslog-ng OSE as a relay that receives log messages from client hosts and forwards them to a central logserver, see *Procedure 4.2, Configuring syslog-ng on server hosts (p. 36)*.

4.1. Procedure – Configuring syslog-ng on client hosts

Purpose:

To configure syslog-ng on a client host, complete the following steps.

Steps:

Step 1. Install the syslog-ng application on the host. For details installing syslog-ng on specific operating systems, see *Chapter 3, Installing syslog-ng (p. 23)*.

Step 2. Configure the local sources to collect the log messages of the host. Starting with version 3.2, syslog-ng OSE automatically collects the log messages that use the native system logging method of the platform, for example, messages from `/dev/log` on Linux, or `/dev/klog` on FreeBSD. For a complete list of messages that are collected automatically, see *Section 6.11, Collecting the system-specific log messages of a platform (p. 95)*.

Add sources to collect the messages from your log files. File sources look like this:

```
source s_myfilesource {
        file("/var/log/myapplication.log" follow-freq(1)); };
```

Name every source uniquely. For details on configuring file sources, see *Section 6.3, Collecting messages from text files (p. 53)*.

> **Tip**
> Many applications send log messages to logfiles by default (for example, the Roundcube webmail client, or the ProFTPD FTP server), but can be configured to send them to syslog instead. If possible, it is recommended to reconfigure the application that way.

> **Note**
> The default configuration file of syslog-ng OSE collects platform-specific log messages and the intenal log messages of syslog-ng OSE.

```
source s_local {
        system();
        internal();
};
```

Step 3. Create a network destination that points directly to the syslog-ng server, or to a local relay. The network destination greatly depends on the protocol that your log server or relay accepts messages. Many systems still use the legacy BSD-syslog protocol (RFC3162) over the unreliable UDP transport:

```
destination d_network { udp("10.1.2.3"); };
```

However, if possible, use the much more reliable IETF-syslog protocol over TCP transport:

```
destination d_network {syslog(ip("10.1.2.3") transport("tcp")); };
```

Step 4. Create a log statement connecting the local sources to the syslog-ng server or relay. For example:

```
log {
        source(s_local); destination(d_network); };
```

Step 5. If the logs will also be stored locally on the host, create local file destinations.

Note
The default configuration of syslog-ng OSE places the collected messages into the `/var/log/messages` file:

```
destination d_local {
    file("/var/log/messages"); };
```

Step 6. Create a log statement connecting the local sources to the file destination.

Note
The default configuration of syslog-ng OSE has only one log statement:

```
log {
    source(s_local); destination(d_local); };
```

Step 7. Set filters, macros and other features and options (for example TLS encryption) as necessary.

Example 4.1. The default configuration file of syslog-ng OSE
The following is the default configuration file of syslog-ng OSE 3.2. It collects local log messages and the log messages of syslog-ng OSE and forwards them to a logserver using the IETF-syslog protocol.

```
@version: 3.6
@include "scl.conf"
source s_local { system(); internal(); };
destination d_local {
        file("/var/log/messages"); };
log { source(s_local); destination(d_local); };
```

Example 4.2. A simple configuration for clients
The following is a simple configuration file that collects local log messages to the `/var/log/messages` file.

```
@version: 3.6
@include "scl.conf"
source s_local { system(); internal(); };
destination d_syslog_tcp {
             syslog("192.168.1.1" transport("tcp") port(2010)); };
log { source(s_local);destination(d_syslog_tcp); };
```

4.2. Procedure – Configuring syslog-ng on server hosts

Purpose:

To configure syslog-ng on a server host, complete the following steps.

Steps:

Step 1. Install the syslog-ng application on the host. For details installing syslog-ng on specific operating systems, see *Chapter 3, Installing syslog-ng (p. 23)*.

Step 2. Starting with version 3.2, syslog-ng OSE automatically collects the log messages that use the native system logging method of the platform, for example, messages from `/dev/log` on Linux, or `/dev/klog` on FreeBSD. For a complete list of messages that are collected automatically, see *Section 6.11, Collecting the system-specific log messages of a platform (p. 95)*.

Step 3. Configure the network sources that collect the log messages sent by the clients and relays. How the network sources should be configured depends also on the capabilities of your client hosts: many older networking devices support only the legacy BSD-syslog protocol (RFC3164) using UDP transport:

```
source s_network { syslog(ip(10.1.2.3) transport("udp")); };
```

However, if possible, use the much more reliable TCP transport:

```
source s_network { syslog(ip(10.1.2.3) transport("tcp")); };
```

For other options, see *Section 6.10, Collecting messages using the IETF syslog protocol (syslog() driver) (p. 87)* and *Section 6.14, Collecting messages from remote hosts using the BSD syslog protocol (p. 100)*.

Note
Starting with syslog-ng OSE version 3.2, the `syslog()` source driver can handle both BSD-syslog (RFC 3164) and IETF-syslog (RFC 5424-26) messages.

Step 4. Create local destinations that will store the log messages, for example file- or program destinations. The default configuration of syslog-ng OSE places the collected messages into the `/var/log/messages` file:

```
destination d_local {
    file("/var/log/messages"); };
```

If you want to create separate logfiles for every client host, use the *${HOST}* macro when specifying the filename, for example:

```
destination d_local {
    file("/var/log/messages_${HOST}"); };
```

For details on further macros and how to use them, see *Chapter 11, Manipulating messages (p. 238)*.

Step 5. Create a log statement connecting the sources to the local destinations.

```
log {
        source(s_local); source(s_network); destination(d_local); };
```

Step 6. Set filters, options (for example TLS encryption) and other advanced features as necessary.

Note
By default, the syslog-ng server will treat the relayed messages as if they were created by the relay host, not the host that originally sent them to the relay. In order to use the original hostname on the syslog-ng server, use the `keep-hostname(yes)` option both on the syslog-ng relay and the syslog-ng server. This option can be set individually for every source if needed.

If you are relaying log messages and want to resolve IP addresses to hostnames, configure the first relay to do the name resolution.

Example 4.3. A simple configuration for servers
The following is a simple configuration file for syslog-ng Open Source Edition that collects incoming log messages and stores them in a text file.

```
@version: 3.6
@include "scl.conf"
    options {
        time-reap(30);
        mark-freq(10);
        keep-hostname(yes);
        };
    source s_local { system(); internal(); };
    source s_network {
        syslog(transport(tcp));
        };
    destination d_logs {
        file(
            "/var/log/syslog-ng/logs.txt"
            owner("root")
            group("root")
            perm(0777)
            ); };
    log { source(s_local); source(s_network); destination(d_logs); };
```

4.3. Configuring syslog-ng relays

This section describes how to configure syslog-ng OSE as a relay.

4.3.1. Procedure – Configuring syslog-ng on relay hosts

Purpose:

To configure syslog-ng on a relay host, complete the following steps:

Steps:

Step 1. Install the syslog-ng application on the host. For details installing syslog-ng on specific operating systems, see *Chapter 3, Installing syslog-ng (p. 23)*.

Step 2. Configure the network sources that collect the log messages sent by the clients.

Step 3. Create a network destination that points to the syslog-ng server.

Step 4. Create a log statement connecting the network sources to the syslog-ng server.

Step 5. Configure the local sources that collect the log messages of the relay host.

Step 6. Create a log statement connecting the local sources to the syslog-ng server.

Step 7. Enable the `keep-hostname()` and disable the `chain-hostnames()` options.

> **Note**
> It is recommended to use these options on your syslog-ng OSE server as well.

Step 8. Set filters and options (for example TLS encryption) as necessary.

> **Note**
> By default, the syslog-ng server will treat the relayed messages as if they were created by the relay host, not the host that originally sent them to the relay. In order to use the original hostname on the syslog-ng server, use the `keep-hostname(yes)` option both on the syslog-ng relay and the syslog-ng server. This option can be set individually for every source if needed.
>
> If you are relaying log messages and want to resolve IP addresses to hostnames, configure the first relay to do the name resolution.

Example 4.4. A simple configuration for relays
The following is a simple configuration file that collects local and incoming log messages and forwards them to a logserver using the IETF-syslog protocol.

```
@version: @techversion;
@include "scl.conf"
    options {
        time-reap(30);
        mark-freq(10);
        keep-hostname(yes);
        chain-hostnames(no);
        };
    source s_local { system(); internal(); };
    source s_network {
        syslog(transport(tcp));
        };
    destination d_syslog_tcp {
        syslog("192.168.1.5" transport("tcp") port(2010));
    };

    log { source(s_local); source(s_network);
        destination(d_syslog_tcp);
        };
```

4.3.2. How relaying log messages works

Depending on your exact needs about relaying log messages, there are many scenarios and syslog-ng OSE options that influence how the log message will look like on the logserver. Some of the most common cases are summarized in the following example.

Consider the following example: *client-host > syslog-ng-relay > syslog-ng-server*, where the IP address of `client-host` is 192.168.1.2. The `client-host` device sends a syslog message to `syslog-ng-relay`. Depending on the settings of `syslog-ng-relay`, the following can happen.

- By default, the `keep-hostname()` option is disabled, so `syslog-ng-relay` writes the IP address of the sender host (in this case, 192.168.1.2) to the HOST field of the syslog message, discarding any IP address or hostname that was originally in the message.

- If the `keep-hostname()` option is enabled on `syslog-ng-relay`, but name resolution is disabled (the `use-dns()` option is set to no), `syslog-ng-relay` uses the HOST field of the message as-is, which is probably 192.168.1.2.

- To resolve the 192.168.1.2 IP address to a hostname on `syslog-ng-relay` using a DNS server, use the keep-hostname(no) and use-dns(yes) options. If the DNS server is properly configured and reverse DNS lookup is available for the 192.168.1.2 address, syslog-ng OSE will rewrite the HOST field of the log message to `client-host`.

> **Note**
> It is also possible to resolve IP addresses locally, without relying on the DNS server. For details on local name resolution, see *Procedure 17.3.1, Resolving hostnames locally (p. 307)*.

- The above points apply to the syslog-ng OSE server (`syslog-ng-server`) as well, so if `syslog-ng-relay` is configured properly, use the keep-hostname(yes) option on `syslog-ng-server` to retain the proper HOST field. Setting keep-hostname(no) on `syslog-ng-server` would result in syslog-ng OSE rewriting the HOST field to the address of the host that sent the message to `syslog-ng-server`, which is `syslog-ng-relay` in this case.

- If you cannot or do not want to resolve the 192.168.1.2 IP address on `syslog-ng-relay`, but want to store your log messages on `syslog-ng-server` using the IP address of the original host (that is, `client-host`), you can enable the `spoof-source()` option on `syslog-ng-relay`. However, `spoof-source()` works only under the following conditions:

  - The syslog-ng OSE binary has been compiled with the `--enable-spoof-source` option.

  - The log messages are sent using the highly unreliable UDP transport protocol. (Extremely unrecommended.)

Chapter 5. The syslog-ng OSE configuration file

5.1. Location of the syslog-ng configuration file

The syslog-ng application is configured by editing the `syslog-ng.conf` file. Use any regular text editor application to modify the file. The location of the configuration file depends on how you installed syslog-ng OSE. Native packages of a platform (like the ones downloaded from Linux repositories) typically place the configuration file under the `/etc/syslog-ng/` directory.

5.2. The configuration syntax in detail

Every syslog-ng configuration file must begin with a line containing the version information of syslog-ng. For syslog-ng version 3.6, this line looks like:

```
@version: 3.6
```

Versioning the configuration file was introduced in syslog-ng 3.0. If the configuration file does not contain the version information, syslog-ng assumes that the file is for syslog-ng version 2.x. In this case it interprets the configuration and sends warnings about the parts of the configuration that should be updated. Version 3.0 and later will correctly operate with configuration files of version 2.x, but the default values of certain parameters have changed since 3.0.

Example 5.1. A simple configuration file
The following is a very simple configuration file for syslog-ng: it collects the internal messages of syslog-ng and the messages from `/dev/log` into the `/var/log/messages_syslog-ng.log` file.

```
@version: 3.6

source s_local { unix-stream("/dev/log"); internal(); };

destination d_file {file("/var/log/messages_syslog-ng.log"); };

log { source(s_local); destination(d_file); };
```

As a syslog-ng user described on a *mailing list*:

> *The syslog-ng's config file format was written by programmers for programmers to be understood by programmers. That may not have been the stated intent, but it is how things turned out. The syntax is exactly that of C, all the way down to braces and statement terminators.*

> —Alan McKinnon

■ The main body of the configuration file consists of object definitions: sources, destinations, logpaths define which log message are received and where they are sent. All identifiers, option names and attributes, and any other strings used in the syslog-ng configuration file are case sensitive. Objects must be defined before they are referenced in another statement. Object definitions (also called statements) have the following syntax:

```
object_type object_id {<options>};
```

- *Type of the object*: One of `source`, `destination`, `log`, `filter`, `parser`, `rewrite` rule, or `template`.

- *Identifier of the object*: A unique name identifying the object. When using a reserved word as an identifier, enclose the identifier in quotation marks.

Tip
Use identifiers that refer to the type of the object they identify. For example, prefix source objects with s_, destinations with d_, and so on.

Note
Repeating a definition of an object (that is, defining the same object with the same id more than once) is not allowed, unless you use the `@define allow-config-dups 1` definition in the configuration file.

- *Parameters*: The parameters of the object, enclosed in braces {`parameters`}.
- *Semicolon*: Object definitions end with a semicolon (`;`).
For example, the following line defines a source and calls it `s_internal`.

```
source s_internal { internal(); };
```

The object can be later referenced in other statements using its ID, for example, the previous source is used as a parameter of the following log statement:

```
log { source(s_internal); destination(d_file); };
```

- The parameters and options within a statement are similar to function calls of the C programming language: the name of the option followed by a list of its parameters enclosed within brackets and terminated with a semicolon.

```
option(parameter1, parameter2); option2(parameter1, parameter2);
```

For example, the `file()` driver in the following source statement has three options: the filename (`/var/log/apache/access.log`), `follow-freq()`, and `flags()`. The `follow-freq()` option also has a parameter, while the `flags()` option has two parameters.

```
source s_tail { file("/var/log/apache/access.log"
    follow-freq(1) flags(no-parse, validate-utf8)); };
```

Objects may have required and optional parameters. Required parameters are positional, meaning that they must be specified in a defined order. Optional parameters can be specified in any order using the `option(value)` format. If a parameter (optional or required) is not specified, its default value is used. The parameters and their default values are listed in the reference section of the particular object.

Example 5.2. Using required and optional parameters
The `unix-stream()` source driver has a single required argument: the name of the socket to listen on. Optional parameters follow the socket name in any order, so the following source definitions have the same effect:

```
source s_demo_stream1 {
        unix-stream("/dev/log" max-connections(10) group(log)); };
source s_demo_stream2 {
        unix-stream("/dev/log" group(log) max-connections(10)); };
```

- Some options are global options, or can be set globally, for example, whether syslog-ng OSE should use DNS resolution to resolve IP addresses. Global options are detailed in *Chapter 9, Global options of syslog-ng OSE (p. 217)*.

```
options { use-dns(no); };
```

- All identifiers, attributes, and any other strings used in the syslog-ng configuration file are case sensitive.

- Objects can be used before definition.

- Objects can be defined inline as well. This is useful if you use the object only once (for example, a filter). For details, see *Section 5.4, Defining configuration objects inline (p. 43)*.

- To add comments to the configuration file, start a line with # and write your comments. These lines are ignored by syslog-ng.

```
# Comment: This is a stream source
source s_demo_stream {
        unix-stream("/dev/log" max-connections(10) group(log)); };
```

Tip
Before activating a new configuration, check that your configuration file is syntactically correct using the `syslog-ng --syntax-only` command.

To activate the configuration, reload the configuration of syslog-ng using the `/etc/init.d/syslog-ng reload` command.

5.3. Notes about the configuration syntax

When you are editing the syslog-ng configuration file, note the following points:

- The configuration file can contain a maximum of 6665 source / destination / log elements.

- When writing the names of options and parameters (or other reserved words), the hyphen (`-`) and underscore (`_`) characters are equivalent, for example `max-connections(10)` and `max_connections(10)` are both correct.

- Numbers can be prefixed with + or - to indicate positive or negative values. Numbers beginning with zero (0) or 0x are treated as octal or hexadecimal numbers, respectively.

  Starting with syslog-ng OSE version 3.5, you can use suffixes for kilo-, mega-, and gigabytes. Use the Kb, Mb, or Gb suffixes for the base-10 version, and Kib, Mib, or Gib for the base-2 version. That

is, 2MB means 2000000, while 2MiB means 2097152. For example, to set the *log-fifo-size()* option to 2000000 bytes, use *log-fifo-size(2Mb)*.

- You can use commas (,) to separate options or other parameters for readability; syslog-ng completely ignores them. The following declarations are equivalent:

```
source s_demo_stream {
        unix-stream("/dev/log" max-connections(10) group(log)); };
```

```
source s_demo_stream {
        unix-stream("/dev/log", max-connections(10), group(log)); };
```

- When enclosing object IDs (for example the name of a destination) between double-quotes ("mydestination"), the ID can include whitespace as well, for example:

```
source "s demo stream" {
        unix-stream("/dev/log" max-connections(10) group(log)); };
```

- For notes on using regular expressions, see *Section 11.3, Regular expressions (p. 263)*.

5.4. Defining configuration objects inline

Starting with syslog-ng OSE 3.4, you can define configuration objects inline, where they are actually used, without having to define them in a separate placement. This is useful if you need an object only once, for example, a filter or a rewrite rule. Every object can be defined inline: sources, destinations, filters, parsers, rewrite rules, and so on.

To define an object inline, use braces instead of parentheses. That is, instead of *<object-type> (<object-id>);*, you use *<object-type> {<object-definition>};*

Example 5.3. Using inline definitions
The following two configuration examples are equivalent. The first one uses traditional statements, while the second uses inline definitions.

```
source s_local {
    system();
    internal();
};
destination d_local {
    file("/var/log/messages");
};
log {
    source(s_local);
    destination(d_local);
};
```

```
log {
    source {
        system();
        internal();
    };
    destination {
        file("/var/log/messages");s
    };
};
```

5.5. Using channels in configuration objects

Starting with syslog-ng OSE 3.4, every configuration object is a log expression. Every configuration object is essentially a configuration block, and can include multiple objects. To reference the block, only the top-level object must be referenced. That way you can use embedded log statements, junctions and in-line object definitions within source, destination, filter, rewrite and parser definitions. For example, a source can include a rewrite rule to modify the messages received by the source, and that combination can be used as a simple source in a log statement. This feature allows you to preprocess the log messages very close to the source itself.

To embed multiple objects into a configuration object, use the following syntax. Note that you must enclose the configuration block between braces instead of parenthesis.

```
<type-of-top-level-object> <name-of-top-level-object> {
    channel {
        <configuration-objects>
    }
}
```

Example 5.4. Using channels
For example, to process a log file in a specific way, you can define the required processing rules (parsers and rewrite expressions) and combine them in a single object:

```
source s_apache {
    channel {
        source { file("/var/log/apache/error.log"); };
        parser(p_apache_parser); };
  };
};

log { source(s_apache); ... };
```

The *s_apache* source uses a file source (the error log of an Apache webserver) and references a specific parser to process the messages of the error log. The log statement references only the *s_apache* source, and any other object in the log statement can already use the results of the *p_apache_parser* parser.

Note
You must start the object definition with a *channel* even if you will use a *junction*, for example:

```
parser demo-parser() {
    channel {
        junction {
            channel { ... }
            channel { ... }
        }
    }
}
```

If you want to embed configuration objects into sources or destinations, always use channels, otherwise the source or destination will not behave as expected. For example, the following configuration is good:

```
source s_filtered_hosts {
    pipe("/dev/pipe");
    syslog(ip(192.168.0.1) transport("tcp"));
    syslog(ip(127.0.0.1) transport("tcp"));
    filter (netmask(10.0.0.0/16));
}
```

5.6. Global and environmental variables

Starting with syslog-ng OSE version 3.2, it is possible to define global variables in the configuration file. Global variables are actually `name-value` pairs; when syslog-ng processes the configuration file during startup, it automatically replaces `name` with `value`. To define a global variable, use the following syntax:

```
@define name "value"
```

The value can be any string, but special characters must be escaped.To use the variable, insert the name of the variable enclosed between backticks (`` ` ``, similarly to using variables in Linux or UNIX shells) anywhere in the configuration file.

The value of the global variable can be also specified using the following methods:

- Without any quotes, as long as the value does not contain any spaces or special characters. In other word, it contains only the following characters: `a-zA-Z0-9_..`
- Between apostrophes, in case the value does not contain apostrophes.
- Between double quotes, in which case special characters must be escaped using backslashes (\).

Tip
The environmental variables of the host are automatically imported and can be used as global variables.

Example 5.5. Using global variables
For example, if an application is creating multiple log files in a directory, you can store the path in a global variable, and use it in your source definitions.

```
@define mypath "/opt/myapp/logs"
        source s_myapp_1 { file("`mypath`/access.log" follow-freq(1)); };
        source s_myapp_2 { file("`mypath`/error.log" follow-freq(1)); };
        source s_myapp_3 { file("`mypath`/debug.log" follow-freq(1)); };
```

The syslog-ng OSE application will interpret this as:

```
@define mypath "/opt/myapp/logs"
        source s_myapp_1 { file("/opt/myapp/logs/access.log" follow-freq(1)); };
        source s_myapp_2 { file("/opt/myapp/logs/error.log" follow-freq(1)); };
        source s_myapp_3 { file("/opt/myapp/logs/debug.log" follow-freq(1)); };
```

5.7. Modules in syslog-ng OSE

Starting with syslog-ng Open Source Edition version 3.3, syslog-ng OSE became modular to increase its flexibility and also to simplify the development of additional modules. Most of the functionality of syslog-ng OSE has been moved to separate modules. That way it becomes also possible to finetune the resource requirements of syslog-ng OSE for example, by loading only the modules that are actually used in the configuration, or simply omitting modules that are not used but require large amount of memory.

Each module contains one or more plugins, which add some functionality to syslog-ng OSE, for example, a destination or a source driver.

- To display the list of available modules, execute the `syslog-ng --version` command.

- To the description of the available modules, execute the `syslog-ng --module-registry` command.

- To customize which modules are loaded automatically when syslog-ng OSE is started, use the `--default-modules` command-line option of syslog-ng OSE.

- To request loading a module from the syslog-ng OSE configuration file, see *Section 5.7.1, Loading modules (p. 46)*.

For details on the command-line parameters of syslog-ng OSE mentioned in the previous list, see the syslog-ng OSE man page at *syslog-ng(8) (p. 321)*.

5.7.1. Loading modules

The syslog-ng Open Source Edition application loads every available module during startup.

To load a module that is not loaded automatically, include the following statement in the syslog-ng OSE configuration file:

```
@module <module-name>
```

Note the following points about the *@module* statement:

- The *@module* statement is a top-level statement, that is, it cannot be nested into any other statement. Usually it is used immediately after the *@version* statement.

- Every *@module* statement loads a single module: loading multiple modules requires a separate *@module* statement for every module.

- In the configuration file, the *@module* statement of a module must be earlier than the module is used.

Note
To disable loading every module automatically, set the *autoload-compiled-modules* global variable to 0 in your configuration file:

```
@define autoload-compiled-modules O
```

Note that in this case, you have to explicitly load the modules you want to use.

5.8. Managing complex syslog-ng configurations

The following sections describe some methods that can be useful to simplify the management of large-scale syslog-ng installations.

5.8.1. Including configuration files

The syslog-ng application supports including external files in its configuration file, so parts of its configuration can be managed separately. To include the contents of a file in the syslog-ng configuration, use the following syntax:

```
include "<filename>";
```

This imports the entire file into the configuration of syslog-ng OSE, at the location of the include statement. The <filename> can be one of the following:

- A filename, optionally with full path. The filename (not the path) can include UNIX-style wildcard characters (*, ?). When using wildcard characters, syslog-ng OSE will include every matching file. For details on using wildcard characters, see *Section glob (p. 265)*.

- A directory. When including a directory, syslog-ng OSE will try to include every file from the directory, except files beginning with a ~ (tilde) or a . (dot) character. Including a directory is not recursive. The files are included in alphabetic order, first files beginning with uppercase characters, then files beginning with lowercase characters. For example, if the directory contains the a.conf, B. conf, c.conf, D.conf files, they will be included in the following order: B.conf, D. conf, a.conf, c.conf.

When including configuration files, consider the following points:

- Defining an object twice is not allowed, unless you use the *@define allow-config-dups 1* definition in the configuration file. If an object is defined twice (for example the original syslog-ng configuration file and the file imported into this configuration file both define the same option, source, or other object), then the object that is defined later in the configuration file will be effective. For example, if you set a global option at the beginning of the configuration file, and later include a file that defines the same option with a different value, then the option defined in the imported file will be used.

- Files can be embedded into each other: the included files can contain include statements as well, up to a maximum depth of 15 levels.

- You cannot include complete configuration files into each other, only configuration snippets can be included. This means that the included file cannot have a *@version* statement.

- Include statements can only be used at top level of the configuration file. For example, the following is correct:

```
@version: 3.6
include "example.conf";
```

But the following is not:

```
source s_example {
    include "example.conf"
     };
```

Warning
The syslog-ng application will not start if it cannot find a file that is to be included in its configuration. Always double-check the filenames, paths, and access rights when including configuration files, and use the --syntax-only command-line option to check your configuration.

5.8.2. Reusing configuration blocks

Starting with syslog-ng OSE 3.2, parts of a configuration file can be easily reused, you have to define the block (for example, a source) once, and reference it later. Any syslog-ng object can be a block. Use the following syntax to define a block:

```
block type name() {<contents of the block>}
```

Type must be one of the following: *destination*, *filter*, *log*, *parser*, *rewrite*, *root*, *source*. The *root* blocks can be used in the "root" context of the configuration file, that is, outside any other statements.

Blocks may be nested into each other, so for example a block can be built from other blocks. Blocks are somewhat similar to C++ templates.

The type and name combination of each block must be unique, that is, two blocks can have the same name if their type is different.

To use a block in your configuration file, you have to do two things:

- Include the file defining the block in the `syslog-ng.conf` file — or a file already included into `syslog-ng.conf`.
- Reference the name of the block in your configuration file. This will insert the block into your configuration. For example, to use a block called `myblock`, include the following line in your configuration:

  ```
  myblock()
  ```

  Blocks may have parameters, but even if they do not, the reference must include opening and closing parentheses like in the previous example.

The contents of the block will be inserted into the configuration when syslog-ng OSE is started or reloaded.

Example 5.6. Reusing configuration blocks
Suppose you are running an application on your hosts that logs into the `/opt/var/myapplication.log` file. Create a file (for example, `myblocks.conf`) that stores a source describing this file and how it should be read:

```
block source myappsource() {
            file("/opt/var/myapplication.log" follow-freq(1) default-facility(syslog));
 };
```

Include this file in your main syslog-ng configuration file,reference the block, and use it in a logpath:

```
@version: 3.6
include "<correct/path>/myblocks.conf";
source s_myappsource { myappsource(); }
...
log { source(s_myappsource); destination(…); };
```

To define a block that defines more than one object, use *root* as the type of the block, and reference the block from the main part of the syslog-ng OSE configuration file.

Example 5.7. Defining blocks with multiple elements
The following example defines a source, a destination, and a log path to connect them.

```
block root mylogs() {
        source s_file { file("/var/log/mylogs.log" follow-freq(1)); };
        destination d_local { file("/var/log/messages"); };
        log { source(s_file); destination(d_local); };
        };
```

Tip
Since the block is inserted into the syslog-ng OSE configuration when syslog-ng OSE is started, the block can be generated dynamically using an external script if needed. This is useful when you are running syslog-ng OSE on different hosts and you want to keep the main configuration identical.

If you want to reuse more than a single configuration object, for example, a logpath and the definitions of its sources and destinations, use the include feature to reuse the entire snippet. For details, see *Section 5.8.1, Including configuration files (p. 46)*.

5.8.2.1. Passing arguments to configuration blocks

Configuration blocks can receive arguments as well. The parameters the block can receive must be specified when the block is defined, using the following syntax:

```
block type block_name(argument1(<default-value-of-the-argument>)
argument2(<default-value-of-the-argument>) argument3())
```

If an argument does not have a default value, use empty parentheses after the name of the argument. To refer the value of the argument in the block, use the name of the argument between backticks (for example, `argument1`).

Example 5.8. Passing arguments to blocks
The following sample defines a file source block, which can receive the name of the file as a parameter. If no parameter is set, it reads messages from the /var/log/messages file.

```
block source s_logfile (filename("messages")) {
  file("/var/log/`filename`" );
};

source s_example {
  s_logfile(filename("logfile.log"));
};
```

Example 5.9. Using arguments in blocks
The following example is the code of the *pacct() source driver*, which is actually a block that can optionally receive two arguments.

```
block source pacct(file("/var/log/account/pacct") follow-freq(1)) {
@module pacctformat
        file("`file`" follow-freq(`follow-freq`) format("pacct") tags(".pacct"));
};
```

Chapter 6. Collecting log messages — sources and source drivers

6.1. How sources work

A source is where syslog-ng receives log messages. Sources consist of one or more drivers, each defining where and how messages are received.

To define a source, add a source statement to the syslog-ng configuration file using the following syntax:

```
source <identifier> { source-driver(params); source-driver(params); ... };
```

Example 6.1. A simple source statement
The following source statement receives messages on the TCP port 1999 of the interface having the 10.1.2.3 IP address.

```
source s_demo_tcp { tcp(ip(10.1.2.3) port(1999)); };
```

Example 6.2. A source statement using two source drivers
The following source statement receives messages on the 1999 TCP port and the 1999 UDP port of the interface having the 10.1.2.3 IP address.

```
source s_demo_two_drivers {
        tcp(ip(10.1.2.3) port(1999));
        udp(ip(10.1.2.3) port(1999)); };
```

Example 6.3. Setting default priority and facility
If the message received by the source does not have a proper syslog header, you can use the *default-facility()* and *default-priority()* options to set the facility and priority of the messages. Note that these values are applied only to messages that do not set these parameters in their header.

```
source headerless_messages { udp(default-facility(syslog) default-priority(emerg)); };
```

Define a source only once. The same source can be used in several log paths. Duplicating sources causes syslog-ng to open the source (TCP/IP port, file, and so on) more than once, which might cause problems. For example, include the /dev/log file source only in one source statement, and use this statement in more than one log path if needed.

Warning
Sources and destinations are initialized only when they are used in a log statement. For example, syslog-ng OSE starts listening on a port or starts polling a file only if the source is used in a log statement. For details on creating log statements, see *Chapter 8, Routing messages: log paths and filters (p. 198)*.

To collect log messages on a specific platform, it is important to know how the native *syslogd* communicates on that platform. The following table summarizes the operation methods of *syslogd* on some of the tested platforms:

Platform	Method
Linux	A *SOCK_STREAM* unix socket named /dev/log; some of the distributions switched over to using *SOCK_DGRAM*, though applications still work with either method. Newer distributions that use systemd collect log messages into a journal file.
BSD flavors	A *SOCK_DGRAM* unix socket named /var/run/log.
Solaris (2.5 or below)	An SVR4 style *STREAMS* device named /dev/log.
Solaris (2.6 or above)	In addition to the *STREAMS* device used in earlier versions, 2.6 uses a new multithreaded IPC method called door. By default the door used by *syslogd* is /etc/.syslog_door.
HP-UX 11 or later	HP-UX uses a named pipe called /dev/log that is padded to 2048 bytes, for example source s_hp-ux {pipe ("/dev/log" pad-size(2048)}.
AIX 5.2 and 5.3	A *SOCK_STREAM* or *SOCK_DGRAM* unix socket called /dev/log.

Table 6.1. Communication methods used between the applications and syslogd

Each possible communication mechanism has a corresponding source driver in syslog-ng. For example, to open a unix socket with *SOCK_DGRAM* style communication use the driver *unix-dgram*. The same socket using the *SOCK_STREAM* style — as used under Linux — is called *unix-stream*.

Example 6.4. Source statement on a Linux based operating system
The following source statement collects the following log messages:

- *internal()*: Messages generated by syslog-ng.

- *udp(ip(0.0.0.0) port(514))*: Messages arriving to the 514/UDP port of any interface of the host.

- *unix-stream("/dev/log");*: Messages arriving to the /dev/log socket.

```
source s_demo {
        internal();
        udp(ip(0.0.0.0) port(514));
        unix-stream("/dev/log"); };
```

The following table lists the source drivers available in syslog-ng.

Name	Description
file()	Opens the specified file and reads messages.
network()	Listens on the specified TCP or UDP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively. (This driver is a unified wrapper for the *tcp()*, *tcp6()*, *udp()*, and *udp6()* drivers.
nodejs()	Receives JSON messages from nodejs applications.

Name	Description
internal()	Messages generated internally in syslog-ng.
pacct()	Reads messages from the process accounting logs on Linux.
pipe()	Opens the specified named pipe and reads messages.
program()	Opens the specified application and reads messages from its standard output.
sun-stream(), sun-streams()	Opens the specified *STREAMS* device on Solaris systems and reads incoming messages.
syslog()	Listens for incoming messages using the new *IETF-standard syslog protocol*.
system()	Automatically detects which platform syslog-ng OSE is running on, and collects the native log messages of that platform.
systemd-journal()	Collects messages directly from the journal of platforms that use systemd.
systemd-syslog()	Collects messages from the journal using a socket on platforms that use systemd.
tcp(), tcp6()	Listens on the specified TCP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively.
udp(), udp6()	Listens on the specified UDP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively.
unix-dgram()	Opens the specified unix socket in *SOCK_DGRAM* mode and listens for incoming messages.
unix-stream()	Opens the specified unix socket in *SOCK_STREAM* mode and listens for incoming messages.

Table 6.2. Source drivers available in syslog-ng

6.2. Collecting internal messages

All messages generated internally by syslog-ng use this special source. To collect warnings, errors and notices from syslog-ng itself, include this source in one of your source statements.

```
internal()
```

The syslog-ng application will issue a warning upon startup if none of the defined log paths reference this driver.

Example 6.5. Using the internal() driver

```
source s_local { internal(); };
```

The syslog-ng OSE application sends the following message types from the internal() source:

- *fatal*: Priority value: critical (2); Facility value: syslog (5)

- *error*: Priority value: error (3); Facility value: syslog (5)

- *warning*: Priority value: warning (4); Facility value: syslog (5)

- *notice*: Priority value: notice (5); Facility value: syslog (5)

- *info*: Priority value: info (6); Facility value: syslog (5)

6.2.1. internal() source options

The `internal()` driver has the following options:

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

6.3. Collecting messages from text files

Collects log messages from plain-text files, for example from the logfiles of an Apache webserver.

The syslog-ng application notices if a file is renamed or replaced with a new file, so it can correctly follow the file even if logrotation is used. When syslog-ng is restarted, it records the position of the last sent log message in the `/opt/syslog-ng/var/syslog-ng.persist` file, and continues to send messages from this position after the restart.

The file driver has a single required parameter specifying the file to open. For the list of available optional parameters, see *Section 6.3.2, file() source options (p. 54)*.

```
Declaration:
          file(filename);
```

Example 6.6. Using the file() driver

```
source s_file { file("/var/log/messages"); };
```

Example 6.7. Tailing files
The following source checks the `access.log` file every second for new messages.

```
source s_tail { file("/var/log/apache/access.log"
             follow-freq(1) flags(no-parse)); };
```

Note
If the message does not have a proper syslog header, syslog-ng treats messages received from files as sent by the `kern` facility. Use the `default-facility` and `default-priority` options in the source definition to assign a different facility if needed.

6.3.1. Notes on reading kernel messages

Note the following points when reading kernel messages on various platforms.

- The kernel usually sends log messages to a special file (`/dev/kmsg` on BSDs, `/proc/kmsg` on Linux). The `file()` driver reads log messages from such files. The syslog-ng application can periodically check the file for new log messages if the `follow-freq()` option is set.

- On Linux, the `klogd` daemon can be used in addition to syslog-ng to read kernel messages and forward them to syslog-ng. `klogd` used to preprocess kernel messages to resolve symbols and so on, but as this is deprecated by `ksymoops` there is really no point in running both `klogd` and syslog-ng in parallel. Also note that running two processes reading `/proc/kmsg` at the same time might result in dead-locks.

- When using syslog-ng to read messages from the `/proc/kmsg` file, syslog-ng automatically disables the `follow-freq()` parameter to avoid blocking the file.

- To read the kernel messages on HP-UX platforms, use the following options in the source statement:

```
file("/dev/klog" program-override("kernel") flags(kernel) follow-freq(0));
```

6.3.2. file() source options

The `file()` driver has the following options:

default-facility()

Type:	facility string
Default:	kern

Description: This parameter assigns a facility value to the messages received from the file source, if the message does not specify one.

default-priority()

Type:	priority string
Default:	

Description: This parameter assigns an emergency level to the messages received from the file source, if the message does not specify one.

file()

Type:	filename with path
Default:	

Description: The file to read messages from, including the path.

encoding()

Type:	string
Default:	

Description: Specifies the characterset (encoding, for example UTF-8) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the `iconv -l` command.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the `LOG_KERN | LOG_NOTICE` priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

```
source s_dell { udp(port(2000) flags(no-hostname)); };
```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually

supports multi-line messages. Currently the , `syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before `msg` in the following message: `Jan 22 10:06:11 host program:msg`). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

follow-freq()

Type:	number
Default:	1

Description: Indicates that the source should be checked periodically. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use `poll()` on the file, but checks whether the file changed every time the `follow-freq()` interval (in seconds) has elapsed. Floating-point numbers (for example `1.5`) can be used as well.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if `log-fetch-limit()` is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. Make sure that `log-iw-size()` is larger than the value of `log-fetch-limit()`.

log-msg-size()

Type:	number
Default:	Use the global `log-msg-size()` option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the _global option_ if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use `program-override()` instead.

multi-line-garbage()

Type:	regular expression
Default:	empty string

Description: Use the `multi-line-garbage()` option when processing multi-line messages that contain unneeded parts between the messages. Specify a string or regular expression that matches the beginning of the unneeded message parts. If the `multi-line-garbage()` option is set, syslog-ng OSE ignores the lines between the line matching the `multi-line-garbage()` and the next line matching `multi-line-prefix()`. See also the `multi-line-prefix()` option.

When receiving multi-line messages from a source when the `multi-line-garbage()` option is set, but no matching line is received between two lines that match `multi-line-prefix()`, syslog-ng OSE will continue to process the incoming lines as a single message until a line matching `multi-line-garbage()` is received.

To use the `multi-line-garbage()` option, set the `multi-line-mode()` option to `prefix-garbage`.

Warning

If the `multi-line-garbage()` option is set, syslog-ng OSE discards lines between the line matching the `multi-line-garbage()` and the next line matching `multi-line-prefix()`.

Note

A message is considered complete if no new lines arrive to the message for 10 seconds, even if no line matching the `multi-line-garbage()` or the `multi-line-suffix()` option is received.

multi-line-mode()

Type:	indented|regexp
Default:	empty string

Description: Use the `multi-line-mode()` option when processing multi-line messages. The syslog-ng OSE application provides the following methods to process multi-line messages: `multi-line-mode(indented)`, and `multi-line-mode(prefix-garbage)`.

- The *indented* mode can process messages where each line that belongs to the previous line is indented by whitespace, and the message continues until the first non-indented line. For example, the Linux kernel (starting with version 3.5) uses this format for `/dev/log`, as well as several applications, like Apache Tomcat.

Example 6.8. Processing indented multi-line messages

```
source s_tomcat {
    file("/var/log/tomcat/xxx.log" multi-line-mode(indented));
};
```

- The *prefix-garbage* mode uses a string or regular expression (set in `multi-line-prefix()`) that matches the beginning of the log messages, ignores newline characters from the source until a line matches the regular expression again, and treats the lines between the matching lines as a single message. For details on using `multi-line-mode(prefix-garbage)`, see the `multi-line-prefix()` and `multi-line-garbage()` options.

- The `prefix-suffix` mode uses a string or regular expression (set in `multi-line-prefix()`) that matches the beginning of the log messages, ignores newline characters from the source until a line matches the regular expression set in `multi-line-suffix()`, and treats the lines between `multi-line-prefix()` and `multi-line-suffix()` as a single message. Any other lines between the end of the message and the beginning of a new message (that is, a line that matches the `multi-line-prefix()` expression) is discarded. For details on using `multi-line-mode(prefix-suffix)`, see the `multi-line-prefix()` and `multi-line-suffix()` options.

  The `prefix-suffix` mode is similar to the `prefix-garbage` mode, but it appends the garbage part to the message instead of discarding it.

multi-line-prefix()

Type:	regular expression
Default:	empty string

Description: Use the `multi-line-prefix()` option to process multi-line messages, that is, log messages that contain newline characters (for example, Tomcat logs). Specify a string or regular expression that matches the beginning of the log messages. Use as simple regular expressions as possible, because complex regular expressions can severely reduce the rate of processing multi-line messages. If the `multi-line-prefix()` option is set, syslog-ng OSE ignores newline characters from the source until a line matches the regular expression again, and treats the lines between the matching lines as a single message. See also the `multi-line-garbage()` option.

Note

A message is considered complete if no new lines arrive to the message for 10 seconds, even if no line matching the `multi-line-garbage()` or the `multi-line-suffix()` option is received.

Tip

- To make multi-line messages more readable when written to a file, use a template in the destination and instead of the `${MESSAGE}` macro, use the following: `$(indent-multi-line ${MESSAGE})`. This expression inserts a tab after every newline character (except when a tab is already present), indenting every line of the message after the first. For example:

```
destination d_file {
    file ("/var/log/messages"
        template("${ISODATE} ${HOST} $(indent-multi-line ${MESSAGE})\n") );
};
```

For details on using templates, see *Section 11.1.2, Templates and macros (p. 239)*.

- To actually convert the lines of multi-line messages to single line (by replacing the newline characters with whitespaces), use the `flags(no-multi-line)` option in the source.

Example 6.9. Processing Tomcat logs
The log messages of the Apache Tomcat server are a typical example for multi-line log messages. The messages start with the date and time of the query in the `YYYY.MM.DD HH:MM:SS` format, as you can see in the following example.

```
2010.06.09. 12:07:39 org.apache.catalina.startup.Catalina start
SEVERE: Catalina.start:
LifecycleException:  service.getName(): "Catalina";  Protocol handler start failed:
java.net.BindException: Address already in use<null>:8080
        at org.apache.catalina.connector.Connector.start(Connector.java:1138)
        at org.apache.catalina.core.StandardService.start(StandardService.java:531)
        at org.apache.catalina.core.StandardServer.start(StandardServer.java:710)
        at org.apache.catalina.startup.Catalina.start(Catalina.java:583)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:39)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:25)
        at java.lang.reflect.Method.invoke(Method.java:597)
        at org.apache.catalina.startup.Bootstrap.start(Bootstrap.java:288)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:39)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:25)
        at java.lang.reflect.Method.invoke(Method.java:597)
        at org.apache.commons.daemon.support.DaemonLoader.start(DaemonLoader.java:177)
2010.06.09. 12:07:39 org.apache.catalina.startup.Catalina start
INFO: Server startup in 1206 ms
2010.06.09. 12:45:08 org.apache.coyote.http11.Http11Protocol pause
INFO: Pausing Coyote HTTP/1.1 on http-8080
2010.06.09. 12:45:09 org.apache.catalina.core.StandardService stop
INFO: Stopping service Catalina
```

To process these messages, specify a regular expression matching the timestamp of the messages in the `multi-line-prefix()` option. Such an expression is the following:

```
source s_file{file("/var/log/tomcat6/catalina.2010-06-09.log" follow-freq(0)
multi-line-mode(regexp) multi-line-prefix("[0-9]{4}\.[0-9]{2}\.[0-9]{2}\.")
flags(no-parse));};
};
```

Note that the `flags(no-parse)` is needed to avoid syslog-ng OSE trying to interpret the date in the message.

optional()

Type:	yes or no
Default:	

Description: Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the `pipe()`, `unix-dgram`, and `unix-stream` drivers.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in `pad-size()`. Mostly used on HP-UX where `/dev/log` is a named pipe and every write is padded to 2048 bytes. If `pad-size()` was given and the incoming message does not fit into `pad-size()`, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

6.4. Collecting messages using the RFC3164 protocol (network() driver)

The network() source driver can receive syslog messages conforming to RFC3164 from the network using the TCP, TLS, and UDP networking protocols.

Example 6.10. Using the network() driver
TCP source listening on the localhost on port 2222 without using the network() driver.

```
source s_tcp6 {
        tcp6(
            ip("::1")
            port(2222)
        );
};
```

TCP source listening on the localhost on port 2222 using the network() driver.

```
source s_network6 {
        network(
            ip("::1")
            transport("tcp")
            port(2222)
            ip-protocol(6)
        );
};
```

Note
For details on the `tcp()`, `tcp6()`, `udp()`, `udp6()` drivers, see *Section 6.14, Collecting messages from remote hosts using the BSD syslog protocol (p. 100).*

6.4.1. network() source options

The *network()* driver has the following options.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The *assume-utf8* flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the *validate-utf8* flag.

- *empty-lines*: Use the *empty-lines* flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the *expect-hostname* flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the *no-hostname* flag by default.

- *kernel*: The *kernel* flag makes the source default to the LOG_KERN | LOG_NOTICE priority if not specified otherwise.

- *no-hostname*: Enable the *no-hostname* flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

```
source s_dell { udp(port(2000) flags(no-hostname)); };
```

- *no-multi-line*: The *no-multi-line* flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the *, syslog(), udp(), unix-dgram()* drivers support multi-line messages; other drivers, for example, the *tcp()* driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The *no-parse* flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before `msg` in the following message: `Jan 22 10:06:11 host program:msg`). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

- *threaded*: The `threaded` flag enables multithreading for the source. For details on multithreading, see *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)*.

> **Note**
> The `syslog` source uses multiple threads only if the source uses the `tls` or `tcp` transport protocols.

host-override()

Type:	string
Default:	

Description: Replaces the ${HOST} part of the message with the parameter string.

ip() or localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to. Note that this is not the address where messages are accepted from.

ip-protocol()

Type:	number
Default:	4

Description: Determines the internet protocol version of the given driver (`network()` or `syslog()`). The possible values are 4 and 6, corresponding to IPv4 and IPv6. The default value is 4.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

ip-tos()

Type:	number
Default:	0

Description: Specifies the Type-of-Service value of outgoing packets.

ip-ttl()

Type:	number
Default:	0

Description: Specifies the Time-To-Live value of outgoing packets.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to sources should be closed when syslog-ng is forced to reload its configuration (upon the receipt of a SIGHUP signal). Note that this applies to the server (source) side of the syslog-ng connections, client-side (destination) connections are always reopened after receiving a HUP signal unless the *keep-alive* option is enabled for the destination.

keep-hostname()

Type:	yes or no
Default:	no

Description: Enable or disable hostname rewriting.

- If enabled (`keep-hostname(yes)`), syslog-ng OSE assumes that the incoming log message was sent by the host specified in the *HOST* field of the message.

- If disabled (`keep-hostname(no)`), syslog-ng OSE rewrites the *HOST* field of the message, either to the IP address (if the *use-dns()* parameter is set to `no`), or to the hostname (if the *use-dns()* parameter is set to `yes` and the IP address can be resolved to a hostname) of the host sending the message to syslog-ng OSE. For details on using name resolution in syslog-ng OSE, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

Note

If the log message does not contain a hostname in its *HOST* field, syslog-ng OSE automatically adds a hostname to the message.

- For messages received from the network, this hostname is the address of the host that sent the message (this means the address of the last hop if the message was transferred via a relay).
- For messages received from the local host, syslog-ng OSE adds the name of the host.

This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Note
When relaying messages, enable this option on the syslog-ng OSE server and also on every relay, otherwise syslog-ng OSE will treat incoming messages as if they were sent by the last relay.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if `log-fetch-limit()` is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the `max-connections()` option is set, the `log-iw-size()` will be divided by the number of connections, otherwise `log-iw-size()` is divided by 10 (the default value of the `max-connections()` option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the `flush-lines()` option set in the destination of your clients.

Example 6.11. Initial window size of a connection
If `log-iw-size(1000)` and `max-connections(10)`, then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global `log-msg-size()` option, which defaults to `8192`.

Description: Specifies the maximum length of incoming log messages. Uses the value of the _global option_ if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use _program-override()_ instead.

max-connections()

Type:	number
Default:	10

Description: Specifies the maximum number of simultaneous connections.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in _pad-size()_. Mostly used on HP-UX where `/dev/log` is a named pipe and every write is padded to 2048 bytes. If _pad-size()_ was given and the incoming message does not fit into _pad-size()_, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

port() or localport()

Type:	number
Default:	In case of TCP transport: 601
	In case of UDP transport: 514

Description: The port number to bind to.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

> **Warning**
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to `2 097 152` bytes.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

tcp-keep-alive()

Type:	yes or no
Default:	no

Description: This is an obsolete alias of the *so-keepalive()* option.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

transport()

Type:	udp, tcp, or tls
Default:	tcp

Description: Specifies the protocol used to receive messages from the source.

> **Warning**
>
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to `2 097 152` bytes.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

use-dns()

Type:	yes, no, persist_only
Default:	yes

Description: Enable or disable DNS usage. The `persist_only` option attempts to resolve hostnames locally from file (for example from `/etc/hosts`). The syslog-ng OSE application blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

use-fqdn()

Type:	yes or no
Default:	no

Description: Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

6.5. Receiving JSON messages from nodejs applications

Using the `nodejs()` driver, syslog-ng OSE can receive application logs directly from nodejs applications that use the widespread _Winston_ logging API. The syslog-ng OSE application automatically adds the `.nodejs.winston.` prefix to the name of the fields the extracted from the message.

To use the `nodejs()` driver, the `scl.conf` file must be included in your syslog-ng configuration:

```
@include "scl.conf"
```

The `nodejs()` driver is actually a reusable configuration snippet configured to receive log messages using the `tcp()` driver, and process its JSON contents. For details on using or writing such configuration snippets, see _Section 5.8.2, Reusing configuration blocks (p. 48)_. You can find the source of the nodejs configuration snippet on _GitHub_.

Example 6.12. Using the nodejs() driver
The following example uses the default settings of the driver, listening for messages on port 9003 of every IP address of the syslog-ng OSE host.

```
@include "scl.conf"
source apps { nodejs(); };
```

The following example listens only on IP address `192.168.1.1`, port 9999.

```
@include "scl.conf"
source apps { nodejs(
            localip(192.168.1.1)
            port(9999)
            )
};
```

Note
For details on the parameters of the `nodejs()` driver, see _Section 6.5.1, nodejs() source options (p. 70)_.

6.5.1. nodejs() source options

The *nodejs()* driver has the following options.

ip() or localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to. Note that this is not the address where messages are accepted from.

port() or localport()

Type:	number
Default:	9003

Description: The port number to bind to.

6.6. Collecting messages from named pipes

The pipe driver opens a named pipe with the specified name and listens for messages. It is used as the native message delivery protocol on HP-UX.

The pipe driver has a single required parameter, specifying the filename of the pipe to open. For the list of available optional parameters, see *Section 6.6.1, pipe() source options (p. 71)*.

```
Declaration:
            pipe(filename);
```

> **Note**
> As of syslog-ng Open Source Edition 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the mkfifo(1) command.

Pipe is very similar to the *file()* driver, but there are a few differences, for example *pipe()* opens its argument in read-write mode, therefore it is not recommended to be used on special files like /proc/kmsg.

> **Warning**
> - It is not recommended to use *pipe()* on anything else than real pipes.
> - By default, syslog-ng OSE uses the flags(no-hostname) option for pipes, meaning that syslog-ng OSE assumes that the log messages received from the pipe do not contain the hostname field. If your messages do contain the hostname field, use flags(expect-hostname). For details, see *Section flags() (p. 71)*.

> **Example 6.13. Using the pipe() driver**
> ```
> source s_pipe { pipe("/dev/pipe" pad-size(2048)); };
> ```

6.6.1. pipe() source options

The *pipe* driver has the following options:

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the LOG_KERN | LOG_NOTICE priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

  ```
  source s_dell { udp(port(2000) flags(no-hostname)); };
  ```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the , `syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before msg

in the following message: Jan 22 10:06:11 host program:msg). If you do not want to store the original header of the message, enable the *dont-store-legacy-msghdr* flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

follow-freq()

Type:	number
Default:	1

Description: Indicates that the source should be checked periodically. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use *poll()* on the file, but checks whether the file changed every time the *follow-freq()* interval (in seconds) has elapsed. Floating-point numbers (for example 1.5) can be used as well.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if *log-fetch-limit()* is too high.

log-iw-size()

Type:	number
Default:	1000

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

Description: The size of the initial window, this value is used during flow control. If the `max-connections()` option is set, the `log-iw-size()` will be divided by the number of connections, otherwise `log-iw-size()` is divided by 10 (the default value of the `max-connections()` option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the `flush-lines()` option set in the destination of your clients.

> **Example 6.14. Initial window size of a connection**
> If `log-iw-size(1000)` and `max-connections(10)`, then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global `log-msg-size()` option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the _global option_ if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use `program-override()` instead.

multi-line-garbage()

Type:	regular expression
Default:	empty string

Description: Use the `multi-line-garbage()` option when processing multi-line messages that contain unneeded parts between the messages. Specify a string or regular expression that matches the beginning of the unneeded message parts. If the `multi-line-garbage()` option is set, syslog-ng OSE ignores the lines between the line matching the `multi-line-garbage()` and the next line matching `multi-line-prefix()`. See also the `multi-line-prefix()` option.

When receiving multi-line messages from a source when the `multi-line-garbage()` option is set, but no matching line is received between two lines that match `multi-line-prefix()`, syslog-ng OSE will continue to process the incoming lines as a single message until a line matching `multi-line-garbage()` is received.

To use the `multi-line-garbage()` option, set the `multi-line-mode()` option to `prefix-garbage`.

Warning

If the `multi-line-garbage()` option is set, syslog-ng OSE discards lines between the line matching the `multi-line-garbage()` and the next line matching `multi-line-prefix()`.

Note

A message is considered complete if no new lines arrive to the message for 10 seconds, even if no line matching the `multi-line-garbage()` or the `multi-line-suffix()` option is received.

multi-line-mode()

Type:	indented|regexp
Default:	empty string

Description: Use the `multi-line-mode()` option when processing multi-line messages. The syslog-ng OSE application provides the following methods to process multi-line messages: `multi-line-mode(indented)`, and `multi-line-mode(prefix-garbage)`.

- The *indented* mode can process messages where each line that belongs to the previous line is indented by whitespace, and the message continues until the first non-indented line. For example, the Linux kernel (starting with version 3.5) uses this format for `/dev/log`, as well as several applications, like Apache Tomcat.

Example 6.15. Processing indented multi-line messages

```
source s_tomcat {
    file("/var/log/tomcat/xxx.log" multi-line-mode(indented));
};
```

- The *prefix-garbage* mode uses a string or regular expression (set in `multi-line-prefix()`) that matches the beginning of the log messages, ignores newline characters from the source until a line matches the regular expression again, and treats the lines between the matching lines as a single message. For details on using `multi-line-mode(prefix-garbage)`, see the `multi-line-prefix()` and `multi-line-garbage()` options.

- The `prefix-suffix` mode uses a string or regular expression (set in `multi-line-prefix()`) that matches the beginning of the log messages, ignores newline characters from the source until a line matches the regular expression set in `multi-line-suffix()`, and treats the lines between `multi-line-prefix()` and `multi-line-suffix()` as a single message. Any other lines between the end of the message and the beginning of a new message (that is, a line that matches the `multi-line-prefix()` expression) is discarded. For details on using `multi-line-mode(prefix-suffix)`, see the `multi-line-prefix()` and `multi-line-suffix()` options.

  The `prefix-suffix` mode is similar to the `prefix-garbage` mode, but it appends the garbage part to the message instead of discarding it.

multi-line-prefix()

Type:	regular expression
Default:	empty string

Description: Use the *multi-line-prefix()* option to process multi-line messages, that is, log messages that contain newline characters (for example, Tomcat logs). Specify a string or regular expression that matches the beginning of the log messages. Use as simple regular expressions as possible, because complex regular expressions can severely reduce the rate of processing multi-line messages. If the *multi-line-prefix()* option is set, syslog-ng OSE ignores newline characters from the source until a line matches the regular expression again, and treats the lines between the matching lines as a single message. See also the *multi-line-garbage()* option.

Note

A message is considered complete if no new lines arrive to the message for 10 seconds, even if no line matching the *multi-line-garbage()* or the *multi-line-suffix()* option is received.

Tip

■ To make multi-line messages more readable when written to a file, use a template in the destination and instead of the *${MESSAGE}* macro, use the following: *$(indent-multi-line ${MESSAGE})*. This expression inserts a tab after every newline character (except when a tab is already present), indenting every line of the message after the first. For example:

```
destination d_file {
    file ("/var/log/messages"
        template("${ISODATE} ${HOST} $(indent-multi-line ${MESSAGE})\n") );
};
```

For details on using templates, see *Section 11.1.2, Templates and macros (p. 239)*.

■ To actually convert the lines of multi-line messages to single line (by replacing the newline characters with whitespaces), use the *flags(no-multi-line)* option in the source.

Example 6.16. Processing Tomcat logs
The log messages of the Apache Tomcat server are a typical example for multi-line log messages. The messages start with the date and time of the query in the YYYY.MM.DD HH:MM:SS format, as you can see in the following example.

```
2010.06.09. 12:07:39 org.apache.catalina.startup.Catalina start
SEVERE: Catalina.start:
LifecycleException:  service.getName(): "Catalina";  Protocol handler start failed:
java.net.BindException: Address already in use<null>:8080
        at org.apache.catalina.connector.Connector.start(Connector.java:1138)
        at org.apache.catalina.core.StandardService.start(StandardService.java:531)
        at org.apache.catalina.core.StandardServer.start(StandardServer.java:710)
        at org.apache.catalina.startup.Catalina.start(Catalina.java:583)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:39)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:25)
        at java.lang.reflect.Method.invoke(Method.java:597)
        at org.apache.catalina.startup.Bootstrap.start(Bootstrap.java:288)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:39)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:25)
        at java.lang.reflect.Method.invoke(Method.java:597)
        at org.apache.commons.daemon.support.DaemonLoader.start(DaemonLoader.java:177)
2010.06.09. 12:07:39 org.apache.catalina.startup.Catalina start
INFO: Server startup in 1206 ms
2010.06.09. 12:45:08 org.apache.coyote.http11.Http11Protocol pause
INFO: Pausing Coyote HTTP/1.1 on http-8080
2010.06.09. 12:45:09 org.apache.catalina.core.StandardService stop
INFO: Stopping service Catalina
```

To process these messages, specify a regular expression matching the timestamp of the messages in the `multi-line-prefix()` option. Such an expression is the following:

```
source s_file{file("/var/log/tomcat6/catalina.2010-06-09.log" follow-freq(0)
multi-line-mode(regexp) multi-line-prefix("[0-9]{4}\.[0-9]{2}\.[0-9]{2}\.")
flags(no-parse));};
};
```

Note that the flags(no-parse) is needed to avoid syslog-ng OSE trying to interpret the date in the message.

optional()

Type:	yes or no
Default:	

Description: Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the *pipe()*, *unix-dgram*, and *unix-stream* drivers.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where /dev/log is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

pipe()

Type:	filename with path
Default:	

Description: The filename of the pipe to read messages from.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

6.7. Collecting process accounting logs on Linux

Starting with version 3.2, syslog-ng OSE can collect process accounting logs on Linux systems. Process accounting is the method of recording and summarizing commands executed on Linux, for example, the commands being run, the user who executed the command, CPU time used by the process, exit code, and so on. When process accounting (also called pacct) is enabled on a system, the kernel writes accounting records to the `/var/log/account/pacct` file (might vary between different Linux distributions).

To use the *pacct()* driver, the following conditions must be met:

- The syslog-ng OSE application must be compiled with the *--enable-pacct* option. Execute the `syslog-ng -V` command to check if your binary supports process accounting.

- The *pacctformat* plugin must be loaded. By default, syslog-ng OSE automatically loads the available modules.

- The `scl.conf` file must be included in your syslog-ng configuration:

  ```
  @include "scl.conf"
  ```

- Process accounting must be running on the host. You can enable it with the `accton` command.

The *pacct()* driver parses the fields of the accounting logs and transforms them into name-value pairs. The fields are defined in the manual page of the accounting log file (`man acct`), syslog-ng OSE prepends every field with the `.pacct.` prefix. For example, the `ac_uid` field that contains the id of the user who started the process will be available under the `$.pacct.ac_uid` name. These can be used as macros in templates, in filters to select specific messages, and so on.

To use the *pacct()* driver, use the following syntax.

```
@version: @techversion;
@include "scl.conf"
source s_pacct { pacct(); };
...
log { source(s_pacct); destination(...); };
```

The *pacct()* driver is actually a reusable configuration snippet configured to handle Linux accounting logs. For details on using or writing such configuration snippets, see *Section 5.8.2, Reusing configuration blocks (p. 48)*. You can find the source of the pacct configuration snippet on *GitHub*.

6.7.1. pacct() options

The *pacct()* driver has the following options:

file

Type:	filename with path
Default:	/var/log/account/pacct

Description: The file where the process accounting logs are stored — syslog-ng OSE reads accounting messages from this file.

follow-freq()

Type:	number
Default:	1

Description: Indicates that the source should be checked periodically. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use *poll()* on the file, but checks whether the file changed every time the *follow-freq()* interval (in seconds) has elapsed. Floating-point numbers (for example `1.5`) can be used as well.

6.8. Receiving messages from external applications

The program driver starts an external application and reads messages from the standard output (stdout) of the application. It is mainly useful to receive log messages from daemons that accept incoming messages and convert them to log messages.

The program driver has a single required parameter, specifying the name of the application to start.

```
Declaration:
program(filename);
```

Example 6.17. Using the program() driver

```
source s_program { program("/etc/init.d/mydaemon"); };
```

Note
The program is restarted automatically if it exits.

6.8.1. program() source options

The *program* driver has the following options:

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the `LOG_KERN | LOG_NOTICE` priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

```
source s_dell { udp(port(2000) flags(no-hostname)); };
```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the , `syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before msg in the following message: Jan 22 10:06:11 host program:msg). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if `log-fetch-limit()` is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the `max-connections()` option is set, the `log-iw-size()` will be divided by the number of connections, otherwise `log-iw-size()` is divided by 10 (the default value of the `max-connections()` option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the `flush-lines()` option set in the destination of your clients.

Example 6.18. Initial window size of a connection
If log-iw-size(1000) and max-connections(10), then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global `log-msg-size()` option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the _global option_ if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use `program-override()` instead.

optional()

Type:	yes or no
Default:	

Description: Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the `pipe()`, `unix-dgram`, and `unix-stream` drivers.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where `/dev/log` is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

program

Type:	filename with path
Default:	

Description: The name of the application to start and read messages from.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

6.9. Collecting messages on Sun Solaris

Solaris uses its *STREAMS* framework to send messages to the *syslogd* process. Solaris 2.5.1 and above uses an IPC called *door* in addition to *STREAMS*, to confirm the delivery of a message. The syslog-ng application supports the IPC mechanism via the *door()* option (see below).

> **Note**
> The *sun-streams()* driver must be enabled when the syslog-ng application is compiled (see `./configure --help`).

The *sun-streams()* driver has a single required argument specifying the *STREAMS* device to open, and the *door()* option. For the list of available optional parameters, see *Section 6.9.1, sun-streams() source options (p. 83)*.

```
Declaration:
    sun-streams(name_of_the_streams_device door(filename_of_the_door));
```

> **Example 6.19. Using the sun-streams() driver**
> ```
> source s_stream { sun-streams("/dev/log" door("/etc/.syslog_door")); };
> ```

6.9.1. sun-streams() source options

The *sun-streams()* driver has the following options.

door()

Type:	string
Default:	none

Description: Specifies the filename of a door to open, needed on Solaris above 2.5.1.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The *assume-utf8* flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the *validate-utf8* flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the `LOG_KERN | LOG_NOTICE` priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

  ```
  source s_dell { udp(port(2000) flags(no-hostname)); };
  ```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the `, syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before msg in the following message: Jan 22 10:06:11 host program:msg). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

follow-freq()

Type:	number
Default:	1

Description: Indicates that the source should be checked periodically. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use *poll()* on the file, but checks whether the file changed every time the *follow-freq()* interval (in seconds) has elapsed. Floating-point numbers (for example 1.5) can be used as well.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if *log-fetch-limit()* is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the *max-connections()* option is set, the *log-iw-size()* will be divided by the number of connections, otherwise *log-iw-size()* is divided by 10 (the default value of the *max-connections()* option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the *flush-lines()* option set in the destination of your clients.

Example 6.20. Initial window size of a connection
If log-iw-size(1000) and max-connections(10), then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global *log-msg-size()* option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the *global option* if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use *program-override()* instead.

optional()

Type:	yes or no
Default:	

Description: Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the *pipe()*, *unix-dgram*, and *unix-stream* drivers.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where `/dev/log` is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

6.10. Collecting messages using the IETF syslog protocol (syslog() driver)

The `syslog()` driver can receive messages from the network using the standard IETF-syslog protocol (as described in RFC5424-26). UDP, TCP, and TLS-encrypted TCP can all be used to transport the messages.

> **Note**
>
> The `syslog()` driver can also receive BSD-syslog-formatted messages (described in RFC 3164, see *Section 2.8.1, BSD-syslog or legacy-syslog messages (p. 11)*) if they are sent using the IETF-syslog protocol.
>
> In syslog-ng OSE versions 3.1 and earlier, the `syslog()` driver could handle only messages in the IETF-syslog (RFC 5424-26) format.

For the list of available optional parameters, see *Section 6.10.1, syslog() source options (p. 88)*.

```
Declaration:
        syslog(ip() port() transport() options());
```

> **Example 6.21. Using the syslog() driver**
> TCP source listening on the localhost on port 1999.
>
> ```
> source s_syslog { syslog(ip(127.0.0.1) port(1999) transport("tcp")); };
> ```
>
> UDP source with defaults.
>
> ```
> source s_udp { syslog( transport("udp")); };
> ```
>
> Encrypted source where the client is also authenticated. For details on the encryption settings, see *Section 10.4, TLS options (p. 235)*.
>
> ```
> source s_syslog_tls{ syslog(
>     ip(10.100.20.40)
>     transport("tls")
>     tls(
>     peer-verify(required-trusted)
>     ca-dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
>     key-file('/opt/syslog-ng/etc/syslog-ng/keys/server_privatekey.pem')
>     cert-file('/opt/syslog-ng/etc/syslog-ng/keys/server_certificate.pem')
>     )
>     );};
> ```

> **Warning**
>
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the `so-rcvbuf()` option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the `so-rcvbuf()` so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the `so-rcvbuf()` at least to `2 097 152` bytes.

6.10.1. syslog() source options

The `syslog()` driver has the following options.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the `LOG_KERN | LOG_NOTICE` priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

```
source s_dell { udp(port(2000) flags(no-hostname)); };
```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the , `syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before `msg` in the following message: `Jan 22 10:06:11 host program:msg`). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

- *threaded*: The `threaded` flag enables multithreading for the source. For details on multithreading, see *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)*.

> **Note**
> The `syslog` source uses multiple threads only if the source uses the `tls` or `tcp` transport protocols.

host-override()

Type:	string
Default:	

Description: Replaces the ${HOST} part of the message with the parameter string.

ip() or localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to. Note that this is not the address where messages are accepted from.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

ip-protocol()

Type:	number
Default:	4

Description: Determines the internet protocol version of the given driver (`network()` or `syslog()`). The possible values are 4 and 6, corresponding to IPv4 and IPv6. The default value is 4.

ip-tos()

Type:	number
Default:	0

Description: Specifies the Type-of-Service value of outgoing packets.

ip-ttl()

Type:	number
Default:	0

Description: Specifies the Time-To-Live value of outgoing packets.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to sources should be closed when syslog-ng is forced to reload its configuration (upon the receipt of a SIGHUP signal). Note that this applies to the server (source) side of the syslog-ng connections, client-side (destination) connections are always reopened after receiving a HUP signal unless the `keep-alive` option is enabled for the destination.

keep-hostname()

Type:	yes or no
Default:	no

Description: Enable or disable hostname rewriting.

- If enabled (`keep-hostname(yes)`), syslog-ng OSE assumes that the incoming log message was sent by the host specified in the *HOST* field of the message.

- If disabled (`keep-hostname(no)`), syslog-ng OSE rewrites the *HOST* field of the message, either to the IP address (if the `use-dns()` parameter is set to no), or to the hostname (if the `use-dns()` parameter is set to `yes` and the IP address can be resolved to a hostname) of the host sending the message to syslog-ng OSE. For details on using name resolution in syslog-ng OSE, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

Note

If the log message does not contain a hostname in its *HOST* field, syslog-ng OSE automatically adds a hostname to the message.

- For messages received from the network, this hostname is the address of the host that sent the message (this means the address of the last hop if the message was transferred via a relay).

- For messages received from the local host, syslog-ng OSE adds the name of the host.

This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Note

When relaying messages, enable this option on the syslog-ng OSE server and also on every relay, otherwise syslog-ng OSE will treat incoming messages as if they were sent by the last relay.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if `log-fetch-limit()` is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the `max-connections()` option is set, the `log-iw-size()` will be divided by the number of connections, otherwise `log-iw-size()` is divided by 10 (the default value of the `max-connections()` option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the `flush-lines()` option set in the destination of your clients.

Example 6.22. Initial window size of a connection
If `log-iw-size(1000)` and `max-connections(10)`, then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global *log-msg-size()* option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the _global option_ if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use *program-override()* instead.

max-connections()

Type:	number
Default:	10

Description: Specifies the maximum number of simultaneous connections.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where `/dev/log` is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

port() or localport()

Type:	number
Default:	In case of TCP transport: 601
	In case of UDP transport: 514

Description: The port number to bind to.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the `program-override("kernel")` option in the source containing `/proc/kmsg`. NOTE: This option replaces the deprecated *log-prefix()* option.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

> **Warning**
>
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to `2 097 152` bytes.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

tcp-keep-alive()

Type:	yes or no
Default:	no

Description: This is an obsolete alias of the *so-keepalive()* option.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

transport()

Type:	udp, tcp, or tls
Default:	tcp

Description: Specifies the protocol used to receive messages from the source.

> **Warning**
>
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to `2 097 152` bytes.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

use-dns()

Type:	yes, no, persist_only
Default:	yes

Description: Enable or disable DNS usage. The *persist_only* option attempts to resolve hostnames locally from file (for example from `/etc/hosts`). The syslog-ng OSE application blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

use-fqdn()

Type:	yes or no
Default:	no

Description: Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

6.11. Collecting the system-specific log messages of a platform

Starting with version 3.2, syslog-ng OSE can automatically collect the system-specific log messages of the host on a number of platforms using the *system()* driver. If the *system()* driver is included in the syslog-ng OSE configuration file, syslog-ng OSE automatically adds the following sources to the syslog-ng OSE configuration.

> **Note**
> syslog-ng OSE versions 3.2-3.3 used an external script to generate the *system()* source, but this was problematic in certain situations, for example, when the host used a strict AppArmor profile. Therefore, the *system()* source is now generated internally in syslog-ng OSE.

The *system()* driver is also used in the default configuration file of syslog-ng OSE. For details on the default configuration file, see *Example 4.1, The default configuration file of syslog-ng OSE (p. 35)*. Starting with syslog-ng OSE version 3.6, you can use the `system-expand` command-line utility (which is a shell script, located in the `modules/system-source/` directory) to display the configuration that the *system()* source will use.

> **Warning**
> If syslog-ng OSE does not recognize the platform it is installed on, it does not add any sources.

Starting with version 3.6, syslog-ng OSE parses messages complying with the *Common Information Model (CIM)* and marked with @cim as JSON messages (for example, the ulogd from the netfilter project can emit such messages). That way, you can forward such messages without losing any information to CIM-aware applications (for example, Splunk).

Platform	Message source
AIX and Tru64	`unix-dgram("/dev/log");`

FreeBSD	`unix-dgram("/var/run/log");`
	
	`unix-dgram("/var/run/logpriv" perm(0600));`
	
	`file("/dev/klog" follow-freq(0) program-override("kernel") flags(no-parse));`
	
	For FreeBSD versions earlier than 9.1, `follow-freq(1)` is used.

GNU/kFreeBSD	`unix-dgram("/var/run/log");`
	
	`file("/dev/klog" follow-freq(0) program-override("kernel"));`

HP-UX	`pipe("/dev/log" pad-size(2048));`

Linux	`unix-dgram("/dev/log");`
	
	`file("/proc/kmsg" program-override("kernel") flags(kernel));`
	
	Note that on Linux, the `so-rcvbuf()` option of the `system()` source is automatically set to 8192.
	
	If the host is running under systemd, syslog-ng OSE reads directly from the systemd journal file using the `systemd-journal()` source.
	
	If the kernel of the host is version 3.5 or newer, and /dev/kmsg is seekable, syslog-ng OSE will use that instead of /proc/kmsg, using the `multi-line-mode(indented)` and the `format(linux-kmsg)` options.
	
	If syslog-ng OSE is running in a jail or a Linux Container (LXC), it will not read from the /dev/kmsg or /proc/kmsg files.

Solaris 8	`sun-streams("/dev/log");`

Platform	Message source
Solaris 9	```sun-streams("/dev/log" door("/etc/.syslog_door"));```
Solaris 10	```sun-streams("/dev/log" door("/var/run/syslog_door"));```

Table 6.3. Sources automatically added by syslog-ng Open Source Edition

6.12. Collecting messages from the systemd-journal system log storage

The *systemd-journal()* source collects messages directly from the journal of platforms that use systemd. If you use the *systemd-journal()* source directly (instead of using the *system()* source), make sure that your other syslog-ng OSE sourcer do not read from /dev/kmsg or /proc/kmsg to avoid duplicating log messages. (The messages from /dev/kmsg or /proc/kmsg are automatically available in the journal as well.)

```
Declaration:
systemd-journal(options);
```

Example 6.23. Sending all fields through syslog protocol using the systemd-journal() driver
To send all fields through the syslog protocol, enter the prefix in the following format: ".SDATA.<name>".

```
@version: 3.6

source s_journald {
 systemd-journal(prefix(".SDATA.journald."));
};

destination d_network {
 syslog("server.host");
};

log {
  source(s_journald);
  destination(d_network);
};
```

Example 6.24. Filtering for a specific field using the systemd-journal() driver

```
@version: 5.2

source s_journald {
 systemd-journal(prefix(".SDATA.journald."));
};

filter f_uid {"${.SDATA.journald._UID}" eq "1000"};

destination d_network {
 syslog("server.host");
};

log {
  source(s_journald);
  filter(f_uid);
  destination(d_network);
};
```

Example 6.25. Sending all fields in value-pairs using the systemd-journal() driver

```
@version: 5.2

source s_local {
 systemd-journal(prefix("journald."));
};

destination d_network {
 network("server.host" template("$(format_json --scope rfc5424 --key journald.*)\n"));
};

log {
 source(s_local);
 destination(d_network);
};
```

The journal contains credential information about the process that sent the log message. The syslog-ng OSE application makes this information available in the following macros:

Journald field	syslog-ng predefined macro
MESSAGE	$MESSAGE
_HOSTNAME	$HOST
_PID	$PID
_COMM if does not exist SYSLOG_IDENTIFIER	$PROGRAM
SYSLOG_FACILITY	$FACILITY_NUM
PRIORITY	$LEVEL_NUM

6.12.1. systemd-journal() source options

The *systemd-journal()* driver has the following options:

default-facility()

Type:	facility string
Default:	local0

Description: The default facility value if the SYSLOG_FACILITY entry does not exist.

default-level()

Type:	string
Default:	notice

Description: The default level value if the PRIORITY entry does not exist.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if `log-fetch-limit()` is too high.

max-field-size()

Type:	number (characters)
Default:	65536

Description: The maximum length of a field's value.

prefix()

Type:	string
Default:	

Description: If this option is set, every non-built-in mapped names get a prefix (for example: `".SDATA.journald."`).

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

6.13. Collecting systemd messages using a socket

On platforms running systemd, the `systemd-syslog()` driver reads the log messages of systemd using the `/run/systemd/journal/syslog` socket. Note the following points about this driver:

- If possible, use the more reliable _systemd-journal()_ driver instead.
- The socket activation of systemd is buggy, causing some log messages to get lost during system startup.
- If syslog-ng OSE is running in a jail or a Linux Container (LXC), it will not read from the `/dev/kmsg` or `/proc/kmsg` files.

```
Declaration:
systemd-syslog();
```

Example 6.26. Using the systemd-syslog() driver

```
@version: 3.6

source s_systemdd {
 systemd-syslog();
};

destination d_network {
 syslog("server.host");
};
```

```
log {
  source(s_journald);
  destination(d_network);
};
```

6.14. Collecting messages from remote hosts using the BSD syslog protocol

Note
The `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers will be deprecated in later versions, use the `network()` driver instead.

The `tcp()`, `tcp6()`, `udp()`, `udp6()` drivers can receive syslog messages conforming to RFC3164 from the network using the TCP and UDP networking protocols. The `tcp6()` and `udp6()` drivers use the IPv6 network protocol, while `tcp()` and `udp()` use IPv4.

UDP is a simple datagram oriented protocol, which provides "best effort service" to transfer messages between hosts. It may lose messages, and no attempt is made at the protocol level to retransmit such lost messages. The _BSD-syslog_ protocol traditionally uses UDP.

TCP provides connection-oriented service, which basically means that the path of the messages is flow-controlled. Along this path, each message is acknowledged, and retransmission is done for lost packets. Generally it is safer to use TCP, because lost connections can be detected, and no messages get lost, assuming that the TCP connection does not break. When a TCP connection is broken the 'in-transit' messages that were sent by syslog-ng but not yet received on the other side are lost. (Basically these messages are still sitting in the socket buffer of the sending host and syslog-ng has no information about the fate of these messages).

The `tcp()` and `udp()` drivers do not have any required parameters. By default they bind to the `0.0.0.0:514` address, which means that syslog-ng will listen on all available interfaces, port 514. To limit accepted connections to only one interface, use the `localip()` parameter as described below. For the list of available optional parameters, see _Section 6.14.1, tcp(), tcp6(), udp() and udp6() source options (p. 101)_.

```
Declaration:
            tcp([options]);
            udp([options]);
```

Note
The TCP port 514 is reserved for use with `rshell`, so select a different port if syslog-ng and `rshell` is used at the same time.

If you specify a multicast bind address to `udp()` and `udp6()`, syslog-ng will automatically join the necessary multicast group. TCP does not support multicasting.

The syslog-ng application supports TLS (Transport Layer Security, also known as SSL) for the tcp() and tcp6() drivers. For details, see the TLS-specific options below and _Section 10.2, Encrypting log messages with TLS (p. 230)_. For the list of available optional parameters, see _Section 6.14.1, tcp(), tcp6(), udp() and udp6() source options (p. 101)_.

Tip
The *syslog()* driver also supports TLS-encrypted connections.

Example 6.27. Using the udp() and tcp() drivers
A simple udp() source with default settings.

```
source s_udp { udp(); };# An UDP source with default settings.
```

A TCP source listening on the localhost interface, with a limited number of connections allowed.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999) max-connections(10)); };
```

A TCP source listening on a TLS-encrypted channel.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999)
          tls(peer-verify('required-trusted')
             key-file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.key')
          cert-file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.crt')));
};
```

A TCP source listening for messages using the IETF-syslog message format. Note that for transferring IETF-syslog messages, generally you are recommended to use the *syslog()* driver on both the client and the server, as it uses both the IETF-syslog message format and the protocol. For details, see *Section 6.10, Collecting messages using the IETF syslog protocol (syslog() driver) (p. 87)*.

```
source s_tcp_syslog { tcp(ip(127.0.0.1) port(1999) flags(syslog-protocol) ); };
```

6.14.1. tcp(), tcp6(), udp() and udp6() source options

The *tcp()*, *tcp6()*, *udp()*, *udp6()* drivers can receive messages conforming to RFC3164 from the network using the TCP and UDP networking protocols.

The following options are valid for *tcp()*, *tcp6()*, *udp()*, and *udp6()* drivers:

Warning
When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the net.core.rmem_max parameter of the host (for example, to 1024000), but do not modify net.core.rmem_default parameter.

As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to 2 097 152 bytes.

encoding()

Type:	string
Default:	

Description: Specifies the characterset (encoding, for example UTF-8) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the iconv -l command.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the `LOG_KERN | LOG_NOTICE` priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

  ```
  source s_dell { udp(port(2000) flags(no-hostname)); };
  ```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the , `syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before msg in the following message: `Jan 22 10:06:11 host program:msg`). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame

header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

- *threaded*: The `threaded` flag enables multithreading for the destination. For details on multithreading, see *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)*.

> **Note**
> Only the `tcp` and `tcp6` sources can use multiple threads.

host-override()

Type:	string
Default:	

Description: Replaces the ${HOST} part of the message with the parameter string.

ip() or localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to. Note that this is not the address where messages are accepted from.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to sources should be closed when syslog-ng is forced to reload its configuration (upon the receipt of a SIGHUP signal). Note that this applies to the server (source) side of the syslog-ng connections, client-side (destination) connections are always reopened after receiving a HUP signal unless the `keep-alive` option is enabled for the destination.

keep-hostname()

Type:	yes or no
Default:	no

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

Description: Enable or disable hostname rewriting.

- If enabled (`keep-hostname(yes)`), syslog-ng OSE assumes that the incoming log message was sent by the host specified in the *HOST* field of the message.

- If disabled (`keep-hostname(no)`), syslog-ng OSE rewrites the *HOST* field of the message, either to the IP address (if the *use-dns()* parameter is set to no), or to the hostname (if the *use-dns()* parameter is set to yes and the IP address can be resolved to a hostname) of the host sending the message to syslog-ng OSE. For details on using name resolution in syslog-ng OSE, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

> **Note**
> If the log message does not contain a hostname in its *HOST* field, syslog-ng OSE automatically adds a hostname to the message.
> - For messages received from the network, this hostname is the address of the host that sent the message (this means the address of the last hop if the message was transferred via a relay).
> - For messages received from the local host, syslog-ng OSE adds the name of the host.

This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

> **Note**
> When relaying messages, enable this option on the syslog-ng OSE server and also on every relay, otherwise syslog-ng OSE will treat incoming messages as if they were sent by the last relay.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if *log-fetch-limit()* is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the *max-connections()* option is set, the *log-iw-size()* will be divided by the number of connections, otherwise *log-iw-size()* is divided by 10 (the default value of the *max-connections()* option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE clients, make sure that the window size is larger than the *flush-lines()* option set in the destination of your clients.

Example 6.28. Initial window size of a connection
If log-iw-size(1000) and max-connections(10), then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global *log-msg-size()* option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the *global option* if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding kernel: to the kernel messages on Linux. NOTE: This option is deprecated. Use *program-override()* instead.

max-connections()

Type:	number
Default:	10

Description: Specifies the maximum number of simultaneous connections.

Note that the *udp()* and *udp6()* drivers do not support this option.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where /dev/log is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

port() or localport()

Type:	number
Default:	514

Description: The port number to bind to.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the program-override("kernel") option in the source containing /proc/kmsg. NOTE: This option replaces the deprecated *log-prefix()* option.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the socket(7) manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the socket(7) manual page.

Warning

When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the net.core.rmem_max parameter of the host (for example, to 1024000), but do not modify net.core.rmem_default parameter.

As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to 2 097 152 bytes.

tcp-keep-alive()

Type:	yes or no
Default:	no

Description: This is an obsolete alias of the `so-keepalive()` option.

tcp-keepalive-intvl()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (number of seconds) between subsequential keepalive probes, regardless of the traffic exchanged in the connection. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_intvl`. The default value is `0`, which means using the kernel default.

> **Warning**
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-probes()

Type:	number
Default:	0

Description: Specifies the number of unacknowledged probes to send before considering the connection dead. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_probes`. The default value is `0`, which means using the kernel default.

> **Warning**
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-time()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (in seconds) between the last data packet sent and the first keepalive probe. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_time`. The default value is *0,* which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

use-dns()

Type:	yes, no, persist_only
Default:	yes

Description: Enable or disable DNS usage. The *persist_only* option attempts to resolve hostnames locally from file (for example from `/etc/hosts`). The syslog-ng OSE application blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

use-fqdn()

Type:	yes or no
Default:	no

Description: Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

6.15. Collecting messages from UNIX domain sockets

The *unix-stream()* and *unix-dgram()* drivers open an *AF_UNIX* socket and start listening on it for messages. The *unix-stream()* driver is primarily used on Linux and uses *SOCK_STREAM* semantics (connection oriented, no messages are lost); while *unix-dgram()* is used on BSDs and uses *SOCK_DGRAM* semantics: this may result in lost local messages if the system is overloaded.

To avoid denial of service attacks when using connection-oriented protocols, the number of simultaneously accepted connections should be limited. This can be achieved using the *max-connections()* parameter. The default value of this parameter is quite strict, you might have to increase it on a busy system.

Both unix-stream and unix-dgram have a single required argument that specifies the filename of the socket to create. For the list of available optional parameters, see *Section 6.15.2, unix-stream() and unix-dgram() source options (p. 110)*

```
Declaration:
          unix-stream(filename [options]);
          unix-dgram(filename [options]);
```

Note

syslogd on Linux originally used *SOCK_STREAM* sockets, but some distributions switched to *SOCK_DGRAM* around 1999 to fix a possible DoS problem. On Linux you can choose to use whichever driver you like as syslog clients automatically detect the socket type being used.

Example 6.29. Using the unix-stream() and unix-dgram() drivers

```
source s_stream { unix-stream("/dev/log" max-connections(10)); };
```

```
source s_dgram { unix-dgram("/var/run/log"); };
```

6.15.1. UNIX credentials and other metadata

Starting with syslog-ng OSE 3.6, the *unix-stream()* and *unix-dgram()* sources automatically extract the available UNIX credentials and other metainformation from the received log messages. The syslog-ng OSE application can extract the following information on Linux and FreeBSD platforms (examples show the value of the macro for the su - myuser command). Similar information is available for the *systemd-journal* source.

Macro	Description
`${.unix.cmdline}`	The name (without the path) and command-line options of the executable belonging to the PID that sent the message. For example, `su - myuser`
`${.unix.exe}`	The path of the executable belonging to the PID that sent the message. For example, `/usr/bin/su`
`${.unix.gid}`	The group ID (GID) corresponding to the UID of the application that sent the log message. Note that this is the ID number of the group, not its human-readable name. For example, `0`
`${.unix.pid}`	The process ID (PID) of the application that sent the log message. For example, `774`. Note that on every UNIX platforms, if the `system()` source uses sockets, it will overwrite the PID macro with the value of `${.unix.pid}`, if it is available.
`${.unix.uid}`	The user ID (UID) of the application that sent the log message. Note that this is the ID number of the user, not its human-readable name. For example, `0`

Table 6.4. UNIX credentials available via UNIX domain sockets

6.15.2. unix-stream() and unix-dgram() source options

These two drivers behave similarly: they open an *AF_UNIX* socket and start listening on it for messages. The following options can be specified for these divers:

encoding()

Type:	string
Default:	

Description: Specifies the characterset (encoding, for example UTF-8) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the iconv -l command.

flags()

Type:	assume-utf8, empty-lines, expect-hostname, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8
Default:	empty set

Description: Specifies the log parsing options of the source.

- *assume-utf8*: The `assume-utf8` flag assumes that the incoming messages are UTF-8 encoded, but does not verify the encoding. If you explicitly want to validate the UTF-8 encoding of the incoming message, use the `validate-utf8` flag.

- *empty-lines*: Use the `empty-lines` flag to keep the empty lines of the messages. By default, syslog-ng OSE removes empty lines automatically.

- *expect-hostname*: If the `expect-hostname` flag is enabled, syslog-ng OSE will assume that the log message contains a hostname and parse the message accordingly. This is the default behavior for TCP sources. Note that pipe sources use the `no-hostname` flag by default.

- *kernel*: The `kernel` flag makes the source default to the LOG_KERN | LOG_NOTICE priority if not specified otherwise.

- *no-hostname*: Enable the `no-hostname` flag if the log message does not include the hostname of the sender host. That way syslog-ng OSE assumes that the first part of the message header is ${PROGRAM} instead of ${HOST}. For example:

```
source s_dell { udp(port(2000) flags(no-hostname)); };
```

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line. Note that this happens only if the underlying transport method actually supports multi-line messages. Currently the `, syslog()`, `udp()`, `unix-dgram()` drivers support multi-line messages; other drivers, for example, the `tcp()` driver do not.

- *no-parse*: By default, syslog-ng OSE parses incoming messages as syslog messages. The `no-parse` flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. The syslog-ng OSE application will generate a new syslog header (timestamp, host, and so on) automatically and put the entire incoming message into the MSG part of the syslog message. This flag is useful for parsing messages not complying to the syslog format.

- *dont-store-legacy-msghdr*: By default, syslog-ng stores the original incoming header of the log message. This is useful of the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, for example, adds a whitespace before msg in the following message: Jan 22 10:06:11 host program:msg). If you do not want to store the original header of the message, enable the `dont-store-legacy-msghdr` flag.

- *syslog-protocol*: The `syslog-protocol` flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard (RFC5424), but without the frame header. Note that this flag is not needed for the `syslog` driver, which handles only messages that have a frame header.

- *validate-utf8*: The `validate-utf8` flag enables encoding-verification for messages formatted according to the new IETF syslog standard (for details, see *Section 2.8.2, IETF-syslog messages (p. 13)*). If the *BOM* character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.

The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.

group()

Type:	string
Default:	root

Description: Set the gid of the socket.

host-override()

Type:	string
Default:	

Description: Replaces the ${HOST} part of the message with the parameter string.

keep-alive()

Type:	yes or no
Default:	yes

Description: Selects whether to keep connections open when syslog-ng is restarted; cannot be used with *unix-dgram()*.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fetch-limit()

Type:	number
Default:	10

Description: The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if *log-fetch-limit()* is too high.

log-iw-size()

Type:	number
Default:	1000

Description: The size of the initial window, this value is used during flow control. If the *max-connections()* option is set, the *log-iw-size()* will be divided by the number of connections, otherwise *log-iw-size()* is divided by 10 (the default value of the *max-connections()* option). The resulting number is the initial window size of each connection. For optimal performance when receiving messages from syslog-ng OSE

clients, make sure that the window size is larger than the `flush-lines()` option set in the destination of your clients.

> **Example 6.30. Initial window size of a connection**
> If `log-iw-size(1000)` and `max-connections(10)`, then each connection will have an initial window size of 100.

log-msg-size()

Type:	number
Default:	Use the global `log-msg-size()` option, which defaults to 8192.

Description: Specifies the maximum length of incoming log messages. Uses the value of the *global option* if not specified.

log-prefix() (DEPRECATED)

Type:	string
Default:	

Description: A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding `kernel:` to the kernel messages on Linux. NOTE: This option is deprecated. Use `program-override()` instead.

max-connections()

Type:	number
Default:	256

Description: Limits the number of simultaneously open connections. Cannot be used with `unix-dgram()`.

optional()

Type:	yes or no
Default:	

Description: Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the `pipe()`, `unix-dgram`, and `unix-stream` drivers.

owner()

Type:	string
Default:	root

Description: Set the uid of the socket.

pad-size()

Type:	number
Default:	0

Description: Specifies input padding. Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). The syslog-ng OSE application will pad reads from the associated device to the number of bytes set in *pad-size()*. Mostly used on HP-UX where /dev/log is a named pipe and every write is padded to 2048 bytes. If *pad-size()* was given and the incoming message does not fit into *pad-size()*, syslog-ng will not read anymore from this pipe and displays the following error message:

```
Padding was set, and couldn't read enough bytes
```

perm()

Type:	number
Default:	0666

Description: Set the permission mask. For octal numbers prefix the number with '0', for example: use 0755 for rwxr-xr-x.

program-override()

Type:	string
Default:	

Description: Replaces the ${PROGRAM} part of the message with the parameter string. For example, to mark every message coming from the kernel, include the program-override("kernel") option in the source containing /proc/kmsg. NOTE: This option replaces the deprecated *log-prefix()* option.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the socket(7) manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the socket(7) manual page.

> **Warning**
>
> When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the *so-rcvbuf()* option of the source is increased. In such cases, you will need to increase the `net.core.rmem_max` parameter of the host (for example, to `1024000`), but do not modify `net.core.rmem_default` parameter.
>
> As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to `2 097 152` bytes.

tags()

Type:	string
Default:	

Description: Label the messages received from the source with custom tags. Tags must be unique, and enclosed between double quotes. When adding multiple tags, separate them with comma, for example `tags("dmz", "router")`. This option is available only in syslog-ng 3.1 and later.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	

Description: The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Chapter 7. Sending and storing log messages — destinations and destination drivers

A destination is where a log message is sent if the filtering rules match. Similarly to sources, destinations consist of one or more drivers, each defining where and how messages are sent.

> **Tip**
> If no drivers are defined for a destination, all messages sent to the destination are discarded. This is equivalent to omitting the destination from the log statement.

To define a destination, add a destination statement to the syslog-ng configuration file using the following syntax:

```
destination <identifier> {
          destination-driver(params); destination-driver(params); ... };
```

> **Example 7.1. A simple destination statement**
> The following destination statement sends messages to the TCP port 1999 of the 10.1.2.3 host.
>
> ```
> destination d_demo_tcp { tcp("10.1.2.3" port(1999)); };
> ```
>
> If name resolution is configured, the hostname of the target server can be used as well.
>
> ```
> destination d_tcp { tcp("target_host" port(1999); localport(999)); };
> ```

> **Warning**
> - Do not define the same drivers with the same parameters more than once, because it will cause problems. For example, do not open the same file in multiple destinations.
> - Do not use the same destination in different log paths, because it can cause problems with most destination types. Instead, use filters and log paths to avoid such situations.
> - Sources and destinations are initialized only when they are used in a log statement. For example, syslog-ng OSE starts listening on a port or starts polling a file only if the source is used in a log statement. For details on creating log statements, see *Chapter 8, Routing messages: log paths and filters (p. 198)*.

The following table lists the destination drivers available in syslog-ng.

Name	Description
amqp()	Publishes messages using the AMQP (Advanced Message Queuing Protocol).
file()	Writes messages to the specified file.
graphite()	Sends metrics to a *Graphite* server to store numeric time-series data.

Name	Description
mongodb()	Sends messages to a _MongoDB_ database.
pipe()	Writes messages to the specified named pipe.
program()	Forks and launches the specified program, and sends messages to its standard input.
redis()	Sends messages as name-value pairs to a _Redis_ key-value store.
riemann()	Sends metrics or events to a _Riemann_ monitoring system.
smtp()	Sends e-mail messages to the specified recipients.
sql()	Sends messages into an SQL database. In addition to the standard syslog-ng packages, the `sql()` destination requires database-specific packages to be installed. Refer to the section appropriate for your platform in *Chapter 3, Installing syslog-ng (p. 23)*.
stomp()	Sends messages to a STOMP server.
syslog()	Sends messages to the specified remote host using the _IETF-syslog protocol_. The IETF standard supports message transport using the UDP, TCP, and TLS networking protocols.
tcp() and tcp6()	Sends messages to the specified TCP port of a remote host using the _BSD-syslog protocol_ over IPv4 and IPv6, respectively.
udp() and udp6()	Sends messages to the specified UDP port of a remote host using the _BSD-syslog protocol_ over IPv4 and IPv6, respectively.
unix-dgram()	Sends messages to the specified unix socket in `SOCK_DGRAM` style (BSD).
unix-stream()	Sends messages to the specified unix socket in `SOCK_STREAM` style (Linux).
usertty()	Sends messages to the terminal of the specified user, if the user is logged in.

Table 7.1. Destination drivers available in syslog-ng

7.1. Publishing messages using AMQP

The `amqp()` driver publishes messages using the _AMQP (Advanced Message Queuing Protocol)_. syslog-ng OSE supports AMQP versions 0.9.1 and 1.0. The syslog-ng OSE `amqp()` driver supports persistence, and every available exchange types.

The name-value pairs selected with the `value-pairs()` option will be sent as AMQP headers, while the body of the AMQP message is empty by default (but you can add custom content using the `body()` option). Publishing

the name-value pairs as headers makes it possible to use the Headers exchange-type and subscribe only to interesting log streams. This solution is more flexible than using the *routing-key()* option.

For the list of available parameters, see *Section 7.1.1, amqp() destination options (p. 118)*.

```
Declaration:
amqp( host("<amqp-server-address>") )
```

Example 7.2. Using the amqp() driver
The following example shows the default values of the available options.

```
destination d_amqp {
    amqp(
        vhost("/")
        host("127.0.0.1")
        port(5672)
        exchange("syslog")
        exchange-type("header")
        routing-key("")
        body("")
        persistent(yes)
        value-pairs(
            scope("selected-macros" "nv-pairs" "sdata")
        )
    );
};
```

7.1.1. amqp() destination options

The *amqp()* driver publishes messages using the AMQP (Advanced Message Queuing Protocol).

The *amqp()* destination has the following options:

body()

Type:	string
Default:	empty string

Description: The body of the AMQP message. You can also use macros and templates.

exchange()

Type:	string
Default:	syslog

Description: The name of the AMQP exchange where syslog-ng OSE sends the message. Exchanges take a message and route it into zero or more queues.

exchange-declare()

Type:	yes|no
Default:	no

Description: By default, syslog-ng OSE does not create non-existing exchanges. Use the `exchange-declare(yes)` option to automatically create exchanges.

exchange-type()

Type:	direct|fanout|topic|headers
Default:	header

Description: The type of the AMQP exchange.

host()

Type:	hostname or IP address
Default:	127.0.0.1

Description: The hostname or IP address of the AMQP server.

password()

Type:	string
Default:	n/a

Description: The password used to authenticate on the AMQP server.

persistent()

Type:	yes|no
Default:	yes

Description: If this option is enabled, the AMQP server or broker will store the messages on its hard disk. That way, the messages will be retained if the AMQP server is restarted, if the message queue is set to be durable on the AMQP server.

port()

Type:	number
Default:	5672

Description: The port number of the AMQP server.

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches *retries*, then drops the message.

routing-key()

Type:	string
Default:	empty string

Description: Specifies a routing key for the exchange. The routing key selects certain messages published to an exchange to be routed to the bound queue. In other words, the routing key acts like a filter. The routing key can include macros and templates.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

username()

Type:	string
Default:	empty string

Description: The username used to authenticate on the AMQP server.

value-pairs()

Type:	parameter list of the *value-pairs()* option
Default:	`scope("selected-macros" "nv-pairs")`

Description: The *value-pairs()* option creates structured name-value pairs from the data and metadata of the log message. For details on using *value-pairs()*, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*.

> **Note**
> Empty keys are not logged.

vhost()

Type:	string
Default:	/

Description: The name of the AMQP virtual host to send the messages to.

7.2. Storing messages in plain-text files

The file driver is one of the most important destination drivers in syslog-ng. It allows to output messages to the specified text file, or to a set of files.

The destination filename may include macros which get expanded when the message is written, thus a simple `file()` driver may create several files: for example, syslog-ng OSE can store the messages of client hosts in a separate file for each host. For more information on available macros see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*.

If the expanded filename refers to a directory which does not exist, it will be created depending on the `create-dirs()` setting (both global and a per destination option).

The `file()` has a single required parameter that specifies the filename that stores the log messages. For the list of available optional parameters, see *Section 7.2.1, file() destination options (p. 122)*.

```
Declaration:
            file(filename options());
```

Example 7.3. Using the file() driver

```
destination d_file { file("/var/log/messages" ); };
```

Example 7.4. Using the file() driver with macros in the file name and a template for the message

```
destination d_file {
        file("/var/log/${YEAR}.${MONTH}.${DAY}/messages"
            template("${HOUR}:${MIN}:${SEC} ${TZ} ${HOST} [${LEVEL}] ${MSG} ${MSG}\n")
            template-escape(no));
};
```

Note
When using this destination, update the configuration of your log rotation program to rotate these files. Otherwise, the log files can become very large.

Also, after rotating the log files, reload syslog-ng OSE using the `syslog-ng-ctl reload` command, or use another method to send a SIGHUP to syslog-ng OSE.

Warning
Since the state of each created file must be tracked by syslog-ng, it consumes some memory for each file. If no new messages are written to a file within 60 seconds (controlled by the `time-reap()` global option), it is closed, and its state is freed.

Exploiting this, a DoS attack can be mounted against the system. If the number of possible destination files and its needed memory is more than the amount available on the syslog-ng server.

The most suspicious macro is `${PROGRAM}`, where the number of possible variations is rather high. Do not use the `${PROGRAM}` macro in insecure environments.

7.2.1. file() destination options

The `file()` driver outputs messages to the specified text file, or to a set of files. The `file()` destination has the following options:

> **Warning**
> When creating several thousands separate log files, syslog-ng might not be able to open the required number of files. This might happen for example when using the `${HOST}` macro in the filename while receiving messages from a large number of hosts. To overcome this problem, adjust the `--fd-limit` command-line parameter of syslog-ng or the global ulimit parameter of your host. For setting the `--fd-limit` command-line parameter of syslog-ng see the *syslog-ng(8) (p. 321)* manual page. For setting the `ulimit` parameter of the host, see the documentation of your operating system.

create-dirs()

Type:	yes or no
Default:	no

Description: Enable creating non-existing directories.

dir-group()

Type:	string
Default:	Use the global settings

Description: The group of the directories created by syslog-ng. To preserve the original properties of an existing directory, use the option without specifying an attribute: `dir-group()`.

dir-owner()

Type:	string
Default:	Use the global settings

Description: The owner of the directories created by syslog-ng. To preserve the original properties of an existing directory, use the option without specifying an attribute: `dir-owner()`.

dir-perm()

Type:	number
Default:	Use the global settings

Description: The permission mask of directories created by syslog-ng. Log directories are only created if a file after macro expansion refers to a non-existing directory, and directory creation is enabled (see also the `create-dirs()` option). For octal numbers prefix the number with 0, for example use 0755 for rwxr-xr-x.

To preserve the original properties of an existing directory, use the option without specifying an attribute: `dir-perm()`. Note that when creating a new directory without specifying attributes for `dir-perm()`, the default permission of the directories is masked with the umask of the parent process (typically 0022).

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The `syslog-protocol` flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the `syslog` driver, and that the `syslog` driver automatically adds the frame header to the messages.

- *threaded*: The `threaded` flag enables multithreading for the destination. For details on multithreading, see *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)*.

> **Note**
> The `file` destination uses multiple threads only if the destination filename contains macros.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

fsync()

Type:	yes or no
Default:	no

Description: Forces an `fsync()` call on the destination fd after each write. Note: enabling this option may seriously degrade performance.

group()

Type:	string
Default:	Use the global settings

Description: Set the group of the created file to the one specified. To preserve the original properties of an existing file, use the option without specifying an attribute: `group()`.

local-time-zone()

Type:	name of the timezone or the timezone offset
Default:	The local timezone.

Description: Sets the timezone used when expanding filename and tablename templates. The timezone can be specified as using the name of the (for example `time-zone("Europe/Budapest")`), or as the timezone offset (for example `+01:00`). The valid timezone names are listed under the `/usr/share/zoneinfo` directory.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The `mark-mode()` option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  `file(),pipe(),unix-stream(),unix-dgram(),program()`

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(),udp(),syslog(),program(),` `file(),pipe(),unix-stream(),unix-dgram()`.

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(),udp(),syslog(),program(),` `file(),pipe(),unix-stream(),unix-dgram()`.

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(),udp(),syslog(),program(),` `file(),pipe(),unix-stream(),unix-dgram()`.

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where `none` is the default value, then `none` will be used.

- `global`: Destination driver uses the global `mark-mode()` setting. The syslog-ng interprets syntax error if the global `mark-mode()` is global.

> **Note**
> In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

overwrite-if-older()

Type:	number
Default:	0

Description: If set to a value higher than 0, syslog-ng OSE checks when the file was last modified before starting to write into the file. If the file is older than the specified amount of time (in seconds), then syslog-ng removes the existing file and opens a new file with the same name. In combination with for example the *${WEEKDAY}* macro, this can be used for simple log rotation, in case not all history has to be kept. (Note that in this weekly log rotation example if its Monday 00:01, then the file from last Monday is not seven days old, because it was probably last modified shortly before 23:59 last Monday, so it is actually not even six days old. So in this case, set the *overwrite-if-older()* parameter to a-bit-less-than-six-days, for example, to 518000 seconds.

owner()

Type:	string
Default:	Use the global settings

Description: Set the owner of the created file to the one specified. To preserve the original properties of an existing file, use the option without specifying an attribute: *owner()*.

pad-size()

Type:	number
Default:	0

Description: If set, syslog-ng OSE will pad output messages to the specified size (in bytes). Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes).

> **Warning**
> Hazard of data loss! If the size of the incoming message is larger than the previously set pad-size() value, syslog-ng will truncate the message to the specified size. Therefore, all message content above that size will be lost.

perm()

Type:	number
Default:	Use the global settings

Description: The permission mask of the file if it is created by syslog-ng. For octal numbers prefix the number with 0, for example use 0755 for rwxr-xr-x.

To preserve the original properties of an existing file, use the option without specifying an attribute: *perm()*.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the Last message repeated n times. message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the ', ", and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.3. Sending metrics to Graphite

The `graphite()` destination can send metrics to a *Graphite* server to store numeric time-series data. There are many ways to feed the Graphite template function with name value pairs. The syslog-ng OSE CSV and PatternDB parsers (for details, see *Section 13.5.1, Using pattern parsers (p. 286)*) can parse log messages and generate name value pairs based on message content. The CSV parser (for details, see *Section 12.2, Parsing messages (p. 268)*) can be used for logs which have a constant field based structure, like the Apache web server access logs. The patterndb parser can parse information and can extract important fields from free form log messages, as long as patterns describing the log messages are available. Another way is to send JSON-based log messages (for details, see *Section 12.3, The JSON parser (p. 272)*) to syslog-ng OSE, like running a simple shell script collecting metrics and running it from cron regularly.

```
Declaration:
        graphite(payload());
```

Example 7.5. Using the graphite() driver
To use the graphite() destination, the only mandatory parameter is payload, which specifies the value pairs to send to graphite. In the following example any value pairs starting with "monitor." are forwarded to graphite.

```
destination d_graphite { graphite(payload("--key monitor.*")); };
```

Note
The `graphite()` destination is only a wrapper around the `tcp()` destination and the `graphite-output` template function. If you want to fine-tune the TCP parameters, use the `tcp()` destination instead, as described in *Section graphite-output (p. 250)*.

7.3.1. graphite() destination options

The `graphite()` destination has the following options:

host()

Type:	hostname or IP address
Default:	localhost

Description: The hostname or IP address of the Graphite server.

port()

Type:	number
Default:	2003

Description: The port number of the Graphite server.

payload()

Type:	parameter list of the `payload()` option
Default:	`empty string`

Description: The payload() option allows you to select which value pairs to forward to graphite.

The syntax of `payload` is different from the syntax of `value-pairs()`: use the command-line syntax used in the *format-json template function*. For details on using the payload() option, see *Section graphite-output (p. 250)*.

> **Note**
> If left empty, there is no data to be forwarded to Graphite.

7.4. Storing messages in a MongoDB database

The `mongodb()` driver sends messages to a *MongoDB* database. MongoDB is a schema-free, document-oriented database. For the list of available optional parameters, see *Section 7.4.2, mongodb() destination options (p. 131)*.

```
Declaration:
                mongodb(parameters);
```

The `mongodb()` driver does not support creating indexes, as that can be a very complex operation in MongoDB. If needed, the administrator of the MongoDB database must ensure that indexes are created on the collections.

The `mongodb()` driver does not add the `_id` field to the message: the MongoDB server will do that automatically, if none is present. If you want to override this field from syslog-ng OSE, use the `key()` parameter of the `value-pairs()` option.

The syslog-ng OSE `mongodb()` driver is compatible with MongoDB server version 1.4 and newer.

> **Note**
> By default, syslog-ng OSE handles every message field as a string. For details on how to send selected fields as other types of data (for example, handle the PID as a number), see *Section 2.10.1, Specifying data types in value-pairs (p. 17)*.

> **Example 7.6. Using the mongodb() driver**
> The following example creates a `mongodb()` destination using only default values.
>
> ```
> destination d_mongodb {
>     mongodb();
> };
> ```
>
> The following example displays the default values, and is equivalent with the previous example.
>
> ```
> destination d_mongodb {
>     mongodb(
>         servers("localhost:27017")
>         database("syslog")
>         collection("messages")
>         value-pairs(
>             scope("selected-macros" "nv-pairs" "sdata")
>         )
> ```

```
        );
};
```

7.4.1. Procedure – How syslog-ng OSE connects the MongoDB server

When syslog-ng OSE connects the MongoDB server during startup, it completes the following steps.

Step 1.   The syslog-ng OSE application connects the first address listed in the *servers()* option.

Step 2.   ■ If the server is accessible and it is a master MongoDB server, syslog-ng OSE authenticates on the server (if needed), then starts sending the log messages to the server.

■ If the server is not accessible, or it is not a master server in a MongoDB replicaset and it does not send the address of the master server, syslog-ng OSE connects the next address listed in the *servers()* option.

■ If the server is not a master server in a MongoDB replicaset, but it sends the address of the master server, syslog-ng OSE connects the received address.

Step 3.   When syslog-ng OSE connects the master MongoDB server, it retrieves the list of replicas (from the `replSet` option of the server), and appends this list to the *servers()* option.

> **Warning**
>
> ■ This means that syslog-ng OSE can send log messages to addresses that are not listed in its configuration.
>
> ■ Addresses retrieved from the MongoDB servers are not stored, and can be lost when syslog-ng OSE is restarted. The retrieved addresses are not lost if the *server()* option of the destination was not changed in the configuration file since the last restart.
>
> ■ The failover mechanism used in the *mongodb()* driver is different from the client-side failover used in other drivers.

Step 4.   The syslog-ng OSE application attempts to connect another server if the *servers()* list contains at least two addresses, and one of the following events happens:

■ The `safe-mode()` option is set to `no`, and the MongoDB server becomes unreachable.

■ The `safe-mode()` option is set to `yes`, and syslog-ng OSE cannot insert a log message into the database because of an error.

In such case, syslog-ng OSE starts to connect the addresses in from the *servers()* list (starting from the first address) to find the new master server, authenticates on the new server (if needed), then continues to send the log messages to the new master server.

During this failover step, one message can be lost if the *safe-mode()* option is disabled.

Step 5.   If the original master becomes accessible again, syslog-ng OSE will automatically connect to the original master.

7.4.2. mongodb() destination options

The *mongodb()* driver sends messages to a MongoDB database. MongoDB is a schema-free, document-oriented database.

The *mongodb()* destination has the following options:

collection()

Type:	string
Default:	messages

Description: The name of the MongoDB collection where the log messages are stored (collections are similar to SQL tables). Note that the name of the collection must not start with a dollar sign ($), and that it may contain dot (.) characters.

> **Warning**
> Hazard of data loss! The syslog-ng OSE application does not verify that the specified collection name does not contain invalid characters. If you specify a collection with an invalid name, the log messages sent to the MongoDB database will be irrevocably lost without any warning.

database()

Type:	string
Default:	syslog

Description: The name of the MongoDB database where the log messages are stored. Note that the name of the database must not start with a dollar sign ($) and it cannot contain dot (.) characters.

> **Warning**
> Hazard of data loss! The syslog-ng OSE application does not verify that the specified database name does not contain invalid characters. If you specify a database with an invalid name, the log messages sent to the MongoDB database will be irrevocably lost without any warning.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The *frac-digits()* parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

local-time-zone()

Type:	name of the timezone or the timezone offset
Default:	The local timezone.

Description: Sets the timezone used when expanding filename and tablename templates. The timezone can be specified as using the name of the (for example `time-zone("Europe/Budapest")`), or as the timezone offset (for example `+01:00`). The valid timezone names are listed under the `/usr/share/zoneinfo` directory.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

on-error()

Accepted values:	drop-message|drop-property|fallback-to-string|silently-drop-message|silently-drop-property|silently-fallback-to-string
Default:	Use the global setting (which defaults to `drop-message`)

Description: Controls what happens when type-casting fails and syslog-ng OSE cannot convert some data to the specified type. By default, syslog-ng OSE drops the entire message and logs the error. Currently the `value-pairs()` option uses the settings of `on-error()`.

- `drop-message`: Drop the entire message and log an error message to the `internal()` source. This is the default behavior of syslog-ng OSE.

- `drop-property`: Omit the affected property (macro, template, or message-field) from the log message and log an error message to the `internal()` source.

- `fallback-to-string`: Convert the property to string and log an error message to the `internal()` source.

- `silently-drop-message`: Drop the entire message silently, without logging the error.

- `silently-drop-property`: Omit the affected property (macro, template, or message-field) silently, without logging the error.

- `silently-fallback-to-string`: Convert the property to string silently, without logging the error.

password()

Type:	string
Default:	n/a

Description: Password of the database user.

path()

Type:	string
Default:	empty

Description: If the *path()* option is set, syslog-ng OSE will connect to the database using the specified UNIX domain socket. Note that you cannot set the *path()* and the *servers()* options at the same time.

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches *retries*, then drops the message.

For MongoDB operations, syslog-ng OSE uses a one-minute timeout: if an operation times out, syslog-ng OSE assumes the operation has failed.

safe-mode()

Type:	yes or no
Default:	yes

Description: If *safe-mode()* is enabled, syslog-ng OSE performs an extra check after each insert to verify that the insert succeeded. The insert is successful only if this second check is successful. Note that enabling this option reduces the performance of the driver.

servers()

Type:	list of hostname:port pairs
Default:	127.0.0.1:27017

Description: Specifies the hostname or IP address and the port number of the database server. When specifying an IP address, IPv4 (for example, 192.168.0.1) or IPv6 (for example, [::1]) can be used as well.

To send the messages to a MongoDB replicaset, specify the addresses of the database servers as a comma-separated list, for example: servers(192.168.1.1:27017,192.168.3.3:27017)

For details on how syslog-ng OSE connects the MongoDB server, see *Procedure 7.4.1, How syslog-ng OSE connects the MongoDB server (p. 130)*.

To connect to the server using a UNIX domain socket, use *path* option. Note that you cannot set the *path()* and the *servers()* options at the same time.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying `0` or a lower value sets the output limit to unlimited.

username()

Type:	string
Default:	n/a

Description: Name of the database user. Note that the *mongodb()* driver currently does not support TLS-encrypted authentication.

value-pairs()

Type:	parameter list of the *value-pairs()* option
Default:	`scope("selected-macros" "nv-pairs")`

Description: The *value-pairs()* option creates structured name-value pairs from the data and metadata of the log message. For details on using *value-pairs()*, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17).*

> **Note**
> Empty keys are not logged.

7.5. Sending messages to a remote log server using the RFC3164 protocol (network() driver)

The network() destination driver can send syslog messages conforming to RFC3164 from the network using the TCP, TLS, and UDP networking protocols.

Example 7.7. Using the network() driver
TCP source listening on the localhost on port 2222 without using the network() driver.

```
destination d_tcp6 {
        tcp6(
            ip("::1")
            port(2222)
        );
        };
```

TCP source listening on the localhost on port 2222 using the network() driver.

```
destination d_network6 {
        network(
            ip("::1")
            port(2222)
            transport(tcp)
            ip-protocol(6)
        );
};
```

Note

For details on the `tcp()`, `tcp6()`, `udp()`, `udp6()` drivers, see *Section 7.15, Sending messages to a remote log server using the legacy BSD-syslog protocol (tcp(), udp() drivers) (p. 184)*.

7.5.1. network() destination options

The `network()` driver sends messages to a remote host (for example a syslog-ng server or relay) on the local intranet or internet using the RFC3164 syslog protocol (for details about the protocol, see *Section 2.8.1, BSD-syslog or legacy-syslog messages (p. 11)*). The `network()` driver supports sending messages using the UDP, TCP, or the encrypted TLS networking protocols.

These destinations have the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The `syslog-protocol` flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the `syslog` driver, and that the `syslog` driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the *flush-lines()* option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The *frac-digits()* parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

ip-protocol()

Type:	number
Default:	4

Description: Determines the internet protocol version of the driver. The possible values are *4* and *6*, corresponding to IPv4 and IPv6. The default value is *4*.

ip-tos()

Type:	number
Default:	0

Description: Specifies the Type-of-Service value of outgoing packets.

ip-ttl()

Type:	number
Default:	0

Description: Specifies the Time-To-Live value of outgoing packets.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to destinations should be closed when syslog-ng is reloaded. Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always reopened after receiving a HUP signal unless the *keep-alive* option is enabled for the source.

localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to before connecting to target.

localport()

Type:	number
Default:	0

Description: The port number to bind to. Messages are sent from this port.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

mark-freq()

Accepted values:	number
Default:	1200

Description: An alias for the obsolete *mark()* option, retained for compatibility with syslog-ng version 1.6.x. The number of seconds between two *MARK* messages. *MARK* messages are generated when there was no message traffic to inform the receiver that the connection is still alive. If set to zero (0), no *MARK* messages are sent. The *mark-freq()* can be set for global option and/or every *MARK* capable destination driver if *mark-mode()* is periodical or dst-idle or host-idle. If *mark-freq()* is not defined in the destination, then the *mark-freq()* will be inherited from the global options. If the destination uses internal *mark-mode()*, then the global *mark-freq()* will be valid (does not matter what *mark-freq()* set in the destination side).

mark-mode()

Accepted values:	`internal` | `dst-idle` | `host-idle` | `periodical` | `none` | `global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The *mark-mode()* option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  file(), pipe(), unix-stream(), unix-dgram(), program()

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where none is the default value, then none will be used.

- `global`: Destination driver uses the global *mark-mode()* setting. The syslog-ng interprets syntax error if the global *mark-mode()* is global.

> **Note**
> In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

port() or destport()

Type:	number
Default:	601

Description: The port number to connect to. Note that the default port numbers used by syslog-ng do not comply with the latest RFC which was published after the release of syslog-ng 3.0.2, therefore the default port numbers will change in the future releases.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

spoof-source()

Type:	yes or no
Default:	no

Description: Enables source address spoofing. This means that the host running syslog-ng generates UDP packets with the source IP address matching the original sender of the message. It is useful when you want to perform some kind of preprocessing via syslog-ng then forward messages to your central log management solution with the source address of the original sender. This option only works for UDP destinations though the original message can be received by TCP as well. This option is only available if syslog-ng was compiled using the *--enable-spoof-source* configuration option.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

tcp-keepalive-intvl()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (number of seconds) between subsequential keepalive probes, regardless of the traffic exchanged in the connection. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_intvl`. The default value is *0*, which means using the kernel default.

> **Warning**
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the *TCP_KEEPCNT*, *TCP_KEEPIDLE*, and *TCP_KEEPINTVL* setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-probes()

Type:	number
Default:	0

Description: Specifies the number of unacknowledged probes to send before considering the connection dead. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_probes`. The default value is *0*, which means using the kernel default.

> **Warning**
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the *TCP_KEEPCNT*, *TCP_KEEPIDLE*, and *TCP_KEEPINTVL* setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-time()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (in seconds) between the last data packet sent and the first keepalive probe. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_time`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

> **Note**
>
> If a message uses the IETF-syslog format (RFC5424), only the text of the message can be customized (that is, the $MESSAGE part of the log), the structure of the header is fixed.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

transport()

Type:	udp, tcp, or tls
Default:	tcp

Description: Specifies the protocol used to send messages to the destination server.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.6. Sending messages to named pipes

The `pipe()` driver sends messages to a named pipe like `/dev/xconsole`.

The pipe driver has a single required parameter, specifying the filename of the pipe to open. The filename can include macros. For the list of available optional parameters, see *Section 7.6.1, pipe() destination options (p. 143)*.

```
Declaration:
          pipe(filename);
```

Warning
Starting with syslog-ng OSE 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the mkfifo(1) command.

Example 7.8. Using the pipe() driver

```
destination d_pipe { pipe("/dev/xconsole"); };
```

7.6.1. pipe() destination options

This driver sends messages to a named pipe like `/dev/xconsole`.

The *pipe()* destination has the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The *no-multi-line* flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The *syslog-protocol* flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the *syslog* driver, and that the *syslog* driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the *flush-timeout()* option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the *flush-timeout()* time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the *flush-lines()* is smaller than the window size set using the *log-iw-size()* option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

group()

Type:	string
Default:	Use the global settings

Description: Set the group of the created file to the one specified. To preserve the original properties of an existing file, use the option without specifying an attribute: `group()`.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The `mark-mode()` option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- **internal**: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  file(), pipe(), unix-stream(), unix-dgram(), program()

- **dst-idle**: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- **host-idle**: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- **periodical**: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().*

- **none**: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where none is the default value, then none will be used.

- **global**: Destination driver uses the global *mark-mode()* setting. The syslog-ng interprets syntax error if the global *mark-mode()* is global.

> **Note**
> In case of dst-idle, host-idle and periodical; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

owner()

Type:	string
Default:	Use the global settings

Description: Set the owner of the created file to the one specified. To preserve the original properties of an existing file, use the option without specifying an attribute: *owner().*

pad-size()

Type:	number
Default:	0

Description: If set, syslog-ng OSE will pad output messages to the specified size (in bytes). Some operating systems (such as HP-UX) pad all messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes).

> **Warning**
> Hazard of data loss! If the size of the incoming message is larger than the previously set pad-size() value, syslog-ng will truncate the message to the specified size. Therefore, all message content above that size will be lost.

perm()

Type:	number
Default:	0600

Description: The permission mask of the pipe. For octal numbers prefix the number with '0', for example: use 0755 for rwxr-xr-x.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.7. Sending messages to external applications

The `program()` driver starts an external application or script and sends the log messages to its standard input (`stdin`).

The `program()` driver has a single required parameter, specifying a program name to start. The program is executed with the help of the current shell, so the command may include both file patterns and I/O redirections. For the list of available optional parameters, see *Section 7.7.1, program() destination options (p. 148)*.

```
Declaration:
            program(command_to_run);
```

> **Note**
> The syslog-ng OSE application executes program destinations through the standard system shell. If the system shell is not bash and you experience problems with the program destination, try changing the `/bin/sh` link to `/bin/bash`.

> **Note**
> - The syslog-ng OSE application automatically restarts the external program if it exits for reliability reasons. However it is not recommended to launch programs for single messages, because if the message rate is high, launching several instances of an application might overload the system, resulting in Denial of Service.
> - Certain external applications buffer the log messages, which might cause unexpected latency and other problems. For example, if you send the log messages to an external Perl script, Perl uses a line buffer for terminal output and block buffer otherwise. You might want to disable buffering in the external application.

> **Warning**
> The syslog-ng OSE application must be able to start and restart the external program, and have the necessary permissions to do so. For example, if your host is running AppArmor, you might have to modify your AppArmor configuration to enable syslog-ng OSE to execute external applications.

Note that the message format does not include the priority and facility values by default. To add these values, specify a template for the program destination, as shown in the following example.

Example 7.9. Using the program() destination driver

```
destination d_prog { program("/bin/script" template("<${PRI}>${DATE} ${HOST} ${MSG}\n")
); };
```

7.7.1. program() destination options

This driver starts an external application or script and sends the log messages to its standard input (*stdin*).

The *program()* destination has the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The *no-multi-line* flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The *syslog-protocol* flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the *syslog* driver, and that the *syslog* driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The `mark-mode()` option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  `file(), pipe(), unix-stream(), unix-dgram(), program()`

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where `none` is the default value, then `none` will be used.

- `global`: Destination driver uses the global `mark-mode()` setting. The syslog-ng interprets syntax error if the global `mark-mode()` is global.

Note
In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying `0` or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.8. pseudofile()

The `pseudofile()` destination driver is a very simple driver, aimed at delivering messages to special files such as files in the `/proc`, `/dev` or `sys` directories. It opens and closes the file after each write operation, instead of keeping it open. It does not append further data. It does not support templates in the filename, and does not have a queue, processing is performed immediately as read by the source. Therefore, no loss is possible, but it takes CPU time from the source, so it is not adequate in high traffic situations.

```
Declaration:
        pseudofile(filename options());
```

7.8.1. pseudofile() destination options

The `pseudofile()` destination has the following options:

file()

Type:	filename with path
Default:	

Description: The file to write messages to, including the path.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like `syslogd` or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

7.9. Storing name-value pairs in Redis

The `redis()` driver sends messages as name-value pairs to a _Redis_ key-value store.

For the list of available parameters, see *Section 7.9.1, redis() destination options (p. 153)*.

```
Declaration:
redis(
    host("<redis-server-address>")
    port("<redis-server-port>")
    command("<redis-command>", "<first-command-parameter>",
"<second-command-parameter>", "<third-command-parameter>")
)
```

Example 7.10. Using the redis() driver
The following destination counts the number of log messages received per host.

```
destination d_redis {
    redis(
        host("localhost")
        port(6379)
        command("HINCRBY", "hosts", "$HOST", "1")
    );
};
```

The following example creates a statistic from Apache webserver logs about the browsers that the visitors use (per minute)

```
@version: 3.6

source s_apache {
    file("/var/log/apache2/access.log");
};

parser p_apache {
    csv-parser(columns("APACHE.CLIENT_IP", "APACHE.IDENT_NAME", "APACHE.USER_NAME",
                       "APACHE.TIMESTAMP", "APACHE.REQUEST_URL", "APACHE.REQUEST_STATUS",
                       "APACHE.CONTENT_LENGTH", "APACHE.REFERER", "APACHE.USER_AGENT",
                       "APACHE.PROCESS_TIME", "APACHE.SERVER_NAME")
                flags(escape-double-char,strip-whitespace)
    delimiters(" ")
    quote-pairs('""[]')
    );
};

destination d_redis {
  redis( command("HINCRBY" "${MONTH_ABBREV} ${DAY} ${HOUR}:${MIN}"  "${APACHE.USER_AGENT}"
 "1"));
};

log {
    source(s_apache);
    parser(p_apache);
    destination(d_redis);
};
```

7.9.1. redis() destination options

The *redis()* driver sends messages as name-value pairs to a *Redis* key-value store.

The *redis()* destination has the following options:

command()

Type:	comma-separated list of strings ("<redis-command>", "<first-command-parameter>", "<second-command-parameter>", "<third-command-parameter>")
Default:	empty string

Description: The *Redis command* to execute, for example, LPUSH, INCR, or HINCRBY. Using the HINCRBY command with an increment value of 1 allows you to create various statistics. For example, the `command("HINCRBY" "${HOST}/programs" "${PROGRAM}" "1")` command counts the number of log messages on each host for each program.

Note the following points when using the `redis()` destination:

- You can use macros and templates in the parameters of the Redis command.

- Currently you can use only one command in a redis() destination.

- The syslog-ng OSE application ignores the return value of the command. If the Redis server returns an error, syslog-ng OSE closes the connection.

host()

Type:	hostname or IP address
Default:	127.0.0.1

Description: The hostname or IP address of the Redis server.

port()

Type:	number
Default:	6379

Description: The port number of the Redis server.

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches `retries`, then drops the message.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

7.10. Monitoring your data with Riemann

The `riemann()` driver sends your data (for example, metrics or events) to a _Riemann_ monitoring system.

For the list of available parameters, see _Section 7.10.1, riemann() destination options (p. 155)_.

```
Declaration:
riemann(
 server("<riemann-server-address>")
 port("<riemann-server-port>")
 metric("<the-metric-or-data-to-send-to-riemann>")
)
```

Example 7.11. Using the riemann() driver
The following destination sends the value of the SEQNUM macro (the number of messages sent to this destination) as a metric to the Riemann server.

```
@version: 3.6

source s_network {
 tcp(port(12345));
};

destination d_riemann {
 riemann(
  server("localhost")
  port(5555)
  ttl("300.5")
  metric(int("$SEQNUM"))
  description("syslog-ng riemann test")
  state("ok")
  attributes(x-ultimate-answer("$(+ $PID 42)")
      key("MESSAGE", rekey(add-prefix("x-")) )
      )
 );
};

log {
 source(s_network);
 destination(d_riemann);
 flags(flow-control);
};
```

For a detailed use-case on using syslog-ng OSE with the Riemann monitoring system, see the article of Fabien Wernli at *http://devops.com/features/guide-modern-monitoring-alerting/*.

7.10.1. riemann() destination options

The *riemann()* driver sends metrics or events to a *Riemann* monitoring system.

The *riemann()* destination has the following options:

attributes()

Type:	parameter list of the *value-pairs()* option
Default:	

Description: The *attributes()* option adds extra metadata to the Riemann event, that can be displayed on the Riemann dashboard. To specify the metadata to add, use the syntax of the *value-pairs()* option. For details on using *value-pairs()*, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*.

description()

Type:	template, macro, or string
Default:	

Description: The value to add as the description field of the Riemann event.

host()

Type:	template, macro, or string
Default:	${HOST}

Description: The value to add as the host field of the Riemann event.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

metric()

Type:	template, macro, or string
Default:	

Description: The numeric value to add as the metric field of the Riemann event. If possible, include type-hinting as well, otherwise the Riemann server will interpret the value as a floating-point number. The following example specifies the SEQNUM macro as an integer.

```
metric(int("$SEQNUM"))
```

port()

Type:	number
Default:	5555

Description: The port number of the Riemann server.

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches *retries*, then drops the message.

server()

Type:	hostname or IP address
Default:	127.0.0.1

Description: The hostname or IP address of the Riemann server.

service()

Type:	template, macro, or string
Default:	${PROGRAM}

Description: The value to add as the service field of the Riemann event.

state()

Type:	template, macro, or string
Default:	

Description: The value to add as the state field of the Riemann event.

tags()

Type:	string list
Default:	the tags already assigned to the message

Description: The list of tags to add as the tags field of the Riemann event. If not specified syslog-ng OSE automatically adds the tags already assigned to the message. If you set the *tags()* option, only the tags you specify will be added to the event.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

ttl()

Type:	template, macro, or number
Default:	

Description: The value (in seconds) to add as the ttl (time-to-live) field of the Riemann event.

type()

Type:	tcp	udp
Default:	tcp

Description: The type of the network connection to the Riemann server: UDP or TCP.

7.11. Generating SMTP messages (e-mail) from logs

The *smtp()* driver sends e-mail messages triggered by log messages. The *smtp()* driver uses SMTP, without needing external applications. You can customize the main fields of the e-mail, add extra headers, send the e-mail to multiple recipients, and so on.

The *subject()*, *body()*, and *header()* fields may include macros which get expanded in the e-mail. For more information on available macros see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*.

The *smtp()* driver has the following required parameters: *host()*, *port()*, *from()*, *to()*, *subject()*, and *body()*. For the list of available optional parameters, see *Section 7.11.1, smtp() destination options (p. 159)*.

> **Note**
> The *smtp()* destination driver is available only in syslog-ng OSE 3.4 and later.

Declaration:
```
    smtp(host() port() from() to() subject() body() options());
```

Example 7.12. Using the smtp() driver
The following example defines an *smtp()* destination using only the required parameters.

```
destination d_smtp {
    smtp(
        host("localhost")
        port(25)
        from("syslog-ng alert service" "noreply@example.com")
        to("Admin #1" "admin1@example.com")
        subject("[ALERT] Important log message of $LEVEL condition received from
$HOST/$PROGRAM!")
        body("Hi!\nThe syslog-ng alerting service detected the following important log
message:\n $MSG\n-- \nsyslog-ng\n")
    );
};
```

The following example sets some optional parameters as well.

```
destination d_smtp {
    smtp(
        host("localhost")
        port(25)
        from("syslog-ng alert service" "noreply@example.com")
        to("Admin #1" "admin1@example.com")
        to("Admin #2" "admin2@example.com")
        cc("Admin BOSS" "admin.boss@example.com")
        bcc("Blind CC" "blindcc@example.com")
        subject("[ALERT] Important log message of $LEVEL condition received from
$HOST/$PROGRAM!")
        body("Hi!\nThe syslog-ng alerting service detected the following important log
message:\n $MSG\n-- \nsyslog-ng\n")
        header("X-Program", "$PROGRAM")
        );
};
```

Example 7.13. Simple e-mail alerting with the *smtp()* driver
The following example sends an e-mail alert if the eth0 network interface of the host is down.

```
filter f_linkdown {
    match("ethO: link down" value("MESSAGE"));
};
destination d_alert {
    smtp(
        host("localhost") port(25)
        from("syslog-ng alert service" "syslog@localhost")
        reply-to("Admins" "root@localhost")
        to("Ennekem" "me@localhost")
        subject("[SYSLOG ALERT]: ethO link down")
        body("Syslog received an alert:\n$MSG")
        );
};

log {
    source(s_local);
    filter(f_linkdown);
    destination(d_alert);
};
```

7.11.1. smtp() destination options

The *smtp()* sends e-mail messages using SMTP, without needing external applications. The *smtp()* destination has the following options:

body()

Type:	string
Default:	n/a

Description: The BODY field of the e-mail. You can also use macros in the string. Use \n to start a new line. For example:

```
body("syslog-ng OSE received the following alert from $HOST:\n$MSG")
```

bcc()

Type:	string
Default:	n/a

Description: The BCC recipient of the e-mail (contents of the BCC field). You can specify the e-mail address, or the name and the e-mail address. Set the *bcc()* option multiple times to send the e-mail to multiple recipients. For example:

```
bcc("admin@example.com")
```

or

```
bcc("Admin" "admin@example.com")
```

or

```
bcc("Admin" "admin@example.com")
bcc("Admin2" "admin2@example.com")
```

cc()

Type:	string
Default:	n/a

Description: The CC recipient of the e-mail (contents of the CC field). You can specify the e-mail address, or the name and the e-mail address. Set the *cc()* option multiple times to send the e-mail to multiple recipients. For example:

```
cc("admin@example.com")
```

or

```
cc("Admin" "admin@example.com")
```

or

```
cc("Admin" "admin@example.com")
cc("Admin2" "admin2@example.com")
```

from()

Type:	string
Default:	n/a

Description: The sender of the e-mail (contents of the FROM field). You can specify the e-mail address, or the name and the e-mail address. For example:

```
from("admin@example.com")
```

or

```
from("Admin" "admin@example.com")
```

If you specify the *from()* option multiple times, the last value will be used. Instead of the *from()* option, you can also use *sender()*, which is just an alias of the *from()* option.

header()

Type:	string
Default:	n/a

Description: Adds an extra header to the e-mail with the specified name and content. The first parameter sets the name of the header, the second one its value. The value of the header can contain macros. Set the *header()* option multiple times to add multiple headers. For example:

```
header("X-Program", "$PROGRAM")
```

When using the header option, note the following points:

- Do not use the *header()* option to set the values of headers that have dedicated options. Use it only to add extra headers.

- If you set the same custom header multiple times, only the first will be added to the e-mail, other occurrences will be ignored.
- It is not possible to set the DATE header.

host()

Type:	hostname or IP address
Default:	n/a

Description: Hostname or IP address of the SMTP server.

> **Note**
> If you specify `host="localhost"`, syslog-ng OSE will use a socket to connect to the local SMTP server. Use `host="127.0.0.1"` to force TCP communication between syslog-ng OSE and the local SMTP server.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

port()

Type:	number
Default:	25

Description: The port number of the SMTP server.

reply-to()

Type:	string
Default:	n/a

Description: Replies of the recipient will be sent to this address (contents of the REPLY-TO field). You can specify the e-mail address, or the name and the e-mail address. Set the *reply-to()* option multiple times to send the e-mail to multiple recipients. For example:

```
reply-to("admin@example.com")
```

or

```
reply-to("Admin" "admin@example.com")
```

or

```
reply-to("Admin" "admin@example.com")
reply-to("Admin2" "admin2@example.com")
```

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches *retries*, then drops the message.

subject()

Type:	string
Default:	n/a

Description: The SUBJECT field of the e-mail. You can also use macros. For example:

```
subject("[SYSLOG ALERT]: Critical error message received from $HOST")
```

If you specify the *subject()* option multiple times, the last value will be used.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

to()

Type:	string
Default:	localhost

Description: The recipient of the e-mail (contents of the TO field). You can specify the e-mail address, or the name and the e-mail address. Set the *to()* option multiple times to send the e-mail to multiple recipients. For example:

```
to("admin@example.com")
```

or

```
to("Admin" "admin@example.com")
```

or

```
to("Admin" "admin@example.com")
to("Admin2" "admin2@example.com")
```

7.12. Storing messages in an SQL database

The *sql()* driver sends messages into an SQL database. Currently the Microsoft SQL (MSSQL), MySQL, Oracle, PostgreSQL, and SQLite databases are supported.

```
Declaration:
    sql(database_type host_parameters database_parameters [options]);
```

The *sql()* driver has the following required parameters: *type()*, *database()*, *table*, *columns()*, and *values*.

Warning
The syslog-ng application requires read and write access to the SQL table, otherwise it cannot verify that the destination table exists.

Currently the syslog-ng application has default schemas for the different databases and uses these defaults if the database schema (for example columns and column types) is not defined in the configuration file. However, these schemas will be deprecated and specifying the exact database schema will be required in later versions of syslog-ng.

The *table* and *value* parameters can include macros to create tables and columns dynamically (for details, see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*).

Warning
When using macros in table names, note that some databases limit the maximum allowed length of table names. Consult the documentation of the database for details.

Inserting the records into the database is performed by a separate thread. The syslog-ng application automatically performs the escaping required to insert the messages into the database.

Example 7.14. Using the sql() driver
The following example sends the log messages into a PostgreSQL database running on the `logserver` host. The messages are inserted into the `logs` database, the name of the table includes the exact date and the name of the host sending the messages. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
  sql(type(pgsql)
  host("logserver") username("syslog-ng") password("password")
  database("logs")
  table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
  columns("datetime", "host", "program", "pid", "message")
  values("{$R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
  indexes("datetime", "host", "program", "pid", "message"));
  };
```

The following example specifies the type of the database columns as well:

```
destination d_sql {
  sql(type(pgsql)
  host("logserver") username("syslog-ng") password("password")
  database("logs")
  table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
  columns("datetime varchar(16)", "host varchar(32)", "program  varchar(20)", "pid
varchar(8)", "message  varchar(200)")
  values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
```

```
    indexes("datetime", "host", "program", "pid", "message"));
};
```

7.12.1. Using the sql() driver with an Oracle database

The Oracle sql destination has some special aspects that are important to note.

■ The hostname of the database server is set in the `tnsnames.ora` file, not in the *host* parameter of the *sql()* destination.

 If the `tnsnames.ora` file is not located in the /etc directory (or in the /var/opt/oracle directory on Solaris), set the following Oracle-related environment variables, so syslog-ng OSE will find the file: *ORACLE_BASE*, *ORACLE_HOME*, and *ORACLE_SID*. For details, see the documentation of the Oracle Instant Client.

■ You cannot use the same *database()* settings in more than one destination, because the *database()* option of the SQL driver is just a reference to the connection string of the `tnsnames.ora` file. To overcome this problem, you can duplicate the connections in the `tnsnames.ora` file under a different name, and use a different table in each Oracle destination in syslog-ng OSE.

■ As certain database versions limit the maximum length of table names, macros in the table names should be used with care.

■ In the current version of syslog-ng OSE, the types of database columns must be explicitly set for the Oracle destination. The column used to store the text part of the syslog messages should be able to store messages as long as the longest message permitted by syslog-ng, therefore it is usually recommended to use the `varchar2` or `clob` column type. (The maximum length of the messages can be set using the *log-msg-size()* option.) For details, see the following example.

■ The Oracle Instant Client used by syslog-ng OSE supports only the following character sets:

 • Single-byte character sets: `US7ASCII`, `WE8DEC`, `WE8MSWIN1252`, and `WE8ISO8859P1`

 • Unicode character sets: `UTF8`, `AL16UTF16`, and `AL32UTF8`

Example 7.15. Using the sql() driver with an Oracle database
The following example sends the log messages into an Oracle database running on the `logserver` host, which must be set in the `/etc/tnsnames.ora` file. The messages are inserted into the `LOGS` database, the name of the table includes the exact date when the messages were sent. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
  sql(type(oracle)
  username("syslog-ng") password("password")
  database("LOGS")
  table("msgs_${R_YEAR}${R_MONTH}${R_DAY}")
  columns("datetime varchar(16)", "host varchar(32)", "program varchar(32)", "pid
varchar(8)", "message varchar2")
  values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
  indexes("datetime", "host", "program", "pid", "message"));
  };
```

The Oracle Instant Client retrieves the address of the database server from the `/etc/tnsnames.ora` file. Edit or create this file as needed for your configuration. A sample is provided below.

```
LOGS =
(DESCRIPTION =
(ADDRESS_LIST =
(ADDRESS = (PROTOCOL = TCP)
(HOST = logserver)
(PORT = 1521))
)
(CONNECT_DATA =
(SERVICE_NAME = EXAMPLE.SERVICE)
)
)
```

7.12.2. Using the sql() driver with a Microsoft SQL database

The *mssql* database driver can access Microsoft SQL (MSSQL) destinations. This driver has some special aspects that are important to note.

- The date format used by the MSSQL database must be explicitly set in the /etc/locales.conf file of the syslog-ng server. For details, see the following example.

- As certain database versions limit the maximum length of table names, macros in the table names should be used with care.

- In the current version of syslog-ng OSE, the types of database columns must be explicitly set for the MSSQL destination.

Warning
The following column types cannot be used in MSSQL destinations: nchar, nvarchar, ntext, and xml.

- The column used to store the text part of the syslog messages should be able to store messages as long as the longest message permitted by syslog-ng. The *varchar* column type can store maximum 4096 bytes-long messages. The maximum length of the messages can be set using the *log-msg-size()* option. For details, see the following example.

- Remote access for SQL users must be explicitly enabled on the Microsoft Windows host running the Microsoft SQL Server. For details, see *Procedure 3.3, Configuring Microsoft SQL Server to accept logs from syslog-ng (p. 27)*.

Example 7.16. Using the sql() driver with an MSSQL database
The following example sends the log messages into an MSSQL database running on the logserver host. The messages are inserted into the syslogng database, the name of the table includes the exact date when the messages were sent. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_mssql {
sql(type(mssql) host("logserver") port("1433")
  username("syslogng") password("syslogng") database("syslogng")
  table("msgs_${R_YEAR}${R_MONTH}${R_DAY}")columns("datetime varchar(16)", "host
varchar(32)",
  "program varchar(32)", "pid varchar(8)", "message varchar(4096)")
  values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
  indexes("datetime", "host", "program", "pid"));
  };
```

The date format used by the MSSQL database must be explicitly set in the `/etc/locales.conf` file of the syslog-ng server. Edit or create this file as needed for your configuration. A sample is provided below.

```
[default]
date = "%Y-%m-%d %H:%M:%S"
```

7.12.3. The way syslog-ng interacts with the database

Used SQL operations by syslog-ng.

Create table:

- If the given table does not exist, syslog-ng tries to create it with the given column types.
- The syslog-ng OSE application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.
- If syslog-ng cannot create or alter a table, it tries to do it again when reach the next time-reopen().

Alter table:

- If the table structure is different from given structure in an existing table, syslog-ng tries to add columns in this table but never will delete or modify an existing column.
- If syslog-ng OSE cannot create or alter a table, it tries to do it again when reach the next time-reopen().
- The syslog-ng OSE application requires read and write access to the SQL table, otherwise it cannot verify that the destination table exists.

Insert table:

- Insert new records in a table.
- Inserting the records into the database is performed by a separate thread.
- The syslog-ng OSE application automatically performs the escaping required to insert the messages into the database.
- If insert returns with error, syslog-ng tries to insert the message +two times by default, then drops it. Retrying time is the value of time-reopen().

Encoding.

The syslog-ng OSE application uses UTF-8 by default when writes logs into database.

Start/stop and reload.

Start:

- The syslog-ng OSE application will connect to database automatically after starting regardless existing incoming messages.

Stop:

- The syslog-ng OSE application will close the connection to database before shutting down.

Possibility of losing logs:

- The syslog-ng OSE application cannot lose logs during shutting down if disk buffer was given and it is not full yet.

- The syslog-ng OSE application cannot lose logs during shutting down if disk buffer was not given.

Reload:

- The syslog-ng OSE application will close the connection to database if it received SIGHUP signal (reload).

- It will reconnect to the database when it tries to send a new message to this database again.

Macros:

The value of ${SEQNUM} macro will be overrided by sql driver regardless of local or relayed incoming message.

It will be grown continously.

7.12.3.1. MySQL-specific interaction methods

To specify the socket to use, set and export the *MYSQL_UNIX_PORT* environment variable, for example `MYSQL_UNIX_PORT=/var/lib/mysql/mysql.sock; export MYSQL_UNIX_PORT`.

7.12.3.2. MsSQL-specific interaction methods

In SQL Server 2005 this restriction is lifted - kind of. The total length of all key columns in an index cannot exceed 900 bytes.

If you are using *null()* in your configuration, be sure that the columns allow *NULL* to insert. Give the column as the following example: `"datetime varchar(16) NULL"`.

The date format used by the MSSQL database must be explicitly set in the `/etc/locales.conf` file of the syslog-ng server. `[default] date = "%Y-%m-%d %H:%M:%S"`.

7.12.4. sql() destination options

This driver sends messages into an SQL database. The *sql()* destination has the following options:

columns()

Type:	string list
Default:	"date", "facility", "level", "host", "program", "pid", "message"

Description: Name of the columns storing the data in `fieldname [dbtype]` format. The `[dbtype]` parameter is optional, and specifies the type of the field. By default, syslog-ng OSE creates `text` columns. Note that not every database engine can index text fields.

> **Warning**
> The following column types cannot be used in MSSQL destinations: `nchar`, `nvarchar`, `ntext`, and `xml`.

database()

Type:	string
Default:	logs

Description: Name of the database that stores the logs. Macros cannot be used in database name. Also, when using an Oracle database, you cannot use the same *database()* settings in more than one destination.

dbd-option()

Type:	string
Default:	empty string

Description: Specify database options that are set whenever syslog-ng OSE connects to the database server. Consult the documentation of your database server for details on the available options. Syntax:

```
dbd-option(OPTION_NAME VALUE)
```

OPTION_NAME is always a string, VALUE is a string or a number. For example:

```
dbd-option("null.sleep.connect" 1)
dbd-option("null.sleep.query" 5)
```

flags()

Type:	list of flags
Default:	empty string

Description: Flags related to the *sql()* destination.

- *dont-create-tables*: Enable this flag to prevent syslog-ng OSE from creating non-existing database tables automatically. The syslog-ng OSE application typically has to create tables if you use macros in the table names. Available in syslog-ng OSE version 3.2 and later.

- *explicit-commits*: By default, syslog-ng OSE commits every log message to the target database individually. When the *explicit-commits* option is enabled, messages are committed in batches. This improves the performance, but results in some latency, as the messages are not immediately sent to the database. The size and frequency of batched commits can be set using the *flush-lines()* and *flush-timeout()* parameters. The *explicit-commits* option is available in syslog-ng OSE version 3.2 and later.

Example 7.17. Setting flags for SQL destinations
The following example sets the *dont-create-tables* and *explicit-commits* flags for an *sql()* destination.

```
flags(dont-create-tables,explicit-commits)
```

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

host()

Type:	hostname or IP address
Default:	n/a

Description: Hostname of the database server. Note that Oracle destinations do not use this parameter, but retrieve the hostname from the `/etc/tnsnames.ora` file.

> **Note**
> If you specify host="localhost", syslog-ng will use a socket to connect to the local database server. Use host="127.0.0.1" to force TCP communication between syslog-ng and the local database server.
>
> To specify the socket to use, set and export the *MYSQL_UNIX_PORT* environment variable, for example MYSQL_UNIX_PORT=/var/lib/mysql/mysql.sock; export MYSQL_UNIX_PORT.

indexes()

Type:	string list
Default:	"date", "facility", "host", "program"

Description: The list of columns that are indexed by the database to speed up searching. To disable indexing for the destination, include the empty *indexes()* parameter in the destination, simply omitting the *indexes* parameter will cause syslog-ng to request indexing on the default columns.

The syslog-ng OSE application will create the name of indexes automaticaly with the following method:

- In case of MsSQL, PostgreSQL, MySQL or SQLite or (Oracle but tablename < 30 characters): {table}_{column}_idx.
- In case of Oracle and tablename > 30 characters: md5sum of {table}_{column}-1 and the first character will be replaced by "i" character and the md5sum will be truncated to 30 characters.

local-time-zone()

Type:	name of the timezone or the timezone offset
Default:	The local timezone.

Description: Sets the timezone used when expanding filename and tablename templates. The timezone can be specified as using the name of the (for example time-zone("Europe/Budapest")), or as the timezone offset (for example +01:00). The valid timezone names are listed under the /usr/share/zoneinfo directory.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

null()

Type:	string
Default:	

Description: If the content of a column matches the string specified in the *null()* parameter, the contents of the column will be replaced with an SQL NULL value. If unset (by default), the option does not match on any string. For details, see the *Example 7.18, Using SQL NULL values (p. 171)*.

Example 7.18. Using SQL NULL values
The *null()* parameter of the SQL driver can be used to replace the contents of a column with a special SQL NULL value. To replace every column that contains an empty string with NULL, use the null("") option, for example

```
destination d_sql {
                sql(type(pgsql)
                host("logserver") username("syslog-ng") password("password")
                database("logs")
                table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
                columns("datetime", "host", "program", "pid", "message")
                values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
                indexes("datetime", "host", "program", "pid", "message")
                null(""));
                };
```

To replace only a specific column (for example *pid*) if it is empty, assign a default value to the column, and use this default value in the *null()* parameter:

```
destination d_sql {
                sql(type(pgsql)
                host("logserver") username("syslog-ng") password("password")
                database("logs")
                table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
                columns("datetime", "host", "program", "pid", "message")
                values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID:-@@NULL@@}",
"${MSGONLY}")

                indexes("datetime", "host", "program", "pid", "message")
                null("@@NULL@@"));
                };
```

Ensure that the default value you use does not appear in the actual log messages, because other occurrences of this string will be replaced with NULL as well.

password()

Type:	string
Default:	n/a

Description: Password of the database user.

port()

Type:	number
Default:	1433 TCP for MSSQL, 3306 TCP for MySQL, 1521 for Oracle, and 5432 TCP for PostgreSQL

Description: The port number to connect to.

retries

Type:	number
Default:	3

Description: The number of insertion attempts. If syslog-ng OSE could not insert a message into the database, it will repeat the attempt until the number of attempts reaches *retries*, then drops the connection to the database. For example, syslog-ng OSE will try to insert a message maximum three times by default (once for first insertion and twice if the first insertion was failed).

session-statements()

Type:	comma-separated list of SQL statements
Default:	empty string

Description: Specifies one or more SQL-like statement which is executed after syslog-ng OSE has successfully connected to the database. For example:

```
session-statements("SET COLLATION_CONNECTION='utf8_general_ci'")
```

> **Warning**
> The syslog-ng OSE application does not validate or limit the contents of customized queries. Consequently, queries performed with a user with write-access can potentially modify or even harm the database. Use customized queries with care, and only for your own responsibility.

table()

Type:	string
Default:	messages

Description: Name of the database table to use (can include macros). When using macros, note that some databases limit the length of table names.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

type()

Type:	mssql, mysql, oracle, pgsql, or sqlite3
Default:	mysql

Description: Specifies the type of the database, that is, the DBI database driver to use. Use the *mssql* option to send logs to an MSSQL database. For details, see the examples of the databases on the following sections.

username()

Type:	string
Default:	n/a

Description: Name of the database user.

values()

Type:	string list
Default:	"${R_YEAR}-${R_MONTH}-${R_DAY},${R_HOUR}:${R_MIN}:${R_SEC}", "${FACILITY}", "${LEVEL}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}"

Description: The parts of the message to store in the fields specified in the `columns` parameter.

It is possible to give a special value calling: default (without quotation marks).It means that the value will be used that is the default of the column type of this value.

Example 7.19. Value: default

```
columns("date datetime", "host varchar(32)", "row_id serial")
 values("${R_DATE}", "${HOST}", default)
```

7.13. Publishing messages using STOMP

The `stomp()` driver sends messages to servers (message brokers) using the *Simple (or Streaming) Text Oriented Message Protocol (STOMP)*, formerly known as TTMP. syslog-ng OSE supports version 1.0 of the STOMP protocol. The syslog-ng OSE `stomp()` driver supports persistence.

The name-value pairs selected with the `value-pairs()` option will be sent as STOMP headers, while the body of the STOMP message is empty by default (but you can add custom content using the `body()` option). Publishing the name-value pairs as headers makes it possible to use the Headers exchange-type and subscribe only to interesting log streams. This solution is more flexible than using the `routing-key()` option.

For the list of available parameters, see *Section 7.13.1, stomp() destination options (p. 174)*.

```
Declaration:
stomp( host("<stomp-server-address>") )
```

Example 7.20. Using the stomp() driver
The following example shows the default values of the available options.

```
destination d_stomp {
    stomp(
        host("localhost")
        port(61613)
        destination("/topic/syslog")
        routing-key("")      # optional, empty by default
        body("")             # optional, empty by default
        persistent(yes)
        ack(no)
        username("user")     # optional, empty by default
        password("password") # optional, empty by default
        value-pairs(scope(selected-macros, nv-pairs, sdata))
    );
};
```

7.13.1. stomp() destination options

The *stomp()* driver publishes messages using the Simple (or Streaming) Text Oriented Message Protocol (STOMP).

The *stomp()* destination has the following options:

ack()

Type:	string
Default:	no

Description: Request the STOMP server to acknowledge the receipt of the messages. If you enable this option, then after sending a message, syslog-ng OSE waits until the server confirms that it has received the message. This delay can seriously limit the performance of syslog-ng OSE if the message rate is high, and the server cannot acknowledge the messages fast enough.

body()

Type:	string
Default:	empty string

Description: The body of the STOMP message. You can also use macros and templates.

destination()

Type:	string
Default:	/topic/syslog

Description: The name of the destination (message queue) on the STOMP server. It can include macros and templates.

host()

Type:	hostname or IP address
Default:	127.0.0.1

Description: The hostname or IP address of the STOMP server.

password()

Type:	string
Default:	n/a

Description: The password used to authenticate on the STOMP server.

persistent()

Type:	yes|no
Default:	yes

Description: If this option is enabled, the STOMP server or broker will store the messages on its hard disk. That way, the messages will be retained if the STOMP server is restarted, if the message queue is set to be durable on the STOMP server.

port()

Type:	number
Default:	61613

Description: The port number of the STOMP server.

retries()

Type:	number (of attempts)
Default:	3

Description: The number of times syslog-ng OSE attempts to send a message to this destination. If syslog-ng OSE could not send a message, it will try again until the number of attempts reaches *retries*, then drops the message.

routing-key()

Type:	string
Default:	empty string

Description: Specifies a routing key for the exchange. The routing key selects certain messages published to an exchange to be routed to the bound queue. In other words, the routing key acts like a filter. The routing key can include macros and templates.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

username()

Type:	string
Default:	empty string

Description: The username used to authenticate on the STOMP server.

value-pairs()

Type:	parameter list of the *value-pairs()* option
Default:	`scope("selected-macros" "nv-pairs")`

Description: The *value-pairs()* option creates structured name-value pairs from the data and metadata of the log message. For details on using *value-pairs()*, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*.

> **Note**
> Empty keys are not logged.

7.14. Sending messages to a remote logserver using the IETF-syslog protocol

The *syslog()* driver sends messages to a remote host (for example a syslog-ng server or relay) on the local intranet or internet using the new standard syslog protocol developed by IETF (for details about the new protocol, see *Section 2.8.2, IETF-syslog messages (p. 13)*). The protocol supports sending messages using the UDP, TCP, or the encrypted TLS networking protocols.

The required arguments of the driver are the address of the destination host (where messages should be sent). The transport method (networking protocol) is optional, syslog-ng uses the TCP protocol by default. For the list of available optional parameters, see *Section 7.14.1, syslog() destination options (p. 177)*.

```
Declaration:
          syslog(host transport [options]);
```

> **Note**
> Note that the *syslog* destination driver has required parameters, while the source driver defaults to the local bind address, and every parameter is optional.

The *udp* transport method automatically sends multicast packets if a multicast destination address is specified. The *tcp* and *tls* methods do not support multicasting.

> **Note**
> The default ports for the different transport protocols are as follows: UDP — 514; TLS — 6514.

> **Example 7.21. Using the syslog() driver**
>
> `destination d_tcp { syslog("10.1.2.3" transport("tcp") port(1999) localport(999)); };`
>
> If name resolution is configured, the hostname of the target server can be used as well.
>
> `destination d_tcp { syslog("target_host" transport("tcp") port(1999) localport(999)); };`

Send the log messages using TLS encryption and use mutual authentication. For details on the encryption and authentication options, see *Section 10.4, TLS options (p. 235)*.

```
destination d_syslog_tls{
            syslog("10.100.20.40"
            transport("tls")
            port(6514)
            tls(peer-verify(required-trusted)
            ca-dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
            key-file('/opt/syslog-ng/etc/syslog-ng/keys/client_key.pem')
            cert-file('/opt/syslog-ng/etc/syslog-ng/keys/client_certificate.pem'))
            );};
```

Note
If a message uses the IETF-syslog format (RFC5424), only the text of the message can be customized (that is, the $MESSAGE part of the log), the structure of the header is fixed.

7.14.1. syslog() destination options

The `syslog()` driver sends messages to a remote host (for example a syslog-ng server or relay) on the local intranet or internet using the RFC5424 syslog protocol developed by IETF (for details about the protocol, see *Section 2.8.2, IETF-syslog messages (p. 13)*). The protocol supports sending messages using the UDP, TCP, or the encrypted TLS networking protocols.

These destinations have the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

■ *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line.

■ *syslog-protocol*: The `syslog-protocol` flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the `syslog` driver, and that the `syslog` driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number

increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

ip-tos()

Type:	number
Default:	0

Description: Specifies the Type-of-Service value of outgoing packets.

ip-ttl()

Type:	number
Default:	0

Description: Specifies the Time-To-Live value of outgoing packets.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to destinations should be closed when syslog-ng is reloaded. Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always reopened after receiving a HUP signal unless the *keep-alive* option is enabled for the source.

localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to before connecting to target.

localport()

Type:	number
Default:	0

Description: The port number to bind to. Messages are sent from this port.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

mark-freq()

Accepted values:	number
Default:	1200

Description: An alias for the obsolete *mark()* option, retained for compatibility with syslog-ng version 1.6.x. The number of seconds between two *MARK* messages. *MARK* messages are generated when there was no message traffic to inform the receiver that the connection is still alive. If set to zero (0), no *MARK* messages are sent. The *mark-freq()* can be set for global option and/or every *MARK* capable destination driver if *mark-mode()* is periodical or dst-idle or host-idle. If *mark-freq()* is not defined in the destination, then the *mark-freq()* will be inherited from the global options. If the destination uses internal *mark-mode()*, then the global *mark-freq()* will be valid (does not matter what *mark-freq()* set in the destination side).

mark-mode()

Accepted values:	internal|dst-idle|host-idle|periodical|none|global
Default:	internal for pipe, program drivers
	none for file, unix-dgram, unix-stream drivers
	global for syslog, tcp, udp destinations
	host-idle for global option

Description: The `mark-mode()` option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  `file()`, `pipe()`, `unix-stream()`, `unix-dgram()`, `program()`

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp()`, `udp()`, `syslog()`, `program()`, `file()`, `pipe()`, `unix-stream()`, `unix-dgram()`.

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp()`, `udp()`, `syslog()`, `program()`, `file()`, `pipe()`, `unix-stream()`, `unix-dgram()`.

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp()`, `udp()`, `syslog()`, `program()`, `file()`, `pipe()`, `unix-stream()`, `unix-dgram()`.

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where `none` is the default value, then `none` will be used.

- `global`: Destination driver uses the global `mark-mode()` setting. The syslog-ng interprets syntax error if the global `mark-mode()` is global.

> **Note**
> In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

port() or destport()

Type:	number
Default:	601

Description: The port number to connect to. Note that the default port numbers used by syslog-ng do not comply with the latest RFC which was published after the release of syslog-ng 3.0.2, therefore the default port numbers will change in the future releases.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

spoof-source()

Type:	yes or no
Default:	no

Description: Enables source address spoofing. This means that the host running syslog-ng generates UDP packets with the source IP address matching the original sender of the message. It is useful when you want to perform some kind of preprocessing via syslog-ng then forward messages to your central log management solution with the source address of the original sender. This option only works for UDP destinations though the original message can be received by TCP as well. This option is only available if syslog-ng was compiled using the *--enable-spoof-source* configuration option.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

tcp-keepalive-intvl()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (number of seconds) between subsequential keepalive probes, regardless of the traffic exchanged in the connection. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_intvl`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-probes()

Type:	number
Default:	0

Description: Specifies the number of unacknowledged probes to send before considering the connection dead. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_probes`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-time()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (in seconds) between the last data packet sent and the first keepalive probe. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_time`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`, and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

> **Note**
>
> If a message uses the IETF-syslog format (RFC5424), only the text of the message can be customized (that is, the $MESSAGE part of the log), the structure of the header is fixed.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

transport()

Type:	udp, tcp, or tls
Default:	tcp

Description: Specifies the protocol used to send messages to the destination server.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.15. Sending messages to a remote log server using the legacy BSD-syslog protocol (tcp(), udp() drivers)

> **Note**
> The `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers will be deprecated in later versions, use the `network()` driver instead.

The `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers send messages to another host (for example a syslog-ng server or relay) on the local intranet or internet using the UDP or TCP protocol. The `tcp6()` and `udp6()` drivers use the IPv6 network protocol.

All four drivers have a single required parameter specifying the destination host address, where messages should be sent. For the list of available optional parameters, see *Section 7.15.1, tcp(), tcp6(), udp(), and udp6() destination options (p. 185)*.

The *udp()* and *udp6()* drivers automatically send multicast packets if a multicast destination address is specified. The *tcp()* and *tcp6()* drivers do not support multicasting.

```
Declaration:
            tcp(host [options]);
            udp(host [options]);
            tcp6(host [options]);
            udp6(host [options]);
```

Example 7.22. Using the tcp() driver

```
destination d_tcp { tcp("10.1.2.3" port(1999) localport(999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { tcp("target_host" port(1999) localport(999)); };
```

To send messages using the IETF-syslog message format without using the IETF-syslog protocol, enable the *syslog-protocol* flag:

```
destination d_tcp { tcp("10.1.2.3" port(1999) flags(syslog-protocol) ); };
```

(For details on how to use the IETF-syslog protocol, see *Section 7.14.1, syslog() destination options (p. 177)*.)

Using an IPv6 address:

```
tcp6("fd00:abcd:4321:2:20c:29ff:fea8:9671" port(514));
```

7.15.1. tcp(), tcp6(), udp(), and udp6() destination options

This driver sends messages to another host on the local intranet or internet according to RFC3164 using the UDP or TCP protocol. The *tcp6()* and *udp6()* drivers use the IPv6 network protocol.

Note
When using IPv6 destination addresses, always enclose the address between double-quotes:

```
tcp6("fd00:abcd:4321:2:20c:29ff:fea8:9671" port(514));
```

These destinations have the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The *no-multi-line* flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The `syslog-protocol` flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the `syslog` driver, and that the `syslog` driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The `frac-digits()` parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

ip-tos()

Type:	number
Default:	0

Description: Specifies the Type-of-Service value of outgoing packets.

ip-ttl()

Type:	number
Default:	0

Description: Specifies the Time-To-Live value of outgoing packets.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to destinations should be closed when syslog-ng is reloaded. Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always reopened after receiving a HUP signal unless the *keep-alive* option is enabled for the source.

localip()

Type:	string
Default:	0.0.0.0

Description: The IP address to bind to before connecting to target.

localport()

Type:	number
Default:	0

Description: The port number to bind to. Messages are sent from this port.

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The *mark-mode()* option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- ■ `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark

messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

`file(), pipe(), unix-stream(), unix-dgram(), program()`

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().`

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().`

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram().`

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where `none` is the default value, then `none` will be used.

- `global`: Destination driver uses the global *mark-mode()* setting. The syslog-ng interprets syntax error if the global *mark-mode()* is global.

Note
In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

port() or destport()

Type:	number
Default:	514

Description: The port number to connect to. Note that the default port numbers used by syslog-ng do not comply with the latest RFC which was published after the release of syslog-ng 3.0.2, therefore the default port numbers will change in the future releases.

> **Note**
> The TCP port 514 is reserved for use with `rshell`, so select a different port if syslog-ng and `rshell` is used at the same time.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

spoof-source()

Type:	yes or no
Default:	no

Description: Enables source address spoofing. This means that the host running syslog-ng generates UDP packets with the source IP address matching the original sender of the message. It is useful when you want to perform some kind of preprocessing via syslog-ng then forward messages to your central log management solution with the source address of the original sender. This option only works for UDP destinations though the original message can be received by TCP as well. This option is only available if syslog-ng was compiled using the *--enable-spoof-source* configuration option.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

tcp-keepalive-intvl()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (number of seconds) between subsequential keepalive probes, regardless of the traffic exchanged in the connection. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_intvl`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time(), tcp-keepalive-probes(),` and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-probes()

Type:	number
Default:	0

Description: Specifies the number of unacknowledged probes to send before considering the connection dead. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_probes`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time(), tcp-keepalive-probes(),` and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

tcp-keepalive-time()

Type:	number [seconds]
Default:	0

Description: Specifies the interval (in seconds) between the last data packet sent and the first keepalive probe. This option is equivalent to `/proc/sys/net/ipv4/tcp_keepalive_time`. The default value is *0*, which means using the kernel default.

> **Warning**
>
> The `tcp-keepalive-time()`, `tcp-keepalive-probes()`, and `tcp-keepalive-intvl()` options only work on platforms which support the `TCP_KEEPCNT`, `TCP_KEEPIDLE`,and `TCP_KEEPINTVL` setsockopts. Currently, this is Linux.
>
> A connection that has no traffic is closed after `tcp-keepalive-time()` + `tcp-keepalive-intvl()` * `tcp-keepalive-probes()` seconds.

Available in syslog-ng OSE version 3.4 and later.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

tls()

Type:	tls options
Default:	n/a

Description: This option sets various options related to TLS encryption, for example, key/certificate files and trusted CA locations. TLS can be used only with tcp-based transport protocols. For details, see *Section 10.4, TLS options (p. 235)*.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global `ts-format()` parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.16. Sending messages to UNIX domain sockets

The `unix-stream()` and `unix-dgram()` drivers send messages to a UNIX domain socket in either `SOCK_STREAM` or `SOCK_DGRAM` mode.

Both drivers have a single required argument specifying the name of the socket to connect to. For the list of available optional parameters, see *Section 7.16.1, unix-stream() and unix-dgram() destination options (p. 192)*.

```
Declaration:
          unix-stream(filename [options]);
          unix-dgram(filename [options]);
```

Example 7.23. Using the unix-stream() driver

```
destination d_unix_stream { unix-stream("/var/run/logs"); };
```

7.16.1. unix-stream() and unix-dgram() destination options

These drivers send messages to a unix socket in either `SOCK_STREAM` or `SOCK_DGRAM` mode. The `unix-stream()` and `unix-dgram()` destinations have the following options:

flags()

Type:	no_multi_line, syslog-protocol
Default:	empty set

Description: Flags influence the behavior of the destination driver.

- *no-multi-line*: The `no-multi-line` flag disables line-breaking in the messages; the entire message is converted to a single line.

- *syslog-protocol*: The `syslog-protocol` flag instructs the driver to format the messages according to the new IETF syslog protocol standard (RFC5424), but without the frame header. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the `syslog` driver, and that the `syslog` driver automatically adds the frame header to the messages.

flush-lines()

Type:	number
Default:	Use global setting.

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the `flush-timeout()` option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the `flush-timeout()` time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

For optimal performance when sending messages to an syslog-ng OSE server, make sure that the `flush-lines()` is smaller than the window size set using the `log-iw-size()` option in the source of your server.

flush-timeout() (DEPRECATED)

Type:	time in milliseconds
Default:	Use global setting.

Description: This is a deprecated option. Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For details, see the `flush-lines()` option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The *frac-digits()* parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

log-fifo-size()

Type:	number
Default:	Use global setting.

Description: The number of messages that the output queue can store.

keep-alive()

Type:	yes or no
Default:	yes

Description: Specifies whether connections to destinations should be closed when syslog-ng is reloaded. Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always reopened after receiving a HUP signal unless the *keep-alive* option is enabled for the source.

so-broadcast()

Type:	yes or no
Default:	no

Description: This option controls the *SO_BROADCAST* socket option required to make syslog-ng send messages to a broadcast address. For details, see the `socket(7)` manual page.

so-keepalive()

Type:	yes or no
Default:	no

Description: Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. For details, see the `socket(7)` manual page.

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The `mark-mode()` option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

  `file(), pipe(), unix-stream(), unix-dgram(), program()`

- `dst-idle`: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `host-idle`: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `periodical`: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: `tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()`.

- `none`: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where `none` is the default value, then `none` will be used.

- `global`: Destination driver uses the global `mark-mode()` setting. The syslog-ng interprets syntax error if the global `mark-mode()` is global.

> **Note**
> In case of `dst-idle`, `host-idle` and `periodical`; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

so-rcvbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket receive buffer in bytes. For details, see the `socket(7)` manual page.

so-sndbuf()

Type:	number
Default:	0

Description: Specifies the size of the socket send buffer in bytes. For details, see the `socket(7)` manual page.

suppress()

Type:	seconds
Default:	0 (disabled)

Description: If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the `Last message repeated n times.` message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

template()

Type:	string
Default:	A format conforming to the default logfile format.

Description: Specifies a template defining the logformat to be used in the destination. Macros are described in *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like *syslogd* or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.

template-escape()

Type:	yes or no
Default:	no

Description: Turns on escaping for the `'`, `"`, and backspace characters in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.

throttle()

Type:	number
Default:	0

Description: Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using large enough buffers as well to avoid the risk of losing messages. Specifying `0` or a lower value sets the output limit to unlimited.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

ts-format()

Type:	rfc3164, bsd, rfc3339, iso
Default:	Use the global option (which defaults to rfc3164)

Description: Override the global timestamp format (set in the global *ts-format()* parameter) for the specific destination. For details, see *Section ts-format() (p. 227)*.

7.17. Sending messages to a user terminal — usertty() destination

This driver writes messages to the terminal of a logged-in user.

The *usertty()* driver has a single required argument, specifying a username who should receive a copy of matching messages. Use the asterisk * to specify every user currently logged in to the system.

```
Declaration:
usertty(username);
```

The *usertty()* does not have any further options nor does it support templates.

Example 7.24. Using the usertty() driver

```
destination d_usertty { usertty("root"); };
```

Chapter 8. Routing messages: log paths and filters

8.1. Log paths

Log paths determine what happens with the incoming log messages. Messages coming from the sources listed in the log statement and matching all the filters are sent to the listed destinations.

To define a log path, add a log statement to the syslog-ng configuration file using the following syntax:

```
log {
    source(s1); source(s2); ...
    optional_element(filter1|parser1|rewrite1);...
    optional_element(filter2|parser2|rewrite2);...
    destination(d1); destination(d2); ...
    flags(flag1[, flag2...]);
    };
```

Warning
Log statements are processed in the order they appear in the configuration file, thus the order of log paths may influence what happens to a message, especially when using filters and log flags.

Example 8.1. A simple log statement
The following log statement sends all messages arriving to the localhost to a remote server.

```
source s_localhost { tcp(ip(127.0.0.1) port(1999) ); };
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
log { source(s_localhost); destination(d_tcp); };
```

All matching log statements are processed by default, and the messages are sent to *every* matching destination by default. So a single log message might be sent to the same destination several times, provided the destination is listed in several log statements, and it can be also sent to several different destinations.

This default behavior can be changed using the `flags()` parameter. Flags apply to individual log paths; they are not global options. The following flags available in syslog-ng:

- *final*: Do not send the messages processed by this log path to any further destination.
- *fallback*: Process messages that were not processed by other log paths.
- *catchall*: Process every message, regardless of its source or if it was already processed by other log paths.
- *flow-control*: Stop reading messages from the source if the destination cannot accept them. For details, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

For details on the individual flags, see *Section 8.1.3, Log path flags (p. 202)*. The effect and use of the `flow-control` flag is detailed in *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

8.1.1. Embedded log statements

Starting from version 3.0, syslog-ng can handle embedded log statements (also called log pipes). Embedded log statements are useful for creating complex, multi-level log paths with several destinations and use filters, parsers, and rewrite rules.

For example, if you want to filter your incoming messages based on the facility parameter, and then use further filters to send messages arriving from different hosts to different destinations, you would use embedded log statements.
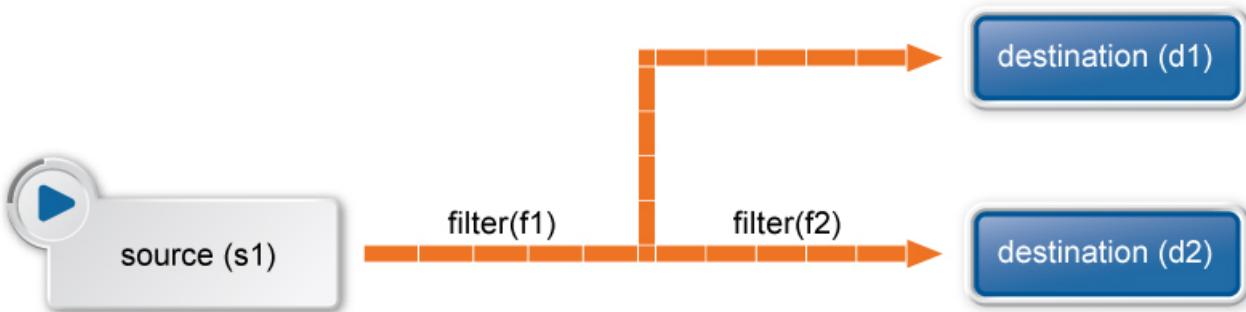
Figure 8.1. Embedded log statement

Embedded log statements include sources — and usually filters, parsers, rewrite rules, or destinations — and other log statements that can include filters, parsers, rewrite rules, and destinations. The following rules apply to embedded log statements:

- Only the beginning (also called top-level) log statement can include sources.
- Embedded log statements can include multiple log statements on the same level (that is, a top-level log statement can include two or more log statements).
- Embedded log statements can include several levels of log statements (that is, a top-level log statement can include a log statement that includes another log statement, and so on).
- After an embedded log statement, you can write either another log statement, or the `flags()` option of the original log statement. You cannot use filters or other configuration objects.
- Embedded log statements that are on the same level receive the same messages from the higher-level log statement. For example, if the top-level log statement includes a filter, the lower-level log statements receive only the messages that pass the filter.
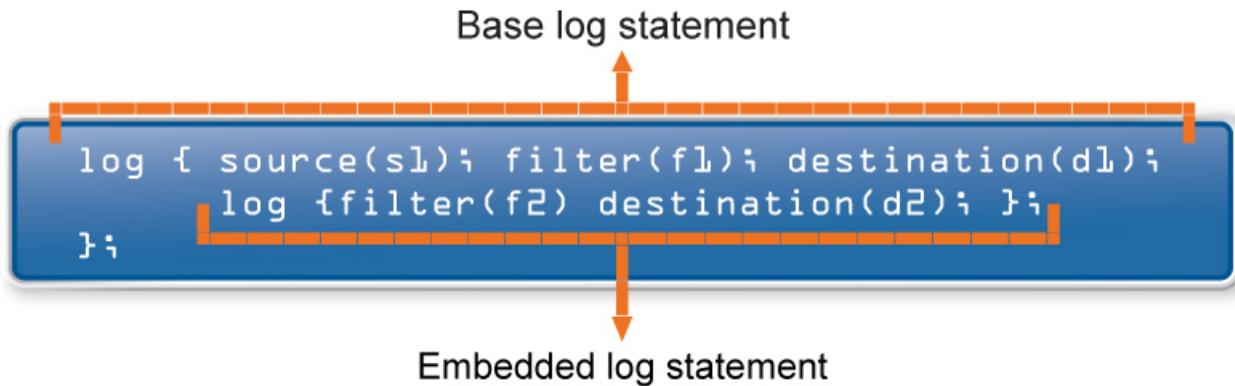
Figure 8.2. Embedded log statements

Embedded log filters can be used to optimize the processing of log messages, for example, to re-use the results of filtering and rewriting operations.

8.1.1.1. Using embedded log statements

Embedded log statements (for details, see *Section 8.1.1, Embedded log statements (p. 199)*) re-use the results of processing messages (for example the results of filtering or rewriting) to create complex log paths. Embedded log statements use the same syntax as regular log statements, but they cannot contain additional sources. To define embedded log statements, use the following syntax:

```
log {
    source(s1); source(s2); ...

    optional_element(filter1|parser1|rewrite1);
    optional_element(filter2|parser2|rewrite2);...

    destination(d1); destination(d2); ...

    #embedded log statement
    log
        {
        optional_element(filter1|parser1|rewrite1);
        optional_element(filter2|parser2|rewrite2);
        ...
        destination(d1); destination(d2); ...
        #another embedded log statement
        log
        {
            optional_element(filter1|parser1|rewrite1);
            optional_element(filter2|parser2|rewrite2);
            ...
            destination(d1); destination(d2); ...};
        };
    #set flags after the embedded log statements
    flags(flag1[, flag2...]);
        };
```

Example 8.2. Using embedded log paths
The following log path sends every message to the configured destinations: both the `d_file1` and the `d_file2` destinations receive every message of the source.

```
log { source(s_localhost); destination(d_file1); destination(d_file2); };
```

The next example is equivalent with the one above, but uses an embedded log statement.

```
log { source(s_localhost); destination(d_file1);
                      log {destination(d_file2); };
};
```

The following example uses two filters:

- messages coming from the host `192.168.1.1` are sent to the `d_file1` destination; and
- messages coming from the host `192.168.1.1` and containing the string `example` are sent to the `d_file2` destination.

```
log { source(s_localhost); host(192.168.1.); destination(d_file1);
                      log {message("example"); destination(d_file2); };
};
```

The following example collects logs from multiple source groups and uses the *source()* filter in the embedded log statement to select messages of the `s_network` source group.

```
log { source(s_localhost); source(s_network); destination(d_file1);
                      log {source(s_network); destination(d_file2); };
                      };
```

8.1.2. Junctions and channels

Junctions make it possible to send the messages to different channels, process the messages differently on each channel, and then join every channel together again. You can define any number of channels in a junction: every channel receives a copy of every message that reaches the junction. Every channel can process the messages differently, and at the end of the junction, the processed messages of every channel return to the junction again, where further processing is possible.

A junction includes one or more channels. A channel usually includes at least one filter, though that is not enforced. Otherwise, channels are identical to log statements, and can include any kind of objects, for example, parsers, rewrite rules, destinations, and so on. (For details on using channels, as well as on using channels outside junctions, see *Section 5.5, Using channels in configuration objects (p. 44)*.)

Note
Certain parsers can also act as filters:

- The JSON parser automatically discards messages that are not valid JSON messages.
- The *csv-parser()* discards invalid messages if the `flags(drop-invalid)` option is set.

You can also use log-path flags in the channels of the junction. Within the junction, a message is processed by every channel, in the order the channels appear in the configuration file. Typically if your channels have filters, you also set the `flags(final)` option for the channel. However, note that the log-path flags of the channel apply only within the junction, for example, if you set the `final` flag for a channel, then the subsequent channels of the junction will not receive the message, but this does not affect any other log path or junction of the configuration. The only exception is the `flow-control` flag: if you enable flow-control in a junction, it affects the entire log path. For details on log-path flags, see *Section 8.1.3, Log path flags (p. 202)*.

```
junction {
    channel { <other-syslog-ng-objects> <log-path-flags>}
    channel { <other-syslog-ng-objects> <log-path-flags>}
    ...
    }
```

Example 8.3. Using junctions

For example, suppose that you have a single network source that receives log messages from different devices, and some devices send messages that are not RFC-compliant (some routers are notorious for that). To solve this problem in earlier versions of syslog-ng OSE, you had to create two different network sources using different IP addresses or ports: one that received the RFC-compliant messages, and one that received the improperly formatted messages (for example, using the `flags(no-parse)` option). Using junctions this becomes much more simple: you can use a single network source to receive every message, then use a junction and two channels. The first channel processes the RFC-compliant messages, the second everything else. At the end, every message is stored in a single file. The filters used in the example can be *host()* filters (if you have a list of the IP addresses of the devices sending non-compliant messages), but that depends on your environment.

```
log {
    source s_network { syslog(ip(10.1.2.3) transport("tcp") flags(no-parse)); };
    junction {
        channel { filter(f_compliant_hosts); parser { syslog-parser(); }; };
        channel { filter(f_noncompliant_hosts); };
    };
    destination { file("/var/log/messages"); };
};
```

Since every channel receives every message that reaches the junction, use the `flags(final)` option in the channels to avoid the unnecessary processing the messages multiple times:

```
log {
    source s_network { syslog(ip(10.1.2.3) transport("tcp") flags(no-parse)); };
    junction {
        channel { filter(f_compliant_hosts); parser { syslog-parser(); }; flags(final);
};
        channel { filter(f_noncompliant_hosts); flags(final); };
    };
    destination { file("/var/log/messages"); };
};
```

Note

Junctions differ from embedded log statements, because embedded log statements are like branches: they split the flow of messages into separate paths, and the different paths do not meet again. Messages processed on different embedded log statements cannot be combined together for further processing. However, junctions split the messages to channels, then combine the channels together.

8.1.3. Log path flags

Flags influence the behavior of syslog-ng, and the way it processes messages. The following flags may be used in the log paths, as described in *Section 8.1, Log paths (p. 198)*.

Flag	Description
catchall	This flag means that the source of the message is ignored, only the filters are taken into account when matching messages. A log statement using the *catchall* flag processes every message that arrives to any of the defined sources.
fallback	This flag makes a log statement 'fallback'. Fallback log statements process messages that were not processed by other, 'non-fallback' log statements.

Flag	Description
final	This flag means that the processing of log messages processed by the log statement ends here, other log statements appearing later in the configuration file will not process the messages processed by the log statement labeled as 'final'. Note that this does not necessarily mean that matching messages will be stored only once, as there can be matching log statements processed prior the current one.
flow-control	Enables flow-control to the log path, meaning that syslog-ng will stop reading messages from the sources of this log statement if the destinations are not able to process the messages at the required speed. If disabled, syslog-ng will drop messages if the destination queues are full. If enabled, syslog-ng will only drop messages if the destination queues/window sizes are improperly sized. For details, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

Table 8.1. Log statement flags

Warning

The *final*, *fallback*, and *catchall* flags apply only for the top-level log paths, they have no effect on embedded log paths.

Example 8.4. Using log path flags

Let's suppose that you have two hosts (myhost_A and myhost_B) that run two applications each (application_A and application_B), and you collect the log messages to a central syslog-ng server. On the server, you create two log paths:

- one that processes only the messages sent by myhost_A; and

- one that processes only the messages sent by application_A.

This means that messages sent by application_A running on myhost_A will be processed by both log paths, and the messages of application_B running on myhost_B will not be processed at all.

- If you add the *final* flag to the first log path, then only this log path will process the messages of myhost_A, so the second log path will receive only the messages of application_A running on myhost_B.

- If you create a third log path that includes the *fallback* flag, it will process the messages not processed by the first two log paths, in this case, the messages of application_B running on myhost_B.

- Adding a fourth log path with the *catchall* flag would process every message received by the syslog-ng server.

```
log { source(s_localhost); destination(d_file); flags(catchall); };
```

8.2. Managing incoming and outgoing messages with flow-control

This section describes the internal message-processing model of syslog-ng, as well as the flow-control feature that can prevent message losses. To use flow-control, the *flow-control* flag must be enabled for the particular log path.

The syslog-ng application monitors (polls) the sources defined in its configuration file, periodically checking each source for messages. When a log message is found in one of the sources, syslog-ng polls every source and reads the available messages. These messages are processed and put into the output buffer of syslog-ng (also called fifo). From the output buffer, the operating system sends the messages to the appropriate destinations.

In large-traffic environments many messages can arrive during a single poll loop, therefore syslog-ng reads only a fixed number of messages from each source. The `log-fetch-limit()` option specifies the number of messages read during a poll loop from a single source.
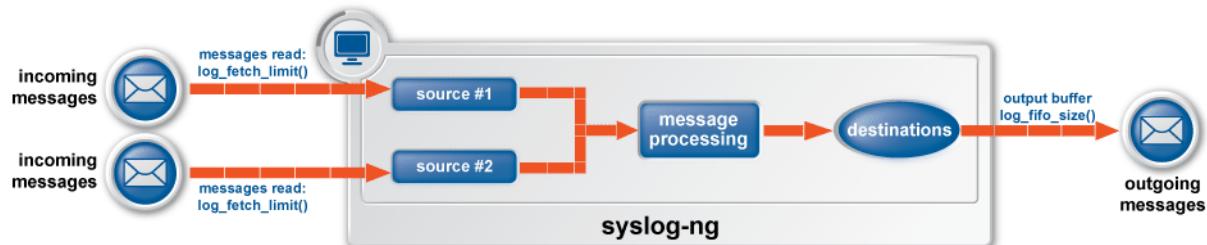
Figure 8.3. Managing log messages in syslog-ng

Every destination has its own output buffer. The output buffer is needed because the destination might not be able to accept all messages immediately. The `log-fifo-size()` parameter sets the size of the output buffer. The output buffer must be larger than the `log-fetch-limit()` of the sources, to ensure that every message read during the poll loop fits into the output buffer. If the log path sends messages to a destination from multiple sources, the output buffer must be large enough to store the incoming messages of every source.

TCP and unix-stream sources can receive the logs from several incoming connections (for example many different clients or applications). For such sources, syslog-ng reads messages from every connection, thus the `log-fetch-limit()` parameter applies individually to every connection of the source.
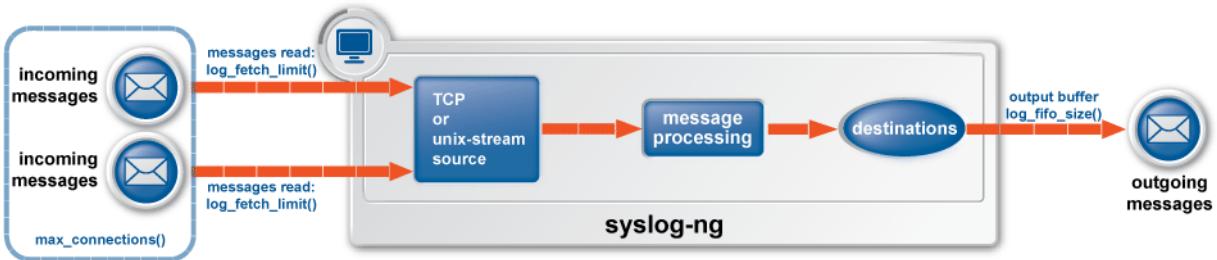
Figure 8.4. Managing log messages of TCP sources in syslog-ng

The flow-control of syslog-ng introduces a control window to the source that tracks how many messages can syslog-ng accept from the source. Every message that syslog-ng reads from the source lowers the window size by one; every message that syslog-ng successfully sends from the output buffer increases the window size by one. If the window is full (that is, its size decreases to zero), syslog-ng stops reading messages from the source. The initial size of the control window is by default 1000: the `log-fifo-size()` must be larger than this value in order for flow-control to have any effect. If a source accepts messages from multiple connections, all messages use the same control window.

> **Note**
> Starting with syslog-ng OSE version 3.3, if the source can handle multiple connections (for example, `tcp()`), the size of the control window is divided by the value of the `max-connections()` parameter and this smaller control window is applied to each connection of the source.

When flow-control is used, every source has its own control window. As a worst-case situation, the output buffer of the destination must be set to accommodate all messages of every control window, that is, the `log-fifo-size()` of the destination must be greater than number_of_sources*`log-iw-size()`. This applies to every source that sends logs to the particular destination. Thus if two sources having several connections and heavy traffic send logs to the same destination, the control window of both sources must fit into the output buffer of the destination. Otherwise, syslog-ng does not activate the flow-control, and messages may be lost.

The syslog-ng application handles outgoing messages the following way:
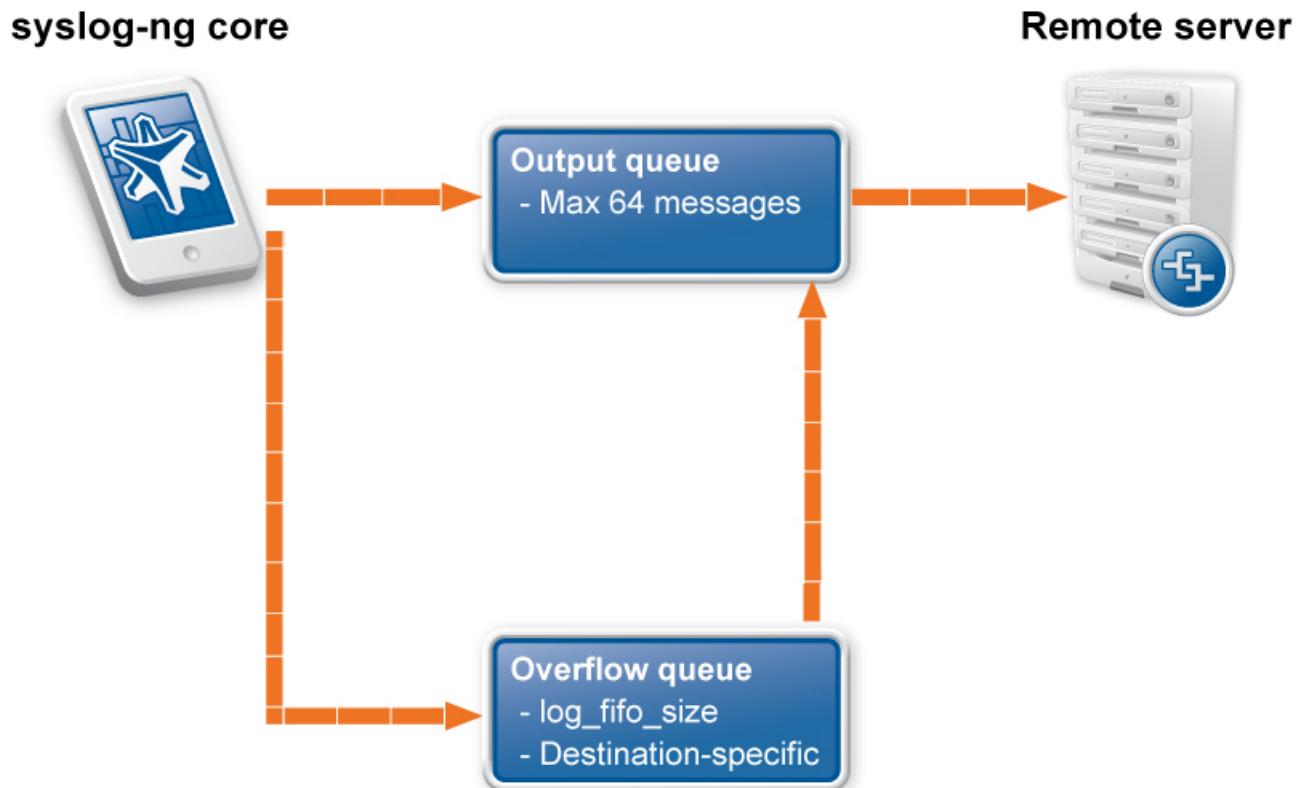
Figure 8.5. Handling outgoing messages in syslog-ng OSE

- *Output queue*: Messages from the output queue are sent to the target syslog-ng server. The syslog-ng application puts the outgoing messages directly into the output queue, unless the output queue is full. The output queue can hold 64 messages, this is a fixed value and cannot be modified.

- *Overflow queue*: If the output queue is full, syslog-ng puts the outgoing messages into the overflow queue of the destination. (The overflow queue is identical to the output buffer used by other destinations.) The `log-fifo-size()` parameter specifies the number of messages stored in the overflow queue. For details on sizing the `log-fifo-size()` parameter, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

There are two types of flow-control: Hard flow-control and soft flow-control.

- *Soft flow-control:* In case of soft flow-control there is no message lost if the destination can accept messages, but it is possible to lose messages if it cannot accept messages (for example non-writeable

file destination, or the disk becomes full), and all buffers are full. Soft flow-control cannot be configured, it is automatically available for file destinations.

Example 8.5. Soft flow-control

```
source s_file {file("/tmp/input_file.log");};
destination d_file {file("/tmp/output_file.log");};
destination d_tcp { tcp("127.0.0.1" port(2222) log-fifo-size(1000)); };
log{ source(s_file); destination(d_file); destination(d_tcp);};
```

Warning
Hazard of data loss! For destinations other than file, soft flow-control is not available. Thus, it is possible to lose log messages on those destinations. To avoid data loss on those destinations, use hard flow-control.

■ *Hard flow-control:* In case of hard flow-control there is no message lost. To use hard flow-control, enable the `flow_control` flag in the logpath. Hard flow-control is available for all destinations.

Example 8.6. Hard flow-control

```
source s_file {file("/tmp/input_file.log");};
destination d_file {file("/tmp/output_file.log");};
destination d_tcp { tcp("127.0.0.1" port(2222) log-fifo-size(1000)); };
log{ source(s_file); destination(d_file); destination(d_tcp)
flags(flow_control);};
```

8.2.1. Flow-control and multiple destinations

Using flow-control on a source has an important side-effect if the messages of the source are sent to multiple destinations. If flow-control is in use and one of the destinations cannot accept the messages, the other destinations do not receive any messages either, because syslog-ng stops reading the source. For example, if messages from a source are sent to a remote server and also stored locally in a file, and the network connection to the server becomes unavailable, neither the remote server nor the local file will receive any messages.

Note
Creating separate log paths for the destinations that use the same flow-controlled source does not avoid the problem.

8.2.2. Configuring flow-control

For details on how flow-control works, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*. The summary of the main points is as follows:

■ The syslog-ng application normally reads a maximum of `log-fetch-limit()` number of messages from a source.

- From TCP and unix-stream sources, syslog-ng reads a maximum of `log-fetch-limit()` from every connection of the source. The number of connections to the source is set using the `max-connections()` parameter.

- Every destination has an output buffer (`log-fifo-size()`).

- Flow-control uses a control window to determine if there is free space in the output buffer for new messages. Every source has its own control window; `log-iw-size()` parameter sets the size of the control window.

- When a source accepts multiple connections, the size of the control window is divided by the value of the `max-connections()` parameter and this smaller control window is applied to each connection of the source.

- The output buffer must be larger than the control window of every source that logs to the destination.

- If the control window is full, syslog-ng stops reading messages from the source until some messages are successfully sent to the destination.

- If the output buffer becomes full, and flow-control is not used, messages may be lost.

Warning

If you modify the `max-connections()` or the `log-fetch-limit()` parameter, do not forget to adjust the `log-iw-size()` and `log-fifo-size()` parameters accordingly.

Example 8.7. Sizing parameters for flow-control

Suppose that syslog-ng has a source that must accept up to 300 parallel connections. Such situation can arise when a network source receives connections from many clients, or if many applications log to the same socket. Therefore, set the `max-connections()` parameter of the source to 300. However, the `log-fetch-limit()` (default value: 10) parameter applies to every connection of the source individually, while the `log-iw-size()` (default value: 1000) parameter applies to the source. In a worst-case scenario, the destination does not accept any messages, while all 300 connections send at least `log-fetch-limit()` number of messages to the source during every poll loop. Therefore, the control window must accommodate at least `max-connections()*log-fetch-limit()` messages to be able to read every incoming message of a poll loop. In the current example this means that (`log-iw-size()` should be greater than 300*10=3000. If the control window is smaller than this value, the control window might fill up with messages from the first connections — causing syslog-ng to read only one message of the last connections in every poll loop.

The output buffer of the destination must accommodate at least `log-iw-size()` messages, but use a greater value: in the current example 3000*10=30000 messages. That way all incoming messages of ten poll loops fit in the output buffer. If the output buffer is full, syslog-ng does not read any messages from the source until some messages are successfully sent to the destination.

```
source s_localhost {
          tcp(ip(127.0.0.1) port(1999) max-connections(300)); };
destination d_tcp {
          tcp("10.1.2.3" port(1999); localport(999)); log-fifo-size(30000); };
log { source(s_localhost); destination(d_tcp); flags(flow-control); };
```

If other sources send messages to this destination, than the output buffer must be further increased. For example, if a network host with maximum 100 connections also logs into the destination, than increase the `log-fifo-size()` by 10000.

```
source s_localhost {
          tcp(ip(127.0.0.1) port(1999) max-connections(300)); };
source s_tcp {
          tcp(ip(192.168.1.5) port(1999) max-connections(100)); };
destination d_tcp {
```

```
                       tcp("10.1.2.3" port(1999); localport(999)); log-fifo-size(40000); };
log { source(s_localhost); destination(d_tcp); flags(flow-control); };
```

8.3. Filters

The following sections describe how to select and filter log messages.

- *Section 8.3.1, Using filters (p. 208)* describes how to configure and use filters.
- *Section 8.3.2, Combining filters with boolean operators (p. 208)* shows how to create complex filters using boolean operators.
- *Section 8.3.3, Comparing macro values in filters (p. 209)* explains how to evaluate macros in filters.
- *Section 8.3.4, Using wildcards, special characters, and regular expressions in filters (p. 210)* provides tips on using regular expressions.
- *Section 8.3.5, Tagging messages (p. 211)* explains how to tag messages and how to filter on the tags.
- *Section 8.3.6, Filter functions (p. 211)* is a detailed description of the filter functions available in syslog-ng OSE.

8.3.1. Using filters

Filters perform log routing within syslog-ng: a message passes the filter if the filter expression is true for the particular message. If a log statement includes filters, the messages are sent to the destinations only if they pass all filters of the log path. For example, a filter can select only the messages originating from a particular host. Complex filters can be created using filter functions and logical boolean expressions.

To define a filter, add a filter statement to the syslog-ng configuration file using the following syntax:

```
filter <identifier> { <filter_type>("<filter_expression>"); };
```

Example 8.8. A simple filter statement
The following filter statement selects the messages that contain the word deny and come from the host example.

```
filter demo_filter { host("example") and match("deny" value("MESSAGE")) };
```

For the filter to have effect, include it in a log statement:

```
log {
        source(s1);
        filter(demo_filter);
        destination(d1);};
```

8.3.2. Combining filters with boolean operators

When a log statement includes multiple filter statements, syslog-ng sends a message to the destination only if all filters are true for the message. In other words, the filters are connected with the logical AND operator. In the following example, no message arrives to the destination, because the filters are exclusive (the hostname of a client cannot be example1 and example2 at the same time):

```
filter demo_filter1 { host("example1"); };
filter demo_filter2 { host("example2"); };
```

```
log {
    source(s1); source(s2);
    filter(demo_filter1); filter(demo_filter2);
    destination(d1); destination(d2); };
```

To select the messages that come from either host `example1` or `example2`, use a single filter expression:

```
filter demo_filter { host("example1") or host("example2"); };
log {
    source(s1); source(s2);
    filter(demo_filter);
    destination(d1); destination(d2); };
```

Use the `not` operator to invert filters, for example, to select the messages that were not sent by host `example1`:

```
filter demo_filter { not host("example1"); };
```

However, to select the messages that were not sent by host `example1` or `example2`, you have to use the `and` operator (that's how boolean logic works):

```
filter demo_filter { not host("example1") and not host("example2"); };
```

Alternatively, you can use parentheses to avoid this confusion:

```
filter demo_filter { not (host("example1") or host("example2")); };
```

For a complete description on filter functions, see *Section 8.3.6, Filter functions (p. 211)*.

The following filter statement selects the messages that contain the word `deny` and come from the host `example`.

```
filter demo_filter { host("example") and match("deny" value("MESSAGE")); };
```

The *value()* parameter of the *match* function limits the scope of the function to the text part of the message (that is, the part returned by the *${MESSAGE}* macro). For details on using the *match()* filter function, see *Section match() (p. 214)*.

> **Tip**
> Filters are often used together with log path flags. For details, see *Section 8.1.3, Log path flags (p. 202)*.

8.3.3. Comparing macro values in filters

Starting with syslog-ng OSE version 3.2, it is also possible to compare macro values and templates as numerical and string values. String comparison is alphabetical: it determines if a string is alphabetically greater or equal to another string. Use the following syntax to compare macro values or templates. For details on macros and templates, see *Section 11.1, Customizing message format (p. 238)*.

```
filter <filter-id>
        {"<macro-or-template>" operator "<value-or-macro-or-template>"};
```

Example 8.9. Comparing macro values in filters

The following expression selects log messages containing a PID (that is, *${PID}* macro is not empty):

```
filter f_pid {"${PID}" !=""};
```

The following expression accomplishes the same, but uses a template as the left argument of the operator and compares the values as strings:

```
filter f_pid {"${HOST}${PID}" eq "${HOST}"};
```

The following example selects messages with priority level 4 or higher.

```
filter f_level {"${LEVEL_NUM}" > "5"};
```

Note that:

- The macro or template must be enclosed in double-quotes.

- The $ character must be used before macros.

- Using comparator operators can be equivalent to using filter functions, but is somewhat slower. For example, using "${HOST}" eq "myhost" is equivalent to using host("myhost" type(string)).

- You can use any macro in the expression, including user-defined macros from parsers and results of pattern database classifications.

- The results of filter functions are boolean values, so they cannot be compared to other values.

- You can use boolean operators to combine comparison expressions.

The following operators are available:

Numerical operator	String operator	Meaning
==	eq	Equals
!=	ne	Not equal to
>	gt	Greater than
<	lt	Less than
>=	ge	Greater than or equal
=<	le	Less than or equal

Table 8.2. Numerical and string comparison operators

8.3.4. Using wildcards, special characters, and regular expressions in filters

The *host()*, *match()*, and *program()* filter functions accept regular expressions as parameters. The exact type of the regular expression to use can be specified with the *type()* option. By default, syslog-ng OSE uses PCRE regular expressions.

In regular expressions, the asterisk (*) character means 0, 1 or any number of the previous expression. For example, in the f*ilter expression the asterisk means 0 or more f letters. This expression matches for the following strings: ilter, filter, ffilter, and so on. To achieve the wildcard functionality commonly represented by the asterisk character in other applications, use .* in your expressions, for example f.*ilter.

Alternatively, if you do not need regular expressions, only wildcards, use type(glob) in your filter:

Example 8.10. Filtering with widcards
The following filter matches on hostnames starting with the `myhost` string, for example, on `myhost-1`, `myhost-2`, and so on.

```
filter f_wildcard {host("myhost*" type(glob));};
```

For details on using regular expressions in syslog-ng OSE, see *Section 8.3.4, Using wildcards, special characters, and regular expressions in filters (p. 210)*.

To filter for special control characters like the carriage return (CR), use the `\r` escape prefix in syslog-ng OSE version 3.0 and 3.1. In syslog-ng OSE 3.2 and later, you can also use the `\x` escape prefix and the ASCII code of the character. For example, to filter on carriage returns, use the following filter:

```
filter f_carriage_return {match("\x0d" value ("MESSAGE"));};
```

8.3.5. Tagging messages

Starting with syslog-ng 3.1, it is also possible to label the messages with custom tags. Tags are simple labels, identified by their names, which must be unique. Currently syslog-ng can tag a message at two different places:

- at the source when the message is received; and
- when the message matches a pattern in the pattern database. For details on using the pattern database, see *Section 13.2, Using pattern databases (p. 278)*, for details on creating tags in the pattern database, see *Section 13.5.3, The syslog-ng pattern database format (p. 289)*.
- Tags can be also added and deleted using rewrite rules. For details, see *Section 11.2.6, Adding and deleting tags (p. 262)*.

When syslog-ng receives a message, it automatically adds the `.source.<id_of_the_source_statement>` tag to the message. Use the *tags()* option of the source to add custom tags, and the *tags()* option of the filters to select only specific messages.

Note

- Tagging messages and also filtering on the tags is very fast, much faster than other types of filters.
- Tags are available locally, that is, if you add tags to a message on the client, these tags will not be available on the server.
- To include the tags in the message, use the *${TAGS}* macro in a template. Alternatively, if you are using the IETF-syslog message format, you can include the *${TAGS}* macro in the *.SDATA.meta* part of the message. Note that the *${TAGS}* macro is available only in syslog-ng OSE 3.1.1 and later.

For an example on tagging, see *Example 8.12, Adding tags and filtering messages with tags (p. 215)*.

8.3.6. Filter functions

The following functions may be used in the filter statement, as described in *Section 8.3, Filters (p. 208)*.

Name	Description
facility()	Filter messages based on the sending facility.

Name	Description
filter()	Call another filter function.
host()	Filter messages based on the sending host.
inlist()	File-based whitelisting and blacklisting.
level() or priority()	Filter messages based on their priority.
match()	Use a regular expression to filter messages based on a specified header or content field.
message()	Use a regular expression to filter messages based on their content.
netmask()	Filter messages based on the IP address of the sending host.
program()	Filter messages based on the sending application.
source()	Select messages of the specified syslog-ng OSE source statement.
tags()	Select messages having the specified tag.

Table 8.3. Filter functions available in syslog-ng OSE

facility()

Synopsis:	facility(<facility-name>) or facility(<facility-code>) or facility(<facility-name>..<facility-name>)

Description: Match messages having one of the listed facility codes.

The `facility()` filter accepts both the name and the numerical code of the facility or the importance level. Facility codes 0-23 are predefined and can be referenced by their usual name. Facility codes above 24 are not defined.

You can use the facility filter the following ways:

- Use a single facility name, for example, `facility(user)`
- Use a single facility code, for example, `facility(1)`
- Use a facility range (works only with facility names), for example, `facility(local0..local5)`

The syslog-ng application recognizes the following facilities: (Note that some of these facilities are available only on specific platforms.)

Numerical Code	Facility name	Facility
0	kern	kernel messages
1	user	user-level messages
2	mail	mail system
3	daemon	system daemons
4	auth	security/authorization messages

Numerical Code	Facility name	Facility
5	syslog	messages generated internally by syslogd
6	lpr	line printer subsystem
7	news	network news subsystem
8	uucp	UUCP subsystem
9	cron	clock daemon
10	authpriv	security/authorization messages
11	ftp	FTP daemon
12	ntp	NTP subsystem
13	security	log audit
14	console	log alert
15	solaris-cron	clock daemon
16-23	local0..local7	locally used facilities (local0-local7)

Table 8.4. syslog Message Facilities recognized by the facility() filter

filter()

Synopsis:	filter(filtername)

Description: Call another filter rule and evaluate its value.

host()

Synopsis:	host(regexp)

Description: Match messages by using a regular expression against the hostname field of log messages.

inlist()

Synopsis:	in-list("</path/to/file.list>", value("<field-to-filter>"));

Description: Matches the value of the specified field to a list stored in a file, allowing you to do simple, file-based black- and whitelisting. The file must be a plain-text file, containing one entry per line. The syslog-ng OSE application loads the entire file, and compares the value of the specified field (for example, ${PROGRAM}) to entries in the file. When you use the *in-list filter*, note the following points:

- Comparing the values is case-sensitive.
- Only exact matches are supported, partial and substring matches are not.
- If you modify the list file, reload the configuration of syslog-ng OSE for the changes to take effect.

Available in syslog-ng OSE 3.5 and later.

Example 8.11. Selecting messages using the in-list filter
Create a text file that contains the programs (as in the ${PROGRAM} field of their log messages) you want to select. For example, you want to forward only the logs of a few applications from a host: kernel, sshd, and sudo. Create the `/etc/syslog-ng/programlist.list` file with the following contents:

```
kernel
sshd
sudo
```

The following filter selects only the messages of the listed applications:

```
filter f_whitelist { in-list("/etc/syslog-ng/programlist.list", value("PROGRAM")); };
```

Create the appropriate sources and destinations for your environment, then create a log path that uses the previous filter to select only the log messages of the applications you need:

```
log {
source(s_all);
filter(f_whitelist);
destination(d_logserver);};
```

To create a blacklist filter, simply negate the `in-list` filter:

```
filter f_blacklist { not in-list("/etc/syslog-ng/programlist.list", value("PROGRAM")); };
```

level() or priority()

Synopsis:	level(<priority-level>) or level(<priority-level>..<priority-level>)

Description: The `level()` filter selects messages corresponding to a single importance level, or a level-range. To select messages of a specific level, use the name of the level as a filter parameter, for example use the following to select warning messages:

```
level(warning)
```

To select a range of levels, include the beginning and the ending level in the filter, separated with two dots (`..`). For example, to select every message of error or higher level, use the following filter:

```
level(err..emerg)
```

The `level()` filter accepts the following levels: *emerg*, *alert*, *crit*, *err*, *warning*, *notice*, *info*, *debug*.

match()

Synopsis:	match(regexp)

Description: Match a regular expression to the headers and the message itself (that is, the values returned by the *MSGHDR* and *MSG* macros). Note that in syslog-ng version 2.1 and earlier, the `match()` filter was applied only to the text of the message, excluding the headers. This functionality has been moved to the `message()` filter.

To limit the scope of the match to a specific part of the message (identified with a macro), use the `match(regexp value("MACRO"))` syntax. Do not include the $ sign in the parameter of the `value()` option.

The `value()` parameter accepts both built-in macros and user-defined ones created with a parser or using a pattern database. For details on macros and parsers, see *Section 11.1.2, Templates and macros (p. 239)*, *Section 12.2, Parsing messages (p. 268)*, and *Section 13.2.1, Using parser results in filters and templates (p. 279)*.

message()

Synopsis:	message(regexp)

Description: Match a regular expression to the text of the log message, excluding the headers (that is, the value returned by the *MSG* macros). Note that in syslog-ng version 2.1 and earlier, this functionality was performed by the *match()* filter.

netmask()

Synopsis:	netmask(ip/mask)

Description: Select only messages sent by a host whose IP address belongs to the specified IP subnet. Note that this filter checks the IP address of the last-hop relay (the host that actually sent the message to syslog-ng), not the contents of the *HOST* field of the message. You can use both the dot-decimal and the CIDR notation to specify the netmask. For example, `192.168.5.0/255.255.255.0` or `192.168.5.0/24`.

program()

Synopsis:	program(regexp)

Description: Match messages by using a regular expression against the program name field of log messages.

source()

Synopsis:	source id

Description: Select messages of a source statement. This filter can be used in embedded log statements if the parent statement contains multiple source groups — only messages originating from the selected source group are sent to the destination of the embedded log statement.

tags()

Synopsis:	tag

Description: Select messages labeled with the specified tag. Every message automatically has the tag of its source in `.source.<id_of_the_source_statement>` format. This option is available only in syslog-ng 3.1 and later.

Example 8.12. Adding tags and filtering messages with tags

```
source s_tcp {
    tcp(ip(192.168.1.1) port(1514) tags("tcp", "router"));
    };
```

Use the *tags()* option of the filters to select only specific messages:

```
filter f_tcp {
    tags(".source.s_tcp");
    };

    filter f_router {
    tags("router");
    };
```

Note
Starting with version 3.2, syslog-ng OSE automatically adds the class of the message as a tag using the *.classifier.<message-class>* format. For example, messages classified as "system" receive the *.classifier.system* tag. Use the *tags()* filter function to select messages of a specific class.

```
filter f_tag_filter {tags(".classifier.system");};
```

8.4. Dropping messages

To skip the processing of a message without sending it to a destination, create a log statement with the appropriate filters, but do not include any destination in the statement, and use the *final* flag.

Example 8.13. Skipping messages
The following log statement drops all *debug* level messages without any further processing.

```
filter demo_debugfilter { level(debug); };
log { source(s_all); filter(demo_debugfilter); flags(final); };
```

Chapter 9. Global options of syslog-ng OSE

9.1. Configuring global syslog-ng options

The syslog-ng application has a number of global options governing DNS usage, the timestamp format used, and other general points. Each option may have parameters, similarly to driver specifications. To set global options, add an option statement to the syslog-ng configuration file using the following syntax:

```
options { option1(params); option2(params); ... };
```

Example 9.1. Using global options
To disable domain name resolving, add the following line to the syslog-ng configuration file:

```
options { use-dns(no); };
```

For a detailed list of the available options, see *Section 9.2, Global options (p. 217)*. For important global options and recommendations on their use, see *Chapter 17, Best practices and examples (p. 306)*.

9.2. Global options

The following options can be specified in the options statement, as described in *Section 9.1, Configuring global syslog-ng options (p. 217)*.

bad-hostname()

Accepted values:	regular expression
Default:	no

Description: A regexp containing hostnames which should not be handled as hostnames.

chain-hostnames()

Accepted values:	yes | no
Default:	no

Description: Enable or disable the chained hostname format. If the log message is forwarded to the log server via a relay, and the *chain-hostnames()* option is enabled, the relay adds its own hostname to the hostname of the client, separated with a */* character. For example, consider a client-relay-server scenario with the following hostnames: *client-hostname*, *relay-hostname*, *server-hostname*. The hostname of the log message received by the server will look like: `client-hostname/relay-hostname`. If the client sends a hostname in the message that is different from its real hostname (as resolved from DNS), the relay can add the resolved hostname to the message, resulting in two different client hostnames in the message, for examle, `client-hostname-from-the-message/client-hostname-resolved-on-the-relay/relay-hostname`.

check-hostname()

Accepted values:	yes | no
Default:	no

Description: Enable or disable checking whether the hostname contains valid characters.

create-dirs()

Accepted values:	yes | no
Default:	no

Description: Enable or disable directory creation for destination files.

custom-domain()

Accepted values:	string
Default:	empty string

Description: Use this option to specify a custom domain name that is appended after the short hostname to receive the FQDN. This option affects every outgoing message: eventlog sources, file sources, MARK messages and internal messages of syslog-ng OSE.

- If the hostname is a short hostname, the custom domain name is appended after the hostname (for example mypc becomes mypc.customcompany.local).

- If the hostname is an FQDN, the domain name part is replaced with the custom domain name (for example if the FQDN in the forwarded message is mypc.mycompany.local and the custom domain name is customcompany.local, the hostname in the outgoing message becomes mypc.customcompany.local).

dir-group()

Accepted values:	groupid
Default:	root

Description: The default group for newly created directories.

dir-owner()

Accepted values:	userid
Default:	root

Description: The default owner of newly created directories.

dir-perm()

Accepted values:	permission value
Default:	0700

Description: The permission mask of directories created by syslog-ng. Log directories are only created if a file after macro expansion refers to a non-existing directory, and directory creation is enabled (see also the `create-dirs()` option). For octal numbers prefix the number with 0, for example use `0755` for `rwxr-xr-x`.

To preserve the original properties of an existing directory, use the option without specifying an attribute: `dir-perm()`. Note that when creating a new directory without specifying attributes for `dir-perm()`, the default permission of the directories is masked with the umask of the parent process (typically `0022`).

dns-cache()

Accepted values:	`yes` | `no`
Default:	`yes`

Description: Enable or disable DNS cache usage.

dns-cache-expire()

Accepted values:	number
Default:	3600

Description: Number of seconds while a successful lookup is cached.

dns-cache-expire-failed()

Accepted values:	number
Default:	60

Description: Number of seconds while a failed lookup is cached.

dns-cache-hosts()

Accepted values:	filename
Default:	unset

Description: Name of a file in `/etc/hosts` format that contains static IP->hostname mappings. Use this option to resolve hostnames locally without using a DNS. Note that any change to this file triggers a reload in syslog-ng and is instantaneous.

dns-cache-size()

Accepted values:	number
Default:	1007

Description: Number of hostnames in the DNS cache.

file-template()

Accepted values:	time offset (for example: `+03:00`)
Default:	local timezone

Description: Specifies a template that file-like destinations use by default. For example:

```
template t_isostamp { template("$ISODATE $HOST $MSGHDR$MSG\n"); };

options { file-template(t_isostamp); };
```

flush-lines()

Accepted values:	number
Default:	100

Description: Specifies how many lines are flushed to a destination at a time. The syslog-ng OSE application waits for this number of lines to accumulate and sends them off in a single batch. Increasing this number increases throughput as more messages are sent in a single batch, but also increases message latency. To limit the latency, adjust the *flush-timeout()* option.

By default, syslog-ng OSE waits for 100 lines. If the source sends a few messages, but less than 100 messages, syslog-ng OSE will not send the messages to the destination until the *flush-timeout()* time passes (which is 10 seconds by default). If you stop or reload syslog-ng OSE or in case of network sources, the connection with the client is closed, syslog-ng OSE automatically sends the unsent messages to the destination.

flush-timeout()

Accepted values:	time in milliseconds
Default:	10000

Description: Specifies the time syslog-ng waits for lines to accumulate in its output buffer. For more information, see the *flush-lines()* option.

frac-digits()

Type:	number
Default:	Value of the global option (which defaults to 0)

Description: The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The *frac-digits()* parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

group()

Accepted values:	groupid
Default:	root

Description: The default group of output files. By default, syslog-ng changes the privileges of accessed files (for example /dev/null) to root.root 0600. To disable modifying privileges, use this option with the -1 value.

keep-hostname()

Type:	yes or no
Default:	no

Description: Enable or disable hostname rewriting.

- If enabled (`keep-hostname(yes)`), syslog-ng OSE assumes that the incoming log message was sent by the host specified in the *HOST* field of the message.

- If disabled (`keep-hostname(no)`), syslog-ng OSE rewrites the *HOST* field of the message, either to the IP address (if the *use-dns()* parameter is set to no), or to the hostname (if the *use-dns()* parameter is set to yes and the IP address can be resolved to a hostname) of the host sending the message to syslog-ng OSE. For details on using name resolution in syslog-ng OSE, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

> **Note**
> If the log message does not contain a hostname in its *HOST* field, syslog-ng OSE automatically adds a hostname to the message.
> - For messages received from the network, this hostname is the address of the host that sent the message (this means the address of the last hop if the message was transferred via a relay).
> - For messages received from the local host, syslog-ng OSE adds the name of the host.

This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

> **Note**
> When relaying messages, enable this option on the syslog-ng OSE server and also on every relay, otherwise syslog-ng OSE will treat incoming messages as if they were sent by the last relay.

keep-timestamp()

Type:	yes or no
Default:	yes

Description: Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

log-fifo-size()

Accepted values:	number
Default:	10000

Description: The number of messages that the output queue can store.

log-msg-size()

Accepted values:	number
Default:	8192

Description: Maximum length of a message in bytes. This length includes the entire message (the data structure and individual fields). The maximal value that can be set is 268435456 bytes (256MB). For messages using the IETF-syslog message format (RFC5424), the maximal size of the value of an SDATA field is 64kB.

mark() (DEPRECATED)

Accepted values:	number
Default:	1200

Description: The *mark-freq()* option is an alias for the deprecated *mark()* option. This is retained for compatibility with syslog-ng version 1.6.x.

mark-freq()

Accepted values:	number
Default:	1200

Description: An alias for the obsolete *mark()* option, retained for compatibility with syslog-ng version 1.6.x. The number of seconds between two *MARK* messages. *MARK* messages are generated when there was no message traffic to inform the receiver that the connection is still alive. If set to zero (0), no *MARK* messages are sent. The *mark-freq()* can be set for global option and/or every *MARK* capable destination driver if *mark-mode()* is periodical or dst-idle or host-idle. If *mark-freq()* is not defined in the destination, then the *mark-freq()* will be inherited from the global options. If the destination uses internal *mark-mode()*, then the global *mark-freq()* will be valid (does not matter what *mark-freq()* set in the destination side).

mark-mode()

Accepted values:	`internal`|`dst-idle`|`host-idle`|`periodical`|`none`|`global`
Default:	`internal` for pipe, program drivers
	`none` for file, unix-dgram, unix-stream drivers
	`global` for syslog, tcp, udp destinations
	`host-idle` for global option

Description: The *mark-mode()* option can be set for the following destination drivers: file(), program(), unix-dgram(), unix-stream(), udp(), udp6(), tcp(), tcp6(), pipe(), syslog() and in global option.

- ■ `internal`: When internal mark mode is selected, internal source should be placed in the log path as this mode does not generate mark by itself at the destination. This mode only yields the mark messages from internal source. This is the mode as syslog-ng OSE 3.3 worked. *MARK* will be generated by internal source if there was NO traffic on local sources:

file(), pipe(), unix-stream(), unix-dgram(), program()

- **dst-idle**: Sends *MARK* signal if there was NO traffic on destination drivers. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()*.

- **host-idle**: Sends *MARK* signal if there was NO local message on destination drivers. For example *MARK* is generated even if messages were received from tcp. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()*.

- **periodical**: Sends *MARK* signal perodically, regardless of traffic on destination driver. *MARK* signal from internal source will be dropped.

  MARK signal can be sent by the following destination drivers: *tcp(), udp(), syslog(), program(), file(), pipe(), unix-stream(), unix-dgram()*.

- **none**: Destination driver drops all *MARK* messages. If an explicit mark-mode() is not given to the drivers where none is the default value, then none will be used.

- **global**: Destination driver uses the global *mark-mode()* setting. The syslog-ng interprets syntax error if the global *mark-mode()* is global.

Note
In case of dst-idle, host-idle and periodical; *MARK* message will not be written in the destination, if it is not open yet.

Available in syslog-ng OSE 3.4 and later.

normalize-hostnames()

Accepted values:	yes | no
Default:	no

Description: If enabled (normalize-hostnames(yes)), syslog-ng OSE converts the hostnames to lowercase.

Note
This setting applies only to hostnames resolved from DNS. It has no effect if the *keep-hostname()* option is enabled, and the message contains a hostname.

on-error()

Accepted values:	drop-message|drop-property|fallback-to-string|silently-drop-message|silently-drop-property|silently-fallback-to-string
Default:	drop-message

Description: Controls what happens when type-casting fails and syslog-ng OSE cannot convert some data to the specified type. By default, syslog-ng OSE drops the entire message and logs the error. Currently the `value-pairs()` option uses the settings of `on-error()`.

- *`drop-message`*: Drop the entire message and log an error message to the `internal()` source. This is the default behavior of syslog-ng OSE.

- *`drop-property`*: Omit the affected property (macro, template, or message-field) from the log message and log an error message to the `internal()` source.

- *`fallback-to-string`*: Convert the property to string and log an error message to the `internal()` source.

- *`silently-drop-message`*: Drop the entire message silently, without logging the error.

- *`silently-drop-property`*: Omit the affected property (macro, template, or message-field) silently, without logging the error.

- *`silently-fallback-to-string`*: Convert the property to string silently, without logging the error.

owner()

Accepted values:	userid
Default:	root

Description: The default owner of output files. By default, syslog-ng changes the privileges of accessed files (for example `/dev/null`) to `root.root 0600`. To disable modifying privileges, use this option with the `-1` value.

perm()

Accepted values:	permission value
Default:	0600

Description: The default permission for output files. By default, syslog-ng changes the privileges of accessed files (for example `/dev/null`) to `root.root 0600`. To disable modifying privileges, use this option with the `-1` value.

proto-template()

Accepted values:	name of a template
Default:	The default message format of the used protocol

Description: Specifies a template that protocol-like destinations (for example, tcp() and syslog()) use by default. For example:

```
template t_isostamp { template("$ISODATE $HOST $MSGHDR$MSG\n"); };

options { proto-template(t_isostamp); };
```

recv-time-zone()

Accepted values:	time offset (for example: +03:00)
Default:	local timezone

Description: Specifies the time zone associated with the incoming messages, if not specified otherwise in the message or in the source driver. For details, see also *Section 2.5, Timezones and daylight saving (p. 8)* and *Section 2.5.2, A note on timezones and timestamps (p. 10)*.

send-time-zone()

Accepted values:	time offset (for example: +03:00)
Default:	local timezone

Description: Specifies the time zone associated with the messages sent by syslog-ng, if not specified otherwise in the message or in the destination driver. For details, see *Section 2.5, Timezones and daylight saving (p. 8)*.

stats-freq()

Accepted values:	number
Default:	600

Description: The period between two STATS messages in seconds. STATS are log messages sent by syslog-ng, containing statistics about dropped log messages. Set to 0 to disable the STATS messages.

stats-level()

Accepted values:	0 | 1 | 2 | 3
Default:	0

Description: Specifies the detail of statistics syslog-ng collects about the processed messages.

- Level 0 collects only statistics about the sources and destinations
- Level 1 contains details about the different connections and log files, but has a slight memory overhead
- Level 2 contains detailed statistics based on the hostname.
- Level 3 contains detailed statistics based on various message parameters like facility, severity, or tags.

Note that level 2 and 3 increase the memory requirements and CPU load. For details on message statistics, see *Chapter 14, Statistics of syslog-ng (p. 297)*.

stats-lifetime()

Accepted values:	number (minutes)
Default:	10

Description: Controls how often dynamic counters are expired. The timer is not exact, some timers may live a little bit longer than the specified time. Dynamic counters are now cleaned up repeatedly in time intervals specified in `stats-lifetime()` instead of only on reloads. This intends to reduce the memory used by dynamic counters.

sync() or sync-freq() (DEPRECATED)

Accepted values:	number
Default:	0

Description: Obsolete aliases for `flush-lines()`

threaded()

Accepted values:	yes|no
Default:	yes

Description: Enable syslog-ng OSE to run in multithreaded mode and use multiple CPUs. Available only in syslog-ng Open Source Edition 3.3 and later. See *Chapter 15, Multithreading and scaling in syslog-ng OSE (p. 300)* for details.

time-reap()

Accepted values:	number
Default:	60

Description: The time to wait in seconds before an idle destination file is closed. Note that only destination files having macros in their filenames are closed automatically.

time-reopen()

Accepted values:	number
Default:	60

Description: The time to wait in seconds before a dead connection is reestablished.

time-sleep() (DEPRECATED)

Accepted values:	number
Default:	0

Description: The time to wait in milliseconds between each invocation of the `poll()` iteration.

time-zone()

Type:	timezone in +/-HH:MM format
Default:	unspecified

Description: Convert timestamps (and all date-related macros derived from the timestamp) to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.

ts-format()

Accepted values:	`rfc3164` | `bsd` | `rfc3339` | `iso`
Default:	`rfc3164`

Description: Specifies the timestamp format used when syslog-ng itself formats a timestamp and nothing else specifies a format (for example: *STAMP* macros, internal messages, messages without original timestamps). For details, see also *Section 2.5.2, A note on timezones and timestamps (p. 10)*.

By default, timestamps include only seconds. To include fractions of a second (for example, milliseconds) use the `frac-digits()` option. For details, see *Section frac-digits() (p. 220)*.

> **Note**
> This option applies only to file and file-like destinations. Destinations that use specific protocols (for example, `tcp()`, or `syslog()`) ignore this option. For protocol-like destinations, use a template locally in the destination, or use the `proto-template` option.

use-dns()

Type:	yes, no, persist_only
Default:	yes

Description: Enable or disable DNS usage. The `persist_only` option attempts to resolve hostnames locally from file (for example from `/etc/hosts`). The syslog-ng OSE application blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

use-fqdn()

Type:	yes or no
Default:	no

Description: Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

use-rcptid()

Accepted values:	yes | no
Default:	no

Description: When the *use-rcptid* global option is set to yes, syslog-ng OSE automatically assigns a unique reception ID to every received message. You can access this ID and use it in templates via the *${RCPTID}* macro. The reception ID is a monotonously increasing 48-bit integer number, that can never be zero (if the counter overflows, it restarts with 1).

Chapter 10. TLS-encrypted message transfer

10.1. Secure logging using TLS

The syslog-ng application can send and receive log messages securely over the network using the Transport Layer Security (TLS) protocol. TLS is an encryption protocol over the TCP/IP network protocol, so it can be used only with TCP-based sources and destinations ($tcp()$ and $tcp6()$).

TLS uses certificates to authenticate and encrypt the communication, as illustrated on the following figure:

Figure 10.1. Certificate-based authentication

The client authenticates the server by requesting its certificate and public key. Optionally, the server can also request a certificate from the client, thus mutual authentication is also possible.

In order to use TLS encryption in syslog-ng, the following elements are required:

- A certificate on the syslog-ng server that identifies the syslog-ng server.
- The certificate of the Certificate Authority that issued the certificate of the syslog-ng server (or the self-signed certificate of the syslog-ng server) must be available on the syslog-ng client.

When using mutual authentication to verify the identity of the clients, the following elements are required:

- A certificate must be available on the syslog-ng client. This certificate identifies the syslog-ng client.
- The certificate of the Certificate Authority that issued the certificate of the syslog-ng client must be available on the syslog-ng server.

Mutual authentication ensures that the syslog-ng server accepts log messages only from authorized clients.

For details on configuring TLS communication in syslog-ng, see *Section 10.2, Encrypting log messages with TLS (p. 230)*.

10.2. Encrypting log messages with TLS

This section describes how to configure TLS encryption in syslog-ng. For the concepts of using TLS in syslog-ng, see *Section 10.1, Secure logging using TLS (p. 229)*.

Create an X.509 certificate for the syslog-ng server.

> **Note**
> The *subject_alt_name* parameter (or the *Common Name* parameter if the *subject_alt_name* parameter is empty) of the server's certificate must contain the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server (for example `syslog-ng.example.com`).
>
> Alternatively, the *Common Name* or the *subject_alt_name* parameter can contain a generic hostname, for example `*.example.com`.
>
> Note that if the *Common Name* of the certificate contains a generic hostname, do not specify a specific hostname or an IP address in the *subject_alt_name* parameter.

10.2.1. Procedure – Configuring TLS on the syslog-ng clients

Purpose:

Complete the following steps on every syslog-ng client host. Examples are provided using both the legacy BSD-syslog protocol (using the *tcp()* driver) and the new IETF-syslog protocol standard (using the *syslog()* driver):

Steps:

Step 1. Copy the CA certificate (for example `cacert.pem`) of the Certificate Authority that issued the certificate of the syslog-ng server (or the self-signed certificate of the syslog-ng server) to the syslog-ng client hosts, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.
Issue the following command on the certificate: `openssl x509 -noout -hash -in cacert.pem`
The result is a hash (for example `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

`ln -s cacert.pem 6d2962a8.0`

Step 2. Add a destination statement to the syslog-ng configuration file that uses the `tls( ca-dir(path_to_ca_directory) )` option and specify the directory using the CA certificate. The destination must use the *tcp()* or *tcpv6()* destination driver, and the IP address and port parameters of the driver must point to the syslog-ng server.

> **Example 10.1. A destination statement using TLS**
> The following destination encrypts the log messages using TLS and sends them to the `6514/TCP` port of the syslog-ng server having the `10.1.2.3` IP address.
>
> ```
> destination demo_tls_destination {
>             tcp("10.1.2.3" port(6514)
>                 tls( ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); };
> ```

A similar statement using the IETF-syslog protocol and thus the *syslog()* driver:

```
destination demo_tls_syslog_destination { syslog("10.1.2.3" port(6514)
                            transport("tls")
                            port(3214)
                            tls(ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) );
};
```

Step 3. Include the destination created in Step 2 in a log statement.

> **Warning**
>
> The encrypted connection between the server and the client fails if the *Common Name* or the *subject_alt_name* parameter of the server certificate does not contain the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server.
>
> Do not forget to update the certificate files when they expire.

10.2.2. Procedure – Configuring TLS on the syslog-ng server

Purpose:

Complete the following steps on the syslog-ng server:

Steps:

Step 1. Copy the certificate (for example `syslog-ng.cert`) of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/cert.d` directory. The certificate must be a valid X.509 certificate in PEM format.

Step 2. Copy the private key (for example `syslog-ng.key`) matching the certificate of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/key.d` directory. The key must be in PEM format, and must not be password-protected.

Step 3. Add a source statement to the syslog-ng configuration file that uses the `tls( key-file(key_file_fullpathname) cert-file(cert_file_fullpathname) )` option and specify the key and certificate files. The source must use the source driver (*tcp()* or *tcpv6()*) matching the destination driver used by the syslog-ng client.

> **Example 10.2. A source statement using TLS**
> The following source receives log messages encrypted using TLS, arriving to the `1999/TCP` port of any interface of the syslog-ng server.
>
> ```
> source demo_tls_source {
>         tcp(ip(0.0.0.0) port(1999)
>             tls( key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
>
> cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")) ); };
> ```
>
> A similar source for receiving messages using the IETF-syslog protocol:
>
> ```
> source demo_tls_syslog_source {
>                     syslog(ip(0.0.0.0) port(1999)
>                     transport("tls")
>                     tls(
> key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
> ```

```
cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")) ); };
```

Step 4. Disable mutual authentication for the source by setting the following TLS option in the source statement:

```
tls( peer-verify(optional-untrusted);
```

For details on how to configure mutual authentication, see *Section 10.3, Mutual authentication using TLS (p. 232)*.

For the details of the available *tls()* options, see *Section 10.4, TLS options (p. 235)*.

Example 10.3. Disabling mutual authentication
The following source receives log messages encrypted using TLS, arriving to the 1999/TCP port of any interface of the syslog-ng server. The identity of the syslog-ng client is not verified.

```
source demo_tls_source {
        tcp(ip(0.0.0.0) port(1999)
            tls( key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")

                cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")

                  peer-verify(optional-untrusted))  ); };
```

A similar source for receiving messages using the IETF-syslog protocol:

```
source demo_tls_syslog_source {
                        syslog(ip(0.0.0.0) port(1999)
                        transport("tls")
                        tls(
key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")

cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
                        peer-verify(optional-untrusted))  ); };
```

Warning
Do not forget to update the certificate and key files when they expire.

10.3. Mutual authentication using TLS

This section describes how to configure mutual authentication between the syslog-ng server and the client. Configuring mutual authentication is similar to configuring TLS (for details, see *Section 10.2, Encrypting log messages with TLS (p. 230)*), but the server verifies the identity of the client as well. Therefore, each client must have a certificate, and the server must have the certificate of the CA that issued the certificate of the clients. For the concepts of using TLS in syslog-ng, see *Section 10.1, Secure logging using TLS (p. 229)*.

10.3.1. Procedure – Configuring TLS on the syslog-ng clients

Purpose:

Complete the following steps on every syslog-ng client host. Examples are provided using both the legacy BSD-syslog protocol (using the *tcp()* driver) and the new IETF-syslog protocol standard (using the *syslog()* driver):

Steps:

Step 1. Create an X.509 certificate for the syslog-ng client.

Step 2. Copy the certificate (for example `client_cert.pem`) and the matching private key (for example `client.key`) to the syslog-ng client host, for example into the `/opt/syslog-ng/etc/syslog-ng/cert.d` directory. The certificate must be a valid X.509 certificate in PEM format and must not be password-protected.

Step 3. Copy the CA certificate of the Certificate Authority (for example `cacert.pem`) that issued the certificate of the syslog-ng server (or the self-signed certificate of the syslog-ng server) to the syslog-ng client hosts, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.
Issue the following command on the certificate: `openssl x509 -noout -hash -in cacert.pem`
The result is a hash (for example `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

```
ln -s cacert.pem 6d2962a8.0
```

Step 4. Add a destination statement to the syslog-ng configuration file that uses the `tls( ca-dir(path_to_ca_directory) )` option and specify the directory using the CA certificate. The destination must use the *tcp()* or *tcpv6()* destination driver, and the IP address and port parameters of the driver must point to the syslog-ng server. Include the client's certificate and private key in the *tls()* options.

> **Example 10.4. A destination statement using mutual authentication**
> The following destination encrypts the log messages using TLS and sends them to the `1999/TCP` port of the syslog-ng server having the `10.1.2.3` IP address. The private key and the certificate file authenticating the client is also specified.
>
> ```
> destination demo_tls_destination {
>     tcp("10.1.2.3" port(1999)
>         tls( ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")
>             key-file("/opt/syslog-ng/etc/syslog-ng/key.d/client.key")
>             cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/client_cert.pem"))
> ); };
> ```
>
> ```
> destination demo_tls_syslog_destination {
>     syslog("10.1.2.3" port(1999)
>         transport("tls")
>             tls( ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")
>                 key-file("/opt/syslog-ng/etc/syslog-ng/key.d/client.key")
>                 cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/client_cert.pem"))
>  ); };
> ```

Step 5. Include the destination created in Step 2 in a log statement.

> **Warning**
>
> The encrypted connection between the server and the client fails if the *Common Name* or the *subject_alt_name* parameter of the server certificate does not the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server.
>
> Do not forget to update the certificate files when they expire.

10.3.2. Procedure – Configuring TLS on the syslog-ng server

Purpose:

Complete the following steps on the syslog-ng server:

Steps:

Step 1.   Copy the certificate (for example `syslog-ng.cert`) of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/cert.d` directory. The certificate must be a valid X.509 certificate in PEM format.

Step 2.   Copy the CA certificate (for example `cacert.pem`) of the Certificate Authority that issued the certificate of the syslog-ng clients to the syslog-ng server, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.
Issue the following command on the certificate: `openssl x509 -noout -hash -in cacert.pem`
The result is a hash (for example `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

```
ln -s cacert.pem 6d2962a8.0
```

Step 3.   Copy the private key (for example `syslog-ng.key`) matching the certificate of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/key.d` directory. The key must be in PEM format, and must not be password-protected.

Step 4.   Add a source statement to the syslog-ng configuration file that uses the `tls( key-file(key_file_fullpathname) cert-file(cert_file_fullpathname) )` option and specify the key and certificate files. The source must use the source driver (*tcp()* or *tcpv6()*) matching the destination driver used by the syslog-ng client. Also specify the directory storing the certificate of the CA that issued the client's certificate.

For the details of the available *tls()* options, see *Section 10.4, TLS options (p. 235)*.

Example 10.5. A source statement using TLS
The following source receives log messages encrypted using TLS, arriving to the `1999/TCP` port of any interface of the syslog-ng server.

```
source demo_tls_source {
        tcp(ip(0.0.0.0) port(1999)
            tls( key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")

                cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")

                  ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); };
```

A similar source for receiving messages using the IETF-syslog protocol:

```
source demo_tls_syslog_source {
                        syslog(ip(0.0.0.0) port(1999)
                        transport("tls")
                        tls(
key-file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")

cert-file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
                        ca-dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); };
```

Warning
Do not forget to update the certificate and key files when they expire.

10.4. TLS options

The syslog-ng application is able to encrypt incoming and outgoing syslog message flows using SSL/TLS, if the TCP transport protocol (the *tcp()* or *tcp6()* sources or destination) is used.

Note
The format of the TLS connections used by syslog-ng is similar to using syslog-ng and stunnel, but the source IP information is not lost.

To encrypt connections, use the *tls()* option in the source and destination statements.

The tls() option can include the following settings:

ca-dir()

Accepted values:	Directory name
Default:	none

Description: Name of a directory, that contains a set of trusted CA certificates in PEM format. The CA certificate files have to be named after the 32-bit hash of the subject's name. This naming can be created using the c_rehash utility in openssl.

cert-file()

Accepted values:	Filename
Default:	none

Description: Name of a file, that contains an X.509 certificate (or a certificate chain) in PEM format, suitable as a TLS certificate, matching the private key. If the file contains a certificate chain, the file must begin with the certificate of the host, followed by the CA certificate that signed the certificate of the host, and any other signing CAs in order.

cipher-suite()

Accepted values:	Cipher name
Default:	aes-128-cbc

Description: Specifies the cipher, hash, and key-exchange algorithms used for the encryption. The following values are accepted: `aes-128-cbc`, `aes-128-ecb`, `aes-192-cbc`, `aes-192-ecb`, `aes-256-cbc`, `aes-256-ecb`, `base64`, `bf`, `bf-cbc`, `bf-cfb`, `bf-ecb`, `bf-ofb`, `cast`, `cast-cbc`, `cast5-cbc`, `cast5-cfb`, `cast5-ecb`, `cast5-ofb`, `des`, `des-cbc`, `des-cfb`, `des-ecb`, `des-ede`, `des-ede-cbc`, `des-ede-cfb`, `des-ede-ofb`, `des-ede3`, `des-ede3-cbc` `des-ede3-cfb` `des-ede3-ofbdes-ofb`, `des3`, `desx`, `rc2`, `rc2-40-cbc`, `rc2-64-cbc`, `rc2-cbc`, `rc2-cfb`, `rc2-ecb`, `rc2-ofb`, `rc4`, `rc4-40`, `md2`, `md4`, `md5`, `rmd160`, `sha`, `sha1`.

crl-dir()

Accepted values:	Directory name
Default:	none

Description: Name of a directory that contains the Certificate Revocation Lists for trusted CAs. Similarly to `ca-dir()` files, use the 32-bit hash of the name of the issuing CAs as filenames. The extension of the files must be `.r0`.

key-file()

Accepted values:	Filename
Default:	none

Description: Name of a file, that contains an unencrypted private key in PEM format, suitable as a TLS key.

peer-verify()

Accepted values:	`optional-trusted` | `optional-untrusted` | `required-trusted` | `required-untrusted`
Default:	`required-trusted`

Description: Verification method of the peer, the four possible values is a combination of two properties of validation:

- whether the peer is required to provide a certificate (required or optional prefix), and
- whether the certificate provided needs to be valid or not.

The following table summarizes the possible options and their results depending on the certificate of the peer.

		The remote peer has:		
		no certificate	**invalid certificate**	**valid certificate**
Local peer-verify() setting	**optional-untrusted**	TLS-encryption	TLS-encryption	TLS-encryption
	optional-trusted	TLS-encryption	rejected connection	TLS-encryption
	required-untrusted	rejected connection	TLS-encryption	TLS-encryption
	required-trusted	rejected connection	rejected connection	TLS-encryption

For untrusted certificates only the existence of the certificate is checked, but it does not have to be valid — syslog-ng accepts the certificate even if it is expired, signed by an unknown CA, or its CN and the name of the machine mismatches.

Warning
When validating a certificate, the entire certificate chain must be valid, including the CA certificate. If any certificate of the chain is invalid, syslog-ng OSE will reject the connection.

trusted-dn()

Accepted values:	list of accepted distinguished names
Default:	none

Description: To accept connections only from hosts using certain certificates signed by the trusted CAs, list the distinguished names of the accepted certificates in this parameter. For example using `trusted-dn("*, O=Example Inc, ST=Some-State, C=*")` will accept only certificates issued for the `Example Inc` organization in `Some-State` state.

trusted-keys()

Accepted values:	list of accepted SHA-1 fingerprints
Default:	none

Description: To accept connections only from hosts using certain certificates having specific SHA-1 fingerprints, list the fingerprints of the accepted certificates in this parameter. For example `trusted-keys("SHA1:00:EF:ED:A4:CE:00:D1:14:A4:AB:43:00:EF:00:91:85:FF:89:28:8F", "SHA1:0C:42:00:3E:B2:60:36:64:00:E2:83:F0:80:46:AD:00:A8:9D:00:15")`.

To find the fingerprint of a certificate, you can use the following command: `openssl x509 -in <certificate-filename> -sha1 -noout -fingerprint`

Note
When using the `trusted-keys()` and `trusted-dn()` parameters, note the following:

- First, the `trusted-keys()` parameter is checked. If the fingerprint of the peer is listed, the certificate validation is performed.

- If the fingerprint of the peer is not listed in the `trusted-keys()` parameter, the `trusted-dn()` parameter is checked. If the DN of the peer is not listed in the `trusted-dn()` parameter, the authentication of the peer fails and the connection is closed.

Chapter 11. Manipulating messages

This chapter explains the methods that you can use to customize, reformat, and modify log messages using syslog-ng Open Source Edition.

- *Section 11.1, Customizing message format (p. 238)* explains how to use templates and macros to change the format of log messages, or the names of logfiles and database tables.

- *Section 11.2, Modifying messages (p. 257)* describes how to use rewrite rules to search and replace certain parts of the message content.

- *Section 11.3, Regular expressions (p. 263)* lists the different types of regular expressions that can be used in various syslog-ng OSE objects like filters and rewrite rules.

11.1. Customizing message format

The following sections describe how to customize the names of logfiles, and also how to use templates, macros, and template functions.

- *Section 11.1.1, Formatting messages, filenames, directories, and tablenames (p. 238)* explains how macros work.

- *Section 11.2, Modifying messages (p. 257)* describes how to use macros and templates to format log messages or change the names of logfiles and database tables.

- *Section 11.1.5, Macros of syslog-ng OSE (p. 241)* lists the different types of macros available in syslog-ng OSE.

- *Section 11.1.6, Using template functions (p. 248)* explains what template functions are and how to use them.

- *Section 11.1.7, Template functions of syslog-ng OSE (p. 248)* lists the template functions available in syslog-ng OSE.

11.1.1. Formatting messages, filenames, directories, and tablenames

The syslog-ng OSE application can dynamically create filenames, directories, or names of database tables using macros that help you organize your log messages. Macros refer to a property or a part of the log message, for example, the `${HOST}` macro refers to the name or IP address of the client that sent the log message, while `${DAY}` is the day of the month when syslog-ng has received the message. Using these macros in the path of the destination log files allows you for example to collect the logs of every host into separate files for every day.

A set of macros can be defined as a template object and used in multiple destinations.

Another use of macros and templates is to customize the format of the syslog message, for example, to add elements of the message header to the message text.

Note

If a message uses the IETF-syslog format (RFC5424), only the text of the message can be customized (that is, the $MESSAGE part of the log), the structure of the header is fixed.

- For details on using templates and macros, see *Section 11.1.2, Templates and macros (p. 239)*.

- For a list and description of the macros available in syslog-ng OSE, see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*.

- For details on using custom macros created with CSV parsers and pattern databases, see *Chapter 12, Parsing and segmenting structured messages (p. 267)* and *Section 13.2.1, Using parser results in filters and templates (p. 279)*, respectively.

11.1.2. Templates and macros

The syslog-ng OSE application allows you to define message templates, and reference them from every object that can use a template. Templates can be used for example to create standard message formats or filenames. Templates can reference one or more macros (for example date, the hostname, and so on). For a list of macros available in syslog-ng Open Source Edition, see *Section 11.1.5, Macros of syslog-ng OSE (p. 241)*. Fields from the structured data (SD) part of messages using the new IETF-syslog standard can also be used as macros.

Template objects have a single option called `template-escape()`, which is disabled by default (`template-escape(no)`). This behavior is useful when the messages are passed to an application that cannot handle escaped characters properly. Enabling template escaping (`template-escape(yes)`) causes syslog-ng to escape the `'`, `"`, and backslash characters from the messages.

Note

In versions 2.1 and earlier, the `template-escape()` option was enabled by default.

Macros can be included by prefixing the macro name with a `$` sign, just like in Bourne compatible shells. Although using braces around macro names is not mandatory, and the `"$MSG"` and `"${MSG}"` formats are equivalent, using the `"${MSG}"` format is recommended for clarity.

To use a literal `$` character in a template, you have to escape it. In syslog-ng OSE versions 3.4 and earlier, use a backslash (`\$`). In version 3.5 and later, use $$.

Note

To use a literal @ character in a template, use @@.

Default values for macros can also be specified by appending the `:-` characters and the default value of the macro. If a message does not contain the field referred to by the macro, or it is empty, the default value will be used when expanding the macro. For example, if a message does not contain a hostname, the following macro can specify a default hostname.

```
${HOST:-default_hostname}
```

Warning
The hostname-related macros (*${FULLHOST}*, *${FULLHOST_FROM}*, *${HOST}*, and *${HOST_FROM}*) do not have any effect if the `keep-hostname()` option is disabled.

By default, syslog-ng sends messages using the following template: ${ISODATE} ${HOST} ${MSGHDR}${MSG}\n. (The ${MSGHDR}${MSG} part is written together because the *${MSGHDR}* macro includes a trailing whitespace.)

Note
Earlier versions of syslog-ng used templates and scripts to send log messages into SQL databases. Starting from version 2.1, syslog-ng natively supports direct database access using the `sql()` destination. For details, see *Section 7.12.4, sql() destination options (p. 167)*.

Example 11.1. Using templates and macros
The following template (`t_demo_filetemplate`) adds the date of the message and the name of the host sending the message to the beginning of the message text. The template is then used in a file destination: messages sent to this destination (`d_file`) will use the message format defined in the template.

```
template t_demo_filetemplate {
            template("${ISODATE} ${HOST} ${MSG}\n"); template-escape(no); };
destination d_file {
            file("/var/log/messages" template(t_demo_filetemplate)); };
```

Templates can also be used inline, if they are used only at a single location. The following destination is equivalent with the previous example:

```
destination d_file {
                file ("/var/log/messages"
                        template("${ISODATE} ${HOST} ${MSG}\n") template-escape(no) );
                };
```

The following file destination uses macros to daily create separate logfiles for every client host.

```
destination d_file {
        file("/var/log/${YEAR}.${MONTH}.${DAY}/${HOST}.log");
};
```

Note
Macros can be used to format messages, and also in the name of destination files or database tables. However, they cannot be used in sources as wildcards, for example, to read messages from files or directories that include a date in their name.

11.1.3. Date-related macros

The macros related to the date of the message (for example: *${ISODATE}*, *${HOUR}*, and so on) have three further variants each:

- S_ prefix, for example, *${S_DATE}*: The *${S_DATE}* macro represents the date found in the log message, that is, when the message was sent by the original application.

- R_ prefix, for example, *${R_DATE}*: *${R_DATE}* is the date when syslog-ng OSE has received the message.

- C_ prefix, for example, *C_DATE*: *C_DATE* is the current date, that is when syslog-ng OSE processes the message and resolves the macro.

The *${DATE}* macro equals the *${S_DATE}* macro.

The values of the date-related macros if calculated based on the original timezone information of the message. To convert it to a different timezone, use the *time-zone()* option. (You can set the *time-zone()* option as a global option, or per destination. For sources, it applies only if the original message does not contain timezone information.

11.1.4. Hard vs. soft macros

Hard macros contain data that is directly derived from the log message, for example, the ${MONTH} macro derives its value from the timestamp. Hard macros are read-only. Soft macros (sometimes also called name-value pairs) are either built-in macros automatically generated from the log message (for example, ${HOST}), or custom user-created macros generated by using the syslog-ng pattern database or a CSV-parser. In contrast to hard macros, soft macros are writable and can be modified within syslog-ng OSE, for example, using rewrite rules.

Hard and soft macros are rather similar and often treated as equivalent. Macros are most commonly used in filters and templates, which does not modify the value of the macro, so both soft and hard macros can be used. However, it is not possible to change the values of hard macros in rewrite rules or via any other means.

The following macros in syslog-ng OSE are hard macros and cannot be modified: *BSDTAG*, *CONTEXT_ID*, *DATE*, *DAY*, *FACILITY_NUM*, *FACILITY*, *FULLDATE*, *HOUR*, *ISODATE*, *LEVEL_NUM*, *LEVEL*, *MIN*, *MONTH_ABBREV*, *MONTH_NAME*, *MONTH*, *MONTH_WEEK*, , *PRIORITY*, *PRI*, *RCPTID*, *SDATA*, *SEC*, *SEQNUM*, *SOURCEIP*, *STAMP*, *TAG*, *TAGS*, *TZOFFSET*, *TZ*, *UNIXTIME*, *WEEK_DAY_ABBREV*, *WEEK_DAY_NAME*, *WEEK_DAY*, *WEEK*, *YEAR_DAY*, *YEAR*.

The following macros can be modified:*FULLHOST_FROM*, *FULLHOST*, *HOST_FROM*, *HOST*, *LEGACY_MSGHDR*, *MESSAGE*, *MSG*,*MSGID*, *MSGONLY*, *PID*, *PROGRAM*, *SOURCE*. Custom values created using rewrite rules or parsers can be modified as well, just like stored matches of regular expressions ($0 ... $255).

11.1.5. Macros of syslog-ng OSE

The following macros are available in syslog-ng OSE.

AMPM
Description: Typically used together with the *${HOUR12}* macro, *${AMPM}* returns the period of the day: AM for hours before mid day and PM for hours after mid day. In reference to a 24-hour clock format, AM is between 00:00-12:00 and PM is between 12:00-24:00. 12AM is midnight. Available in syslog-ng OSE 3.4 and later.

BSDTAG
Description: Facility/priority information in the format used by the FreeBSD syslogd: a priority number followed by a letter that indicates the facility. The priority number can range from 0 to 7. The facility letter can

range from A to Y, where A corresponds to facility number zero (LOG_KERN), B corresponds to facility 1 (LOG_USER), and so on.

Custom macros

Description: CSV parsers and pattern databases can also define macros from the content of the messages, for example, a pattern database rule can extract the username from a login message and create a macro that references the username. For details on using custom macros created with CSV parsers and pattern databases, see *Chapter 12, Parsing and segmenting structured messages (p. 267)* and *Section 13.2.1, Using parser results in filters and templates (p. 279)*, respectively.

DATE, C_DATE, R_DATE, S_DATE

Description: Date of the message using the BSD-syslog style timestamp format (month/day/hour/minute/second, each expressed in two digits). This is the original syslog time stamp without year information, for example: `Jun 13 15:58:00`.

DAY, C_DAY, R_DAY, S_DAY

Description: The day the message was sent.

FACILITY

Description: The name of the facility (for example, *kern*) that sent the message.

FACILITY_NUM

Description: The numerical code of the facility (for example, 0) that sent the message.

FULLDATE, C_FULLDATE, R_FULLDATE, S_FULLDATE

Description: A nonstandard format for the date of the message using the same format as `${DATE}`, but including the year as well, for example: `2006 Jun 13 15:58:00`.

FULLHOST

Description: The full FQDN of the host name chain (without trimming chained hosts), including the domain name.

FULLHOST_FROM

Description: FQDN of the host that sent the message to syslog-ng as resolved by syslog-ng using DNS. If the message traverses several hosts, this is the last host in the chain.

The syslog-ng OSE application uses the following procedure to determine the value of the `$FULLHOST_FROM` macro:

1. The syslog-ng OSE application takes the IP address of the host sending the message.

2. If the `use-dns()` option is enabled, syslog-ng OSE attempts to resolve the IP address to a hostname. If it succeeds, the returned hostname will be the value of the `$FULLHOST_FROM` macro. This value will be the FQDN of the host if the `use-fqdn()` option is enabled, but only the hostname if `use-fqdn()` is disabled.

3. If the *use-dns()* option is disabled, or the address resolution fails, the *${FULLHOST_FROM}* macro will return the IP address of the sender host.

HOUR, C_HOUR, R_HOUR, S_HOUR
Description: The hour of day the message was sent.

HOUR12, C_HOUR12, R_HOUR12, S_HOUR12
Description: The hour of day the message was sent in 12-hour clock format. See also the *${AMPM}* macro. 12AM is midnight. Available in syslog-ng OSE 3.4 and later.

HOST
Description: The name of the source host where the message originates from.

- If the message traverses several hosts and the *chain-hostnames()* option is on, the first host in the chain is used.

- If the *keep-hostname()* option is disabled (keep-hostname(no)), the value of the $HOST macro will be the DNS hostname of the host that sent the message to syslog-ng OSE (that is, the DNS hostname of the last hop). In this case the $HOST and $HOST_FROM macros will have the same value.

- If the *keep-hostname()* option is enabled (keep-hostname(yes)), the value of the $HOST macro will be the hostname retrieved from the log message. That way the name of the original sender host can be used, even if there are log relays between the sender and the server.

HOST_FROM
Description: Name of the host that sent the message to syslog-ng, as resolved by syslog-ng using DNS. If the message traverses several hosts, this is the last host in the chain.

ISODATE, C_ISODATE, R_ISODATE, S_ISODATE
Description: Date of the message in the ISO 8601 compatible standard timestamp format (yyyy-mm-ddThh:mm:ss+-ZONE), for example: 2006-06-13T15:58:00.123+01:00. If possible, it is recommended to use *${ISODATE}* for timestamping. Note that syslog-ng can produce fractions of a second (for example milliseconds) in the timestamp by using the *frac-digits()* global or per-destination option.

LEVEL_NUM
Description: The priority (also called severity) of the message, represented as a numeric value, for example, 3. For the textual representation of this value, use the *${LEVEL}* macro. See *Section PRIORITY or LEVEL (p. 245)* for details.

LOGHOST
Description: The hostname of the computer running syslog-ng OSE — it returns the same result as the hostname command.

MIN, C_MIN, R_MIN, S_MIN

Description: The minute the message was sent.

MONTH, C_MONTH, R_MONTH, S_MONTH

Description: The month the message was sent as a decimal value, prefixed with a zero if smaller than 10.

MONTH_ABBREV, C_MONTH_ABBREV, R_MONTH_ABBREV, S_MONTH_ABBREV

Description: The English abbreviation of the month name (3 letters).

MONTH_NAME, C_MONTH_NAME, R_MONTH_NAME, S_MONTH_NAME

Description: The English name of the month name.

MONTH_WEEK, C_MONTH_WEEK, R_MONTH_WEEK, S_MONTH_WEEK

Description: The number of the week in the given month (0-5). The week with numerical value 1 is the first week containing a Monday. The days of month before the first Monday are considered week 0. For example, if a 31-day month begins on a Sunday, then the 1st of the month is week 0, and the end of the month (the 30th and 31st) is week 5.

MSEC, C_MSEC, R_MSEC, S_MSEC

Description: The millisecond the message was sent.

Available in syslog-ng OSE version 3.4 and later.

MSG or MESSAGE

Description: Text contents of the log message without the program name and pid. Note that this has changed in syslog-ng version 3.0; in earlier versions this macro included the program name and the pid. In syslog-ng 3.0, the *${MSG}* macro became equivalent with the *${MSGONLY}* macro. The program name and the pid together are available in the *${MSGHDR}* macro.

MSGHDR

Description: The name and the PID of the program that sent the log message in `PROGRAM[PID]:` format. Includes a trailing whitespace. Note that the macro returns an empty value if both the PROGRAM and PID fields of the message are empty.

MSGID

Description: A string specifying the type of the message in IETF-syslog (RFC5424-formatted) messages. For example, a firewall might use the ${MSGID} "TCPIN" for incoming TCP traffic and the ${MSGID} "TCPOUT" for outgoing TCP traffic. By default, syslog-ng OSE does not specify this value, but uses a dash (-) character instead. If an incoming message includes the ${MSGID} value, it is retained and relayed without modification.

MSGONLY

Description: Message contents without the program name or pid.

PID

Description: The PID of the program sending the message.

PRI

Description: The priority and facility encoded as a 2 or 3 digit decimal number as it is present in syslog messages.

PRIORITY or LEVEL

Description: The priority (also called severity) of the message, for example, *error*. For the textual representation of this value, use the *${LEVEL}* macro. See *Section PRIORITY or LEVEL (p. 245)* for details.

PROGRAM

Description: The name of the program sending the message. Note that the content of the ${PROGRAM} variable may not be completely trusted as it is provided by the client program that constructed the message.

RCPTID

Description: When the *use-rcptid* global option is set to yes, syslog-ng OSE automatically assigns a unique reception ID to every received message. You can access this ID and use it in templates via the *${RCPTID}* macro. The reception ID is a monotonously increasing 48-bit integer number, that can never be zero (if the counter overflows, it restarts with 1).

RUNID

Description: An ID that changes its value every time syslog-ng OSE is restarted, but not when reloaded.

SDATA, .SDATA.SDID.SDNAME

Description: The syslog-ng application automatically parses the STRUCTURED-DATA part of IETF-syslog messages, which can be referenced in macros. The *${SDATA}* macro references the entire STRUCTURED-DATA part of the message, while structured data elements can be referenced using the *${.SDATA.SDID.SDNAME}* macro.

> **Note**
> When using STRUCTURED-DATA macros, consider the following:
>
> - When referencing an element of the structured data, the macro must begin with the dot (.) character. For example, ${.SDATA.timeQuality.isSynced}.
> - The SDID and SDNAME parts of the macro names are case sensitive: ${.SDATA.timeQuality.isSynced} is not the same as ${.SDATA.TIMEQUALITY.ISSYNCED}.

> **Example 11.2. Using SDATA macros**
> For example, if a log message contains the following structured data: [exampleSDID@0 iut="3" eventSource="Application" eventID="1011"][examplePriority@0 class="high"] you can use macros like: ${.SDATA.exampleSDID@0.eventSource} — this would return the Application string in this case.

SEC, C_SEC, R_SEC, S_SEC

Description: The second the message was sent.

SEQNUM

Description: The *${SEQNUM}* macro contains a sequence number for the log message. The value of the macro depends on the scenario, and can be one of the following:

- If syslog-ng OSE receives a message via the IETF-syslog protocol that includes a sequence ID, this ID is automatically available in the *${SEQNUM}* macro.

- If the message is a Cisco IOS log message using the extended timestamp format, then syslog-ng OSE stores the sequence number from the message in this macro. If you forward this message the IETF-syslog protocol, syslog-ng OSE includes the sequence number received from the Cisco device in the `${.SDATA.meta.sequenceId}` part of the message.

  Note
  To enable sequence numbering of log messages on Cisco devices, use the following command on the device (available in IOS 10.0 and later): `service sequence-numbers`. For details, see the manual of your Cisco device.

- For locally generated messages (that is, for messages that are received from a local source, and not from the network), syslog-ng OSE calculates a sequence number when sending the message to a destination (it is not calculated for relayed messages).

  - The sequence number is not global, but per-destination. Essentially, it counts the number of messages sent to the destination.

  - This sequence number increases by one for every message sent to the destination. It not lost when syslog-ng OSE is reloaded, but it is reset when syslog-ng OSE is restarted.

  - This sequence number is added to every message that uses the IETF-syslog protocol (`${.SDATA.meta.sequenceId}`), and can be added to BSD-syslog messages using the *${SEQNUM}* macro.

Note
If you need a sequence number for every log message that syslog-ng OSE receives, use the *RCPTID* macro.

SOURCE
Description: The identifier of the source statement in the syslog-ng OSE configuration file that received the message. For example, if syslog-ng OSE received the log message from the `source s_local { internal(); };` source statement, the value of the ${SOURCE} macro is `s_local`. This macro is mainly useful for debugging and troubleshooting purposes.

SOURCEIP
Description: IP address of the host that sent the message to syslog-ng. (That is, the IP address of the host in the *${FULLHOST_FROM}* macro.) Please note that when a message traverses several relays, this macro contains the IP of the last relay.

STAMP, R_STAMP, S_STAMP
Description: A timestamp formatted according to the *ts-format()* global or per-destination option.

SYSUPTIME

Description: The time elapsed since the computer running syslog-ng OSE has booted. If this data is not available, the macro contains the time elapsed since syslog-ng OSE was started. The value of this macro is an integer containing the time in 1/100th of the second.

Available in syslog-ng OSE version 3.4 and later.

TAG

Description: The priority and facility encoded as a 2 digit hexadecimal number.

TAGS

Description: A comma-separated list of the tags assigned to the message.

> **Note**
>
> Note that the tags are not part of the log message and are not automatically transferred from a client to the server. For example, if a client uses a pattern database to tag the messages, the tags are not transferred to the server. A way of transferring the tags is to explicitly add them to the log messages using a template and the *${TAGS}* macro, or to add them to the structured metadata part of messages when using the IETF-syslog message format.
>
> When sent as structured metadata, it is possible to reference to the list of tags on the central server, and for example, to add them to a database column.

TZ, C_TZ, R_TZ, S_TZ

Description: An alias of the *${TZOFFSET}* macro.

TZOFFSET, C_TZOFFSET, R_TZOFFSET, S_TZOFFSET

Description: The time-zone as hour offset from GMT; for example: -07:00. In syslog-ng 1.6.x this used to be -0700 but as *${ISODATE}* requires the colon it was added to *${TZOFFSET}* as well.

UNIXTIME, C_UNIXTIME, R_UNIXTIME, S_UNIXTIME

Description: Standard UNIX timestamp, represented as the number of seconds since 1970-01-01T00:00:00.

USEC, C_USEC, R_USEC, S_USEC

Description: The microsecond the message was sent.

Available in syslog-ng OSE version 3.4 and later.

YEAR, C_YEAR, R_YEAR, S_YEAR

Description: The year the message was sent.

WEEK, C_WEEK, R_WEEK, S_WEEK

Description: The week number of the year, prefixed with a zero for the first nine week of the year. (The first Monday in the year marks the first week.)

WEEK_ABBREV, C_WEEK_ABBREV, R_WEEK_ABBREV, S_WEEK_ABBREV

Description: The 3-letter English abbreviation of the name of the day the message was sent, for example *Thu*.

WEEK_DAY, C_WEEK_DAY, R_WEEK_DAY, S_WEEK_DAY
Description: The day of the week as a numerical value (1-7).

WEEKDAY, C_WEEKDAY, R_WEEKDAY, S_WEEKDAY
Description: These macros are deprecated, use *${WEEK_ABBREV}, ${R_WEEK_ABBREV}, ${S_WEEK_ABBREV}* instead. The 3-letter name of the day of week the message was sent, for example Thu.

WEEK_DAY_NAME, C_WEEK_DAY_NAME, R_WEEK_DAY_NAME, S_WEEK_DAY_NAME
Description: The English name of the day.

11.1.6. Using template functions

A template function is a transformation: it modifies the way macros or name-value pairs are expanded. Template functions can be used in template definitions, or when macros are used in the configuration of syslog-ng OSE. Template functions use the following syntax:

```
$(function-name parameter1 parameter2 parameter3 ...)
```

For example, the *$(echo)* template function simply returns the value of the macro it receives as a parameter, thus $(echo ${HOST}) is equivalent to *${HOST}*.

The parameters of template functions are separated by a whitespace character. If you want to use a longer string or multiple macros as a single parameter, enclose the parameter in double-quotes or apostrophes. For example:

```
$(echo "${HOST} ${PROGRAM} ${PID}")
```

Template functions can be nested into each other, so the parameter of a template function can be another template function, like:

```
$(echo $(echo ${HOST}))
```

For details on using template functions, see the descriptions of the individual template functions in *Section 11.1.7, Template functions of syslog-ng OSE (p. 248)*.

11.1.7. Template functions of syslog-ng OSE

The following template functions are available in syslog-ng OSE.

echo
Syntax:

```
$(echo argument)
```

Description: Returns the value of its argument. Using $(echo ${HOST}) is equivalent to *${HOST}*.

env
Syntax:

```
$(env <environment-variable>)
```

Description: Returns the value of the specified environment variable. Available in syslog-ng OSE 3.5 and later.

format-json
Syntax:

```
$(format-json parameters)
```

Description: The *format-json* template function receives value-pairs as parameters and converts them into JavaScript Object Notation (JSON) format. Including the template function in a message template allows you to store selected information about a log message (that is, its content, macros, or other metadata) in JSON format.

For details on selecting value-pairs in syslog-ng OSE and for possibilities to specify which information to convert to JSON format, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*. Note that the syntax of *format-json* is different from the syntax of *value-pairs()*: *format-json* uses a syntax similar to command lines.

> **Note**
> By default, syslog-ng OSE handles every message field as a string. For details on how to send selected fields as other types of data (for example, handle the PID as a number), see *Section 2.10.1, Specifying data types in value-pairs (p. 17)*.

> **Example 11.3. Using the format-json template function**
> The following example selects every available information about the log message, except for the date-related macros (R_* and S_*), selects the .SDATA.meta.sequenceId macro, and defines a new value-pair called MSGHDR that contains the program name and PID of the application that sent the log message (since you will use the template-function in a template, you must escape the double-quotes).
>
> ```
> $(format-json --scope syslog,all_macros,selected_macros \
>   --exclude R_* --exclude S_* --key .SDATA.meta.sequenceId \
>   --pair MSGHDR=\"$PROGRAM[$PID]: \")
> ```
>
> The following example shows how to use this template function to store log messages in JSON format:
>
> ```
> destination d_json {
>   file("/var/log/messages.json" template("$(format-json --scope selected_macros --scope
> nv_pairs)\n"));
> };
> ```

> **Note**
> In case of syslog-ng macros starting with a dot (for example ".SDATA.meta.sequenceID") an empty key name is added at the top level of the JSON structure. You can work around this by adding --shift 1 as a parameter to the template function. For example in case of ".SDATA.meta.sequenceID", an empty key name is added at the top level of the JSON structure:
>
> ```
> {"":
>     {"SDATA" :
>         {"meta" :
>             {"sequenceID": "123"}
>         }
>     }
> }
> ```

geoip

Syntax:

```
$(geoip <IP-address>)
```

Description: This template function returns the 2-letter country code of any IPv4 address or host. Currently only the 2-letter codes are supported, and only from the default database. For example, `$(geoip $HOST)`

> **Note**
> This template function is available only if syslog-ng OSE has been compiled with the `--enable-geoip` compiling option.

graphite-output

Syntax:

```
$(graphite-output parameters)
```

Description: Available in syslog-ng OSE 3.6 and later (Originally appeared in the syslog-ng OSE incubator for syslog-ng 3.5). This template function converts value-pairs from the incoming message to the Graphite plain text protocol format. It is ideal to use with the messages generated by the *monitor-source plugin* (currently available in the syslog-ng incubator project).

For details on selecting value-pairs in syslog-ng OSE and for possibilities to specify which information to convert to Graphite plain text protocol format, see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*. Note that the syntax of `graphite-output` is different from the syntax of `value-pairs()`: `graphite-output` uses a the command-line syntax used in the *format-json template function*.

> **Example 11.4. Using the graphite-output template function**
> The following configuration example shows, how to send value-pairs with names starting with "`vmstat.`" to Graphite running on `localhost`, port 2003:
>
> ```
> destination d_graphite {
>   tcp( host("localhost") port(2003) template("$(graphite-output --key vmstat.*)"));
> };
> ```

grep

Syntax:

```
$(grep condition value-to-select)
```

Description: The *grep* template function is useful when using a pattern database to correlate related log messages. The *grep* template function can be used to filter the messages of the same context when the index of the particular message is not known.

> **Example 11.5. Using the grep template function**
> The following example selects the message of the context that has a `username` name-value pair with the `root` value, and returns the value of the `auth_method` name-value pair.
>
> ```
> $(grep ("${username}" == "root") ${auth_method})
> ```

It is possible to specify multiple name-value pairs as parameters, separated with commas. If multiple messages match the condition of *grep*, these will be returned also separated by commas. This can be used for example to collect the e-mail recipients from postfix messages.

hash

Syntax:

```
$(<method> [opts] $arg1 $arg2 $arg3...)
```

Options:

```
--length N, -l N
```

Truncate the hash to the first N characters.

Description: Calculates a hash of the string or macro received as argument using the specified hashing method. If you specify multiple arguments, effectively you receive the hash of the first argument salted with the subsequent arguments.

<method> can be one of md5, md4, sha1, sha256, sha512 and "hash", which is equivalent to md5. Macros are expected as arguments, and they are concatenated without the use of additional characters.

This template function can be used for anonymizing sensitive parts of the log message (for example username) that were parsed out using PatternDB before storing or forwarding the message. This way, the ability of correlating messages along this value is retained.

Also, using this template, quasi-unique IDs can be generated for data, using the *--length* option. This way, IDs will be shorter than a regular hash, but there is a very small possibility of them not being as unique as a non-truncated hash.

Note
These template functions are available only if syslog-ng OSE has been compiled with the *--enable-ssl* compile option and the `tfhash` module has been loaded. By default, syslog-ng OSE loads every available module.

Example 11.6. Using the $(hash) template function
The following example calculates the SHA1 hash of the hostname of the message:

```
$(sha1 $HOST)
```

The following example calculates the SHA256 hash of the hostname, using the `salted` string to salt the result:

```
$(sha1 $HOST salted)
```

To use shorter hashes, set the *--length*:

```
$(sha1 --length 6 $HOST)
```

To replace the hostname with its hash, use a rewrite rule:

```
rewrite r_rewrite_hostname{set("$(sha1 $HOST)", value("HOST"));};
```

Example 11.7.
The following example replaces every IPv4 address in the MESSAGE part with its SHA-1 hash:

```
rewrite pseudonymize_ip_addresses_in_message {subst
("(([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-5])\.){3}([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-5])",
 "$(sha1)", value("MSG"));};
```

if
Syntax:

```
$(if (<condition>) <true template> <false template>)
```

Description: Returns the value of the `<true template>` parameter if the `<condition>` is true. If the `<condition>` is false, the value of `<false template>` is returned.

Example 11.8. Using pattern databases and the if template function
The following example returns `violation` if the `username` name-value pair of a message processed with pattern database is `root`, and `system` otherwise.

```
$(if ("${username}" == "root") "violation" "system")
```

This can be used to set the class of a message in pattern database rules based on the condition.

```
<value name="username">$(if ("${username}" == "root") "violation" "system")</value>
```

Since template functions can be embedded into each other, it is possible to use another template function as the template of the first one. For example, the following expression returns `root` if the username is `root`, `admin` if the username is `joe`, and `normal user` otherwise.

```
<value name="username">
        $(if ("${username}" == "root")
            "root"
            $(if ("${username}" == "joe") "admin" "normal user")
            "normal user")</value>
```

indent-multi-line
Syntax:

```
$(indent-multi-line parameter)
```

Description: This template function makes it possible to write multi-line log messages into a file. The first line is written like a regular message, subsequent lines are indented with a tab, in compliance with RFC822.

Example 11.9. Using the indent-multi-line template function
The following example writes multi-line messages into a text file.

```
destination d_file {
        file ("/var/log/messages"
                template("${ISODATE} ${HOST} $(indent-multi-line ${MESSAGE})\n") );
};
```

ipv4-to-int
Syntax:

```
$(ipv4-to-int parameter)
```

Description: Converts the specified IPv4 address to its numeric representation. The numerical value of an IPv4 address is calculated by treating the IP address as a 4-byte hexadecimal value. For example, the 192.168.1.1 address equals to: 192=C0, 168=A8, 1=01, 1=01, or C0A80101, which is 3232235777 in decimal representation.

> **Note**
> This template function is available only if the `convertfuncs` module has been loaded. By default, syslog-ng OSE loads every available module.

length
Syntax:

```
$(length "<macro>")
```

Description: Returns the length of the macro in characters, for example, the length of the message. For example, the following filter selects messages that are shorter than 16 characters:

```
f_short {
    match ('-', value ("$(if ($(length "${MSG}") <= 16) "-" "+")"));
};
```

lowercase
Syntax:

```
$(lowercase "<macro>")
```

Description: Returns the lowercase version of the specified string or macro. For example, the following example uses the lowercase version of the hostname in a directory name:

```
destination d_file {
file ("/var/log/${MONTH}/${DAY}/$(lowercase "${HOST}")/messages");
};
```

Available in syslog-ng OSE 3.5 and later.

Numerical operations
Syntax:

```
$(<operation> "<value1>" "<value2>")
```

Description: These template functions allow you to manipulate numbers, that is, to perform addition (+), substraction (-), multiplication (*), division (/), and modulus (%). All of them require two numeric arguments. The result is NaN (Not-a-Number) if the parameters are not numbers, cannot be parsed, or if a division by zero would occur. For example, to add the value of two macros, use the following template function:

```
$(+ "${<MACRO1>}" "${<MACRO2>}");
```

or
Syntax:

```
$(or <macro1> <macro2>)
```

Description: This template function returns the first non-empty argument.

padding
Syntax:

```
$(padding <macro> <prepended characters>)
```

Description: This template function extends a string to a specified width with prepending specified characters. The default padding character is ' '(space). For example:

Example 11.10. Using the padding template function

```
$(padding $MESSAGE 10 X): bar
```

This will result in the following string: XXXXXXXbar

```
$(padding $MESSAGE 10 foo): bar
```

This will result in the following string: foofoofbar

replace-delimiter
Syntax:

```
$(replace-delimiter "<old-delimiters>" "<new-delimiter>" "<macro>")
```

Description: Replaces the delimiter character with a new one. For example, the following example replaces the tabulators (/t) in the message with semicolons (;):

```
$(replace-delimiter "\t" ";" "${MESSAGE}")
```

Available in syslog-ng OSE 3.5 and later.

sanitize
Syntax:

```
$(sanitize <options> "<macro1>" "<macro2> ...")
```

Description: This file replaces the special characters in macro values, for example, it can replace the slash (/) characters in a filename with the underscore (_) character. If you specify multiple arguments, they will be concatenated using the / character, so they can be used as separate directory levels when used in filenames.

The function has the following options:

--ctrl-chars or -c	Filter control characters (characters that have an ASCII code of 32 or lower). This option is used by default.
--invalid-chars <characterlist> or -i <characterlist>	The list of characters to be replaced with underscores (_). The default list contains the / character. The following example replaces the \ and @ characters, so for example, fo\o@bar becomes foobar:

```
$(sanitize -i \@ $PROGRAM)
```

`--no-ctrl-chars or -C`	Do not filter the control characters (characters that have an ASCII code of 32 or lower).

`--replacement`
`<replacement-character> or`
`-r <replacement-character>`
	The character used to replace invalid characters. By default, this is the underscore (_). The following example replaces invalid characters with colons instead of underscores, so for example, foo/bar becomes foo;bar:

```
$(sanitize -r ; $PROGRAM)
```

Example 11.11. Using the sanitize template function
The following example uses the sanitize function on two macros, and the results are used as directory names in a file destination.

```
file("/var/log/$(sanitize $HOST $PROGRAM)/messages");
```

This is equivalent to `file("/var/log/$HOST/$PROGRAM/messages");`, but any slashes in the values of the $HOST and $PROGRAM macros are replaced with underscores.

strip
Syntax:

```
$(strip "<macro>")
```

Description: Deletes whitespaces from the beginning and the end of a macro. You can specify multiple macros separated with whitespace in a single template function, for example:

```
$(strip "${MSG}" "${PROGRAM}");
};
```

substr
Syntax:

```
$(substr "<argument>" "<offset>" "<length>")
```

Description: This function extracts a substring of a string.

argument	The string to extract the substring from, for example, *"${MSG}"*
offset	Specifies where the substring begins (in characters). `0` means to start from the beginning of the string, `5` means to skip the first 5 characters of the string, and so on. Use negative numbers to specify where to start from the end of the string, for example, `-1` means the last character, `-5` means to start five characters before the end of the string.
length	*Optional parameter*: The number of characters to extract. If not specified, the substring will be extracted from the offset to the end of the string. Use negative numbers to stop the substring

before the end of the string, for example, `-5` means the substring ends five characters before the end of the string.

Example 11.12. Using the substr template function
Skip the first 15 characters of the message, and select the rest:

```
$(substr "${MSG}" "15");
```

Select characters 16-30 of the message:

```
$(substr "${MSG}" "15" "30");
```

Select the last 15 characters of the message:

```
$(substr "${MSG}" "-15");
```

A template that converts the message to RFC3164 (BSD-syslog) format and truncates the messages to 1023 characters:

```
template t_truncate_messages {
      template("$(substr \"<$PRI>$DATE $HOST $MSGHDR$MSG\" \"0\" \"1023\")\n");
      template-escape(no);
  };
```

uppercase
Syntax:

```
$(uppercase "<macro>")
```

Description: Returns the uppercase version of the specified string or macro. For example, the following example uses the uppercase version of the hostname in a directory name:

```
destination d_file {
file ("/var/log/${MONTH}/${DAY}/$(uppercase "${HOST}")/messages");
};
```

Available in syslog-ng OSE 3.5 and later.

uuid
Syntax:

```
$(uuid)
```

Description: Generates a Universally Unique IDentifier (UUID) that complies with _RFC4122_. That way, an UUID can be added to the message soon after it is received, so messages stored in multiple destinations can be identified. For example, when storing messages in a database and also in files, the UUID can be used to find a particular message both in the database and the files.

To generate a UUID, you can use a rewrite rule to create a new value-pair for the message.

Example 11.13. Using Universally Unique Identifiers
The following example adds a value-pair called MESSAGE_UUID to the message using a rewrite rule and a template.

```
rewrite r_add_uuid { set("$(uuid)" value("MESSAGE_UUID")); };
destination d_file {
    file ("/var/log/messages"
          template("$MESSAGE_UUID $ISODATE $HOST $MSG\n")
          template-escape(no)
        );
};

log { source(s_network);
      rewrite(r_add_uuid);
      destination(d_file);
};
```

Note
This template function is available only if the `tfuuid` module has been loaded. By default, syslog-ng OSE loads every available module.

11.1.8. Modifying the on-the-wire message format

Macros, templates, and template functions allow you to fully customize the format of the message. This flexibility makes it possible to use syslog-ng OSE in some unexpected way if needed, for example, to emulate simple, plain-text protocols. The following example shows you how to send LPUSH commands to a Redis server.

Note
The purpose of this example is to demonstrate the flexibility of syslog-ng OSE. A dedicated Redis destination is available in syslog-ng OSE version 3.5. For details, see *Section 7.9, Storing name-value pairs in Redis (p. 152)*.

The following template is a valid LPUSH command in accordance with the _Redis protocol_, and puts the $MESSAGE into a separate list for every $PROGRAM:

```
template t_redis_lpush {
    template("*3\r\n$$5\r\nLPUSH\r\n$$$(length
${PROGRAM})\r\n${PROGRAM}\r\n$$$(length ${MESSAGE})\r\n${MESSAGE}\r\n");
};
```

If you use this template in a `tcp()` destination, syslog-ng OSE formats the message according to the template, and sends it to the Redis server.

```
destination d_redis_tcp {
    tcp("127.0.0.1" port(6379) template(t_redis_lpush));
};
```

11.2. Modifying messages

The syslog-ng application can rewrite parts of the messages using rewrite rules. Rewrite rules are global objects similar to parsers and filters and can be used in log paths. The syslog-ng application has two methods to rewrite

---

This example is a summary of a blog post by _Ben Wilder_.

parts of the log messages: substituting (setting) a part of the message to a fix value, and a general search-and-replace mode.

Substitution completely replaces a specific part of the message that is referenced using a built-in or user-defined macro.

General rewriting searches for a string in the entire message (or only a part of the message specified by a macro) and replaces it with another string. Optionally, this replacement string can be a template that contains macros.

Rewriting messages is often used in conjunction with message parsing *Section 12.2, Parsing messages (p. 268)*.

Rewrite rules are similar to filters: they must be defined in the syslog-ng configuration file and used in the log statement.

> **Note**
> The order of filters, rewriting rules, and parsers in the log statement is important, as they are processed sequentially.

11.2.1. Replacing message parts

To replace a part of the log message, you have to:

- define a string or regular expression to find the text to replace
- define a string to replace the original text (macros can be used as well)
- select the field of the message that the rewrite rule should process

Substitution rules can operate on any value available via macros, for example HOST, MESSAGE, PROGRAM, or any user-defined macros created using parsers (for details, see *Chapter 12, Parsing and segmenting structured messages (p. 267)* and *Chapter 13, Processing message content with a pattern database (p. 275)*). The only exceptions are the FACILITY, SEVERITY, TAGS, and the date-related fields, which cannot be rewritten. You can also rewrite the structured-data fields of messages complying to the RFC5424 (IETF-syslog) message format. Substitution rules use the following syntax:

```
Declaration:
rewrite <name_of_the_rule> {
    subst("<string or regular expression to find>",
        "<replacement string>", value(<field name>), flags() );
};
```

The *type()* and *flags()* options are optional. The *type()* specifies the type of regular expression to use; while the *flags()* are the flags of the regular expressions. For details on regular expressions, see *Section 11.3, Regular expressions (p. 263)*.

A single substitution rule can include multiple substitutions that are applied sequentially to the message. Note that rewriting rules must be included in the log statement to have any effect.

Tip

For case-insensitive searches, add the `flags(ignore-case)` option; to replace every occurrence of the string, add `flags(global)` option.

Example 11.14. Using substitution rules

The following example replaces the `IP` in the text of the message with the string `IP-Address`.

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"));};
```

To replace every occurrence, use:

```
rewrite r_rewrite_subst{
    subst("IP", "IP-Address", value("MESSAGE"), flags("global"));
};
```

Multiple substitution rules are applied sequentially; the following rules replace the first occurrence of the string `IP` with the string `IP-Addresses`.

```
rewrite r_rewrite_subst{
    subst("IP", "IP-Address", value("MESSAGE"));
    subst("Address", "Addresses", value("MESSAGE"));
};
```

Example 11.15.

The following example replaces every IPv4 address in the MESSAGE part with its SHA-1 hash:

```
rewrite pseudonymize_ip_addresses_in_message {subst
("(([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-5])\.){3}([0-9]|[1-9][0-9]|1[0-9]{2}|2[0-4][0-9]|25[0-5])",
 "$(sha1)", value("MSG"));};
```

11.2.2. Setting message fields to specific values

To set a field of the message to a specific value, you have to:

- define the string to include in the message, and
- select the field where it should be included.

You can set the value of available macros, for example HOST, MESSAGE, PROGRAM, or any user-defined macros created using parsers (for details, see *Chapter 12, Parsing and segmenting structured messages (p. 267)* and *Chapter 13, Processing message content with a pattern database (p. 275)*). The only exceptions are the FACILITY, SEVERITY, TAGS, and the date-related fields, which cannot be rewritten. Note that the rewrite operation completely replaces any previous value of that field. Use the following syntax:

```
Declaration:
rewrite <name_of_the_rule> {
    set("<string to include>", value(<field name>));
};
```

Example 11.16. Setting message fields to a particular value

The following example sets the HOST field of the message to `myhost`.

```
rewrite r_rewrite_set{set("myhost", value("HOST"));};
```

The following example appends the "suffix" string to the MESSAGE field:

```
rewrite r_rewrite_set{set("$MESSAGE suffix", value("MESSAGE"));};
```

For details on rewriting SDATA fields, see *Section 11.2.3, Creating custom SDATA fields (p. 260)*.

11.2.3. Creating custom SDATA fields

If you use RFC5424-formatted (IETF-syslog) messages, you can also create custom fields in the SDATA part of the message (For details on the SDATA message part, see *Section 2.8.2.3, The STRUCTURED-DATA message part (p. 15)*). According to RFC5424, the name of the field (its SD-ID) must not contain the @ character for reserved SD-IDs. Custom SDATA fields must be in the following format: `name@<private enterprise number>`, for example, `mySDATA-field@18372.4`. (18372.4 is the private enterprise number of BalaBit S.a.r.l., the developer of syslog-ng OSE.)

Example 11.17. Rewriting custom SDATA fields
The following example sets the sequence ID field of the RFC5424-formatted (IETF-syslog) messages to a fixed value. This field is a predefined SDATA field with a reserved SD-ID, therefore its name does not contain the @ character.

```
rewrite r_sd {
    set("55555" value(".SDATA.meta.sequenceId"));
};
```

It is also possible to set the value of a field that does not exist yet, and create a new, custom name-value pair that is associated with the message. The following example creates the `MODIFIED@18372.4` field and sets its value to `yes`. If you use the `${MODIFIED@18372.4}` macro in a template or SQL table, its value will be `yes` for every message that was processed with this rewrite rule, and empty for every other message.

```
rewrite r_rewrite_set {
    set("yes", value("MODIFIED@18372.4"));
};
```

11.2.4. Setting multiple message fields to specific values

The `groupset()` rewrite rule allows you to modify the value of multiple message fields at once, for example, to change the value of sensitive fields extracted using patterndb, or received in a JSON format.

- The first parameter is the new value of the modified fields. This can be a simple string, a macro, or a template (which can include template functions as well).

- The second parameter (`values()`) specifies the fields to modify. You can explicitly list the macros or fields (a space-separated list with the values enclosed in double-quotes), or use wildcards and glob expressions to select multiple fields.

- Note that `groupset()` does not create new fields, it only modifies existing fields.

- You can refer to the old value of the field using the `$_` macro. This is resolved to the value of the current field, and is available only in `groupset()` rules.

```
Declaration:
rewrite <name_of_the_rule> {
    groupset("<new-value-of-the-fields>", values("<field-name-or-glob>"
["<another-field-name-or-glob>"]));
};
```

Example 11.18. Using groupset rewrite rules
The following examples show how to change the values of multiple fields at the same time.

- Change the value of the *HOST* field to *myhost*.

```
groupset ("myhost" values("HOST"))
```

- Change the value of the *HOST* and *FULLHOST* fields to *myhost*.

```
groupset ("myhost" values("HOST" "FULLHOST"))
```

- Change the value of the *HOST FULLHOST* and fields to lowercase.

```
groupset ("$(lowercase "$_")" values("HOST" "FULLHOST"))
```

- Change the value of each field and macro that begins with *.USER* to *nobody*.

```
groupset ("nobody" values(".USER.*"))
```

- Change the value of each field and macro that begins with *.USER* to its SHA-1 hash (truncated to 6 characters).

```
groupset ("$(sha1 --length 6 $_)" values(".USER.*"))
```

11.2.5. Conditional rewrites

Starting with 3.2, it is possible to apply a rewrite rule to a message only if certain conditions are met. The `condition()` option effectively embeds a filter expression into the rewrite rule: the message is modified only if the message passes the filter. If the condition is not met, the message is passed to the next element of the log path (that is, the element following the rewrite rule in the log statement, for example, the destination). Any filter expression normally used in filters can be used as a rewrite condition. Existing filter statements can be referenced using the `filter()` function within the condition. For details on filters, see *Section 8.3, Filters (p. 208)*.

Tip
Using conditions in rewrite rules can simplify your syslog-ng OSE configuration file, as you do not need to create separate log paths to modify certain messages.

11.2.5.1. Procedure – How conditional rewriting works

Purpose:

The following procedure summarizes how conditional rewrite rules (rewrite rules that have the `condition()` parameter set) work. The following configuration snippet is used to illustrate the procedure:

```
rewrite r_rewrite_set{set("myhost", value("HOST")
condition(program("myapplication")));};
log {
    source(s1);
    rewrite(r_rewrite_set);
    destination(d1);};
```

Steps:

Step 1. The log path receives a message from the source (`s1`).

Step 2. The rewrite rule (`r_rewrite_set`) evaluates the condition. If the message matches the condition (the PROGRAM field of the message is "myapplication"), syslog-ng OSE rewrites the log message (sets the value of the HOST field to "myhost"), otherwise it is not modified.

Step 3. The next element of the log path processes the message (`d1`).

Example 11.19. Using conditional rewriting
The following example sets the HOST field of the message to `myhost` only if the message was sent by the `myapplication` program.

```
rewrite r_rewrite_set{set("myhost", value("HOST") condition(program("myapplication")));};
```

The following example is identical to the previous one, except that the condition references an existing filter template.

```
filter f_rewritefilter {program("myapplication");};
rewrite r_rewrite_set{set("myhost", value("HOST") condition(filter(f_rewritefilter)));};
```

11.2.6. Adding and deleting tags

To add or delete a tag, you can use rewrite rules. To add a tag, use the following syntax:

```
rewrite <name_of_the_rule> {
    set-tag("<tag-to-add>");
};
```

To delete a tag, use the following syntax:

```
rewrite <name_of_the_rule> {
    clear-tag("<tag-to-delete>");
};
```

You cannot use macros in the tags.

11.2.7. Anonymizing credit card numbers

Log messages of banking and e-commerce applications might include credit card numbers (Primary Account Number or PAN). According to privacy best practices and the requirements of the Payment Card Industry Data Security Standards (PCI-DSS), PAN must be rendered unreadable. The syslog-ng OSE application uses a regular expression to detect credit card numbers, and provides two ways to accomplish this: you can either mask the credit card numbers, or replace them with a hash. To mask the credit card numbers, use the *credit-card-mask()* or the *credit-card-hash()* rewrite rules in a log path.

```
Usage:
    @include 'scl/rewrite/cc-mask.conf'

    rewrite { credit-card-mask(value("<message-field-to-process>")); };
};
```

By default, these rewrite rules process the MESSAGE part of the log message.

credit-card-hash()

Synopsis:	credit-card-hash(value("\<message-field-to-process\>"))

Description: Process the specified message field (by default, `${MESSAGE}`), and replace any credit card numbers (Primary Account Number or PAN) with its 16-character-long SHA-1 hash.

credit-card-mask()

Synopsis:	credit-card-mask(value("\<message-field-to-process\>"))

Description: Process the specified message field (by default, `${MESSAGE}`), and replace the 7-12th character of any credit card numbers (Primary Account Number or PAN) with asterisks (`*`). For example, syslog-ng OSE replaces the number `5542043004559005` with `554204******9005`.

11.3. Regular expressions

Filters and substitution rewrite rules can use regular expressions. In regular expressions, the characters `()[].*?+^$|\` are used as special symbols. Depending on how you want to use these characters and which quotation mark you use, these characters must be used differently, as summarized below.

- Strings between single quotes (`'string'`) are treated literally and are not interpreted at all, you do not have to escape special characters. For example the output of `'\x41'` is `\x41` (characters as follows: backslash, x(letter), 4(number), 1(number)). This makes writing and reading regular expressions much more simple: it is recommended to use single quotes when writing regular expressions.

- When enclosing strings between double-quotes (`"string"`), the string is interpreted and you have to escape special characters, that is, to precede them with a backslash (\) character if they are meant literally. For example the output of the `"\x41"` is simply the letter `a`. Therefore special characters like \(backslash) "(quotation mark) must be escaped (\\ and \"). The following expressions are interpreted: \a; \n; \r; \t; \v. For example, the \$40 expression matches the $40 string. Backslashes have to be escaped as well if they are meant literally, for example, the \\d expression matches the \d string.

> **Tip**
> If you use single quotes, you do not need to escape the backslash, for example `match("\\.")` is equivalent to `match('\.')`.

- Enclosing alphanumeric strings between double-quotes (`"string"`) is not necessary, you can just omit the double-quotes. For example when writing filters, `match("sometext")` and `match(sometext)` will both match for the `sometext` string.

> **Note**
> Only strings containing alphanumerical characters can be used without quotes or double quotes. If the string contains whitespace or any special characters (`()[].*?+^$|\` or `;:#`), you must use quotes or double quotes.

When using the `;:#` characters, you must use quotes or double quotes, but escaping them is not required.

By default, all regular expressions are case sensitive. To disable the case sensitivity of the expression, add the `flags(ignore-case)` option to the regular expression.

```
filter demo_regexp_insensitive { host("system" flags(ignore-case)); };
```

The regular expressions can use up to 255 regexp matches (`${1} ... ${255}`), but only from the last filter and only if the `flags("store-matches")` flag was set for the filter. For case-insensitive searches, use the `flags("ignore-case")` option.

11.3.1. Types and options of regular expressions

By default, syslog-ng uses PCRE-style regular expressions. To use other expression types, add the *type()* option after the regular expression.

The syslog-ng OSE application supports the following expression types:

- *POSIX regular expressions*
- *Perl Compatible Regular Expressions (PCRE)*
- *Literal string searches*
- *Glob patterns without regular expression support*

posix
Description: Use POSIX regular expressions.

Posix regular expressions have the following flag options:

global: Usable only in rewrite rules; match for every occurrence of the expression, not only the first one.

ignore-case: Disable case-sensitivity.

store-matches: Store the matches of the regular expression into the `$1, ... $255` variables. Matches from the last filter expression can be referenced in regular expressions.

utf8: Use UTF-8 matching.

Example 11.20. Using Posix regular expressions

```
filter f_message { message("keyword" flags("utf8" "ignore-case") ); };
```

pcre
Description: Use Perl Compatible Regular Expressions (PCRE). If the *type()* parameter is not specified, syslog-ng uses PCRE regular expressions by default.

PCRE regular expressions have the following flag options:

global: Usable only in rewrite rules; match for every occurrence of the expression, not only the first one.

ignore-case: Disable case-sensitivity.

store-matches: Store the matches of the regular expression into the `$1, ... $255` variables. Named matches (also called named subpatterns), for example `(?<name>...)`, are stored as well. Matches from the last filter expression can be referenced in regular expressions.

unicode: Use Unicode support for UTF-8 matches: UTF-8 character sequences are handled as single characters.

utf8: An alias for the *unicode* flag.

Example 11.21. Using PCRE regular expressions

```
rewrite r_rewrite_subst
        {subst("a*", "?", value("MESSAGE") type("pcre") flags("utf8" "global"));  };
```

string
Description: Match the strings literally, without regular expression support. By default, only identical strings are matched. For partial matches, use the `flags("prefix")` or the `flags("substring")` flags.

glob
Description: Match the strings against a pattern containing '*' and '?' wildcards, without regular expression and character range support. The advantage of glob patterns to regular expressions is that globs can be processed much faster.

*	matches an arbitrary string, including an empty string
?	matches an arbitrary character

Note
- The wildcards can match the / character.
- You cannot use the * and ? literally in the pattern.

11.3.2. Optimizing regular expressions

The *host()*, *match()*, and *program()* filter functions and some other syslog-ng objects accept regular expressions as parameters. But evaluating general regular expressions puts a high load on the CPU, which can cause problems when the message traffic is very high. Often the regular expression can be replaced with simple filter functions and logical operators. Using simple filters and logical operators, the same effect can be achieved at a much lower CPU load.

Example 11.22. Optimizing regular expressions in filters

Suppose you need a filter that matches the following error message logged by the xntpd NTP daemon:

```
xntpd[1567]: time error -1159.777379 is too large (set clock manually);
```

The following filter uses regular expressions and matches every instance and variant of this message.

```
filter f_demo_regexp {
    program("demo_program") and
    match("time error .* is too large .* set clock manually"); };
```

Segmenting the *match()* part of this filter into separate *match()* functions greatly improves the performance of the filter.

```
filter f_demo_optimized_regexp {
    program("demo_program") and
    match("time error") and
    match("is too large") and
    match("set clock manually"); };
```

Chapter 12. Parsing and segmenting structured messages

The filters and default macros of syslog-ng work well on the headers and metainformation of the log messages, but are rather limited when processing the content of the messages. Parsers can segment the content of the messages into name-value pairs, and these names can be used as user-defined macros. Subsequent filtering or other type of processing of the message can use these custom macros to refer to parts of the message. Parsers are global objects most often used together with filters and rewrite rules.

syslog-ng OSE provides the following possibilities to parse the messages, or parts of the messages:

- By default, syslog-ng OSE parses every message as a syslog message. To disable message parsing, use the *flags(no-parse)* option of the source. To explicitly parse a message as a syslog message, use the *syslog* parser. For details, see *Section 12.1, Parsing syslog messages (p. 267)*.

- To segment a message into columns using a CSV-parser, see *Section 12.2, Parsing messages (p. 268)*.
- To parse JSON-formatted messages, see *Section 12.3, The JSON parser (p. 272)*.
- To identify and parse the messages using a pattern database, see *Chapter 13, Processing message content with a pattern database (p. 275)*.

12.1. Parsing syslog messages

By default, syslog-ng OSE parses every message using the *syslog-parser* as a syslog message, and fills the macros with values of the message. The *syslog-parser* does not discard messages: the message cannot be parsed as a syslog message, the entire message (including its header) is stored in the $MSG macro. If you do not want to parse the message as a syslog message, use the `flags(no-parse)` option of the source.

Example 12.1. Using junctions
For example, suppose that you have a single network source that receives log messages from different devices, and some devices send messages that are not RFC-compliant (some routers are notorious for that). To solve this problem in earlier versions of syslog-ng OSE, you had to create two different network sources using different IP addresses or ports: one that received the RFC-compliant messages, and one that received the improperly formatted messages (for example, using the `flags(no-parse)` option). Using junctions this becomes much more simple: you can use a single network source to receive every message, then use a junction and two channels. The first channel processes the RFC-compliant messages, the second everything else. At the end, every message is stored in a single file. The filters used in the example can be *host()* filters (if you have a list of the IP addresses of the devices sending non-compliant messages), but that depends on your environment.

```
log {
    source s_network { syslog(ip(10.1.2.3) transport("tcp") flags(no-parse)); };
    junction {
        channel { filter(f_compliant_hosts); parser { syslog-parser(); }; };
        channel { filter(f_noncompliant_hosts); };
    };
    destination { file("/var/log/messages"); };
};
```

Since every channel receives every message that reaches the junction, use the `flags(final)` option in the channels to avoid the unnecessary processing the messages multiple times:

```
log {
    source s_network { syslog(ip(10.1.2.3) transport("tcp") flags(no-parse)); };
```

```
      junction {
          channel { filter(f_compliant_hosts); parser { syslog-parser(); }; flags(final);
};
          channel { filter(f_noncompliant_hosts); flags(final); };
      };
      destination { file("/var/log/messages"); };
};
```

12.2. Parsing messages

The syslog-ng application can separate parts of log messages (that is, the contents of the $MSG macro) to named fields (columns). These fields act as user-defined macros that can be referenced in message templates, file- and tablenames, and so on.

Parsers are similar to filters: they must be defined in the syslog-ng configuration file and used in the log statement.

> **Note**
> The order of filters, rewriting rules, and parsers in the log statement is important, as they are processed sequentially.

To create a parser, define the columns of the message, the delimiter or separator characters (for example, semicolon or tabulator), and optionally the characters that are used to escape the delimiter characters (quote-pairs). For the list of parser parameters, see *Section 12.2.1, Options of CSV parsers (p. 270)*.

```
Declaration:
      parser parser_name {
          csv-parser(
          columns(column1, column2, ...)
          delimiters()
          quote-pairs()
          );
      };
```

Column names work like macros. Always use a prefix to identify the columns of the parsers, for example MYPARSER1.COLUMN1, MYPARSER2.COLUMN2, and so on. Column names starting with a dot (for example .HOST) are reserved for use by syslog-ng.

> **Example 12.2. Segmenting hostnames separated with a dash**
> The following example separates hostnames like example-1 and example-2 into two parts.
>
> ```
> parser p_hostname_segmentation {
>     csv-parser(columns("HOSTNAME.NAME", "HOSTNAME.ID")
>     delimiters("-")
>     flags(escape-none)
>     template("${HOST}"));
> };
> destination d_file { file("/var/log/messages-${HOSTNAME.NAME:-examplehost}"); };
> log { source(s_local); parser(p_hostname_segmentation); destination(d_file);};
> ```

Example 12.3. Parsing Apache log files

The following parser processes the log of Apache web servers and separates them into different fields. Apache log messages can be formatted like:

```
"%h %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-Agent}i\" %T %v"
```

Here is a sample message:

```
192.168.1.1 - - [31/Dec/2007:00:17:10 +0100] "GET /cgi-bin/example.cgi HTTP/1.1" 200 2708
 "-" "curl/7.15.5 (i4 86-pc-linux-gnu) libcurl/7.15.5 OpenSSL/0.9.8c zlib/1.2.3 libidn/0.6.5"
 2 example.balabit
```

To parse such logs, the delimiter character is set to a single whitespace (`delimiters(" ")`). Whitespaces between quotes and brackets are ignored (`quote-pairs('""[]')`).

```
parser p_apache {
    csv-parser(columns("APACHE.CLIENT_IP", "APACHE.IDENT_NAME", "APACHE.USER_NAME",
        "APACHE.TIMESTAMP", "APACHE.REQUEST_URL", "APACHE.REQUEST_STATUS",
        "APACHE.CONTENT_LENGTH", "APACHE.REFERER", "APACHE.USER_AGENT",
        "APACHE.PROCESS_TIME", "APACHE.SERVER_NAME")
        flags(escape-double-char,strip-whitespace)
        delimiters(" ")
        quote-pairs('""[]')
        );
};
```

The results can be used for example to separate log messages into different files based on the APACHE.USER_NAME field. If the field is empty, the `nouser` name is assigned.

```
log { source(s_local);
    parser(p_apache); destination(d_file);};
};
destination d_file { file("/var/log/messages-${APACHE.USER_NAME:-nouser}"); };
```

Multiple parsers can be used to split a part of an already parsed message into further segments.

Example 12.4. Segmenting a part of a message

The following example splits the timestamp of a parsed Apache log message into separate fields.

```
parser p_apache_timestamp {
    csv-parser(columns("APACHE.TIMESTAMP.DAY", "APACHE.TIMESTAMP.MONTH",
"APACHE.TIMESTAMP.YEAR", "APACHE.TIMESTAMP.HOUR", "APACHE.TIMESTAMP.MIN",
"APACHE.TIMESTAMP.MIN", "APACHE.TIMESTAMP.ZONE")
    delimiters("/: ")
    flags(escape-none)
    template("${APACHE.TIMESTAMP}"));
    };
log { source(s_local);
    log { parser(p_apache); parser(p_apache_timestamp); destination(d_file);};
};
```

Further examples:

- For an example on using the *greedy* option, see *Example 12.5, Adding the end of the message to the last column (p. 271)*.

12.2.1. Options of CSV parsers

The syslog-ng application can separate parts of log messages (that is, the contents of the ${MSG} macro) to named fields (columns). These fields act as user-defined macros that can be referenced in message templates, file- and tablenames, and so on.

To create a parser, define the columns of the message, the delimiter or separator characters, and optionally the characters that are used to escape the delimiter characters (quote-pairs).

```
Declaration:
    parser parser_name {
        csv-parser(
        columns(column1, column2, ...)
        delimiters()
        quote-pairs()
        );
    };
```

Column names work like macros. Always use a prefix to identify the columns of the parsers, for example MYPARSER1.COLUMN1, MYPARSER2.COLUMN2, and so on. Column names starting with a dot (for example .HOST) are reserved for use by syslog-ng.

csv-parser

Synopsis:	csv-parser(columns("PARSER.COLUMN1", "PARSER.COLUMN2", ...))

Description: Specifies the type of parser to use, and the name of the columns to separate messages to. Currently only the *csv-parser* is implemented, which can separate columns based on delimiter characters and strings.

delimiters

Synopsis:	delimiters("<delimiter_characters>")

Description: The character that separates the columns in the message. If you specify multiple characters, every character will be treated as a delimiter. To separate the columns at the tabulator (tab character), specify \t. For example, to separate the text at evert hyphen (-) and colon (:) character, use delimiters("-:")

flags()

Synopsis:	drop-invalid, escape-none, escape-backslash, escape-double-char, greedy, strip-whitespace

Description: Specifies various options for parsing the message. The following flags are available:

- *drop-invalid*: When the *drop-invalid* option is set, the parser does not process messages that do not match the parser. For example, a message does not match the parser if it has less columns than specified in the parser, or it has more columns but the *greedy* flag is not enabled. Using the *drop-invalid* option practically turns the parser into a special filter, that matches messages that have the predefined number of columns (using the specified delimiters).

Tip
Messages dropped as invalid can be processed by a `fallback` log path. For details on the `fallback` option, see *Section 8.1.3, Log path flags (p. 202)*.

- *escape-backslash*: The parsed message uses the backslash (\) character to escape quote characters.

- *escape-double-char*: The parsed message repeats the quote character when the quote character is used literally. For example, to escape a comma (,), the message contains two commas (,,).

- *escape-none*: The parsed message does not use any escaping for using the quote character literally.

- *greedy*: The `greedy` option assigns the remainder of the message to the last column, regardless of the delimiter characters set. You can use this option to process messages where the number of columns varies.

Example 12.5. Adding the end of the message to the last column
If the `greedy` option is enabled, the syslog-ng application adds the not-yet-parsed part of the message to the last column, ignoring any delimiter characters that may appear in this part of the message.

For example, you receive the following comma-separated message: `example 1, example2, example3`, and you segment it with the following parser:

```
csv-parser(columns("COLUMN1", "COLUMN2", "COLUMN3") delimiters(","));
```

The COLUMN1, COLUMN2, and COLUMN3 variables will contain the strings `example1`, `example2`, and `example3`, respectively. If the message looks like `example 1, example2, example3, some more information`, then any text appearing after the third comma (that is, `some more information`) is not parsed, and possibly lost if you use only the variables to reconstruct the message (for example, to send it to different columns of an SQL table).

Using the `greedy` flag will assign the remainder of the message to the last column, so that the COLUMN1, COLUMN2, and COLUMN3 variables will contain the strings `example1`, `example2`, and `example3, some more information`.

```
csv-parser(columns("COLUMN1", "COLUMN2", "COLUMN3") delimiters(",")
flags(greedy));
```

- *strip-whitespace*: The `strip-whitespace` flag removes trailing whitespaces from the beginning and the end of the columns.

quote-pairs()

Synopsis:	quote-pairs('<quote_pairs>')

Description: List quote-pairs between single quotes. Delimiter characters enclosed between quote characters are ignored. Note that the beginning and ending quote character does not have to be identical, for example [] can also be a quote-pair. For an example of using `quote-pairs()` to parse Apache log files, see *Example 12.3, Parsing Apache log files (p. 269)*.

template()

Synopsis:	template("${<macroname>}")

Description: The macro that contains the part of the message that the parser will process. It can also be a macro created by a previous parser of the log path. By default, this is empty and the parser processes the entire message (`${MESSAGE}`). For examples, see *Example 12.2, Segmenting hostnames separated with a dash (p. 268)* and *Example 12.4, Segmenting a part of a message (p. 269)*.

12.3. The JSON parser

JavaScript Object Notation (JSON) is a text-based open standard designed for human-readable data interchange. It is used primarily to transmit data between a server and web application, serving as an alternative to XML. It is described in <u>RFC 4627</u>. The syslog-ng OSE application can separate parts of JSON-encoded log messages to name-value pairs, using json-c. For details on using value-pairs in syslog-ng OSE see *Section 2.10, Structuring macros, metadata, and other value-pairs (p. 17)*.

You can refer to the separated parts of the JSON message using the key of the JSON object as a macro. For example, if the JSON contains {"KEY1":"value1","KEY2":"value2"}, you can refer to the values as ${KEY1} and ${KEY2}. If the JSON content is structured, syslog-ng OSE converts it to dot-notation-format. For example, to access the value of the following structure {"KEY1": {"KEY2": "VALUE"}}, use the ${KEY1.KEY2} macro.

Warning

If the names of keys in the JSON content are the same as the names of syslog-ng OSE soft macros, the value from the JSON content will overwrite the value of the macro. For example, the {"PROGRAM":"value1","MESSAGE":"value2"} JSON content will overwrite the ${PROGRAM} and ${MESSAGE} macros. To avoid overwriting such macros, use the `prefix()` option.

Hard macros cannot be modified, so they will not be overwritten. For details on the macro types, see *Section 11.1.4, Hard vs. soft macros (p. 241)*.

Note

The JSON parser currently supports only integer, double and string values when interpreting JSON structures. As syslog-ng does not handle different data types internally, the JSON parser converts all JSON data to string values. In case of boolean types, the value is converted to 'TRUE' or 'FALSE' as their string representation.

The JSON parser discards messages if it cannot parse them as JSON messages, so it acts as a JSON-filter as well.

To create a JSON parser, define a parser that has the `json-parser()` option. Defining the prefix and the marker are optional. By default, the parser will process the ${MESSAGE} part of the log message. To process other parts of a log message with the JSON parser, use the `template()` option.

```
Declaration:
    parser parser_name {
        json-parser(
        marker()
        prefix()
    );
};
```

Example 12.6. Using a JSON parser
In the following example, the source is a JSON encoded log message. The syslog parser is disabled, so that syslog-ng OSE does not parse the message: `flags(no-parse)`. The json-parser inserts `".json."` prefix before all extracted name-value pairs. The destination is a file, that uses the *format-json* template function. Every name-value pair that begins with a dot (".") character will be written to the file (`dot-nv-pairs`). The log line connects the source, the destination and the parser.

```
source s_json {
    tcp(port(21514) flags(no-parse));
};

destination d_json {
    file("/tmp/test.json"
        template("$(format-json --scope dot-nv-pairs)\n"));
};

parser p_json {
    json-parser (prefix(".json."));
};

log {
    source(s_json);
    parser(p_json);
    destination(d_json);
};
```

12.3.1. Options of JSON parsers

The JSON parser has the following options.

extract-prefix

Synopsis:	extract-prefix()

Description: Extract only the specified subtree from the JSON message. Use the dot-notation to specify the subtree. The rest of the message will be ignored. For example, assuming that the incoming object is named `msg`, the `json-parser(extract-prefix("foo.bar[5]"));` syslog-ng OSE parser is equivalent to the `msg.foo.bar[5]` javascript code. Note that the resulting expression must be a JSON object, so that syslog-ng OSE can extract its members into name-value pairs.

This feature also works when the top-level object is an array, because you can use an array index at the first indirection level, for example: `json-parser(extract-prefix("[5]"))`, which is equivalent to `msg[5]`.

marker

Synopsis:	marker()

Description: Use a marker in case of mixed log messages, to identify JSON encoded messages for the parser.

Some logging implementations require a marker to be set before the JSON payload. The JSON parser is able to find these markers and parse the message only if it is present.

Example 12.7. Using the marker option in JSON parser
This json parser parses log messages which use the "@cee:" marker in front of the json payload. It inserts ".cee." in front of the name of name-value pairs, so later on it is easier to find name-value pairs that were parsed using this parser. (For details on selecting name-value pairs, see *Section value-pairs() (p. 19)*.)

```
parser {
        json-parser(
            marker("@cee:")
            prefix(".cee.")
        );
    };
```

prefix

Synopsis:	prefix()

Description: Insert a prefix before the name part of the name-value pairs to help further processing. For example, if you forward the parsed messages using the IETF-syslog protocol, you can insert all the parsed data into the SDATA part of the message using the `prefix(.SDATA.json.)`.

template()

Synopsis:	template("${<macroname>}")

Description: The macro that contains the part of the message that the parser will process. It can also be a macro created by a previous parser of the log path. By default, this is empty and the parser processes the entire message (`${MESSAGE}`).

Chapter 13. Processing message content with a pattern database

13.1. Classifying log messages

The syslog-ng application can compare the contents of the received log messages to predefined message patterns. By comparing the messages to the known patterns, syslog-ng is able to identify the exact type of the messages, and sort them into message classes. The message classes can be used to classify the type of the event described in the log message. The message classes can be customized, and for example can label the messages as user login, application crash, file transfer, and so on events.

To find the pattern that matches a particular message, syslog-ng uses a method called longest prefix match radix tree. This means that syslog-ng creates a tree structure of the available patterns, where the different characters available in the patterns for a given position are the branches of the tree.

To classify a message, syslog-ng selects the first character of the message (the text of message, not the header), and selects the patterns starting with this character, other patterns are ignored for the rest of the process. After that, the second character of the message is compared to the second character of the selected patterns. Again, matching patterns are selected, and the others discarded. This process is repeated until a single pattern completely matches the message, or no match is found. In the latter case, the message is classified as unknown, otherwise the class of the matching pattern is assigned to the message.

To make the message classification more flexible and robust, the patterns can contain pattern parsers: elements that match on a set of characters. For example, the NUMBER parser matches on any integer or hexadecimal number (for example 1, 123, 894054, 0xFFFF, and so on). Other pattern parsers match on various strings and IP addresses. For the details of available pattern parsers, see *Section 13.5.1, Using pattern parsers (p. 286)*.

The functionality of the pattern database is similar to that of the logcheck project, but it is much easier to write and maintain the patterns used by syslog-ng, than the regular expressions used by logcheck. Also, it is much easier to understand syslog-ng pattens than regular expressions.

Pattern matching based on regular expressions is computationally very intensive, especially when the number of patterns increases. The solution used by syslog-ng can be performed real-time, and is independent from the number of patterns, so it scales much better. The following patterns describe the same message: `Accepted password for bazsi from 10.50.0.247 port 42156 ssh2`

A regular expression matching this message from the logcheck project: `Accepted (gssapi(-with-mic|-keyex)?|rsa|dsa|password|publickey|keyboard-interactive/pam) for [^[:space:]]+ from [^[:space:]]+ port [0-9]+( (ssh|ssh2))?`

A syslog-ng database pattern for this message: `Accepted @QSTRING:auth_method: @ for@QSTRING:username: @from @QSTRING:client_addr: @port @NUMBER:port:@ ssh2`

For details on using pattern databases to classify log messages, see *Section 13.2, Using pattern databases (p. 278)*.

13.1.1. The structure of the pattern database

The pattern database is organized as follows:

Pattern database

Ruleset 1

Ruleset 2 — Program pattern — Rule 1: Message pattern Class Tags
Rule 2: Message pattern Class Tags
...
Rule n: Message pattern Class Tags
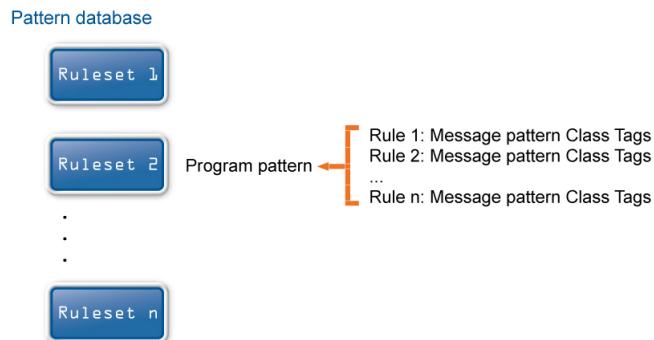
Ruleset n

Figure 13.1. The structure of the pattern database

- The pattern database consists of rulesets. A ruleset consists of a Program Pattern and a set of rules: the rules of a ruleset are applied to log messages if the name of the application that sent the message matches the Program Pattern of the ruleset. The name of the application (the content of the ${PROGRAM} macro) is compared to the Program Patterns of the available rulesets, and then the rules of the matching rulesets are applied to the message.

- The Program Pattern can be a string that specifies the name of the appliation or the beginning of its name (for example, to match for sendmail, the program pattern can be sendmail, or just send), and the Program Pattern can contain pattern parsers. Note that pattern parsers are completely independent from the syslog-ng parsers used to segment messages. Additionally, every rule has a unique identifier: if a message matches a rule, the identifier of the rule is stored together with the message.

- Rules consist of a message pattern and a class. The Message Pattern is similar to the Program Pattern, but is applied to the message part of the log message (the content of the ${MESSAGE} macro). If a message pattern matches the message, the class of the rule is assigned to the message (for example, Security, Violation, and so on).

- Rules can also contain additional information about the matching messages, such as the description of the rule, an URL, name-value pairs, or free-form tags.

- Patterns can consist of literals (keywords, or rather, keycharacters) and pattern parsers.

> **Note**
> If the ${PROGRAM} part of a message is empty, rules with an empty Program Pattern are used to classify the message.
>
> If the same Program Pattern is used in multiple rulesets, the rules of these rulesets are merged, and every rule is used to classify the message. Note that message patterns must be unique within the merged rulesets, but the currently only one ruleset is checked for uniqueness.

13.1.2. How pattern matching works

A sample log message:

`<133>Feb 25 14:09:07 webserver syslogd: restart.`

Program pattern   Message pattern

Figure 13.2. Applying patterns

The followings describe how patterns work. This information applies to program patterns and message patterns alike, even though message patterns are used to illustrate the procedure.

Patterns can consist of literals (keywords, or rather, keycharacters) and pattern parsers. Pattern parsers attempt to parse a sequence of characters according to certain rules.

> **Note**
> Wildcards and regular expressions cannot be used in patterns. The @ character must be escaped, that is, to match for this character, you have to write @@ in your pattern. This is required because pattern parsers of syslog-ng are enclosed between @ characters.

When a new message arrives, syslog-ng attempts to classify it using the pattern database. The available patterns are organized alphabetically into a tree, and syslog-ng inspects the message character-by-character, starting from the beginning. This approach ensures that only a small subset of the rules must be evaluated at any given step, resulting in high processing speed. Note that the speed of classifying messages is practically independent from the total number of rules.

For example, if the message begins with the `Apple` string, only patterns beginning with the character `A` are considered. In the next step, syslog-ng selects the patterns that start with `Ap`, and so on, until there is no more specific pattern left.

Note that literal matches take precedence over pattern parser matches: if at a step there is a pattern that matches the next character with a literal, and another pattern that would match it with a parser, the pattern with the literal match is selected. Using the previous example, if at the third step there is the literal pattern `Apport` and a pattern parser `Ap@STRING@`, the `Apport` pattern is matched. If the literal does not match the incoming string (for example, `Apple`), syslog-ng attempts to match the pattern with the parser. However, if there are two or more parsers on the same level, only the first one will be applied, even if it does not perfectly match the message.

If there are two parsers at the same level (for example, `Ap@STRING@` and `Ap@QSTRING@`), it is random which pattern is applied (technically, the one that is loaded first). However, if the selected parser cannot parse at least one character of the message, the other parser is used. But having two different parsers at the same level is extremely rare, so the impact of this limitation is much less than it appears.

13.1.3. Artificial ignorance

Artificial ignorance is a method to detect anomalies. When applied to log analysis, it means that you ignore the regular, common log messages - these are the result of the regular behavior of your system, and therefore are not too interesting. However, new messages that have not appeared in the logs before can sign important events, and should be therefore investigated. "By definition, something we have never seen before is anomalous" (Marcus J. Ranum).

The syslog-ng application can classify messages using a pattern database: messages that do not match any pattern are classified as unknown. This provides a way to use artificial ignorance to review your log messages. You can periodically review the unknown messages — syslog-ng can send them to a separate destination - and add patterns for them to the pattern database. By reviewing and manually classifying the unknown messages, you can iteratively classify more and more messages, until the only the really anomalous messages show up as unknown.

Obviously, for this to work, a large number of message patterns are required. The radix-tree matching method used for message classification is very effective, can be performed very fast, and scales very well; basically the time required to perform a pattern matching is independent from the number of patterns in the database. For sample pattern databases, see *Section 13.2.2, Downloading sample pattern databases (p. 280)*.

13.2. Using pattern databases

To classify messages using a pattern database, include a *db-parser()* statement in your syslog-ng configuration file using the following syntax:

```
Declaration:
        parser <identifier> {db-parser(file("<database_filename>"));};
```

Note that using the parser in a log statement only performs the classification, but does not automatically do anything with the results of the classification.

Example 13.1. Defining pattern databases
The following statement uses the database located at `/opt/syslog-ng/var/db/patterndb.xml`.

```
parser pattern_db {
        db-parser(
            file("/opt/syslog-ng/var/db/patterndb.xml")
        );
        };
```

To apply the patterns on the incoming messages, include the parser in a log statement:

```
log {
        source(s_all);
        parser(pattern_db);
        destination( di_messages_class);
        };
```

Note
The default location of the pattern database file is `/opt/syslog-ng/var/run/patterndb.xml`. The *file* option of the *db-parser* statement can be used to specify a different file, thus different db-parser statements can use different pattern databases. Later versions of syslog-ng will be able to dynamically generate a main database from separate pattern database files.

Example 13.2. Using classification results
The following destination separates the log messages into different files based on the class assigned to the pattern that matches the message (for example Violation and Security type messages are stored in a separate file), and also adds the ID of the matching rule to the message:

```
destination di_messages_class {
        file("/var/log/messages-${.classifier.class}"

template("${.classifier.rule_id};${S_UNIXTIME};${SOURCEIP};${HOST};${PROGRAM};${PID};${MSG}\n")
```

```
            template-escape(no)
    );
};
```

For details on how to create your own pattern databases see *Section 13.5.3, The syslog-ng pattern database format (p. 289)*.

13.2.1. Using parser results in filters and templates

The results of message classification and parsing can be used in custom filters and templates, for example, in file and database templates. The following built-in macros allow you to use the results of the classification:

- The `.classifier.class` macro contains the class assigned to the message (for example violation, security, or unknown).

- The `.classifier.rule_id` macro contains the identifier of the message pattern that matched the message.

- The `.classifier.context_id` macro contains the identifier of the context for messages that were correlated. For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

Example 13.3. Using classification results for filtering messages
To filter on a specific message class, create a filter that checks the `.classifier_class` macro, and use this filter in a log statement.

```
filter fi_class_violation {
                match("violation"
                value(".classifier.class")
                type("string")
                );
                };
```

```
log {
                source(s_all);
                parser(pattern_db);
                filter(fi_class_violation);
                destination(di_class_violation);
                };
```

Filtering on the *unknown* class selects messages that did not match any rule of the pattern database. Routing these messages into a separate file allows you to periodically review new or unknown messages.

To filter on messages matching a specific classification rule, create a filter that checks the `.classifier.rule_id` macro. The unique identifier of the rule (for example e1e9c0d8-13bb-11de-8293-000c2922ed0a) is the id attribute of the rule in the XML database.

```
filter fi_class_rule {
                match("e1e9c0d8-13bb-11de-8293-000c2922ed0a"
                value(".classifier.rule_id")
                type("string")
                );
                };
```

Pattern database rules can assign tags to messages. These tags can be used to select tagged messages using the *tags()* filter function.

> **Note**
> Starting with version 3.2, syslog-ng OSE automatically adds the class of the message as a tag using the `.classifier.<message-class>` format. For example, messages classified as "system" receive the `.classifier.system` tag. Use the `tags()` filter function to select messages of a specific class.
>
> ```
> filter f_tag_filter {tags(".classifier.system");};
> ```

The message-segments parsed by the pattern parsers can also be used as macros as well. To accomplish this, you have to add a name to the parser, and then you can use this name as a macro that refers to the parsed value of the message.

> **Example 13.4. Using pattern parsers as macros**
> For example, you want to parse messages of an application that look like `"Transaction: <type>."`, where <type> is a string that has different values (for example refused, accepted, incomplete, and so on). To parse these messages, you can use the following pattern:
>
> ```
> 'Transaction: @ESTRING::.@'
> ```
>
> Here the @ESTRING@ parser parses the message until the next full stop character. To use the results in a filter or a filename template, include a name in the parser of the pattern, for example:
>
> ```
> 'Transaction: @ESTRING:TRANSACTIONTYPE:.@'
> ```
>
> After that, add a custom template to the log path that uses this template. For example, to select every `accepted` transaction, use the following custom filter in the log path:
>
> ```
> match("accepted" value("TRANSACTIONTYPE"));
> ```

> **Note**
> The above macros can be used in database columns and filename templates as well, if you create custom templates for the destination or logspace.
>
> Use a consistent naming scheme for your macros, for example, APPLICATIONNAME_MACRONAME.

13.2.2. Downloading sample pattern databases

To simplify the building of pattern databases, BalaBit has released (and will continue to release) sample databases. You can download sample pattern databases from the *BalaBit GitHub page* (older samples are temporarily available *here*.

Note that these pattern databases are only samples and experimental databases. They are not officially supported, and may or may not work in your environment.

The syslog-ng pattern databases are available under the Creative Commons Attribution-Share Alike 3.0 (CC by-SA) license. This includes every pattern database written by community contributors or the BalaBit staff. It means that:

- you are free to use and modify the patterns for your needs;
- if you redistribute the pattern databases, you must distribute your modifications under the same license;
- if you redistribute the pattern databases, you must make it obvious that the source of the original syslog-ng pattern databases is the *BalaBit GitHub page*.

For legal details, the full text of the license is *available here*.

If you create patterns that are not available in the GitHub repository, consider sharing them with us and the syslog-ng community, and send them to the *syslog-ng mailing list*, or to the following e-mail address: <patterndb@balabit.com>

13.3. Correlating log messages

Starting with version 3.2, the syslog-ng OSE application is able to correlate log messages identified using pattern databases.

Log messages are supposed to describe events, but applications often separate information about a single event into different log messages. For example, the Postfix e-mail server logs the sender and recipient addresses into separate log messages, or in case of an unsuccessful login attempt, the OpenSSH server sends a log message about the authentication failure, and the reason of the failure in the next message.

Of course, messages that are not so directly related can be correlated as well, for example, login-logout messages, and so on.

To correlate log messages, syslog-ng OSE uses the pattern database to add messages into message-groups called contexts. A context consists of a series of log messages that are related to each other in some way, for example, the log messages of an SSH session can belong to the same context. As new messages come in, they may be added to a context. Also, when an incoming message is identified it can trigger actions to be performed, for example, generate a new message that contains all the important information that was stored previously in the context. (For details on triggering actions and generating messages, see *Section 13.4, Triggering actions for identified messages (p. 283)*.)

There are two attributes for pattern database rules that determine if a message matching the rule is added to a context: `context-scope` and `context-id`. The `context-scope` attribute acts as an early filter, selecting messages sent by the same process (${HOST}${PROGRAM}${PID} is identical), application (${HOST}${PROGRAM} is identical), or host, while the `context-id` actually adds the message to the context specified in the id. The `context-id` can be a simple string, or can contain macros or values extracted from the log messages for further filtering. Starting with syslog-ng OSE version 3.5, if a message is added to a context, syslog-ng OSE automatically adds the identifier of the context to the `.classifier.context_id` macro of the message.

Note
Message contexts are persistent and are not lost when syslog-ng OSE is reloaded (SIGHUP), but are lost when syslog-ng OSE is restarted.

Another parameter of a rule is the `context-timeout` attribute, which determines how long a context is stored, that is, how long syslog-ng OSE waits for related messages to arrive. Note the following points about timeout values:

- When a new message is added to a context, syslog-ng OSE will restart the timeout using the `context-timeout` set for the new message.

- When calculating if the timeout has already expired or not, syslog-ng OSE uses the timestamps of the incoming messages, not system time elapsed between receiving the two messages (unless the messages do not include a timestamp, or the `keep-timestamp(no)` option is set). That way syslog-ng OSE can be used to process and correlate already existing log messages offline. However, the timestamps of the messages must be in chronological order (that is, a new message cannot be older than the one already processed), and if a message is newer than the current system time (that is, it seems to be coming from the future), syslog-ng OSE will replace its timestamp with the current system time.

Example 13.5. How syslog-ng OSE calculates `context-timeout`
Consider the following two messages:

```
<38>1990-01-01T14:45:25 customhostname program6[1234]: program6 testmessage
<38>1990-01-01T14:46:25 customhostname program6[1234]: program6 testmessage
```

If the `context-timeout` is 10 seconds and syslog-ng OSE receives the messages within 1 sec, the timeout event will occour immediately, because the difference of the two timestamp (60 sec) is larger than the timeout value (10 sec).

- Avoid using unnecessarily long timeout values on high-traffic systems, as storing the contexts for many messages can require considerable memory. For example, if two related messages usually arrive within seconds, it is not needed to set the timeout to several hours.

Example 13.6. Using message correlation

```
<rule xml:id="..." context-id="ssh-session" context-timeout="86400" context-scope="process">

    <patterns>
        <pattern>Accepted @ESTRING:usracct.authmethod: @for @ESTRING:usracct.username:
@from @ESTRING:usracct.device: @port @ESTRING:: @@ANYSTRING:usracct.service@</pattern>
    </patterns>
...
</rule>
```

For details on configuring message correlation, see the description of the *context-id*, *context-timeout*, and *context-scope* attributes of pattern database rules.

13.3.1. Referencing earlier messages of the context

When using the `<value>` element in pattern database rules together with message correlation, you can also refer to fields and values of earlier messages of the context by adding the `@<distance-of-referenced-message-from-the-current>` suffix to the macro. For example, if there are three log messages in a context, and you are creating a generated message for the third log message, the `${HOST}@1` expression refers to the host field of the current (third) message in the context, the `${HOST}@2` expression refers to the host field of the previous (second) message in the context, `${PID}@3` to the PID of the first message, and so on. For example, the following message can be created from SSH login/logout messages (for details on generating new messages, see *Section 13.4, Triggering actions for identified messages (p. 283)*):
`An SSH session for ${SSH_USERNAME}@1 from ${SSH_CLIENT_ADDRESS}@2 closed. Session lasted from ${DATE}@2 to ${DATE}.`

Warning
When referencing an earlier message of the context, always enclose the field name between braces, for example, `${PID}@3`. The reference will not work if you omit the braces.

Note
To use a literal @ character in a template, use @@.

13.4. Triggering actions for identified messages

Starting with version 3.2, the syslog-ng OSE application is able to generate (trigger) messages automatically if certain events occur, for example, a specific log message is received, or the correlation timeout of a message expires. Basically, you can define messages for every pattern database rule that are emitted when a message matching the rule is received. Triggering messages is often used together with message correlation, but can also be used separately.

The generated message is injected into the same place where the *db-parser()* statement is referenced in the log path. To post the generated message into the *internal()* source instead, use the *inject-mode()* option in the definition of the parser.

Example 13.7. Sending triggered messages to the *internal()* source
To send the generated messages to the *internal* source, use the inject-mode(internal) option:

```
parser p_db {db-parser(
    file("mypatterndbfile.xml")
    inject-mode(internal)
);};
```

To inject the generated messages where the pattern database is referenced, use the inject-mode(pass-through) option:

```
parser p_db {db-parser(
    file("mypatterndbfile.xml")
    inject-mode(pass-through)
);};
```

Note
Version 3.2 of syslog-ng OSE was able to send the generated messages only to the *internal* source.

The generated message must be configured in the pattern database rule. It is possible to create an entire message, use macros and values extracted from the original message with pattern database, and so on.

Example 13.8. Generating messages for pattern database matches
When inserted in a pattern database rule, the following example generates a message when a message matching the rule is received.

```
<actions>
    <action>
        <message>
```

```
                <values>
                    <value name="MESSAGE">A log message from ${HOST} matched rule number
$.classifier.rule_id</value>
                </values>
            </message>
        </action>
</actions>
```

To inherit the properties and values of the triggering message, set the *inherit-properties* attribute of the *<message>* element to TRUE. That way the triggering log message is cloned, including name-value pairs and tags. If you set any values for the message in the *<action>* element, they will override the values of the original message.

Example 13.9. Generating messages with inherited values
The following action generates a message that is identical to the original message, but its $PROGRAM field is set to *overriding-original-program-name*

```
<actions>
    <action>
        <message inherit-properties='TRUE'>
            <values>
                <value name="PROGRAM">overriding-original-program-name</value>
            </values>
        </message>
    </action>
</actions>
```

For details on configuring actions, see the description of the *pattern database format*.

13.4.1. Conditional actions

To limit when a message is triggered, use the *condition* attribute and specify a filter expression: the action will be executed only if the condition is met. For example, the following action is executed only if the message was sent by the host called myhost.

```
<action condition="'${HOST}' == 'example'">
```

You can use the same operators in the condition that can be used in filters. For details, see *Section 8.3.3, Comparing macro values in filters (p. 209)*.

The following action can be used to log the length of an SSH session (the time difference between a login and a logout message in the context):

```
<actions>
    <action>
        <message>
            <values>
                <value name="MESSAGE">An SSH session for ${SSH_USERNAME}@1 from
${SSH_CLIENT_ADDRESS}@2 closed. Session lasted from ${DATE}@2 ${DATE} </value>
            </values>
        </message>
    </action>
</actions>
```

Example 13.10. Actions based on the number of messages
The following example triggers different actions based on the number of messages in the context. This way you can check if the context contains enough messages for the event to be complete, and execute a different action if it does not.

```
<actions>
    <action condition='"$(context-length)" >= "4"'>
        <message>
            <values>
                <value name="PROGRAM">event</value>
                <value name="MESSAGE">Event complete</value>
            </values>
        </message>
    </action>
    <action condition='"$(context-length)" < "4"'>
        <message>
            <values>
                <value name="PROGRAM">error</value>
                <value name="MESSAGE">Error detected</value>
            </values>
        </message>
    </action>
</actions>
```

13.4.2. External actions

To perform an external action when a message is triggered, for example, to send the message in an e-mail, you have to route the generated messages to an external application using the *program()* destination.

Example 13.11. Sending triggered messages to external applications
The following sample configuration selects the triggered messages and sends them to an external script.

1. Set a field in the triggered message that is easy to identify and filter. For example:

   ```
   <values>
       <value name="MESSAGE">A log message from ${HOST} matched rule number
   $.classifier.rule_id</value>
       <value name="TRIGGER">yes</value>
   </values>
   ```

2. Create a destination that will process the triggered messages.

   ```
   destination d_triggers { program("/bin/myscript"; ); };
   ```

3. Create a filter that selects the triggered messages from the internal source.

   ```
   filter f_triggers {match("yes" value ("TRIGGER") type(string));};
   ```

4. Create a logpath that selects the triggered messages from the internal source and sends them to the script:

   ```
   log { source(s_local); filter(f_triggers); destination(d_triggers); };
   ```

5. Create a script that will actually process the generated messages, for example:

   ```
   #!/usr/bin/perl
   while (<>) {
           # body of the script to send emails, snmp traps, and so on
   }
   ```

13.4.3. Actions and message correlation

Certain features of generating messages can be used only if message correlation is used as well.

- The syslog-ng OSE application automatically fills the fields for the generated message based on the scope of the context, for example, the HOST and PROGRAM fields if the *context-scope* is program.

- When used together with message correlation, you can also refer to fields and values of earlier messages of the context by adding the @<distance-of-referenced-message-from-the-current> suffix to the macro. For details, see *Section 13.3.1, Referencing earlier messages of the context (p. 282)*.

- It is possible to generate a message when the *context-timeout* of the original message expires and no new message is added to the context during this time. To accomplish this, include the trigger="timeout" attribute in the action element:

```
<action trigger="timeout">
```

For example, the following action can be used to log the length of an SSH session (the time difference between a login and a logout message in the context):

```
<actions>
    <action>
        <message>
            <values>
                <value name="MESSAGE">An SSH session for ${SSH_USERNAME}@1 from
${SSH_CLIENT_ADDRESS}@2 closed. Session lasted from ${DATE}@2 ${DATE} </value>
            </values>
        </message>
    </action>
</actions>
```

For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

13.5. Creating pattern databases

13.5.1. Using pattern parsers

Pattern parsers attempt to parse a part of the message using rules specific to the type of the parser. Parsers are enclosed between @ characters. The syntax of parsers is the following:

- a beginning @ character;
- the type of the parser written in capitals;
- optionally a name;
- parameters of the parser, if any;
- a closing @ character.

Example 13.12. Pattern parser syntax
A simple parser:

```
@STRING@
```

A named parser:

```
@STRING:myparser_name@
```

A named parser with a parameter:

```
@STRING:myparser_name:*@
```

A parser with a parameter, but without a name:

```
@STRING::*@
```

Patterns and literals can be mixed together. For example, to parse a message that begins with the `Host:` string followed by an IP address (for example, `Host: 192.168.1.1`), the following pattern can be used: `Host:@IPv4@`.

Note
Note that using parsers is a CPU-intensive operation. Use the ESTRING and QSTRING parsers whenever possible, as these can be processed much faster than the other parsers.

Example 13.13. Using the STRING and ESTRING parsers
For example, if the message is `user=joe96 group=somegroup`, `@STRING:mytext:@` parses only to the first non-alphanumeric character (=), parsing only `user`. `@STRING:mytext:=@` parses the equation mark as well, and proceeds to the next non-alphanumeric character (the whitespace), resulting in `user=joe96` being parsed. `@STRING:mytext:= @` will parse the whitespace as well, and proceed to the next non-alphanumeric non-equation mark non-whitespace character, resulting in `user=joe96 group=somegroup`.

Of course, usually it is better to parse the different values separately, like this: `"user=@STRING:user@ group=@STRING:group@"`.

If the username or the group may contain non-alphanumeric characters, you can either include these in the second parameter of the parser (as shown at the beginning of this example), or use an ESTRING parser to parse the message till the next whitespace: `"user=@ESTRING:user: @group=@ESTRING:group: @"`.

13.5.1.1. Pattern parsers of syslog-ng OSE

The following parsers are available in syslog-ng OSE.

@ANYSTRING@
Parses everything to the end of the message; you can use it to collect everything that is not parsed specifically to a single macro. In that sense its behavior is similar to the *greedy()* option of the CSV parser.

@DOUBLE@
An obsolete alias of the *@FLOAT@* parser.

@EMAIL@
This parser matches an e-mail address. The parameter is a set of characters to strip from the beginning and the end of the e-mail address. That way e-mail addresses enclosed between other characters can be matched easily (for example, *<user@example.com>* or *"user@example.com"*. Characters that are valid for a hostname are not stripped from the end of the hostname. This includes a trailing period if present.

For example, the *@EMAIL:email:"[<]>@* parser will match any of the following e-mail addresses: *<user@example.com>*, *[user@example.com]*, *"user@example.com"*, and set the value of the *email* macro to *user@example.com*.

@ESTRING@

This parser has a required parameter that acts as the stopcharacter: the parser parses everything until it finds the stopcharacter. For example, to stop by the next " (double quote) character, use @ESTRING::"@. You can use the colon (:) as stopcharacter as well, for example: @ESTRING:::@. You can also specify a stopstring instead of a single character, for example, @ESTRING::stop_here.@. The @ character cannot be a stopcharacter, nor can line-breaks or tabs.

@FLOAT@

A floating-point number that may contain a dot (.) character. (Up to syslog-ng 3.1, the name of this parser was *@DOUBLE@*.)

@HOSTNAME@

Parses a generic hostname. The hostname may contain only alphanumeric characters (A-Z,a-z,0-9), hypen (-), or dot (.).

@IPv4@

Parses an IPv4 IP address (numbers separated with a maximum of 3 dots).

@IPv6@

Parses any valid IPv6 IP address.

@IPvANY@

Parses any IP address.

@LLADDR@

Parses a Link Layer Address in the *xx:xx:xx:...* form, where each xx is a 2 digit HEX number (an octet). The parameter specifies the maximum number of octets to match and defaults to 20. The MACADDR parser is a special wrapper using the LLADDR parser. For example, the following parser parses maximally 10 octets, and stores the results in the *link-level-address* macro:

```
@LLADDR:link-level-address:10@
```

@MACADDR@

Parses the standard format of a MAC-48 address, consisting of is six groups of two hexadecimal digits, separated by colons. For example, *00:50:fc:e3:cd:37*.

@NUMBER@

A sequence of decimal (0-9) numbers (for example, 1, 0687, and so on). Note that if the number starts with the 0x characters, it is parsed as a hexadecimal number, but only if at least one valid character follows 0x. A leading hyphen (–) is accepted for non-hexadecimal numbers, but other separator characters (for example, dot or comma) are not. To parse floating-point numbers, use the @FLOAT@ parser.

@PCRE@

Use Perl-Compatible Regular Expressions (as implemented by the PCRE library), after the identification of the potential patterns has happened by the radix implementation.

Syntax: *@PCRE:name:regexp@*

@QSTRING@

Parse a string between the quote characters specified as parameter. Note that the quote character can be different at the beginning and the end of the quote, for example: @QSTRING::"@ parses everything between two quotation marks ("), while @QSTRING:<>@ parses from an opening bracket to the closing bracket. The @ character cannot be a quote character, nor can line-breaks or tabs.

@SET@

Parse any combination of the specified characters until another character is found. For example, specifying a whitespace character parses any number of whitespaces, and can be used to process paddings (for example, log messages of the Squid application have whitespace padding after the username).

For example, the *@SET:: "@* parser will parse any combination of whitespaces and double-quotes.

Available in syslog-ng OSE 3.4 and later.

@STRING@

A sequence of alphanumeric characters (0-9, A-z), not including any whitespace. Optionally, other accepted characters can be listed as parameters (for example, to parse a complete sentence, add the whitespace as parameter, like: @STRING:: @). Note that the @ character cannot be a parameter, nor can line-breaks or tabs.

13.5.2. What's new in the syslog-ng pattern database format V4

The V4 database format has the following differences compared to the V3 format:

- It is now possible to specify multiple program patterns for a ruleset. For details, see the description of the <u>*patterns tag*</u>.

- The <value> element of name-value pairs can include template functions. For details, see *Section 11.1.6, Using template functions (p. 248)*, for examples, see *Section if (p. 252)*.

- It is now possible to correlate log messages processed with the pattern database. For details, see *Section 13.3, Correlating log messages (p. 281)*.

- It is now possible to generate new messages based on pattern matching and correlation results. For details, see *Section 13.4, Triggering actions for identified messages (p. 283)* and the description of the <u>*actions tag*</u>.

13.5.3. The syslog-ng pattern database format

Pattern databases are XML files that contain rules describing the message patterns. For sample pattern databases, see *Section 13.2.2, Downloading sample pattern databases (p. 280)*.

The following scheme describes the V4 format of the pattern database. This format is used by syslog-ng OSE 3.2 and later, and is backwards-compatible with the earlier V3 format.

For a sample database containing only a single pattern, see *Example 13.14, A V4 pattern database containing a single rule (p. 295)*.

> **Tip**
>
> Use the `pdbtool` utility that is bundled with syslog-ng to test message patterns and convert existing databases to the latest format. For details, see *pdbtool(1) (p. 315)*.
>
> To automatically create an initial pattern database from an existing log file, use the `pdbtool patternize` command. For details, see *the section called "The patternize command" (p. 318)*.

- **\<patterndb\>**: The container element of the pattern database. For example:

  ```
  <patterndb version='4' pub_date='2010-10-25'>
  ```

- *version*: The schema version of the pattern database. The current version is 4.
- *pubdate*: The publication date of the XML file.
- **\<ruleset\>**: A container element to group log patterns for an application or program. For example:

  ```
  <ruleset name='su' id='480de478-d4a6-4a7f-bea4-0c0245d361e1'>
  ```

  A `<patterndb>` element may contain any number of `<ruleset>` elements.

  - *name*: The name of the application. Note that the function of this attribute is to make the database more readable, syslog-ng uses the `<pattern>` element to identify the applications sending log messages.
  - *id*: A unique ID of the application, for example, the md5 sum of the `name` attribute.
  - **description**: OPTIONAL — A description of the ruleset or the application.
  - **url**: OPTIONAL — An URL referring to further information about the ruleset or the application.
  - **\<patterns\>**: A container element storing program names also called `program pattern`. For example:

    ```
    <patterns>
        <pattern>su</pattern>
    </patterns>
    ```

    A `<patterns>` element may contain any number of `<pattern>` elements.

    - **pattern**: The name of the application — syslog-ng matches this value to the ${PROGRAM} header of the syslog message to find the rulesets applicable to the syslog message.

      Specifying multiple patterns is useful if two or more applications have different names (that is, different ${PROGRAM} fields), but otherwise send identical log messages.

      ```
      <patterns>
          <pattern>firstapplication</pattern>
          <pattern>otherapplication</pattern>
      </patterns>
      ```

It is not necessary to use multiple patterns if only the end of the ${PROGRAM} fields is different, use only the beginning of the ${PROGRAM} field as the `pattern`. For example, the Postfix e-mail server sends messages using different process names, but all of them begin with the `postfix` string.

You can also use parsers in the program pattern if needed, and use the parsed results later. For example: *<pattern>postfix\@ESTRING:.postfix.component:[@</pattern>*

> **Note**
> If the `<pattern>` element of a ruleset is not specified, syslog-ng OSE will use this ruleset as a fallback ruleset: it will apply the ruleset to messages that have an empty PROGRAM header, or if none of the program patterns matched the PROGRAM header of the incoming message.

- **<rules>**: A container element for the rules of the ruleset.
  - **<rule>**: An element containing message patterns and how a message that matches these patterns is classified. For example:

    ```
    <rule provider='balabit' id='f57196aa-75fd-11dd-9bba-001e6806451b'
    class='violation'>
    ```

    The following example specifies attributes for correlating messages as well. For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

    ```
    <rule provider='balabit' id='f57196aa-75fd-11dd-9bba-001e6806451b'
    class='violation' context-id='same-session' context-scope='process'
    context-timeout='360'>
    ```

    > **Note**
    > If the following characters appear in the message, they must be escaped in the rule as follows:
    > - @: Use @@, for example `user@@example.com`
    > - <: Use <
    > - >: Use >
    > - &: Use &

    The **<rules>** element may contain any number of **<rule>** elements.

- *provider*: The provider of the rule. This is used to distinguish between who supplied the rule; that is, if it has been created by BalaBit, or added to the XML by a local user.
- *id*: The globally unique ID of the rule.
- *class*: The class of the rule — syslog-ng assigns this class to the messages matching a pattern of this rule.
- *context-id*: OPTIONAL — An identifier to group related log messages when using the pattern database to correlate events. The ID can be a descriptive string describing the events related to the log message (for example, `ssh-sessions` for log messages related to SSH traffic), but can also contain macros to generate IDs dynamically. When using macros in IDs, see also the

context-scope attribute. Starting with syslog-ng OSE version 3.5, if a message is added to a context, syslog-ng OSE automatically adds the identifier of the context to the *.classifier.context_id* macro of the message. For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

> **Note**
> The syslog-ng OSE application determines the context of the message *after* the pattern matching is completed. This means that macros and name-value pairs created by the matching pattern database rule can be used as context-id macros.

- *context-timeout*: OPTIONAL — The number of seconds the context is stored. Note that for high-traffic logservers, storing open contexts for long time can require significant amount of memory. For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

- *context-scope*: OPTIONAL — Specifies which messages belong to the same context. This attribute is used to determine the context of the message if the *context-id* does not specify any macros. Usually, *context-scope* acts a filter for the context, with *context-id* refining the filtering if needed. The *context-scope* attribute has the following possible values:

  - *process*: Only messages that are generated by the same process of a client belong to the same context, that is, messages that have identical ${HOST}, ${PROGRAM} and ${PID} values. This is the default behavior of syslog-ng OSE if *context-scope* is not specified.

  - *program*: Messages that are generated by the same application of a client belong to the same context, that is, messages that have identical ${HOST} and ${PROGRAM} values.

  - *host*: Every message generated by a client belongs to the same context, only the ${HOST} value of the messages must be identical.

  - *global*: Every message belongs to the same context.

> **Note**
> Using the *context-scope* attribute is significantly faster than using macros in the *context-id* attribute.

For details on correlating messages, see *Section 13.3, Correlating log messages (p. 281)*.

- **<patterns>**: An element containing the patterns of the rule. If a **<patterns>** element contains multiple **<pattern>** elements, the class of the **<rule>** is assigned to every syslog message matching any of the patterns.

  - **<pattern>**: A pattern describing a log message. This element is also called `message pattern`. For example:

    ```
    <pattern>+ ??? root-</pattern>
    ```

> **Note**
> Support for XML entities is limited, you can use only the following entities: & < > " '. User-defined entities are not supported.

- **description**: OPTIONAL — A description of the pattern or the log message matching the pattern.

- **urls**: OPTIONAL — An element containing one or more URLs referring to further information about the patterns or the matching log messages.

  - **url**: OPTIONAL — An URL referring to further information about the patterns or the matching log messages.

- **values**: OPTIONAL — Name-value pairs that are assigned to messages matching the patterns, for example, the representation of the event in the message according to the Common Event Format (CEF) or Common Event Exchange (CEE). The names can be used as macros to reference the assigned values.

  - **value**: OPTIONAL — Contains the value of the name-value pair that is assigned to the message. For example:

    ```
    <value name=".classifier.outcome">/Success</value>
    ```

    The <value> element of name-value pairs can include template functions. For details, see *Section 11.1.6, Using template functions (p. 248)*, for examples, see *Section if (p. 252)*.

    When used together with message correlation, the <value> element of name-value pairs can include references to the values of earlier messages from the same context. For details, see *Section 13.3, Correlating log messages (p. 281)*.

  - *name*: The name of the name-value pair. It can also be used as a macro to reference the assigned value.

- **examples**: OPTIONAL — A container element for sample log messages that should be recognized by the pattern. These messages can be used also to test the patterns and the parsers.

  - **example**: OPTIONAL — A container element for a sample log message.

    - **test_message**: OPTIONAL — A sample log message that should match this pattern. For example:

      ```
      <test_message program="myapplication">Content filter has been
      enabled</test_message>
      ```

      - *program*: The program pattern of the test message. For example:

        ```
        <test_message program="proftpd">ubuntu
        (::ffff:192.168.2.179[::ffff:192.168.2.179]) - FTP session
        closed.</test_message>
        ```

    - **test_values**: OPTIONAL — A container element to test the results of the parsers used in the pattern.

- **test_value**: OPTIONAL — The expected value of the parser when matching the pattern to the test message. For example:

  ```
  <test_value name=".dict.ContentFilter">enabled</test_value>
  ```

  - *name*: The name of the parser to test.

- **actions**: OPTIONAL — A container element for actions that are performed if a message is recognized by the pattern. For details on actions, see *Section 13.4, Triggering actions for identified messages (p. 283)*.

  - **action**: OPTIONAL — A container element describing an action that is performed when a message matching the rule is received.

    - *condition*: A syslog-ng filter expression. The action is performed only if the message matches the filter. The filter can include macros and name-value pairs extracted from the message. When using actions together with message-correlation, you can also use the `$(context-length)` macro, which returns the number of messages in the current context. For example, this can be used to determine if the expected number of messages has arrived to the context: `condition='"$(context-length)" >= "5"'`

    - *rate*: Specifies maximum how many messages should be generated in the specified time period in the following format: `<number-of-messages>/<period-in-seconds>`. For example: `1/60` allows 1 message per minute. Rates apply within the scope of the context, that is, if `context-scope="host"` and `rate="1/60"`, then maximum one message is generated per minute for every host that sends a log message matching the rule. Excess messages are dropped. Note that when applying the rate to the generated messages, syslog-ng OSE uses the timestamps of the log messages, similarly to calculating the `context-timeout`. That way `rate` is applied correctly even if the log messages are processed offline.

    - *trigger*: Specifies when the action is executed. The `trigger` attribute has the following possible values:

      - *match*: Execute the action immediately when a message matching the rule is received.

      - *timeout*: Execute the action when the correlation timer (`context-timeout`) expires. This is available only if actions are used together with correlating messages.

    - **message**: A container element storing the message to be sent when the action is executed. Currently syslog-ng OSE sends these messages to the `internal()` destination.

      - *inherit-properties*: If set to TRUE, the original message that triggered the action is cloned, including its name-value pairs and tags. For details, see *Section 13.4, Triggering actions for identified messages (p. 283)*.

    - **values**: A container element for values and fields that are used to create the message generated b the action.

      - **value**: Sets the value of the message field specified in the *name* attribute of the element. For example, to specify the body of the generated message, use the following syntax:

        ```
        <value name="MESSAGE">A log message matched rule number
        $.classifier.rule_id</value>
        ```

Note that currently it is not possible to add DATE, FACILITY, or SEVERITY fields to the message.

When the action is used together with message correlation, the syslog-ng OSE application automatically adds fields to the message based on the `context-scope` parameter. For example, using `context-scope="process"` automatically fills the HOST, PROGRAM, and PID fields of the generated message.

- *name*: Name of the message field set by the `value` element.

- **tags**: OPTIONAL — An element containing custom keywords (tags) about the messages matching the patterns. The tags can be used to label specific events (for example user logons). It is also possible to filter on these tags later (for details, see *Section 8.3.5, Tagging messages (p. 211)*). Starting with syslog-ng Open Source Edition 3.2, the list of tags assigned to a message can be referenced with the `${TAGS}` macro.

  - **tag**: OPTIONAL — A keyword or tags applied to messages matching the rule. For example:

```
<tags><tag>UserLogin</tag></tags>
```

Example 13.14. A V4 pattern database containing a single rule
The following pattern database contains a single rule that matches a log message of the `ssh` application. A sample log message looks like:

```
Accepted password for sampleuser from 10.50.0.247 port 42156 ssh2
```

The following is a simple pattern database containing a matching rule.

```
<patterndb version='4' pub_date='2010-10-17'>
    <ruleset name='ssh' id='123456678'>
        <pattern>ssh</pattern>
            <rules>
                <rule provider='me' id='182437592347598' class='system'>
                    <patterns>
                        <pattern>Accepted @QSTRING:SSH.AUTH_METHOD: @
for@QSTRING:SSH_USERNAME: @from\ @QSTRING:SSH_CLIENT_ADDRESS: @port @NUMBER:SSH_PORT_NUMBER:@
 ssh2</pattern>
                    </patterns>
                </rule>
            </rules>
    </ruleset>
</patterndb>
```

Note that the rule uses macros that refer to parts of the message, for example, you can use the `${SSH_USERNAME}` macro refer to the username used in the connection.

The following is the same example, but with a test message and test values for the parsers.

```
<patterndb version='4' pub_date='2010-10-17'>
    <ruleset name='ssh' id='123456678'>
        <pattern>ssh</pattern>
            <rules>
                <rule provider='me' id='182437592347598' class='system'>
                    <patterns>
                        <pattern>Accepted @QSTRING:SSH.AUTH_METHOD: @
for@QSTRING:SSH_USERNAME: @from\ @QSTRING:SSH_CLIENT_ADDRESS: @port @NUMBER:SSH_PORT_NUMBER:@
 ssh2</pattern>
                    </patterns>
                    <examples>
                        <example>
                          <test_message>Accepted password for sampleuser from 10.50.0.247
 port 42156 ssh2</test_message>
```

```
                              <test_values>
                                      <test_value name="SSH.AUTH_METHOD">password</test_value>
                                      <test_value name="SSH_USERNAME">sampleuser</test_value>
                                      <test_value
name="SSH_CLIENT_ADDRESS">10.50.0.247</test_value>
                                      <test_value name="SSH_PORT_NUMBER">42156</test_value>
                              </test_values>
                      </example>
                  </examples>
              </rule>
          </rules>
    </ruleset>
</patterndb>
```

Chapter 14. Statistics of syslog-ng

Periodically, syslog-ng sends a message containing statistics about the received messages, and about any lost messages since the last such message. It includes a *processed* entry for every source and destination, listing the number of messages received or sent, and a *dropped* entry including the IP address of the server for every destination where syslog-ng has lost messages. The center(received) entry shows the total number of messages received from every configured sources.

The following is a sample log statistics message for a configuration that has a single source (s_local) and a network and a local file destination (d_network and d_local, respectively). All incoming messages are sent to both destinations.

```
Log statistics;
        dropped='tcp(AF_INET(192.168.10.1:514))=6439',
        processed='center(received)=234413',
        processed='destination(d_tcp)=234413',
        processed='destination(d_local)=234413',
        processed='source(s_local)=234413'
```

Log statistics can be also retrieved on-demand using one of the following options:

- Use the socat application: echo STATS | socat -vv UNIX-CONNECT:/opt/syslog-ng/var/run/syslog-ng.ctl -

- If you have an OpenBSD-style netcat application installed, use the echo STATS | nc -U /opt/syslog-ng/var/run/syslog-ng.ctl command. Note that the netcat included in most Linux distributions is a GNU-style version that is not suitable to query the statistics of syslog-ng.

- Starting from syslog-ng Open Source Edition version 3.1, syslog-ng Open Source Edition includes the syslog-ng-ctl utility. Use the syslog-ng-ctl stats command.

The statistics include a list of source groups and destinations, as well as the number of processed messages for each. The verbosity of the statistics can be set using the *stats-level()* option. For details, see *Section 9.2, Global options (p. 217)*. An example output is shown below.

```
src.internal;s_all#0;;a;processed;6445
src.internal;s_all#0;;a;stamp;1268989330
destination;df_auth;;a;processed;404
destination;df_news_dot_notice;;a;processed;0
destination;df_news_dot_err;;a;processed;0
destination;d_ssb;;a;processed;7128
destination;df_uucp;;a;processed;0
source;s_all;;a;processed;7128
destination;df_mail;;a;processed;0
destination;df_user;;a;processed;1
destination;df_daemon;;a;processed;1
destination;df_debug;;a;processed;15
destination;df_messages;;a;processed;54
destination;dp_xconsole;;a;processed;671
dst.tcp;d_network#0;10.50.0.111:514;a;dropped;5080
dst.tcp;d_network#0;10.50.0.111:514;a;processed;7128
```

```
dst.tcp;d_network#0;10.50.0.111:514;a;stored;2048
destination;df_syslog;;a;processed;6724
destination;df_facility_dot_warn;;a;processed;0
destination;df_news_dot_crit;;a;processed;0
destination;df_lpr;;a;processed;0
destination;du_all;;a;processed;0
destination;df_facility_dot_info;;a;processed;0
center;;received;a;processed;0
destination;df_kern;;a;processed;70
center;;queued;a;processed;0
destination;df_facility_dot_err;;a;processed;0
```

The statistics are semicolon separated; every line contains statistics for a particular object (for example source, destination, tag, and so on). The statistics have the following fields:

1. The type of the object (for example `dst.file`, `tag`, `src.facility`)

2. The ID of the object used in the syslog-ng configuration file, for example `d_internal` or `source.src_tcp`. The #0 part means that this is the first destination in the destination group.

3. The instance ID (destination) of the object, for example the filename of a file destination, or the name of the application for a program source or destination.

4. The status of the object. One of the following:

   - `a` - active. At the time of quering the statistics, the source or the destination was still alive (it continuously received statistical data).

   - `d` - dynamic. Such objects may not be continuously available, for example, like statistics based on the sender's hostname.

   - `o` - This object was once active, but stopped receiving messages. (For example a dynamic object may disappear and become orphan.)

   > **Note**
   > The syslog-ng OSE application stores the statistics of the objects when syslog-ng OSE is reloaded. However, if the configuration of syslog-ng OSE was changed since the last reload, the statistics of orphaned objects are deleted.

5. The type of the statistics:

   - *processed*: The number of messages that successfully reached their destination driver.

   - *dropped*: The number of dropped messages — syslog-ng OSE could not send the messages to the destination and the output buffer got full, so messages were dropped by the destination driver.

   - *stored*: The number of messages stored in the message queue of the destination driver, waiting to be sent to the destination.

   - *suppressed*: The number of suppressed messages (if the *suppress()* feature is enabled).

   - *stamp*: The UNIX timestamp of the last message sent to the destination.

6. The number of such messages.

Note
Certain statistics are available only if the *stats-level()* option is set to a higher value.

When receiving messages with non-standard facility values (that is, higher than 23), these messages will be listed as *other* facility instead of their facility number.

Chapter 15. Multithreading and scaling in syslog-ng OSE

Starting with version 3.3, syslog-ng OSE can be run in multithreaded mode to scale to multiple CPUs or cores for increased performance.

> **Note**
> Starting with version 3.6, syslog-ng OSE runs in multithreaded mode by default.

15.1. Multithreading concepts of syslog-ng OSE

This section is a brief overview on how syslog-ng OSE works in multithreaded mode. It is mainly for illustration purposes: the concept has been somewhat simplified and may not completely match reality.

> **Note**
> The way syslog-ng OSE uses multithreading may change in future releases. The current documentation applies to version 3.6.

syslog-ng OSE has a main thread that is always running, and a number of worker threads that process the messages. The maximum number of worker threads syslog-ng OSE uses is the number of CPUs or cores in the host running syslog-ng OSE (up to 64) but can be limited using the `--worker-threads` command-line option.

> **Note**
> The `--worker-threads` command-line option sets the maximum total number of threads syslog-ng OSE can use, including the main syslog-ng OSE thread.

When an event requiring a new thread occurs (for example, syslog-ng OSE receives new messages, or a destination becomes available), syslog-ng OSE tries to start a new thread. If there are no free threads, the task waits until a thread finishes its task and becomes available. There are two types of worker threads:

- Reader threads read messages from a source (as many as possible, but limited by the `log-fetch-limit()` and `log-iw-size()` options. The thread then processes these messages, that is, performs filtering, rewriting and other tasks as necessary, and puts the log message into the queue of the destination. If the destination does not have a queue (for example, usertty), the reader thread sends the message to the destination, without the interaction of a separate writer thread.

- Writer threads take the messages from the queue of the destination and send them to the destination, that is, write the messages into a file, or send them to the syslog server over the network. The writer thread starts to process messages from the queue only if the destination is writable, and there are

enough messages in the queue, as set in the *flush-lines()* and the *flush-timeout()* options. Writer threads stop processing messages when the destination becomes unavailable, or there are no more messages in the queue.

The following list describes which sources and destinations can use multiple threads.

- The *tcp* and syslog(tcp) sources can process independent connections in separate threads. The number of independent connections is limited by the *max-connections()* option of the source. Separate sources are processed by separate thread, for example, if you have two separate *tcp* sources defined that receive messages on different IP addresses or port, syslog-ng OSE will use separate threads for these sources even if they both have only a single active connection.

- The *udp*, *file*, and *pipe* sources use a single thread for every source statement.

- The *tcp*, *syslog*, and *pipe* destinations use a single thread for every destination.

- The *file* destination uses a single thread for writing the destination file, but may use a separate thread for each destination file if the filename includes macros.

- Every *sql* destination uses its own thread. These threads are independent from the setting of the *--worker-threads* command-line option.

15.2. Configuring multithreading

Starting with version 3.6, syslog-ng OSE runs in multithreaded mode by default.

Multithreading in syslog-ng OSE can be enabled using the following methods:

- Globally using the threaded(yes) option.
- Separately for selected sources or destinations using the flags("threaded") option.

Example 15.1. Enabling multithreading
To enable multithreading globally, use the *threaded* option:

```
options {threaded(yes) ; };
```

To enable multithreading only for a selected source or destination, use the flags("threaded") option:

```
source s_tcp_syslog { tcp(ip(127.0.0.1) port(1999) flags("syslog-protocol", "threaded")
); };
```

15.3. Optimizing multithreaded performance

Destinations that have a queue process that queue in a single thread. Multiple sources can send messages to the same queue, so the queue can scale to multiple CPUs. However, when the writer thread writes the queue contents to the destination, it will be single-threaded.

Message parsing, rewrite rules, filters, and other types of message processing is performed by the reader thread in a sequential manner. This means that such operations can scale only if reading messages from the source

can be multithreaded. For example, if a *tcp* source can process messages from different connections (clients) in separate threads. If the source cannot use multiple threads to process the messages, the operations will not scale.

To improve the processing power of syslog-ng OSE and scale to more processors, use the following methods:

- To improve scaling on the source side, use more sources, for example, more source files, or receive the messages from more parallel connections. For network sources, you can also configure a part of your clients to send the messages to a different port of your syslog-ng server, and use separate source definitions for each port.

- On the destination side, when writing the log messages to files, use macros in the filename to split the messages to separate files (for example, using the ${HOST} macro). Files with macros in their filenames are processed in separate writer threads.

- On the destination side, when sending messages to a syslog-ng server, you can use multiple connections to the server if you configure the syslog-ng server to receive messages on multiple ports, and configure the clients to use both ports.

Chapter 16. Troubleshooting syslog-ng

This chapter provides tips and guidelines about troubleshooting problems related to syslog-ng.

Tip
As a general rule, first try to get logging the messages to a local file. Once this is working, you know that syslog-ng is running correctly and receiving messages, and you can proceed to forwarding the messages to the server.

If the syslog-ng server does not receive the messages, use tcpdump or a similar packet sniffer tool on the client to verify that the messages are sent correctly, and on the server to verify that it receives the messages.

If syslog-ng is closing the connections for no apparent reason, be sure to check the log messages of syslog-ng. You might also want to run syslog-ng with the `--verbose` or `--debug` command-line options for more-detailed log messages. Starting from syslog-ng OSE version 3.1, you can enable these messages without restarting syslog-ng using the `syslog-ng-ctl verbose --set=on` command. For details, see the syslog-ng-ctl man page at *syslog-ng-ctl(1) (p. 332)*.

Similarly, build up encrypted connections step-by-step: first create a working unencrypted (for example TCP) connection, then add TLS encryption, and finally client authentication if needed.

16.1. Possible causes of losing log messages

During the course of a message from the sending application to the final destination of the message, there are a number of locations where a message may be lost, even though syslog-ng does its best to avoid message loss. Usually losing messages can be avoided with careful planning and proper configuration of syslog-ng and the hosts running syslog-ng. The following list shows the possible locations where messages may be lost, and provides methods to minimize the risk of losing messages.

Note
The following list covers the main possibilities of losing messages, but does not take into account the possible use of flow-control (for details, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*). This topic will be addressed in more detail in the future releases of this guide.

- *Between the application and the syslog-ng client*: Make sure to use an appropriate source to receive the logs from the application (for example from `/dev/log`). For example, use *unix-stream* instead of *unix-dgram* whenever possible.

- *When syslog-ng is sending messages*: If syslog-ng cannot send messages to the destination and the output buffer gets full, syslog-ng will drop messages. Use flags (flow-control) to avoid it (for details, see *Section 8.2.2, Configuring flow-control (p. 206)*). The number of dropped messages is displayed per destination in the log message statistics of syslog-ng (for details, see *Chapter 14, Statistics of syslog-ng (p. 297)*).

- *On the network*: When transferring messages using the UDP protocol, messages may be lost without any notice or feedback — such is the nature of the UDP protocol. Always use the TCP protocol to transfer messages over the network whenever possible.

- *In the socket receive buffer*: When transferring messages using the UDP protocol, the UDP datagram (that is, the message) that reaches the receiving host placed in a memory area called the socket

receive buffer. If the host receives more messages than it can process, this area overflows, and the kernel drops messages without letting syslog-ng know about it. Using TCP instead of UDP prevents this issue. If you must use the UDP protocol, increase the size of the receive buffer using the *so-rcvbuf()* option.

- *When syslog-ng is receiving messages*: The receiving syslog-ng (for example the syslog-ng server or relay) may drop messages if the fifo of the destination file gets full. The number of dropped messages is displayed per destination in the log message statistics of syslog-ng (for details, see *Chapter 14, Statistics of syslog-ng (p. 297)*).

- *When the destination cannot handle large load*: When syslog-ng is sending messages at a high rate into an SQL database, a file, or another destination, it is possible that the destination cannot handle the load, and processes the messages slowly. As a result, the buffers of syslog-ng fill up, syslog-ng cannot process the incoming messages, and starts to loose messages. For details, see the previous entry. Use the *throttle* parameter to avoid this problem.

- *As a result of an unclean shutdown of the syslog-ng server*: If the host running the syslog-ng server experiences an unclean shutdown, it takes time until the clients realize that the connection to the syslog-ng server is down. Messages that are put into the output TCP buffer of the clients during this period are not sent to the server.

16.2. Procedure – Creating syslog-ng core files

Purpose:

When syslog-ng crashes for some reason, it can create a core file that contains important troubleshooting information. To enable core files, complete the following procedure:

Steps:

Step 1. Core files are produced only if the `maximum core file size` ulimit is set to a high value in the init script of syslog-ng. Add the following line to the init script of syslog-ng:

```
ulimit -c unlimited
```

Step 2. Verify that syslog-ng has permissions to write the directory it is started from, for example `/opt/syslog-ng/sbin/`.

Step 3. If syslog-ng crashes, it will create a core file in the directory syslog-ng was started from.

Step 4. To test that syslog-ng can create a core file, you can create a crash manually. For this, determine the PID of syslog-ng (for example using the `ps -All|grep syslog-ng` command), then issue the following command: `kill -ABRT <syslog-ng pid>`
This should create a core file in the current working directory.

16.3. Collecting debugging information with strace, truss, or tusc

To properly troubleshoot certain situations, it can be useful to trace which system calls syslog-ng OSE performs. How this is performed depends on the platform running syslog-ng OSE. In general, note the following points:

- When syslog-ng OSE is started, a supervisor process might stay in the foreground, while the actual syslog-ng daemon goes to the background. Always trace the background process.

- Apart from the system calls, the time between two system calls can be important as well. Make sure that your tracing tool records the time information as well. For details on how to do that, refer to the manual page of your specific tool (for example, `strace` on Linux, or `truss` on Solaris and BSD).

- Run your tracing tool in verbose mode, and if possible, set it to print long output strings, so the messages are not truncated.

- When using `strace`, also record the output of `lsof` to see which files are accessed.

The following are examples for tracing system calls of syslog-ng on some platforms. The output is saved into the `/tmp/syslog-ng-trace.txt` file, suffced with the PID of the related syslog-ng process.The path of the syslog-ng binary assumes that you have installed syslog-ng OSE from the official syslog-ng OSE binaries available at the BalaBit website — native distribution-specific packages may use different paths.

- *Linux*: `strace -o /tmp/trace.txt -s256 -ff -ttT /opt/syslog-ng/sbin/syslog-ng -f /opt/syslog-ng/etc/syslog-ng.conf -Fdv`

- *HP-UX*: `tusc -f -o /tmp/syslog-ng-trace.txt -T /opt/syslog-ng/sbin/syslog-ng`

- *IBM AIX and Solaris*: `truss -f -o /tmp/syslog-ng-trace.txt -r all -w all -u libc:: /opt/syslog-ng/sbin/syslog-ng -d -d -d`

> **Tip**
> To execute these commands on an already running syslog-ng OSE process, use the `-p <pid_of_syslog-ng>` parameter.

16.4. Running a failure script

When syslog-ng is abnormally terminated, it can execute a user-created failure script. This can be used for example to send an automatic e-mail notification. The script must be located at `/opt/syslog-ng/sbin/syslog-ng-failure`.

16.5. Stopping syslog-ng

To avoid problems, always use the init scripts to stop syslog-ng (`/etc/init.d/syslog-ng stop`), instead of using the `kill` command. This is especially true on Solaris and HP-UX systems, here use `/etc/init.d/syslog stop`.

Chapter 17. Best practices and examples

This chapter discusses some special examples and recommendations.

17.1. General recommendations

This section provides general tips and recommendations on using syslog-ng. Some of the recommendations are detailed in the subsequent sections.

- Do not base the separation of log messages into different files on the `facility` parameter. As several applications and processes can use the same facility, the facility does not identify the application that sent the message. By default, the `facility` parameter is not even included in the log message itself. In general, sorting the log messages into several different files can make finding specific log messages difficult. If you must create separate log files, use the application name.

- Standard log messages include the local time of the sending host, without any time zone information. It is recommended to replace this timestamp with an ISODATE timestamp, because the ISODATE format includes the year and timezone as well. To convert all timestamps to the ISODATE format, include the following line in the syslog-ng configuration file:

  ```
  options {ts-format(iso) ; };
  ```

- Resolving the IP addresses of the clients to domain names can decrease the performance of syslog-ng. For details, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

17.2. Handling large message load

This section provides tips on optimizing the performance of syslog-ng. Optimizing the performance is important for syslog-ng hosts that handle large traffic.

- Disable DNS resolution, or resolve hostnames locally. For details, see *Section 17.3, Using name resolution in syslog-ng (p. 307)*.

- Enable flow-control for the TCP sources. For details, see *Section 8.2, Managing incoming and outgoing messages with flow-control (p. 203)*.

- Do not use the `usertty()` destination driver. Under heavy load, the users are not be able to read the messages from the console, and it slows down syslog-ng.

- Do not use regular expressions in our filters. Evaluating general regular expressions puts a high load on the CPU. Use simple filter functions and logical operators instead. For details, see *Section 11.3, Regular expressions (p. 263)*.

- **Warning**

  When receiving messages using the UDP protocol, increase the size of the UDP receive buffer on the receiver host (that is, the syslog-ng OSE server or relay receiving the messages). Note that on certain platforms, for example, on Red Hat Enterprise Linux 5, even low message load (~200 messages per second) can result in message loss, unless the `so-rcvbuf()` option of the source is increased. In such cases, you

will need to increase the net.core.rmem_max parameter of the host (for example, to 1024000), but do not modify net.core.rmem_default parameter.

As a general rule, increase the *so-rcvbuf()* so that the buffer size in kilobytes is higher than the rate of incoming messages per second. For example, to receive 2000 messages per second, set the *so-rcvbuf()* at least to 2 097 152 bytes.

■ Increase the value of the *flush-lines()* parameter. Increasing *flush-lines()* from 0 to 100 can increase the performance of syslog-ng OSE by 100%.

17.3. Using name resolution in syslog-ng

The syslog-ng application can resolve the hostnames of the clients and include them in the log messages. However, the performance of syslog-ng is severely degraded if the domain name server is unaccessible or slow. Therefore, it is not recommended to resolve hostnames in syslog-ng. If you must use name resolution from syslog-ng, consider the following:

■ Use DNS caching. Verify that the DNS cache is large enough to store all important hostnames. (By default, the syslog-ng DNS cache stores 1007 entries.)

```
options { dns-cache-size(2000); };
```

■ If the IP addresses of the clients change only rarely, set the expiry of the DNS cache large.

```
options { dns-cache-expire(87600); };
```

■ If possible, resolve the hostnames locally. For details, see *Procedure 17.3.1, Resolving hostnames locally (p. 307)*.

Note
Domain name resolution is important mainly in relay and server mode.

17.3.1. Procedure – Resolving hostnames locally

Purpose:

Resolving hostnames locally enables you to display hostnames in the log files for frequently used hosts, without having to rely on a DNS server. The known IP address – hostname pairs are stored locally in a file. In the log messages, syslog-ng will replace the IP addresses of known hosts with their hostnames. To configure local name resolution, complete the following steps:

Steps:

Step 1. Add the hostnames and the respective IP addresses to the file used for local name resolution. On Linux and UNIX systems, this is the /etc/hosts file. Consult the documentation of your operating system for details.

Step 2. Instruct syslog-ng to resolve hostnames locally. Set the *use-dns()* option of syslog-ng to *persist_only.*

Step 3. Set the *dns-cache-hosts()* option to point to the file storing the hostnames.

```
options {
        use-dns(persist_only);
        dns-cache-hosts(/etc/hosts); };
```

17.4. Procedure – Collecting logs from chroot

Purpose:

To collect logs from a chroot using a syslog-ng client running on the host, complete the following steps:
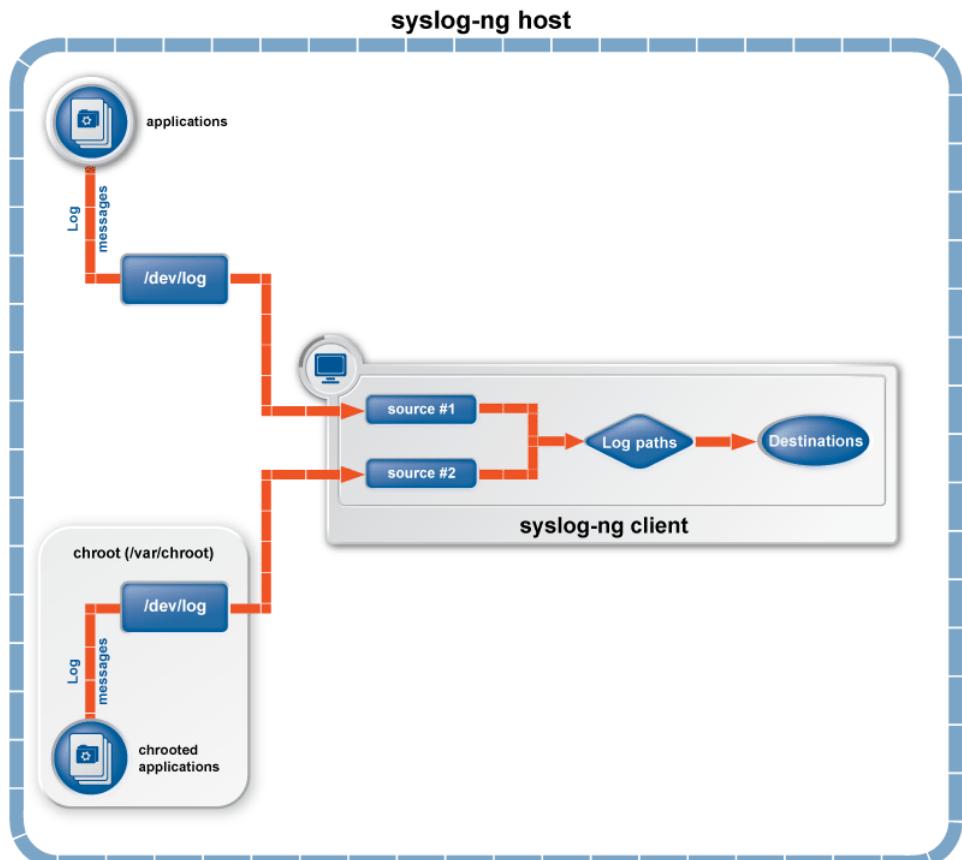
Figure 17.1. Collecting logs from chroot

Steps:

Step 1. Create a /dev directory within the chroot. The applications running in the chroot send their log messages here.

Step 2. Create a local source in the configuration file of the syslog-ng application running outside the chroot. This source should point to the /dev/log file within the chroot (for example to the /chroot/dev/log directory).

Step 3. Include the source in a log statement.

> **Note**
> You need to set up timezone information within your chroot as well. This usually means creating a symlink to `/etc/localtime`.

17.5. Configuring log rotation

The syslog-ng OSE application does not rotate logs by itself. To use syslog-ng OSE for log rotation, consider the following approaches:

Use logrotate together with syslog-ng OSE:

- Ideal for workstations or when processing fewer logs.
- It is included in most distributions by default.
- Less scripting is required, only `logrotate` has to be configured correctly.
- Requires frequent restart (syslog-ng OSE must be reloaded/restarted when the files are rotated).
- The statistics collected by syslog-ng OSE, and the correlation information gathered with Pattern Database is lost with each restart.

Separate incoming logs based on time, host or other information:

- Ideal for central log servers, where regular restart of syslog-ng OSE is unfavorable.
- Requires shell scripts or cron jobs to remove old logs.
- It can be done by using macros in the destination name (in the filename, directory name, or the database table name).

Example 17.1. File destination for log rotation
This sample file destination configuration stores incoming logs in files that are named based on the current year, month and day, and places these files in directories that are named based on the hostname:

```
destination d_sorted { file("/var/log/remote/${HOST}/${YEAR}_${MONTH}_${DAY}.log"
create-dirs(yes)); };
```

Example 17.2. Command for cron for log rotation
This sample command for `cron` removes files older than two weeks from the `/var/log/remote` directory:

```
find /var/log/remote/ -daystart -mtime +14 -type f -exec rm {} \;
```

Appendix A. The syslog-ng manual pages

loggen

loggen — Generate syslog messages at a specified rate

Synopsis

`loggen` [options]`target` [port]

Description

NOTE: The loggen application is distributed with the syslog-ng system logging application, and is usually part of the syslog-ng package. The latest version of the syslog-ng application is available at _the official syslog-ng website_.

This manual page is only an abstract; for the complete documentation of syslog-ng, see _The syslog-ng Administrator Guide_.

The `loggen` application is tool to test and stress-test your syslog server and the connection to the server. It can send syslog messages to the server at a specified rate, using a number of connection types and protocols, including TCP, UDP, and unix domain sockets. The messages can be generated automatically (repeating the _PADD_string over and over), or read from a file or the standard input.

When `loggen` finishes sending the messages, it displays the following statistics:

- _average rate_: Average rate the messages were sent in messages/second.
- _count_: The total number of messages sent.
- _time_: The time required to send the messages in seconds.
- _average message size_: The average size of the sent messages in bytes.
- _bandwidth_: The average bandwidth used for sending the messages in kilobytes/second.

Options

`--active-connections` `<number-of-connections>`	Number of connections `loggen` will use to send messages to the destination. This option is usable only when using TCP or TLS connections to the destination. Default value: 1
	The `loggen` utility waits until every connection is established before starting to send messages. See also the `--idle-connections` option.
`--csv` or `-C`	Send the statistics of the sent messages to stdout as CSV. This can be used for plotting the message rate.
`--dgram` or `-D`	Use datagram socket (UDP or unix-dgram) to send the messages to the target. Requires the `--inet` option as well.
`--dont-parse` or `-d`	Do not parse the lines read from the input files, send them as received.

`--help` or `-h`	Display a brief help message.
`--idle-connections` `<number-of-connections>`	Number of idle connections `loggen` will establish to the destination. Note that `loggen` will not send any messages on idle connections, but the connection is kept open using keep-alive messages. This option is usable only when using TCP or TLS connections to the destination. See also the `--active-connections` option. Default value: 0
`--inet` or `-i`	Use the TCP (by default) or UDP (when used together with the `--dgram` option) protocol to send the messages to the target.
`--interval <seconds>` or `-I` `<seconds>`	The number of seconds `loggen` will run. Default value: 10

> **Note**
> Note that when the `--interval` and `--number` are used together, `loggen` will send messages until the period set in `--interval` expires or the amount of messages set in `--number` is reached, whichever happens first.

`--ipv6` or `-6`	Specify the destination using its IPv6 address. Note that the destination must have a real IPv6 address.
`--loop-reading` or `-l`	Read the file specified in `--read-file` option in loop: loggen will start reading from the beginning of the file when it reaches the end of the file.

`--number`
`<number-of-messages>` or `-n`
`<number-of-messages>`
	Number of messages to generate.

> **Note**
> Note that when the `--interval` and `--number` are used together, `loggen` will send messages until the period set in `--interval` expires or the amount of messages set in `--number` is reached, whichever happens first.

`--no-framing` or `-F`	Do not use the framing of the IETF-syslog protocol style, even if the `syslog-proto` option is set.
`--quiet` or `-Q`	Output statistics only when the execution of `loggen` is finished. If not set, the statistics are displayed every second.
`--rate <message/second>` or `-r` `<message/second>`	The number of messages generated per second for every active connection. Default value: 1000
`--read-file <filename>` or `-R` `<filename>`	Read the messages from a file and send them to the target. See also the `--skip-tokens` option.

	Specify - as the input file to read messages from the standard input (stdio). Note that when reading messages from the standard input, `loggen` can only use a single thread. The `-R` - parameters must be placed at end of command, like: `loggen 127.0.0.1 1061 --read-file -`
`--sdata <data-to-send>` or `-p <data-to-send>`	Send the argument of the `--sdata` option as the SDATA part of IETF-syslog (RFC5424 formatted) messages. Use it together with the `--syslog-proto` option. For example: `--sdata "[test name=\"value\"]"`
`--size <message-size>` or `-s <message-size>`	The size of a syslog message in bytes. Default value: 256. Minimum value: 127 bytes, maximum value: 8192 bytes.
`--skip-tokens <number>`	Skip the specified number of space-separated tokens (words) at the beginning of every line. For example, if the messages in the file look like `foo bar message`, `--skip-tokens 2` skips the `foo bar` part of the line, and sends only the `message` part. Works only when used together with the `--read-file` parameter. Default value: 3
`--stream` or `-S`	Use a stream socket (TCP or unix-stream) to send the messages to the target.
`--syslog-proto` or `-P`	Use the new IETF-syslog message format as specified in RFC5424. By default, loggen uses the legacy BSD-syslog message format (as described in RFC3164). See also the `--no-framing` option.
`--unix </path/to/socket>` or `-x </path/to/socket>`	Use a UNIX domain socket to send the messages to the target.
`--use-ssl` or `-U`	Use an SSL-encrypted channel to send the messages to the target. Note that it is not possible to check the certificate of the target, or to perform mutual authentication.
`--version` or `-V`	Display version number of syslog-ng.

Examples

The following command generates 100 messages per second for ten minutes, and sends them to port 2010 of the localhost via TCP. Each message is 300 bytes long.

```
loggen --size 300 --rate 100 --interval 600 127.0.0.1 2010
```

The following command is similar to the one above, but uses the UDP protocol.

```
loggen --inet --dgram --size 300 --rate 100 --interval 600 127.0.0.1 2010
```

Send a single message on TCP6 to the `::1` IPv6 address, port `1061`:

```
loggen --ipv6 --number 1 ::1 1061
```

Send a single message on UDP6 to the `::1` IPv6 address, port *1061*:

```
loggen --ipv6 --dgram --number 1 ::1 1061
```

Send a single message using a unix domain-socket:

```
loggen --unix --stream --number 1 </path/to/socket>
```

Read messages from the standard input (stdio) and send them to the localhost:

```
loggen 127.0.0.1 1061 --read-file -
```

Files

```
/opt/syslog-ng/bin/loggen
```

See also

```
syslog-ng.conf(5)
```

Note

For the detailed documentation of syslog-ng OSE see *The syslog-ng OSE 3.6 Administrator Guide*

If you experience any problems or need help with syslog-ng, visit *visit the syslog-ng wiki* or the *syslog-ng mailing list*.

For news and notifications about of syslog-ng, visit the *syslog-ng Insider Blog*.

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

The authors grant permission to copy, distribute and/or modify this manual page under the terms of the GNU General Public License Version 2 or newer (GPL v2+).

pdbtool

pdbtool — An application to test and convert syslog-ng pattern database rules

Synopsis

`pdbtool` [command] [options]

Description

This manual page is only an abstract; for the complete documentation of syslog-ng and pdbtool, see _ `The syslog-ng Administrator Guide`.

The syslog-ng application can match the contents of the log messages to a database of predefined message patterns (also called patterndb). By comparing the messages to the known patterns, syslog-ng is able to identify the exact type of the messages, tag the messages, and sort them into message classes. The message classes can be used to classify the type of the event described in the log message. The functionality of the pattern database is similar to that of the logcheck project, but the syslog-ng approach is faster, scales better, and is much easier to maintain compared to the regular expressions of logcheck.

The `pdbtool` application is a utility that can be used to:

- *test messages*, or *specific rules*;
- convert an older pattern database to the latest database format;
- *merge pattern databases* into a single file;
- *automatically create pattern databases* from a large amount of log messages;
- *dump the RADIX tree* built from the pattern database (or a part of it) to explore how the pattern matching works.

The dictionary command

`dictionary` [options]

Lists every name-value pair that can be set by the rules of the pattern database.

`--dump-tags` or `-T`	List the tags instead of the names of the name-value pairs.
`--pdb <path-to-file>` or `-p <path-to-file>`	Name of the pattern database file to use.
`--program <programname>` or `-P <programname>`	List only the name-value pairs that can be set for the messages of the specified *$PROGRAM* application.

The dump command

`dump` [options]

Display the RADIX tree built from the patterns. This shows how are the patterns represented in syslog-ng and it might also help to track down pattern-matching problems. The dump utility can dump the tree used for matching the PROGRAM or the MSG parts.

`--debug` or `-d`	Enable debug/diagnostic messages on stderr.
`--pdb` or `-p`	Name of the pattern database file to use.
`--program` or `-P`	Displays the RADIX tree built from the patterns belonging to the ${PROGRAM} application.
`--program-tree` or `-T`	Display the ${PROGRAM} tree.
`--verbose` or `-v`	Enable verbose messages on stderr.

Example and sample output:

```
pdbtool dump -p patterndb.xml  -P 'sshd'
```

```
'p'
   'assword for'
     @QSTRING:@
       'from'
        @QSTRING:@
           'port '
            @NUMBER:@ rule_id='fc49054e-75fd-11dd-9bba-001e6806451b'
              ' ssh' rule_id='fc55cf86-75fd-11dd-9bba-001e6806451b'
                  '2' rule_id='fc4b7982-75fd-11dd-9bba-001e6806451b'
   'ublickey for'
     @QSTRING:@
       'from'
        @QSTRING:@
           'port '
            @NUMBER:@ rule_id='fc4d377c-75fd-11dd-9bba-001e6806451b'
              ' ssh' rule_id='fc5441ac-75fd-11dd-9bba-001e6806451b'
                  '2' rule_id='fc44a9fe-75fd-11dd-9bba-001e6806451b'
```

The match command

`match` [options]

Use the `match` command to test the rules in a pattern database. The command tries to match the specified message against the patterns of the database, evaluates the parsers of the pattern, and also displays which part of the message was parsed successfully. The command returns with a *0* (success) or *1* (no match) return code and displays the following information:

- the class assigned to the message (that is, system, violation, and so on),
- the ID of the rule that matched the message, and
- the values of the parsers (if there were parsers in the matching pattern).

The `match` command has the following options:

`--color-out` or `-c`	Color the terminal output to highlight the part of the message that was successfully parsed.
`--debug` or `-d`	Enable debug/diagnostic messages on stderr.
`--debug-csv` or `-C`	Print the debugging information returned by the `--debug-pattern` option as comma-separated values.
`--debug-pattern` or `-D`	Print debugging information about the pattern matching. See also the `--debug-csv` option.
`--file=<filename-with-path>` or `-f`	Process the messages of the specified log file with the pattern database. This option allows to classify messages offline, and to apply the pattern database to already existing logfiles. To read the messages from the standard input (stdin), specify a hyphen (`-`) character instead of a filename.
`--filter=<filter-expression>` or `-F`	Print only messages matching the specified syslog-ng filter expression.
`--message` or `-M`	The text of the log message to match (only the *${MESSAGE}* part without the syslog headers).
`--pdb` or `-p`	Name of the pattern database file to use.
`--program` or `-P`	Name of the program to use, as contained in the *${PROGRAM}* part of the syslog message.
`--template=<template-expression>` or `-T`	A syslog-ng template expression that is used to format the output messages.
`--verbose` or `-v`	Enable verbose messages on stderr.

Example: The following command checks if the `patterndb.xml` file recognizes the *Accepted publickey for myuser from 127.0.0.1 port 59357 ssh2* message:

```
pdbtool match -p patterndb.xml -P sshd -M "Accepted publickey for myuser from
127.0.0.1 port 59357 ssh2"
```

The following example applies the `sshd.pdb` pattern database file to the log messages stored in the `/var/log/messages` file, and displays only the messages that received a *useracct* tag.

```
pdbtool match -p sshd.pdb \
  –file /var/log/messages \
  –filter 'tags("usracct");'
```

The merge command

`merge` [options]

Use the `merge` command to combine separate pattern database files into a single file (pattern databases are usually stored in separate files per applications to simplify maintenance). If a file uses an older database format,

it is automatically updated to the latest format (V3). See the `The syslog-ng Administrator Guide` for details on the different pattern database versions.

`--debug` or `-d`	Enable debug/diagnostic messages on stderr.
`--directory` or `-D`	The directory that contains the pattern database XML files to be merged.
`--glob` or `-G`	Specify filenames to be merged using a glob pattern, for example, using wildcards. For details on glob patterns, see `man glob`. This pattern is applied only to the filenames, and not on directory names.
`--pdb` or `-p`	Name of the output pattern database file.
`--recursive` or `-r`	Merge files from subdirectories as well.
`--verbose` or `-v`	Enable verbose messages on stderr.

Example:

```
pdbtool merge --recursive --directory /home/me/mypatterns/  --pdb
/var/lib/syslog-ng/patterndb.xml
```

Currently it is not possible to convert a file without merging, so if you only want to convert an older pattern database file to the latest format, you have to copy it into an empty directory.

The patternize command

`patternize` [options]

Automatically create a pattern database from a log file containing a large number of log messages. The resulting pattern database is printed to the standard output (stdout). The `pdbtool patternize` command uses a data clustering technique to find similar log messages and replacing the differing parts with *@ESTRING:: @* parsers. For details on pattern databases and message parsers, see the *The syslog-ng Administrator Guide*. The `patternize` command is available only in syslog-ng OSE version 3.2 and later.

`--debug` or `-d`	Enable debug/diagnostic messages on stderr.
`--file=<path>` or `-f`	The logfile containing the log messages to create patterns from. To receive the log messages from the standard input (stdin), use `-`.
`--iterate-outliers` or `-o`	Recursively iterate on the log lines to cover as many log messages with patterns as possible.
`--named-parsers` or `-n`	The number of example log messages to include in the pattern database for every pattern. Default value: *1*
`--no-parse` or `-p`	Do not parse the input file, treat every line as the message part of a log message.
`--samples=<number-of-samples>`	Include a generated name in the parsers, for example, *.dict.string1*, *.dict.string2*, and so on.
`--support=<number>` or `-S`	A pattern is added to the output pattern database if at least the specified percentage of log messages from the input logfile match

the pattern. For example, if the input logfile contains 1000 log messages and the `--support=3.0` option is used, a pattern is created only if the pattern matches at least 3 percent of the log messages (that is, 30 log messages). If patternize does not create enough patterns, try to decrease the support value.

Default value: `4.0`

    `--verbose` or `-v`   Enable verbose messages on stderr.

Example:

```
pdbtool patternize --support=2.5 --file=/var/log/messages
```

The test command

`test` [options]

Use the `test` command to validate a pattern database XML file. Note that you must have the `xmllint` application installed. The `test` command is available only in syslog-ng OSE version 3.2 and later.

`--color-out` or `-c`	Enable coloring in terminal output.
`--debug` or `-d`	Enable debug/diagnostic messages on stderr.
`--debug` or `-D`	Print debugging information on non-matching patterns.
`--rule-id` or `-r`	Test only the patterndb rule (specified by its rule id) against its example.
`--validate`	Validate a pattern database XML file.
`--verbose` or `-v`	Enable verbose messages on stderr.

Example:

```
pdbtool test --validate /home/me/mypatterndb.pdb
```

Files

```
/opt/syslog-ng/
```

```
/opt/syslog-ng/etc/syslog-ng.conf
```

See also

```
The syslog-ng Administrator Guide
```

```
syslog-ng.conf(5)
```

```
syslog-ng(8)
```

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

syslog-ng

syslog-ng — syslog-ng system logger application

Synopsis

`syslog-ng` [options]

Description

This manual page is only an abstract; for the complete documentation of syslog-ng, see *The syslog-ng Open Source Edition Administrator Guide* or *the official syslog-ng website*.

The syslog-ng OSE application is a flexible and highly scalable system logging application. Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices - called syslog-ng clients - all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, where the server sorts and stores them.

Options

`--caps`	Run syslog-ng OSE process with the specified POSIX capability flags.

- If the `--no-caps` option is not set, syslog-ng OSE has been compiled with the `--enable-linux-caps` compile option, and the host supports CAP_SYSLOG, syslog-ng OSE uses the following capabilities: "cap_net_bind_service,cap_net_broadcast,cap_net_raw,cap_dac_read_search,cap_dac_override,cap_chown,cap_fowner=p cap_syslog=ep"

- If the `--no-caps` option is not set, and the host does not support CAP_SYSLOG, syslog-ng OSE uses the following capabilities: "cap_net_bind_service,cap_net_broadcast,cap_net_raw,cap_dac_read_search,cap_dac_override,cap_chown,cap_fowner=p cap_sys_admin=ep"

`--cfgfile <file>` or `-f <file>`	Use the specified configuration file.
`--chroot <dir>` or `-C <dir>`	Change root to the specified directory. The configuration file is read after chrooting so, the configuration file must be available within the chroot. That way it is also possible to reload the syslog-ng configuration after chrooting. However, note that the `--user` and `--group` options are resolved before chrooting.
`--debug` or `-d`	Start syslog-ng in debug mode.
`--default-modules`	A comma-separated list of the modules that are loaded automatically. Modules not loaded automatically can be loaded by including the *@module <modulename>* statement in the

syslog-ng OSE configuration file. The following modules are l o a d e d   b y   d e f a u l t : affile,afprog,afsocket,afuser,basicfuncs,csvparser,dbparser,syslogformat,afsql. Available only in syslog-ng Open Source Edition 3.3 and later.

`--enable-core`	Enable syslog-ng to write core files in case of a crash to help support and debugging.
`--fd-limit <number>`	Set the minimal number of required file descriptors (fd-s); this sets how many files syslog-ng can keep open simultaneously. Default value: *4096*. Note that this does not override the global ulimit setting of the host.
`--foreground` or `-F`	Do not daemonize, run in the foreground.
`--group <group>` or `-g <group>`	Switch to the specified group after initializing the configuration file.
`--help` or `-h`	Display a brief help message.
`--module-registry`	Display the list and description of the available modules. Note that not all of these modules are loaded automatically, only the ones specified in the `--default-modules` option. Available only in syslog-ng Open Source Edition 3.3 and later.
`--no-caps`	Run syslog-ng as root, without capability-support. This is the default behavior. On Linux, it is possible to run syslog-ng as non-root with capability-support if syslog-ng was compiled with the *--enable-linux-caps* option enabled. (Execute `syslog-ng --version` to display the list of enabled build parameters.)
	To run syslog-ng OSE with specific capabilities, use the *--caps* option.
`--persist-file <persist-file>` or `-R <persist-file>`	Set the path and name of the `syslog-ng.persist` file where the persistent options and data are stored.
`--pidfile <pidfile>` or `-p <pidfile>`	Set path to the PID file where the pid of the main process is stored.
`--preprocess-into <output-file>`	After processing the configuration file and resolving included files and variables, write the resulting configuration into the specified output file. Available only in syslog-ng Open Source Edition 3.3 and later.
`--process-mode <mode>`	Sets how to run syslog-ng: in the *foreground* (mainly used for debugging), in the *background* as a daemon, or in *safe-background* mode. By default, syslog-ng runs in *safe-background* mode. This mode creates a supervisor process called *supervising syslog-ng* , that restarts syslog-ng if it crashes.

`--stderr` or `-e`	Log internal messages of syslog-ng to stderr. Mainly used for debugging purposes in conjunction with the `--foreground` option. If not specified, syslog-ng will log such messages to its internal source.
`--syntax-only` or `-s`	Verify that the configuration file is syntactically correct and exit.
`--user <user>` or `-u <user>`	Switch to the specified user after initializing the configuration file (and optionally chrooting). Note that it is not possible to reload the syslog-ng configuration if the specified user has no privilege to create the `/dev/log` file.
`--verbose` or `-v`	Enable verbose logging used to troubleshoot syslog-ng.
`--version` or `-V`	Display version number and compilation information, and also the list and short description of the available modules. For detailed description of the available modules, see the `--module-registry` option. Note that not all of these modules are loaded automatically, only the ones specified in the `--default-modules` option.
`--worker-threads`	Sets the number of worker threads syslog-ng OSE can use, including the main syslog-ng OSE thread. Note that certain operations in syslog-ng OSE can use threads that are not limited by this option. This setting has effect only when syslog-ng OSE is running in multithreaded mode. Available only in syslog-ng Open Source Edition 3.3 and later. See `The syslog-ng Open Source Edition 3.6 Administrator Guide` for details.

Files

`/opt/syslog-ng/`

`/opt/syslog-ng/etc/syslog-ng.conf`

See also

syslog-ng.conf(5)

> **Note**
>
> For the detailed documentation of syslog-ng OSE see _The syslog-ng OSE 3.6 Administrator Guide_
>
> If you experience any problems or need help with syslog-ng, visit _visit the syslog-ng wiki_ or the _syslog-ng mailing list_.
>
> For news and notifications about of syslog-ng, visit the _syslog-ng Insider Blog_.

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

The authors grant permission to copy, distribute and/or modify this manual page under the terms of the GNU General Public License Version 2 or newer (GPL v2+).

syslog-ng.conf

syslog-ng.conf — syslog-ng configuration file

Synopsis

```
syslog-ng.conf
```

Description

This manual page is only an abstract; for the complete documentation of syslog-ng, see `The syslog-ng Open Source Edition Administrator Guide` or *the official syslog-ng website*.

The syslog-ng OSE application is a flexible and highly scalable system logging application. Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices - called syslog-ng clients - all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, where the server sorts and stores them.

Basic concepts of syslog-ng OSE

The syslog-ng application reads incoming messages and forwards them to the selected *destinations*. The syslog-ng application can receive messages from files, remote hosts, and other *sources*.

Log messages enter syslog-ng in one of the defined sources, and are sent to one or more *destinations*.

Sources and destinations are independent objects; *log paths* define what syslog-ng does with a message, connecting the sources to the destinations. A log path consists of one or more sources and one or more destinations; messages arriving from a source are sent to every destination listed in the log path. A log path defined in syslog-ng is called a *log statement*.

Optionally, log paths can include *filters*. Filters are rules that select only certain messages, for example, selecting only messages sent by a specific application. If a log path includes filters, syslog-ng sends only the messages satisfying the filter rules to the destinations set in the log path.

Other optional elements that can appear in log statements are *parsers* and *rewriting rules*. Parsers segment messages into different fields to help processing the messages, while rewrite rules modify the messages by adding, replacing, or removing parts of the messages.

Configuring syslog-ng

- The main body of the configuration file consists of object definitions: sources, destinations, logpaths define which log message are received and where they are sent. All identifiers, option names and attributes, and any other strings used in the syslog-ng configuration file are case sensitive. Objects must be defined before they are referenced in another statement. Object definitions (also called statements) have the following syntax:

```
object_type object_id {<options>};
```

- *Type of the object*: One of `source`, `destination`, `log`, `filter`, `parser`, `rewrite` rule, or `template`.

- *Identifier of the object*: A unique name identifying the object. When using a reserved word as an identifier, enclose the identifier in quotation marks.

Tip
Use identifiers that refer to the type of the object they identify. For example, prefix source objects with s_, destinations with d_, and so on.

Note
Repeating a definition of an object (that is, defining the same object with the same id more than once) is not allowed, unless you use the `@define allow-config-dups 1` definition in the configuration file.

- *Parameters*: The parameters of the object, enclosed in braces {`parameters`}.
- *Semicolon*: Object definitions end with a semicolon (`;`).
  For example, the following line defines a source and calls it `s_internal`.

```
source s_internal { internal(); };
```

The object can be later referenced in other statements using its ID, for example, the previous source is used as a parameter of the following log statement:

```
log { source(s_internal); destination(d_file); };
```

■ The parameters and options within a statement are similar to function calls of the C programming language: the name of the option followed by a list of its parameters enclosed within brackets and terminated with a semicolon.

```
option(parameter1, parameter2); option2(parameter1, parameter2);
```

For example, the `file()` driver in the following source statement has three options: the filename (`/var/log/apache/access.log`), `follow-freq()`, and `flags()`. The `follow-freq()` option also has a parameter, while the `flags()` option has two parameters.

```
source s_tail { file("/var/log/apache/access.log"
    follow-freq(1) flags(no-parse, validate-utf8)); };
```

Objects may have required and optional parameters. Required parameters are positional, meaning that they must be specified in a defined order. Optional parameters can be specified in any order using the `option(value)` format. If a parameter (optional or required) is not specified, its default value is used. The parameters and their default values are listed in the reference section of the particular object.

Example A.1. Using required and optional parameters
The *unix-stream()* source driver has a single required argument: the name of the socket to listen on. Optional parameters follow the socket name in any order, so the following source definitions have the same effect:

```
source s_demo_stream1 {
        unix-stream("/dev/log" max-connections(10) group(log)); };
source s_demo_stream2 {
        unix-stream("/dev/log" group(log) max-connections(10)); };
```

- Some options are global options, or can be set globally, for example, whether syslog-ng OSE should use DNS resolution to resolve IP addresses. Global options are detailed in *Chapter 9, Global options of syslog-ng OSE (p. 217)*.

```
options { use-dns(no); };
```

- All identifiers, attributes, and any other strings used in the syslog-ng configuration file are case sensitive.

- Objects can be used before definition.

- Objects can be defined inline as well. This is useful if you use the object only once (for example, a filter). For details, see *Section 5.4, Defining configuration objects inline (p. 43)*.

- To add comments to the configuration file, start a line with # and write your comments. These lines are ignored by syslog-ng.

```
# Comment: This is a stream source
source s_demo_stream {
        unix-stream("/dev/log" max-connections(10) group(log)); };
```

The syntax of log statements is as follows:

```
log {
    source(s1); source(s2); ...
    optional_element(filter1|parser1|rewrite1);...
    optional_element(filter2|parser2|rewrite2);...
    destination(d1); destination(d2); ...
    flags(flag1[, flag2...]);
    };
```

The following log statement sends all messages arriving to the localhost to a remote server.

```
source s_localhost { tcp(ip(127.0.0.1) port(1999) ); };
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
log { source(s_localhost); destination(d_tcp); };
```

The syslog-ng application has a number of global options governing DNS usage, the timestamp format used, and other general points. Each option may have parameters, similarly to driver specifications. To set global options, add an option statement to the syslog-ng configuration file using the following syntax:

```
options { option1(params); option2(params); ... };
```

Example A.2. Using global options
To disable domain name resolving, add the following line to the syslog-ng configuration file:

```
options { use-dns(no); };
```

The sources, destinations, and filters available in syslog-ng are listed below. For details, see ⎵The syslog-ng Administrator Guide.

Name	Description
file()	Opens the specified file and reads messages.
network()	Listens on the specified TCP or UDP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively. (This driver is a unified wrapper for the `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers.
nodejs()	Receives JSON messages from nodejs applications.
internal()	Messages generated internally in syslog-ng.
pacct()	Reads messages from the process accounting logs on Linux.
pipe()	Opens the specified named pipe and reads messages.
program()	Opens the specified application and reads messages from its standard output.
sun-stream(), sun-streams()	Opens the specified `STREAMS` device on Solaris systems and reads incoming messages.
syslog()	Listens for incoming messages using the new *IETF-standard syslog protocol*.
system()	Automatically detects which platform syslog-ng OSE is running on, and collects the native log messages of that platform.
systemd-journal()	Collects messages directly from the journal of platforms that use systemd.
systemd-syslog()	Collects messages from the journal using a socket on platforms that use systemd.
tcp(), tcp6()	Listens on the specified TCP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively.
udp(), udp6()	Listens on the specified UDP port for incoming messages using the *BSD-syslog protocol* over IPv4 and IPv6 networks, respectively.
unix-dgram()	Opens the specified unix socket in `SOCK_DGRAM` mode and listens for incoming messages.

Name	Description
unix-stream()	Opens the specified unix socket in `SOCK_STREAM` mode and listens for incoming messages.

Name	Description
amqp()	Publishes messages using the AMQP (Advanced Message Queuing Protocol).
file()	Writes messages to the specified file.
graphite()	Sends metrics to a *Graphite* server to store numeric time-series data.
mongodb()	Sends messages to a *MongoDB* database.
pipe()	Writes messages to the specified named pipe.
program()	Forks and launches the specified program, and sends messages to its standard input.
redis()	Sends messages as name-value pairs to a *Redis* key-value store.
riemann()	Sends metrics or events to a *Riemann* monitoring system.
smtp()	Sends e-mail messages to the specified recipients.
sql()	Sends messages into an SQL database. In addition to the standard syslog-ng packages, the `sql()` destination requires database-specific packages to be installed. Refer to the section appropriate for your platform in *Chapter 3, Installing syslog-ng (p. 23)*.
stomp()	Sends messages to a STOMP server.
syslog()	Sends messages to the specified remote host using the *IETF-syslog protocol*. The IETF standard supports message transport using the UDP, TCP, and TLS networking protocols.
tcp() and tcp6()	Sends messages to the specified TCP port of a remote host using the *BSD-syslog protocol* over IPv4 and IPv6, respectively.
udp() and udp6()	Sends messages to the specified UDP port of a remote host using the *BSD-syslog protocol* over IPv4 and IPv6, respectively.
unix-dgram()	Sends messages to the specified unix socket in `SOCK_DGRAM` style (BSD).
unix-stream()	Sends messages to the specified unix socket in `SOCK_STREAM` style (Linux).

Name	Description
usertty()	Sends messages to the terminal of the specified user, if the user is logged in.

Table A.2. Destination drivers available in syslog-ng

Name	Description
facility()	Filter messages based on the sending facility.
filter()	Call another filter function.
host()	Filter messages based on the sending host.
inlist()	File-based whitelisting and blacklisting.
level() or priority()	Filter messages based on their priority.
match()	Use a regular expression to filter messages based on a specified header or content field.
message()	Use a regular expression to filter messages based on their content.
netmask()	Filter messages based on the IP address of the sending host.
program()	Filter messages based on the sending application.
source()	Select messages of the specified syslog-ng OSE source statement.
tags()	Select messages having the specified tag.

Table A.3. Filter functions available in syslog-ng OSE

Files

`/opt/syslog-ng/`

`/opt/syslog-ng/etc/syslog-ng.conf`

See also

`syslog-ng(8)`

> **Note**
> For the detailed documentation of syslog-ng OSE see *The syslog-ng OSE 3.6 Administrator Guide*
>
> If you experience any problems or need help with syslog-ng, visit *visit the syslog-ng wiki* or the *syslog-ng mailing list*.
>
> For news and notifications about of syslog-ng, visit the *syslog-ng Insider Blog*.

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

The authors grant permission to copy, distribute and/or modify this manual page under the terms of the GNU General Public License Version 2 or newer (GPL v2+).

syslog-ng-ctl

syslog-ng-ctl — Display message statistics and enable verbose, debug and trace modes in syslog-ng Open Source Edition

Synopsis

`syslog-ng-ctl` [command] [options]

Description

NOTE: The syslog-ng-ctl application is distributed with the syslog-ng Open Source Edition system logging application, and is usually part of the syslog-ng package. The latest version of the syslog-ng application is available at *the official syslog-ng website*.

This manual page is only an abstract; for the complete documentation of syslog-ng, see ⎵ `The syslog-ng Open Source Edition Administrator Guide`.

The `syslog-ng-ctl` application is a utility that can be used to:

- enable/disable various syslog-ng messages for troubleshooting;

- display statistics about the processed messages;

- reload the configuration of syslog-ng OSE.

Enabling troubleshooting messages

`command` [options]

Use the `syslog-ng-ctl <command> --set=on` command to display verbose, trace, or debug messages. If you are trying to solve configuration problems, the verbose (and occasionally trace) messages are usually sufficient; debug messages are needed mostly for finding software errors. After solving the problem, do not forget to turn these messages off using the `syslog-ng-ctl <command> --set=off`. Note that enabling debug messages does not enable verbose and trace messages.

Use `syslog-ng-ctl <command>` without any parameters to display whether the particular type of messages are enabled or not.

If you need to use a non-standard control socket to access syslog-ng, use the `syslog-ng-ctl <command> --set=on --control=<socket>` command to specify the socket to use.

verbose	Print verbose messages. If syslog-ng was started with the `--stderr` or `-e` option, the messages will be sent to stderr. If not specified, syslog-ng will log such messages to its internal source.
trace	Print trace messages of how messages are processed. If syslog-ng was started with the `--stderr` or `-e` option, the messages will be sent to stderr. If not specified, syslog-ng will log such messages to its internal source.

debug   Print debug messages. If syslog-ng was started with the *--stderr* or *-e* option, the messages will be sent to stderr. If not specified, syslog-ng will log such messages to its internal source.

Example:

```
syslog-ng-ctl verbose --set=on
```

The stats command

`stats` [options]

Use the `stats` command to display statistics about the processed messages. The `stats` command has the following options:

--control=<socket> or -c   Specify the socket to use to access syslog-ng. Only needed when using a non-standard socket.

Example:

```
syslog-ng-ctl stats
```

An example output:

```
src.internal;s_all#0;;a;processed;6445
src.internal;s_all#0;;a;stamp;1268989330
destination;df_auth;;a;processed;404
destination;df_news_dot_notice;;a;processed;0
destination;df_news_dot_err;;a;processed;0
destination;d_ssb;;a;processed;7128
destination;df_uucp;;a;processed;0
source;s_all;;a;processed;7128
destination;df_mail;;a;processed;0
destination;df_user;;a;processed;1
destination;df_daemon;;a;processed;1
destination;df_debug;;a;processed;15
destination;df_messages;;a;processed;54
destination;dp_xconsole;;a;processed;671
dst.tcp;d_network#0;10.50.0.111:514;a;dropped;5080
dst.tcp;d_network#0;10.50.0.111:514;a;processed;7128
dst.tcp;d_network#0;10.50.0.111:514;a;stored;2048
destination;df_syslog;;a;processed;6724
destination;df_facility_dot_warn;;a;processed;0
destination;df_news_dot_crit;;a;processed;0
destination;df_lpr;;a;processed;0
destination;du_all;;a;processed;0
destination;df_facility_dot_info;;a;processed;0
center;;received;a;processed;0
destination;df_kern;;a;processed;70
center;;queued;a;processed;0
destination;df_facility_dot_err;;a;processed;0
```

Reloading the configuration

`command [options]`

Use the `syslog-ng-ctl reload` command to reload the configuration file of syslog-ng OSE without having to restart the syslog-ng OSE application. The `syslog-ng-ctl reload` works like a SIGHUP.

Files

`/opt/syslog-ng/sbin/syslog-ng-ctl`

See also

`The syslog-ng Administrator Guide`

`syslog-ng.conf(5)`

`syslog-ng(8)`

> **Note**
>
> For the detailed documentation of syslog-ng OSE see `The syslog-ng OSE 3.6 Administrator Guide`
>
> If you experience any problems or need help with syslog-ng, visit *visit the syslog-ng wiki* or the `syslog-ng mailing list`.
>
> For news and notifications about of syslog-ng, visit the `syslog-ng Insider Blog`.

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

The authors grant permission to copy, distribute and/or modify this manual page under the terms of the GNU General Public License Version 2 or newer (GPL v2+).

Appendix B. GNU General Public License

Version 2, June 1991
Copyright © 1989, 1991 Free Software Foundation, Inc.

Free Software Foundation, Inc.
  51 Franklin Street, Fifth Floor,
  Boston, MA 02110-1301
  USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Version 2, June 1991

B.1. Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software - to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps:

1. copyright the software, and
2. offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

B.2. TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

B.2.1. Section 0

This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program", below, refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification".) Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

B.2.2. Section 1

You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

B.2.3. Section 2

You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of *Section 1* above, provided that you also meet all of these conditions:

a. You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.

b. You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.

c. If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and

telling the user how to view a copy of this License. (Exception: If the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

B.2.4. Section 3

You may copy and distribute the Program (or a work based on it, under _Section 2_ in object code or executable form under the terms of _Sections 1_ and _2_ above provided that you also do one of the following:

a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,

b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,

c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

B.2.5. Section 4

You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

B.2.6. Section 5

You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.

B.2.7. Section 6

Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.

B.2.8. Section 7

If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

B.2.9. Section 8

If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

B.2.10. Section 9

The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

B.2.11. Section 10

If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

B.2.12. NO WARRANTY Section 11

BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

B.2.13. Section 12

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

B.3. How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.> Copyright (C) <year> <name of author>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

Gnomovision version 69, Copyright (C) year name of author Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type "show w". This is free software, and you are welcome to redistribute it under certain conditions; type "show c" for details.

The hypothetical commands "show w" and "show c" should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than "show w" and "show c"; they could even be mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the program, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the program "Gnomovision" (which makes passes at compilers) written by James Hacker.

<signature of Ty Coon>, 1 April 1989 Ty Coon, President of Vice

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

Appendix C. GNU Lesser General Public License

This is the first released version of the Lesser GPL. It also counts as the successor of the GNU Library Public License, version 2, hence the version number 2.1.
Copyright © 1991, 1999 Free Software Foundation, Inc.

Free Software Foundation, Inc.
    51 Franklin Street, Fifth Floor,
    Boston, MA 02110-1301
    USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Version 2.1, February 1999

C.1. Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public Licenses are intended to guarantee your freedom to share and change free software--to make sure the software is free for all its users.

This license, the Lesser General Public License, applies to some specially designated software packages--typically libraries--of the Free Software Foundation and other authors who decide to use it. You can use it too, but we suggest you first think carefully about whether this license or the ordinary General Public License is the better strategy to use in any particular case, based on the explanations below.

When we speak of free software, we are referring to freedom of use, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish); that you receive source code or can get it if you want it; that you can change the software and use pieces of it in new free programs; and that you are informed that you can do these things.

To protect your rights, we need to make restrictions that forbid distributors to deny you these rights or to ask you to surrender these rights. These restrictions translate to certain responsibilities for you if you distribute copies of the library or if you modify it.

For example, if you distribute copies of the library, whether gratis or for a fee, you must give the recipients all the rights that we gave you. You must make sure that they, too, receive or can get the source code. If you link other code with the library, you must provide complete object files to the recipients, so that they can relink them with the library after making changes to the library and recompiling it. And you must show them these terms so they know their rights.

We protect your rights with a two-step method:

1. we copyright the library, and
2. we offer you this license, which gives you legal permission to copy, distribute and/or modify the library.

To protect each distributor, we want to make it very clear that there is no warranty for the free library. Also, if the library is modified by someone else and passed on, the recipients should know that what they have is not the original version, so that the original author's reputation will not be affected by problems that might be introduced by others.

Finally, software patents pose a constant threat to the existence of any free program. We wish to make sure that a company cannot effectively restrict the users of a free program by obtaining a restrictive license from a patent holder. Therefore, we insist that any patent license obtained for a version of the library must be consistent with the full freedom of use specified in this license.

Most GNU software, including some libraries, is covered by the ordinary GNU General Public License. This license, the GNU Lesser General Public License, applies to certain designated libraries, and is quite different from the ordinary General Public License. We use this license for certain libraries in order to permit linking those libraries into non-free programs.

When a program is linked with a library, whether statically or using a shared library, the combination of the two is legally speaking a combined work, a derivative of the original library. The ordinary General Public License therefore permits such linking only if the entire combination fits its criteria of freedom. The Lesser General Public License permits more lax criteria for linking other code with the library.

We call this license the *Lesser* General Public License because it does Less to protect the user's freedom than the ordinary General Public License. It also provides other free software developers Less of an advantage over competing non-free programs. These disadvantages are the reason we use the ordinary General Public License for many libraries. However, the Lesser license provides advantages in certain special circumstances.

For example, on rare occasions, there may be a special need to encourage the widest possible use of a certain library, so that it becomes a de-facto standard. To achieve this, non-free programs must be allowed to use the library. A more frequent case is that a free library does the same job as widely used non-free libraries. In this case, there is little to gain by limiting the free library to free software only, so we use the Lesser General Public License.

In other cases, permission to use a particular library in non-free programs enables a greater number of people to use a large body of free software. For example, permission to use the GNU C Library in non-free programs enables many more people to use the whole GNU operating system, as well as its variant, the GNU/Linux operating system.

Although the Lesser General Public License is Less protective of the users' freedom, it does ensure that the user of a program that is linked with the Library has the freedom and the wherewithal to run that program using a modified version of the Library.

The precise terms and conditions for copying, distribution and modification follow. Pay close attention to the difference between a "work based on the library" and a "work that uses the library". The former contains code derived from the library, whereas the latter must be combined with the library in order to run.

C.2. TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

C.2.1. Section 0

This License Agreement applies to any software library or other program which contains a notice placed by the copyright holder or other authorized party saying it may be distributed under the terms of this Lesser General Public License (also called "this License"). Each licensee is addressed as "you".

A "library" means a collection of software functions and/or data prepared so as to be conveniently linked with application programs (which use some of those functions and data) to form executables.

The "Library", below, refers to any such software library or work which has been distributed under these terms. A "work based on the Library" means either the Library or any derivative work under copyright law: that is to say, a work containing the Library or a portion of it, either verbatim or with modifications and/or translated straightforwardly into another language. (Hereinafter, translation is included without limitation in the term "modification".)

"Source code" for a work means the preferred form of the work for making modifications to it. For a library, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the library.

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running a program using the Library is not restricted, and output from such a program is covered only if its contents constitute a work based on the Library (independent of the use of the Library in a tool for writing it). Whether that is true depends on what the Library does and what the program that uses the Library does.

C.2.2. Section 1

You may copy and distribute verbatim copies of the Library's complete source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and distribute a copy of this License along with the Library.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

C.2.3. Section 2

You may modify your copy or copies of the Library or any portion of it, thus forming a work based on the Library, and copy and distribute such modifications or work under the terms of *Section 1* above, provided that you also meet all of these conditions:

    a. The modified work must itself be a software library.

    b. You must cause the files modified to carry prominent notices stating that you changed the files and the date of any change.

    c. You must cause the whole of the work to be licensed at no charge to all third parties under the terms of this License.

d. If a facility in the modified Library refers to a function or a table of data to be supplied by an application program that uses the facility, other than as an argument passed when the facility is invoked, then you must make a good faith effort to ensure that, in the event an application does not supply such function or table, the facility still operates, and performs whatever part of its purpose remains meaningful.

(For example, a function in a library to compute square roots has a purpose that is entirely well-defined independent of the application. Therefore, *Subsection 2d* requires that any application-supplied function or table used by this function must be optional: if the application does not supply it, the square root function must still compute square roots.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Library, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Library, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Library.

In addition, mere aggregation of another work not based on the Library with the Library (or with a work based on the Library) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

C.2.4. Section 3

You may opt to apply the terms of the ordinary GNU General Public License instead of this License to a given copy of the Library. To do this, you must alter all the notices that refer to this License, so that they refer to the ordinary GNU General Public License, version 2, instead of to this License. (If a newer version than version 2 of the ordinary GNU General Public License has appeared, then you can specify that version instead if you wish.) Do not make any other change in these notices.

Once this change is made in a given copy, it is irreversible for that copy, so the ordinary GNU General Public License applies to all subsequent copies and derivative works made from that copy.

This option is useful when you wish to copy part of the code of the Library into a program that is not a library.

C.2.5. Section 4

You may copy and distribute the Library (or a portion or derivative of it, under *Section 2*) in object code or executable form under the terms of *Sections 1* and *2* above provided that you accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of *Sections 1* and *2* above on a medium customarily used for software interchange.

If distribution of object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place satisfies the requirement to distribute the source code, even though third parties are not compelled to copy the source along with the object code.

C.2.6. Section 5

A program that contains no derivative of any portion of the Library, but is designed to work with the Library by being compiled or linked with it, is called a "work that uses the Library". Such a work, in isolation, is not a derivative work of the Library, and therefore falls outside the scope of this License.

However, linking a "work that uses the Library" with the Library creates an executable that is a derivative of the Library (because it contains portions of the Library), rather than a "work that uses the library". The executable is therefore covered by this License. *Section 6* states terms for distribution of such executables.

When a "work that uses the Library" uses material from a header file that is part of the Library, the object code for the work may be a derivative work of the Library even though the source code is not. Whether this is true is especially significant if the work can be linked without the Library, or if the work is itself a library. The threshold for this to be true is not precisely defined by law.

If such an object file uses only numerical parameters, data structure layouts and accessors, and small macros and small inline functions (ten lines or less in length), then the use of the object file is unrestricted, regardless of whether it is legally a derivative work. (Executables containing this object code plus portions of the Library will still fall under *Section 6*.)

Otherwise, if the work is a derivative of the Library, you may distribute the object code for the work under the terms of *Section 6*. Any executables containing that work also fall under *Section 6*, whether or not they are linked directly with the Library itself.

C.2.7. Section 6

As an exception to the Sections above, you may also combine or link a "work that uses the Library" with the Library to produce a work containing portions of the Library, and distribute that work under terms of your choice, provided that the terms permit modification of the work for the customer's own use and reverse engineering for debugging such modifications.

You must give prominent notice with each copy of the work that the Library is used in it and that the Library and its use are covered by this License. You must supply a copy of this License. If the work during execution displays copyright notices, you must include the copyright notice for the Library among them, as well as a reference directing the user to the copy of this License. Also, you must do one of these things:

a. Accompany the work with the complete corresponding machine-readable source code for the Library including whatever changes were used in the work (which must be distributed under *Sections 1* and *2* above); and, if the work is an executable linked with the Library, with the complete machine-readable "work that uses the Library", as object code and/or source code, so that the user can modify the Library and then relink to produce a modified executable containing the modified Library. (It is understood that the user who changes the contents of definitions files in the Library will not necessarily be able to recompile the application to use the modified definitions.)

b. Use a suitable shared library mechanism for linking with the Library. A suitable mechanism is one that (1) uses at run time a copy of the library already present on the user's computer system, rather than copying library functions into the executable, and (2) will operate properly with a modified version of the library, if the user installs one, as long as the modified version is interface-compatible with the version that the work was made with.

    c. Accompany the work with a written offer, valid for at least three years, to give the same user the materials specified in *Subsection 6a*, above, for a charge no more than the cost of performing this distribution.

    d. If distribution of the work is made by offering access to copy from a designated place, offer equivalent access to copy the above specified materials from the same place.

    e. Verify that the user has already received a copy of these materials or that you have already sent this user a copy.

For an executable, the required form of the "work that uses the Library" must include any data and utility programs needed for reproducing the executable from it. However, as a special exception, the materials to be distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

It may happen that this requirement contradicts the license restrictions of other proprietary libraries that do not normally accompany the operating system. Such a contradiction means you cannot use both them and the Library together in an executable that you distribute.

C.2.8. Section 7

You may place library facilities that are a work based on the Library side-by-side in a single library together with other library facilities not covered by this License, and distribute such a combined library, provided that the separate distribution of the work based on the Library and of the other library facilities is otherwise permitted, and provided that you do these two things:

    a. Accompany the combined library with a copy of the same work based on the Library, uncombined with any other library facilities. This must be distributed under the terms of the Sections above.

    b. Give prominent notice with the combined library of the fact that part of it is a work based on the Library, and explaining where to find the accompanying uncombined form of the same work.

C.2.9. Section 8

You may not copy, modify, sublicense, link with, or distribute the Library except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, link with, or distribute the Library is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

C.2.10. Section 9

You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Library or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Library (or any work based on the Library), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Library or works based on it.

C.2.11. Section 10

Each time you redistribute the Library (or any work based on the Library), the recipient automatically receives a license from the original licensor to copy, distribute, link with or modify the Library subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties with this License.

C.2.12. Section 11

If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Library at all. For example, if a patent license would not permit royalty-free redistribution of the Library by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Library.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply, and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

C.2.13. Section 12

If the distribution and/or use of the Library is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Library under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

C.2.14. Section 13

The Free Software Foundation may publish revised and/or new versions of the Lesser General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Library specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Library does not specify a license version number, you may choose any version ever published by the Free Software Foundation.

C.2.15. Section 14

If you wish to incorporate parts of the Library into other free programs whose distribution conditions are incompatible with these, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

C.2.16. NO WARRANTY Section 15

BECAUSE THE LIBRARY IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE LIBRARY, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE LIBRARY "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE LIBRARY IS WITH YOU. SHOULD THE LIBRARY PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

C.2.17. Section 16

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE LIBRARY AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE LIBRARY (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE LIBRARY TO OPERATE WITH ANY OTHER SOFTWARE), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

C.3. How to Apply These Terms to Your New Libraries

If you develop a new library, and you want it to be of the greatest possible use to the public, we recommend making it free software that everyone can redistribute and change. You can do so by permitting redistribution under these terms (or, alternatively, under the terms of the ordinary General Public License).

To apply these terms, attach the following notices to the library. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

<one line to give the library's name and a brief idea of what it does.> Copyright (C) <year> <name of author>

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2.1 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the Free Software Foundation, Inc., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA

Also add information on how to contact you by electronic and paper mail.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the library, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the library `Frob' (a library for tweaking knobs) written by James Random Hacker.

<signature of Ty Coon>, 1 April 1990 Ty Coon, President of Vice

That's all there is to it!

Appendix D. Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License

THE WORK (AS DEFINED BELOW) IS PROVIDED UNDER THE TERMS OF THIS CREATIVE COMMONS PUBLIC LICENSE ("CCPL" OR "LICENSE"). THE WORK IS PROTECTED BY COPYRIGHT AND/OR OTHER APPLICABLE LAW. ANY USE OF THE WORK OTHER THAN AS AUTHORIZED UNDER THIS LICENSE OR COPYRIGHT LAW IS PROHIBITED. BY EXERCISING ANY RIGHTS TO THE WORK PROVIDED HERE, YOU ACCEPT AND AGREE TO BE BOUND BY THE TERMS OF THIS LICENSE. TO THE EXTENT THIS LICENSE MAY BE CONSIDERED TO BE A CONTRACT, THE LICENSOR GRANTS YOU THE RIGHTS CONTAINED HERE IN CONSIDERATION OF YOUR ACCEPTANCE OF SUCH TERMS AND CONDITIONS.

1. *Definitions*

    a. "Adaptation" means a work based upon the Work, or upon the Work and other pre-existing works, such as a translation, adaptation, derivative work, arrangement of music or other alterations of a literary or artistic work, or phonogram or performance and includes cinematographic adaptations or any other form in which the Work may be recast, transformed, or adapted including in any form recognizably derived from the original, except that a work that constitutes a Collection will not be considered an Adaptation for the purpose of this License. For the avoidance of doubt, where the Work is a musical work, performance or phonogram, the synchronization of the Work in timed-relation with a moving image ("synching") will be considered an Adaptation for the purpose of this License.

    b. "Collection" means a collection of literary or artistic works, such as encyclopedias and anthologies, or performances, phonograms or broadcasts, or other works or subject matter other than works listed in Section 1(f) below, which, by reason of the selection and arrangement of their contents, constitute intellectual creations, in which the Work is included in its entirety in unmodified form along with one or more other contributions, each constituting separate and independent works in themselves, which together are assembled into a collective whole. A work that constitutes a Collection will not be considered an Adaptation (as defined above) for the purposes of this License.

    c. "Distribute" means to make available to the public the original and copies of the Work through sale or other transfer of ownership.

    d. "Licensor" means the individual, individuals, entity or entities that offer(s) the Work under the terms of this License.

    e. "Original Author" means, in the case of a literary or artistic work, the individual, individuals, entity or entities who created the Work or if no individual or entity can be identified, the publisher; and in addition (i) in the case of a performance the actors, singers, musicians, dancers, and other persons who act, sing, deliver, declaim, play in, interpret or otherwise perform literary or artistic works or expressions of folklore; (ii) in the case of a phonogram the producer being the person or legal entity who first fixes the sounds of a performance or other sounds; and, (iii) in the case of broadcasts, the organization that transmits the broadcast.

f. "Work" means the literary and/or artistic work offered under the terms of this License including without limitation any production in the literary, scientific and artistic domain, whatever may be the mode or form of its expression including digital form, such as a book, pamphlet and other writing; a lecture, address, sermon or other work of the same nature; a dramatic or dramatico-musical work; a choreographic work or entertainment in dumb show; a musical composition with or without words; a cinematographic work to which are assimilated works expressed by a process analogous to cinematography; a work of drawing, painting, architecture, sculpture, engraving or lithography; a photographic work to which are assimilated works expressed by a process analogous to photography; a work of applied art; an illustration, map, plan, sketch or three-dimensional work relative to geography, topography, architecture or science; a performance; a broadcast; a phonogram; a compilation of data to the extent it is protected as a copyrightable work; or a work performed by a variety or circus performer to the extent it is not otherwise considered a literary or artistic work.

g. "You" means an individual or entity exercising rights under this License who has not previously violated the terms of this License with respect to the Work, or who has received express permission from the Licensor to exercise rights under this License despite a previous violation.

h. "Publicly Perform" means to perform public recitations of the Work and to communicate to the public those public recitations, by any means or process, including by wire or wireless means or public digital performances; to make available to the public Works in such a way that members of the public may access these Works from a place and at a place individually chosen by them; to perform the Work to the public by any means or process and the communication to the public of the performances of the Work, including by public digital performance; to broadcast and rebroadcast the Work by any means including signs, sounds or images.

i. "Reproduce" means to make copies of the Work by any means including without limitation by sound or visual recordings and the right of fixation and reproducing fixations of the Work, including storage of a protected performance or phonogram in digital form or other electronic medium.

2. *Fair Dealing Rights.* Nothing in this License is intended to reduce, limit, or restrict any uses free from copyright or rights arising from limitations or exceptions that are provided for in connection with the copyright protection under copyright law or other applicable laws.

3. *License Grant.* Subject to the terms and conditions of this License, Licensor hereby grants You a worldwide, royalty-free, non-exclusive, perpetual (for the duration of the applicable copyright) license to exercise the rights in the Work as stated below:

a. to Reproduce the Work, to incorporate the Work into one or more Collections, and to Reproduce the Work as incorporated in the Collections; and,

b. to Distribute and Publicly Perform the Work including as incorporated in Collections.
The above rights may be exercised in all media and formats whether now known or hereafter devised. The above rights include the right to make such modifications as are technically necessary to exercise the rights in other media and formats, but otherwise you have no rights to make Adaptations. Subject to 8(f), all rights not expressly granted by Licensor are hereby reserved, including but not limited to the rights set forth in Section 4(d).

4. *Restrictions.* The license granted in Section 3 above is expressly made subject to and limited by the following restrictions:

a. You may Distribute or Publicly Perform the Work only under the terms of this License. You must include a copy of, or the Uniform Resource Identifier (URI) for, this License with every copy of the Work You Distribute or Publicly Perform. You may not offer or impose any terms on the Work that restrict the terms of this License or the ability of the recipient of the Work to exercise the rights granted to that recipient under the terms of the License. You may not sublicense the Work. You must keep intact all notices that refer to this License and to the disclaimer of warranties with every copy of the Work You Distribute or Publicly Perform. When You Distribute or Publicly Perform the Work, You may not impose any effective technological measures on the Work that restrict the ability of a recipient of the Work from You to exercise the rights granted to that recipient under the terms of the License. This Section 4(a) applies to the Work as incorporated in a Collection, but this does not require the Collection apart from the Work itself to be made subject to the terms of this License. If You create a Collection, upon notice from any Licensor You must, to the extent practicable, remove from the Collection any credit as required by Section 4(c), as requested.

b. You may not exercise any of the rights granted to You in Section 3 above in any manner that is primarily intended for or directed toward commercial advantage or private monetary compensation. The exchange of the Work for other copyrighted works by means of digital file-sharing or otherwise shall not be considered to be intended for or directed toward commercial advantage or private monetary compensation, provided there is no payment of any monetary compensation in connection with the exchange of copyrighted works.

c. If You Distribute, or Publicly Perform the Work or Collections, You must, unless a request has been made pursuant to Section 4(a), keep intact all copyright notices for the Work and provide, reasonable to the medium or means You are utilizing: (i) the name of the Original Author (or pseudonym, if applicable) if supplied, and/or if the Original Author and/or Licensor designate another party or parties (for example a sponsor institute, publishing entity, journal) for attribution ("Attribution Parties") in Licensor's copyright notice, terms of service or by other reasonable means, the name of such party or parties; (ii) the title of the Work if supplied; (iii) to the extent reasonably practicable, the URI, if any, that Licensor specifies to be associated with the Work, unless such URI does not refer to the copyright notice or licensing information for the Work. The credit required by this Section 4(c) may be implemented in any reasonable manner; provided, however, that in the case of a Collection, at a minimum such credit will appear, if a credit for all contributing authors of Collection appears, then as part of these credits and in a manner at least as prominent as the credits for the other contributing authors. For the avoidance of doubt, You may only use the credit required by this Section for the purpose of attribution in the manner set out above and, by exercising Your rights under this License, You may not implicitly or explicitly assert or imply any connection with, sponsorship or endorsement by the Original Author, Licensor and/or Attribution Parties, as appropriate, of You or Your use of the Work, without the separate, express prior written permission of the Original Author, Licensor and/or Attribution Parties.

d. For the avoidance of doubt:

    i. Non-waivable Compulsory License Schemes. In those jurisdictions in which the right to collect royalties through any statutory or compulsory licensing scheme cannot be waived, the Licensor reserves the exclusive right to collect such royalties for any exercise by You of the rights granted under this License;

    ii. Waivable Compulsory License Schemes. In those jurisdictions in which the right to collect royalties through any statutory or compulsory licensing scheme can be waived, the Licensor reserves the exclusive right to collect such royalties for any exercise by You of the rights

granted under this License if Your exercise of such rights is for a purpose or use which is otherwise than noncommercial as permitted under Section 4(b) and otherwise waives the right to collect royalties through any statutory or compulsory licensing scheme; and,

    iii. Voluntary License Schemes. The Licensor reserves the right to collect royalties, whether individually or, in the event that the Licensor is a member of a collecting society that administers voluntary licensing schemes, via that society, from any exercise by You of the rights granted under this License that is for a purpose or use which is otherwise than noncommercial as permitted under Section 4(b).

  e. Except as otherwise agreed in writing by the Licensor or as may be otherwise permitted by applicable law, if You Reproduce, Distribute or Publicly Perform the Work either by itself or as part of any Collections, You must not distort, mutilate, modify or take other derogatory action in relation to the Work which would be prejudicial to the Original Author's honor or reputation.

5. *Representations, Warranties and Disclaimer* UNLESS OTHERWISE MUTUALLY AGREED BY THE PARTIES IN WRITING, LICENSOR OFFERS THE WORK AS-IS AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE WORK, EXPRESS, IMPLIED, STATUTORY OR OTHERWISE, INCLUDING, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTIBILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, OR THE ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OF ABSENCE OF ERRORS, WHETHER OR NOT DISCOVERABLE. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO SUCH EXCLUSION MAY NOT APPLY TO YOU.

6. *Limitation on Liability.* EXCEPT TO THE EXTENT REQUIRED BY APPLICABLE LAW, IN NO EVENT WILL LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, PUNITIVE OR EXEMPLARY DAMAGES ARISING OUT OF THIS LICENSE OR THE USE OF THE WORK, EVEN IF LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

7. *Termination*

  a. This License and the rights granted hereunder will terminate automatically upon any breach by You of the terms of this License. Individuals or entities who have received Collections from You under this License, however, will not have their licenses terminated provided such individuals or entities remain in full compliance with those licenses. Sections 1, 2, 5, 6, 7, and 8 will survive any termination of this License.

  b. Subject to the above terms and conditions, the license granted here is perpetual (for the duration of the applicable copyright in the Work). Notwithstanding the above, Licensor reserves the right to release the Work under different license terms or to stop distributing the Work at any time; provided, however that any such election will not serve to withdraw this License (or any other license that has been, or is required to be, granted under the terms of this License), and this License will continue in full force and effect unless terminated as stated above.

8. *Miscellaneous*

  a. Each time You Distribute or Publicly Perform the Work or a Collection, the Licensor offers to the recipient a license to the Work on the same terms and conditions as the license granted to You under this License.

  b. If any provision of this License is invalid or unenforceable under applicable law, it shall not affect the validity or enforceability of the remainder of the terms of this License, and without further

action by the parties to this agreement, such provision shall be reformed to the minimum extent necessary to make such provision valid and enforceable.

c. No term or provision of this License shall be deemed waived and no breach consented to unless such waiver or consent shall be in writing and signed by the party to be charged with such waiver or consent.

d. This License constitutes the entire agreement between the parties with respect to the Work licensed here. There are no understandings, agreements or representations with respect to the Work not specified here. Licensor shall not be bound by any additional provisions that may appear in any communication from You. This License may not be modified without the mutual written agreement of the Licensor and You.

e. The rights granted under, and the subject matter referenced, in this License were drafted utilizing the terminology of the Berne Convention for the Protection of Literary and Artistic Works (as amended on September 28, 1979), the Rome Convention of 1961, the WIPO Copyright Treaty of 1996, the WIPO Performances and Phonograms Treaty of 1996 and the Universal Copyright Convention (as revised on July 24, 1971). These rights and subject matter take effect in the relevant jurisdiction in which the License terms are sought to be enforced according to the corresponding provisions of the implementation of those treaty provisions in the applicable national law. If the standard suite of rights granted under applicable copyright law includes additional rights not granted under this License, such additional rights are deemed to be included in the License; this License is not intended to restrict the license of any rights under applicable law.

Glossary

alias IP	An additional IP address assigned to an interface that already has an IP address. The normal and alias IP addresses both refer to the same physical interface.
authentication	The process of verifying the authenticity of a user or client before allowing access to a network system or service.
auditing policy	The auditing policy determines which events are logged on host running Microsoft Windows operating systems.
BOM	The byte order mark (BOM) is a Unicode character used to signal the byte-order of the message text.
BSD-syslog protocol	The old syslog protocol standard described in _RFC 3164_. Sometimes also referred to as the legacy-syslog protocol.
CA	A Certificate Authority (CA) is an institute that issues certificates.
certificate	A certificate is a file that uniquely identifies its owner. Certificates contains information identifying the owner of the certificate, a public key itself, the expiration date of the certificate, the name of the CA that signed the certificate, and some other data.
client mode	In client mode, syslog-ng collects the local logs generated by the host and forwards them through a network connection to the central syslog-ng server or to a relay.
destination	A named collection of configured destination drivers.
destination driver	A communication method used to send log messages.
destination, network	A destination that sends log messages to a remote host (that is, a syslog-ng relay or server) using a network connection.
destination, local	A destination that transfers log messages within the host, for example writes them to a file, or passes them to a log analyzing application.
disk queue	See _disk buffer_.
domain name	The name of a network, for example: `balabit.com`.
embedded log statement	A log statement that is included in another log statement to create a complex log path.
filter	An expression to select messages.

fully qualified domain name (FQDN)	A domain name that specifies its exact location in the tree hierarchy of the Domain Name System (DNS). For example, given a device with a local hostname `myhost` and a parent domain name `example.com`, the fully qualified domain name is `myhost.example.com`.
gateway	A device that connects two or more parts of the network, for example: your local intranet and the external network (the Internet). Gateways act as entrances into other networks.
high availability	High availability uses a second syslog-ng server unit to ensure that the logs are received even if the first unit breaks down.
host	A computer connected to the network.
hostname	A name that identifies a host on the network.
IETF-syslog protocol	The syslog-protocol standard developed by the Internet Engineering Task Force (IETF), described in *RFC 5424-5427*.
key pair	A private key and its related public key. The private key is known only to the owner; the public key can be freely distributed. Information encrypted with the private key can only be decrypted using the public key.
log path	A combination of sources, filters, parsers, rewrite rules, and destinations: syslog-ng examines all messages arriving to the sources of the logpath and sends the messages matching all filters to the defined destinations.
LSH	See *log source host*.
log source host	A host or network device (including syslog-ng clients and relays) that sends logs to the syslog-ng server. Log source hosts can be servers, routers, desktop computers, or other devices capable of sending syslog messages or running syslog-ng.
log statement	See *log path*.
name server	A network computer storing the IP addresses corresponding to domain names.
Oracle Instant Client	The Oracle Instant Client is a small set of libraries, which allow you to connect to an Oracle Database. A subset of the full Oracle Client, it requires minimal installation but has full functionality.
output buffer	A part of the memory of the host where syslog-ng stores outgoing log messages if the destination cannot accept the messages immediately.
output queue	Messages from the output queue are sent to the target syslog-ng server. The syslog-ng application puts the outgoing messages directly into the output queue, unless the output queue is full. The output queue can hold 64 messages, this is a fixed value and cannot be modified.

overflow queue	See *output buffer*.
parser	A set of rules to segment messages into named fields or columns.
ping	A command that sends a message from a host to another host over a network to test connectivity and packet loss.
port	A number ranging from 1 to 65535 that identifies the destination application of the transmitted data. For example: SSH commonly uses port 22, web servers (HTTP) use port 80, and so on.
Public-key authentication	An authentication method that uses encryption key pairs to verify the identity of a user or a client.
regular expression	A regular expression is a string that describes or matches a set of strings. The syslog-ng application supports extended regular expressions (also called POSIX modern regular expressions).
relay mode	In relay mode, syslog-ng receives logs through the network from syslog-ng clients and forwards them to the central syslog-ng server using a network connection.
rewrite rule	A set of rules to modify selected elements of a log message.
template	A user-defined structure that can be used to restructure log messages or automatically generate file names.
server mode	In server mode, syslog-ng acts as a central log-collecting server. It receives messages from syslog-ng clients and relays over the network, and stores them locally in files, or passes them to other applications, for example, log analyzers.
source	A named collection of configured source drivers.
source, network	A source that receives log messages from a remote host using a network connection. The following sources are network sources: `tcp()`, `tcp6()`, `udp()`, `udp6()`.
source, local	A source that receives log messages from within the host, for example, from a file.
source driver	A communication method used to receive log messages.
SSL	See *TLS*.
syslog-ng	The syslog-ng application is a flexible and highly scalable system logging application, typically used to manage log messages and implement centralized logging.
syslog-ng agent	The syslog-ng Agent for Windows is a commercial log collector and forwarder application for the Microsoft Windows platform. It collects the log messages

of the Windows-based host and forwards them to a syslog-ng server using regular or SSL-encrypted TCP connections.

syslog-ng client	A host running syslog-ng in client mode.
syslog-ng Premium Edition	The syslog-ng Premium Edition is the commercial version of the open-source application. It offers additional features, like encrypted message transfer and an agent for Microsoft Windows platforms.
syslog-ng relay	A host running syslog-ng in relay mode.
syslog-ng server	A host running syslog-ng in server mode.
TLS	Transport Layer Security (TLS) and its predecessor, Secure Sockets Layer (SSL), are cryptographic protocols which provide secure communications on the Internet. The syslog-ng Open Source Edition application can encrypt the communication between the clients and the server using TLS to prevent unauthorized access to sensitive log messages.
traceroute	A command that shows all routing steps (the path of a message) between two hosts.
UNIX domain socket	A UNIX domain socket (UDS) or IPC socket (inter-procedure call socket) is a virtual socket, used for inter-process communication.

Index

Symbols

flags(no-multi-line), 59, 75
flags(no-parse), 16, 267
flow-control, 198, 203, 206
   example, 207
   hard, 206
   multiple destinations, 206
   soft, 205
flow_control, 206
flush-lines(), 65, 73, 81, 85, 92, 105, 113, 123-124, 135-136, 143-144, 149, 168-169, 178, 186, 193, 220, 226, 300, 307
flush-timeout(), 123, 135, 143, 149, 168-169, 178, 186, 193, 220, 300
follow-freq(), 41, 54, 56, 72, 78, 85, 326
follow-freq(1), 96
foo bar, 313
foo bar message, 313
foreground, 322
format(linux-kmsg), 96
format-json, 17, 21, 25, 249, 273
   compiling, 25
format-json(), 17
formatting messages, 238
formatting multi-line messages, 59, 75
frac-digits(), 124, 131, 136, 144, 149, 169, 178, 186, 194, 220, 227, 243
from(), 158, 160
fsync(), 124
FULLDATE, 241
FULLHOST, 241, 261
FULLHOST_FROM, 241

G

generating alerts, 283
geoip, 25
   compiling, 25
GEOIP, LENGTH, STRIP, SUBSTR, TFHASH, UUID, xix
glob patterns, 265
global objects, 7
global options, **217**
   reference, 217
global variables, 45
gmake, 24
graphite(), 128
graphite-output, xvii, 128, 250
greedy, 269-271
greedy(), 287

grep, 250
group(), 124, 144
groupset(), 260

H

hard macros, 16, 241
HEADER, 11, 13
header(), 158, 160
HOST, 64, 90-91, 104, 215, 221, 241, 261
host, 164
host(), 158, 202, 210, 265, 267
HOST_FROM, 241
HOUR, 241

I

in-list, 214
in-list filter, 213
indenting multi-line messages, 59, 75
indexes, 170
indexes(), 170
info, 214
inherit-properties, 283
inject-mode(), 283
installing syslog-ng, 23
installing syslog-ng OSE from source, 23
int, 18
int32, 18
int64, 18
internal, 283
internal(), 52-53, 132, 224, 283, 294
ip-protocol(), xviii
ISODATE, 241

J

JavaScript Object Notation, 249
JSON, 249
   Common Information Model (CIM), 96
JSON parsers, 272
json-c, 25
json-parser, 25
   compiling, 25
json-parser(), 272
junction, 44
junctions, 201
   and flags, 201

K

keep-alive, 64, 90, 103, 136, 179, 187, 194
keep-hostname(), 38-39, 223, 240, 243
keep-timestamp(), 9
kern, 54, 242
kernel, 55, 62, 71, 79, 84, 88, 102, 111
key(), 17, 19, 21, 129
klogd, 54
kmsg, 54, 96
ksymoops, 54

L

LEGACY_MSGHDR, 241
LEVEL, 241
level(), 214
LEVEL_NUM, 241
libdbi, 26
libgeoip, 25
libopenssl, 26
libpcre, 23
libsystemd-daemon, 26
libwrap, 26
link-level-address, 288
literal, 18
local time, 12, 15
local-time-zone(), 27
localip(), 100
log, 41, 48, 326
log messages, representation, 16
log messages, structure, 10
    BSD-syslog protocol, 11
    IETF-syslog protocol, 13
    legacy-syslog protocol, 11
    RFC 3164, 11
    RFC 5424, 13
log paths, 4, **198**, 325
    defining, 198
    flags, 198, 203
    flow-control, 203, 206-207
log pipes
    embedded log statements, 199
log statements, 8
    embedded, **199**
    log paths, 4, 325
log statistics, 297
    on unix-socket, 297

log-fetch-limit(), 57, 65, 72, 81, 85, 91, 99, 104, 112, 203, 206-207, 300
log-fifo-size(), 42, 203, 205-207
log-fifo-size(2Mb), 42
log-iw-size(), 57, 65, 73, 81, 85, 92, 105, 113, 123, 135, 143, 149, 169, 178, 186, 193, 203, 206-207, 300
log-msg-size(), 11, 57, 65, 73, 81, 85, 92, 105, 113, 164-165
log-prefix(), 61, 67, 77, 82, 86, 93, 106, 114
logging procedure, 4
logrotate, 121
losing messages, 303

M

macros, 8, 238
    date-related, 240
    default value, 239
    hard, 16
    hard and soft macros, 241
    in filenames, 240
    patterndb tags, 247
    read-only, 16
    reference, 241
    rewritable, 16
    SDATA, 245
    soft, 16
make, 24
manipulating tags (see modifying tags)
MARK, 125-126, 137-138, 145, 150, 179-180, 187-188, 195, 222-223
mark(), 137, 179, 222
mark-freq(), 137, 179, 222
mark-mode(), 125-126, 137-138, 145, 150, 179-180, 188, 195, 222-223
match, 208
match(), 208, 210, 214-215, 265-266
max-connections(), 65, 73, 81, 85, 92, 105, 109, 113, 203-204, 206-207, 301
max-field-size(), 99
maximal message size, 222
max_field_size(), 99
message, 313
    facilities, 12, 14
    ID, 246
    statistics, 297
MESSAGE, 241
message classification, 278-279, 289
message context, 281

message correlation, 281
message counters, 297
message facilities, 213
message filtering
   using parsers, 279
message loss, 303
message parsing, 267, 278-279
message statistics, 297
message templates, 238
message triggers, 283
message(), 214
Microsoft SQL
   sql() configuration, 165
Microsoft SQL Server configuration, 27
MIN, 241
modes of operation, 6
   client mode, 6
   relay mode, 7
   server mode, 7
modifying SDATA, 260
modifying tags, 262
modules, 45-46
mongodb (see type-casting)
   failover, 130
   replicasets, 130
mongodb(), xix, 17, 19, 129-131, 134
   compiling, 25-26
MONTH, 241
MONTH_ABBREV, 241
MONTH_NAME, 241
MONTH_WEEK, 241
MSG, 11, 13, 214-215, 241
MSGHDR, 214
MSGID, 241
MSGONLY, 241
mssql, 165, 172
MSSQL
   sql() configuration, 165
multi-line messages, 57-60, 73-76
multi-line-garbage(), 58-60, 74-76
multi-line-mode, xvii
multi-line-mode(), xviii, 58-59, 74-75
multi-line-mode(indented), 59, 75, 96
multi-line-mode(prefix-garbage), 58-59, 74-75
multi-line-mode(prefix-suffix), 58, 74
multi-line-prefix(), 58, 60, 74, 76
multi-line-suffix(), 58-59, 74-75
multiline

indented-multiline, 96
multiline messages (see multi-line messages)
multithreading in syslog-ng OSE, **300**
mutual authentication, 229, 232
myhost, 261
MYSQL_UNIX_PORT, 167, 170

N
name, 294
name resolution, 306-307
   local, 307
network(), xviii, 62-63, 90, 100, 135, 184
no-hostname, 55, 62, 71, 79-80, 84, 88, 102, 111
no-multi-line, 55, 62, 71, 80, 84, 88, 102, 111, 123, 135, 143, 148, 177, 185, 193
no-parse, 56, 62, 71, 80, 84, 89, 102, 111
nobody, 261
nodejs(), 69-70
notice, 214
NULL, 167
null(), 167, 170-171
number of open files, 122
nv-pairs, 22

O
on-error(), 132, 224
optimizing regular expressions, 265
optimizing syslog-ng performance, 306
   regular expressions, 265
options, 8
   reference, 217
or, xvii
Oracle
   sql() configuration, 164
ORACLE_BASE, 164
ORACLE_HOME, 164
ORACLE_SID, 164
other, 299
output buffer, 204, 206
output queue, 205
overflow queue
   output buffer, 205
overriding facility, 50
overriding-original-program-name, 284
overwrite-if-older(), 126
owner(), 126, 145

P

pacct(), 25, 49, 77-78
    compiling, 25
pacctformat, 78
pad-size(), 61, 66, 77, 82, 86, 92, 106, 114
PADD, 311
padding, xvii
pair(), 17, 19
Parameters, xv
parameters
    log-fetch-limit() , 203, 206
    log-fifo-size() , 203, 206
    log-iw-size() , 204, 206
    max-connections() , 204, 206
parser, 41, 48, 326
parsers, 4, 8, 267, 278-279, 325
parsing messages, 267, 278-279, 286
    concepts of, 267
    filtering parsed messages, 279
password, xvii
path(), 133
pattern database, 278-279, 289, 295
    concepts of, 275
    creating parsers, 286
    pattern matching precedence, 277
    structure of, 276
    using the results, 279
pattern database schema, 289
pattern databases
    correlating messages, 281
pattern matching
    procedure of, 277
patterndb
    download, 280
payload, 129
payload(), 129
pdbtool dictionary, xix
performance
    optimizing multithreading, 301
    using multithreading, 300
perm(), 127
persist_only, 69, 95, 108, 227, 308
pid, 171
PID, 241
pipe, 71, 301
pipe(), 60, 70-71, 76, 81, 86, 113, 125, 138, 142-143, 145, 150, 180, 187-188, 195, 222-223
plugins (see modules)

poll(), 56, 72, 78, 85, 226
port(), 158
PostgreSQL
    sql() configuration, 163
prefix(), 99, 272
preventing message loss
    flow-control, 203, 206
PRI, 11, 13, 241
PRIORITY, 241
priority(), xviii
process accounting, 77
processed, 297-298
processing multi-line messages, 57-59, 73-75
program, 79
PROGRAM, 241
program(), 79, 125, 138, 145, 147-148, 150, 180, 187-188, 195, 210, 222-223, 265, 285
program-override(), 57, 66, 73, 81, 86, 92, 105, 113
proto-template, 227
proto-template(), xix
pseudofile(), 152
pseudonymization, 252, 259
p_apache_parser, 44

Q

quote-pairs(), 271

R

RCPTID, 241
read-only macros, 16
reading messages
    from external applications, 79
recv-time-zone(), 9-10
redis(), 152-153
    compiling, 25-26
regular expressions, **208**, **263**, 265, 306
    case-insensitive, 264
    escaping, 264
    pcre, 264
    posix, 210
rekey, xix
rekey(), 21
relay mode, 7
relay-hostname, 217
reload, xx
removing syslog-ng OSE, 27
replace(), xviii, 21
replace-prefix(), xviii

replacing message text, 257

reply-to(), 161

retries, xvii, 119, 133, 154, 156, 162, 171, 175

reusing snippets, 48

rewritable macros, 16

rewrite, 41, 48, 326

rewrite if, 261

rewrite rules, 4, 8, 257, 325

rewriting

    IP addresses, 252, 259

rewriting messages, 257

   concepts of, 257

   conditional rewrites, 261

rfc3164, 22

rfc5424, 22

riemann(), 154-155

rltp, 55, 62, 71, 80, 84, 88, 102, 111

root, 48-49

rotating log files, 121

routing-key(), 117, 173

R_UNIXTIME, 10

S

safe-background, 322

safe-mode(), 130, 133

scaling to multiple CPUs, 300

scl

   system(), 95

scope(), 17, 19-21

SDATA, 241

SEC, 241

secondary messages, 283

sedding messages, 257

segmenting messages, 268, 270, 272

selected-macros, 22

send-time-zone(), 9

sender(), 160

SEQNUM, 241

sequence ID, 245

sequence number, 246

   Cisco, 246

server mode, 7

server(), 130

server-hostname, 217

servers(), 130, 133

session-statements(), xix

setting facility, 50

setting message fields, 259-260

setting multiple fields, 260

silent building, 24

silent rules (see silent building)

silently-drop-message, 132, 224

silently-drop-property, 132, 224

silently-fallback-to-string, 132, 224

skipping messages, 216

smtp(), 23, 158-159

   compiling, 23

so-keepalive(), 68, 94, 107

so-rcvbuf(), 67-68, 88, 93-94, 96, 101, 106, 115, 303, 306

SOCK_DGRAM, 50-52, 109, 117, 192, 328-329

SOCK_STREAM, 50-52, 109, 117, 192, 329

soft macros, 16, 241

source, 41, 48, 326

SOURCE, 241

source drivers, 7, **50**

   file() driver, 53-54

   internal() driver, 52-53

   list of, 52, 329

   network() driver, 62

   nodejs() driver, 70

   pacct() driver, 77

   pipe() driver, 70-71

   program() driver, 79

   reference, 50

   sun-streams() driver, 83

   syslog() driver, 87-88

   system() driver, 95

   systemd-journal() driver, 97

   systemd-syslog() driver, 99

   tcp() driver, 100-101

   tcp6() driver, 100-101

   udp() driver, 100-101

   udp6() driver, 100-101

   unix-dgram() driver, 110

   unix-stream() driver, 110

source(), 201

SOURCEIP, 241

sources, 4, 8, **50**

   defining, 50

   on different platforms, 51

   pacct(), 25

SO_BROADCAST, 67, 93, 139, 181, 189, 194

splitting messages, 268, 270, 272

spoof-source(), 39

spoof_source

W

warning, 214
WEEK, 241
WEEK_DAY, 241
WEEK_DAY_ABBREV, 241
WEEK_DAY_NAME, 241
Winston API, 70

X

xx:xx:xx:..., 288

Y

YEAR, 241
YEAR_DAY, 241

List of syslog-ng OSE parameters

Symbols

credit-card-mask(), 262

crit, 214

csv-parser, 270

csv-parser(), 8, 201

custom-domain(), xvii, 218

C_DATE, 241

D

database(), 163-164, 168

DATE, 22, 241

datetime, 18

DAY, 241

db-parser, 278

db-parser(), 278, 283

dbd-option(), xix

debug, 214, 216

default-facility, 54

default-facility(), 50, 98

default-priority, 54

default-priority(), 50

delimiter, xx

destination, 41, 48, 326

dir-group(), 122

dir-owner(), 122

dir-perm(), 122, 218

dns-cache-hosts(), 308

dont-create-tables, 168

dont-store-legacy-msghdr, 56, 63, 71, 80, 84, 89, 102, 111

door(), 83

dot-nv-pairs, 22

double, 18

drop-invalid, 270

drop-message, 132, 224

drop-property, 132, 224

dropped, 297-298

dynamic, 25

E

email, 287

emerg, 214

empty-lines, 55, 62, 71, 79, 84, 88, 102, 111

err, 214

error, 245

exclude(), 17, 19

expect-hostname, xix, 55, 62, 71, 79, 84, 88, 102, 111

explicit-commits, 168

extract-prefix, xvii

F

FACILITY, 241

facility, 306

facility(), xviii, 212

FACILITY_NUM, 241

fallback, 203, 271

fallback-to-string, 132, 224

file, 123, 278, 301

file(), 41, 53-54, 70, 121-122, 125, 138, 145, 150, 180, 187-188, 195, 222-223, 326

file-template(), xix

filter, 41, 48, 326

filter(), 261

final, 5, 203, 216

flags(), 41, 198-199, 258, 326

flags(no-multi-line), 59, 75

flags(no-parse), 16, 267

flow-control, 198, 203, 206

flow_control, 206

flush-lines(), 65, 73, 81, 85, 92, 105, 113, 123-124, 135-136, 143-144, 149, 168-169, 178, 186, 193, 220, 226, 300, 307

flush-timeout(), 123, 135, 143, 149, 168-169, 178, 186, 193, 220, 300

follow-freq(), 41, 54, 56, 72, 78, 85, 326

follow-freq(1), 96

foo bar, 313

foo bar message, 313

foreground, 322

format(linux-kmsg), 96

format-json, 17, 21, 25, 249, 273

format-json(), 17

frac-digits(), 124, 131, 136, 144, 149, 169, 178, 186, 194, 220, 227, 243

from(), 158, 160

fsync(), 124

FULLDATE, 241

FULLHOST, 241, 261

FULLHOST_FROM, 241

G

geoip, 25

GEOIP, LENGTH, STRIP, SUBSTR, TFHASH, UUID, xix

gmake, 24

graphite(), 128

graphite-output, xvii, 128, 250

tcp-keepalive-time() + tcp-keepalive-intvl() * tcp-keepalive-probes(), 107-108, 140-141, 182-183, 190-191

tcp6, 103

tcp6(), 51, 62, 100-101, 135, 184-185, 229, 235, 328, 357

tcpv6(), 230-231, 233-234

TCP_KEEPCNT, 107-108, 140-141, 182-183, 190-191

TCP_KEEPIDLE, 107-108, 140-141, 182-183, 190-191

TCP_KEEPINTVL, 107-108, 140-141, 182-183, 190-191

template, 41, 326

template(), 272

template-escape(), 239

threaded, 63, 89, 103, 123, 301

throttle, xvii, 304

Thu, 247

time-reap(), 121

time-zone(), 9, 27, 240

tls, 63, 89, 176

tls(), 232-235

to(), 158, 162

trigger, 294

trusted-dn(), 237

trusted-keys(), 237

ts-format(), 12, 15, 128, 142, 147, 152, 184, 192, 197, 246

type(), 163, 210, 258, 264

TZ, 241

TZOFFSET, 241

tztab, 27

U

udp, 176, 301

udp(), 51, 55, 62, 71, 80, 84, 88, 100-102, 105, 111, 125, 135, 138, 145, 150, 180, 184-185, 188, 195, 223, 328, 357

udp6(), 51, 62, 100-101, 105, 135, 184-185, 328, 357

ulimit, 122

unicode, 265

unix-dgram, 50, 60, 76, 81, 86, 113, 303

unix-dgram(), 55, 62, 71, 80, 84, 88, 102, 109-113, 125, 138, 145, 150, 180, 187-188, 192, 195, 222-223

unix-stream, 50, 60, 76, 81, 86, 113, 303

unix-stream(), 42, 109-110, 125, 138, 145, 150, 180, 187-188, 192, 195, 222-223, 327

UNIXTIME, 241

unknown, 279

use-dns(), 39, 64, 90, 104, 221, 242-243, 308

use-fqdn(), 242

use-rcptid, 228, 245

user@example.com, 287

useracct, 316

username, xvii

usertty(), 197, 306

V

validate-utf8, 55-56, 62-63, 71-72, 79-80, 83-84, 88-89, 102-103, 110-111

value, 163, 295

value(), 208, 214

value-pairs, 17

value-pairs(), xviii-xix, 19, 117, 120, 129, 132, 134, 155, 173, 176, 224, 249-250

values, 163

values(), 260

varchar, 165

W

warning, 214

WEEK, 241

WEEK_DAY, 241

WEEK_DAY_ABBREV, 241

WEEK_DAY_NAME, 241

X

xx:xx:xx:..., 288

Y

YEAR, 241

YEAR_DAY, 241